



# Dipartimento di Elettronica e Informazione

## Politecnico di Milano

20133 Milano (Italia)  
Piazza Leonardo da Vinci, 32  
Tel. (+39) 02-2399.3400  
Fax (+39) 02-2399.3411

---

### Advanced Computer Architecture

June 29, 2022

Prof. Donatella Sciuto and Marco Santambrogio

Name
Last Name
Professor:

You have 90' to complete the exam

Problem 1 (6pts)	
Problem 2 (6pts)	
Problem 3 (6pts)	
Question 1 (5pts)	
Question 2 (5pts)	
Question 3 (5pts)	
Total (32pts)	

To pass the exam (the following conditions are in AND):

- The overall score has to be equal or higher than 18
- At least 1 problem and 1 Question has to be correct

## Problem 1

Design and describe (all the decisions have to be motivated) a 1-BHT and a 2-BHT able to execute the following assembly code.

(R0 is set to 100, R1 is set to 0)

LOOP:	LD	F3	0	R0
	ADDD	F1	F3	F3
	MULTD	F2	F3	F1
	ADDI	R1	R1	1000
LOOP2:	LD	F3	0	R1
	MULTD	F2	F2	F3
	SUBI	R1	R1	10
	BNEZ	R1	LOOP2	
	SUBI	R0	R0	10
	BNE	R0	R1	LOOP

The obtained result, in terms of miss predictions, is inline with theoretical characteristics of the two predictors? Please effectively support your answer.

## Answer

LOOP 10 times, LOOP2 steady 100

Worst case:

1BHT, not colliding initialized with NT and NT. 1+1 misprediction for LOOP and (1+1) \* 10 for LOOP2. Total of 22 miss.

2BHT, NT strong and NT strong. 2 + 1 misprediction@LOOP and 2 + (1)\*10@LOOP2. Total of 15 miss.

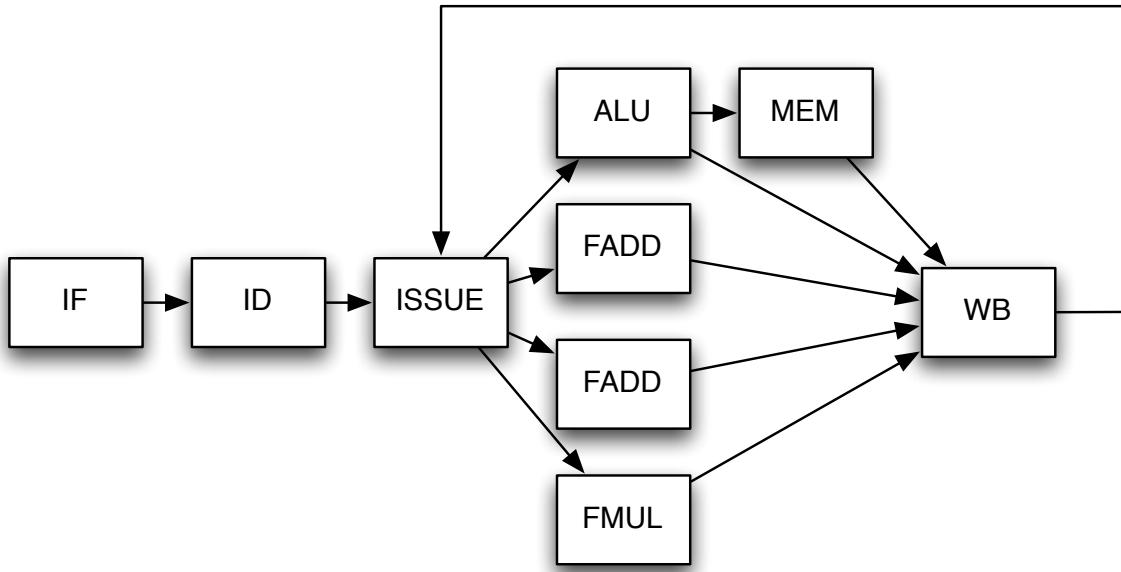
Best case:

1BHT. T and T: 1@LOOP and 1 + 2\*9@LOOP2. Total of 20 miss.

Best of 1BHT is worst than worst-case of 2BHT

## Problem 2

In this problem we will examine the execution of a code segment on the following single-issue out-of-order processor:



You can assume that:

- All functional units are pipelined
- ALU operations take 1 cycle
- Memory operations take 3 cycles (includes time in ALU)
- Floating-point add instructions take 3 cycles
- Floating-point multiply instructions take 5 cycles
- There is no register renaming
- Instructions are fetched, decoded and issued in order
- The issue stage is a buffer of unlimited length that holds instructions waiting to start execution
- An instruction will only enter the issue stage if it does not cause a WAR or WAW hazard
- Only one instruction can be issued at a time, and in the case multiple instructions are ready, the oldest one will go first

To simplify book-keeping for this problem we will only track instructions during a few critical stages:

- I - When an instruction enters the issue stage
- E - When an instruction starts execution (enters FU)
- C - When an instruction completes execution (enters WB, results available this cycle)

For this problem we will be using the following code:

```

I1    L.D   F2, 0(R1)
I2    ADD.D F1, F1, F2
I3    MUL.D F4, F3, F3
I4    ADDI   R1, R1, 8
I5    L.D   F2, 0(R1)
I6    ADD.D F1, F2, F4

```

Fill in the following table to indicate how the given code will execute. Assume all register values are available at the start of execution.

### Problem 3

Consider the following access pattern on a two-processor system with a direct-mapped, write-back cache with one cache block and a two cache block memory. Assume the MESI protocol is used, with write-back caches, write-allocate, and invalidation of other caches on write (instead of updating the value in the other caches).

Time	After Operation	P1 cache block state	P2 cache block state	Memory at block 0 up to date?	Memory at block 1 up to date?
0	P1: read block 0				
1	P2: write block 1				
2	P1: write block 0				
3	P2: read block 1				
4	P1: read block 1				
5	P2 : write block 1				
6	P1: read block 0				
7	P2: write block 1				
8	P1: read block 0				
9	P2: read block 0				

## **Question 1**

The following 5 questions are True/False.  
Each of them will count 1 point

### **Question 1.1**

A superscalar architecture can let instructions from different threads to be executed at the same clock cycle.

Given the previous statement, confirm if it is TRUE or FALSE and **effectively support** your answer.

Circle the **right** answer:    True              False

### **Question 1.2**

The ARM big.LITTLE architecture, considering the fact that is a multicore architecture where the cores are sharing the same ISA, cannot be considered a MIMD architecture but a SIMD one.

Circle the **right** answer:    True              False

### **Question 1.3**

A physically centralized memory can be used to implement a shared-memory architecture

Circle the **right** answer:    True              False

### **Question 1.4**

A VLIW Architecture has to have multiple Program Counters to load the necessary Multiple Data

Circle the **right** answer:    True              False

### **Question 1.5**

In a 5-stage pipeline MIPS architecture, it is possible to forward a data twice from stage-i to stage-j.

As an example, given the following code:

I1 ADDD F2, F4, F6  
I2 ADDD F4, F2, F26  
I3 ADDD F6, F2, F4

F2 will be forwarder from I1 to both I2 and I3 by using the EX->EX forwarding path

Circle the **right** answer:    True              False

## **Question 2**

Explain the difference among fine-grained multithreading, coarse-grained multithreading and simultaneous multithreading.

## **Question 3**

1. Explain the main differences between a VLIW architecture and a superscalar architecture, in terms of instructions scheduling, hardware complexity, and code size.
2. In real superscalar processors, some architectural limits do not allow fully exploitation of the ILP present in the program. Specify these architectural limits and specify how they impact on the potential ILP

1

LOOP: LD F3 0 R0  
 ADDD F1 F3 F3  
 MULTD F2 F3 F1  
 ADDI R1 R1 1000

LOOP2: LD F3 0 R1  
 MULTD F2 F2 F3  
 SUBI R1 R1 10  
 BNEZ R1 LOOP2  
 SUBI R0 R0 10  
 BNE R0 R1 LOOP

R0=100

R1=0

100 ITERATIONS

10 ITERATIONS

1-BHT NO COLLISION

LOOP: LD F3 0 R0  
 ADDD F1 F3 F3  
 MULTD F2 F3 F1  
 ADDI R1 R1 1000

LOOP2: LD F3 0 R1  
 MULTD F2 F2 F3  
 SUBI R1 R1 10  
 BNEZ R1 LOOP2  
 SUBI R0 R0 10  
 BNE R0 R1 LOOP

$$\begin{array}{c} T \rightarrow NT \\ \uparrow \end{array} \quad \begin{array}{c} NT \rightarrow T \\ \uparrow \end{array} \quad \begin{array}{c} T \rightarrow NT \\ \uparrow \end{array}$$

$$1 + (1+1) \cdot (10-1) +$$

$$+ 1 = 20$$

$$\begin{array}{c} T \rightarrow NT \\ \uparrow \end{array} \quad \begin{array}{c} MT \rightarrow T \\ \uparrow \end{array} \quad \begin{array}{c} T \rightarrow NT \\ \uparrow \end{array}$$

$$1 + (1+1) \cdot (10-1) = 19$$

WITH COLLISION T, WE WOULD HAVE GET

2-BHT NO COLLISION

LOOP: LD F3 0 R0  
 ADDD F1 F3 F3  
 MULTD F2 F3 F1  
 ADDI R1 R1 1000

LOOP2: LD F3 0 R1  
 MULTD F2 F2 F3  
 SUBI R1 R1 10  
 BNEZ R1 LOOP2  
 SUBI R0 R0 10  
 BNE R0 R1 LOOP

MT<sub>1</sub> T<sub>1</sub>  
NT<sub>1</sub> T<sub>2</sub>

$$\begin{array}{c} T_1 \rightarrow T_W \\ \uparrow \end{array} \quad \begin{array}{c} T_1 \rightarrow T_W \\ \uparrow \end{array} \quad \text{LAST}$$

$$1 \cdot 10 + 1 = 11$$

$$\begin{array}{c} T \rightarrow NT \\ \uparrow \end{array} \quad \begin{array}{c} T_1 \rightarrow T_W \\ \uparrow \end{array} \quad \begin{array}{c} T \rightarrow NT \\ \uparrow \end{array} \quad \begin{array}{c} T_2 \rightarrow T_W \\ \uparrow \end{array}$$

$$2 + 1 \cdot 10 + 2 + 1 = 15$$

RESULTS MAKE SENSE, AS WORST 2-BHT OUTPERFORMS 1-BHT'S BEST CASE

2

ALU 1cc FP+ 3cc  
MEM 3cc FP. 5cc

I1 L.D F2, 0(R1)  
 I2 ADD.D F1, F1, F2  
 I3 MUL.D F4, F3, F3  
 I4 ADDI R1, R1, 8  
 I5 L.D F2, 0(R1)  
 I6 ADD.D F1, F2, F4

RAW F2 I1-I2

WAW F2 I1-I5

WAR F2 I2-I5 RAW F2 I5-I6

WAW F1 I2-I6 WAW F1 I2-I6 RAW F4 I3-I6

WAR R1 I1-I4 RAW R1 I4-I5

Cycle \ Inst	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
I1	F	D	I	E	E	E	W													
I2		F	D	J	J	J	I	E	E	E	E	W								
I3			F	D	J	J	J	I	E	E	E	E	W							
I4				F	J	D	J	J	I	E	J	W								
I5					F	J	J	J	D	J	I	E	E	E	W					
I6								F	P	J	J	J	J	I	E	E	E	W		

3

Time	After Operation	P1 cache block state	P2 cache block state	Memory at block 0 up to date?	Memory at block 1 up to date?
0	P1: read block 0	E(0)	I	YES	YES
1	P2: write block 1	E(0)	M(1)	YES	NO
2	P1: write block 0	M(0)	M(1)	NO	NO
3	P2: read block 1	M(0)	M(1)	NO	NO
4	P1: read block 1	S(1)	S(1)	YES	YES
5	P2 : write block 1	I	M(1)	YES	NO
6	P1: read block 0	E(0)	M(1)	YES	NO
7	P2: write block 1	E(0)	M(1)	YES	NO
8	P1: read block 0	E(0)	M(1)	YES	NO
9	P2: read block 0	S(0)	S(0)	YES	YES

# THEORY 1

## Question 1.1

A superscalar architecture can let instructions from different threads to be executed at the same clock cycle.

Given the previous statement, confirm if it is TRUE or FALSE and **effectively support** your answer.

Circle the **right** answer: True ✓

False

THANKS TO EXPLOITATION  
OF MULTIPLE P.C

## Question 1.2

The ARM big.LITTLE architecture, considering the fact that is a multicore architecture where the cores are sharing the same ISA, cannot be considered a MIMD architecture but a SIMD one.

Circle the **right** answer: True ✓

False DIFFERENT P.C.

## Question 1.3

A physically centralized memory can be used to implement a shared-memory architecture

Circle the **right** answer: True ✓

False

## Question 1.4

A VLIW Architecture has to have multiple Program Counters to load the necessary Multiple Data

Circle the **right** answer: True ✓

False

VLIW, BY DEFINITION, IS  
A SINGLE INSTRUCTION  
ARCHITECTURE

## Question 1.5

In a 5-stage pipeline MIPS architecture, it is possible to forward a data twice from stage-i to stage-j.

As an example, given the following code:

I1 ADDD F2, F4, F6  
I2 ADDD F4, F2, F26  
I3 ADDD F6, F2, F4

F2 will be forwarder from I1 to both I2 and I3 by using the EX->EX forwarding path

Circle the **right** answer: True ✓

False

E → E TO I3 IBADS TO  
RECEIVE F4

# THEORY 2

## FINE GRAINED PARALLELISM

SWITCH BETWEEN THREADS AT EACH INSTRUCTIONS

- HIDE LONG STALLS BUT SLOWS OVERALL EXECUTION

## COARSE GRAINED MULTITHREADING

SWITCH BETWEEN THREADS WHEN LONG STALLS OCCUR

- OVERALL EXECUTION NOT AFFECTIONED, BUT FOR STEADY STALLS NOTHING IS <sup>MANAGED</sup>

## SIMULTANEOUS MULTITHREADING

DYNAMICALLY ADAPT AND "LEARN" WHEN TO STALL A THREAD, START

ANOTHER ONE AND MOVE BACK TO THE "WAITING LIST THREADS"

3

VLIW

SUPERSCALAR

STATIC SCHEDULING

DYNAMIC SCHEDULING

IN-ORDER ISSUE, PARALLEL EXECUTION

OUT-OF-ORDER ISSUE/EXECUTION

COMPILE-TIME MANAGEMENT

DYNAMIC, RUN-TIME MANAGEMENT

HIGHER COMPLEXITY BECAUSE OF CONTINUOUS ADAPTATION

LONGER CODE, AS INSTRUCTIONS INCLUDES MULTIPLE OPERATIONS AND NOP WHERE "NOTHING HAPPENS"  
THAT MAKES CODE HEAVIER

SIMPLER, PARALLELIZATION  
HANDLED AT RUNTIME

REAL WORLD SUPERSCALAR ARCHITECTURES ARE LIMITED BY:

- LIMIT NUMBER OF F.U.
- RAW, WAR, WAW HAZARDS
- BRANCH PREDICTION AND CONTROL HAZARDS
- LIMITED INSTRUCTION WINDOW SIZES