

Introduction to CUDA: CUDA Programming and Memory Models

Beatrice Branchini
beatrice.branchini@polimi.it

Advanced Computer Architecture Course
T.1.1 - 5/5/2024

Who I Am



Beatrice Branchini



POLITECNICO
MILANO 1863

Ph.D. Information Technology

2022 – Now



Intern DCGPU HPC Solutions

Feb 2025 – Now



Pacific Northwest
NATIONAL LABORATORY

Ph.D. Intern (HPC Group)

Sept 2023 – Mar 2024

Research interests:

- Genomics
- High Performance Computing
- Heterogeneous architectures (GPUs and FPGAs)
- Quantum Computing

More Computing Power

**Financial
Analysis**

**Scientific
Simulation**

**Engineering
Simulation**

**Data Intensive
Analytics**

**Medical
Imaging**

**Digital Audio
Processing**

**Digital Video
Processing**

**Computer
Vision**

**Biomedical
Informatics**

**Electronic
Design
Automation**

**Statistical
Modeling**

**Numerical
Methods**

**Ray Tracing
Rendering**

**Interactive
Physics**

How Can We Program GPUs?

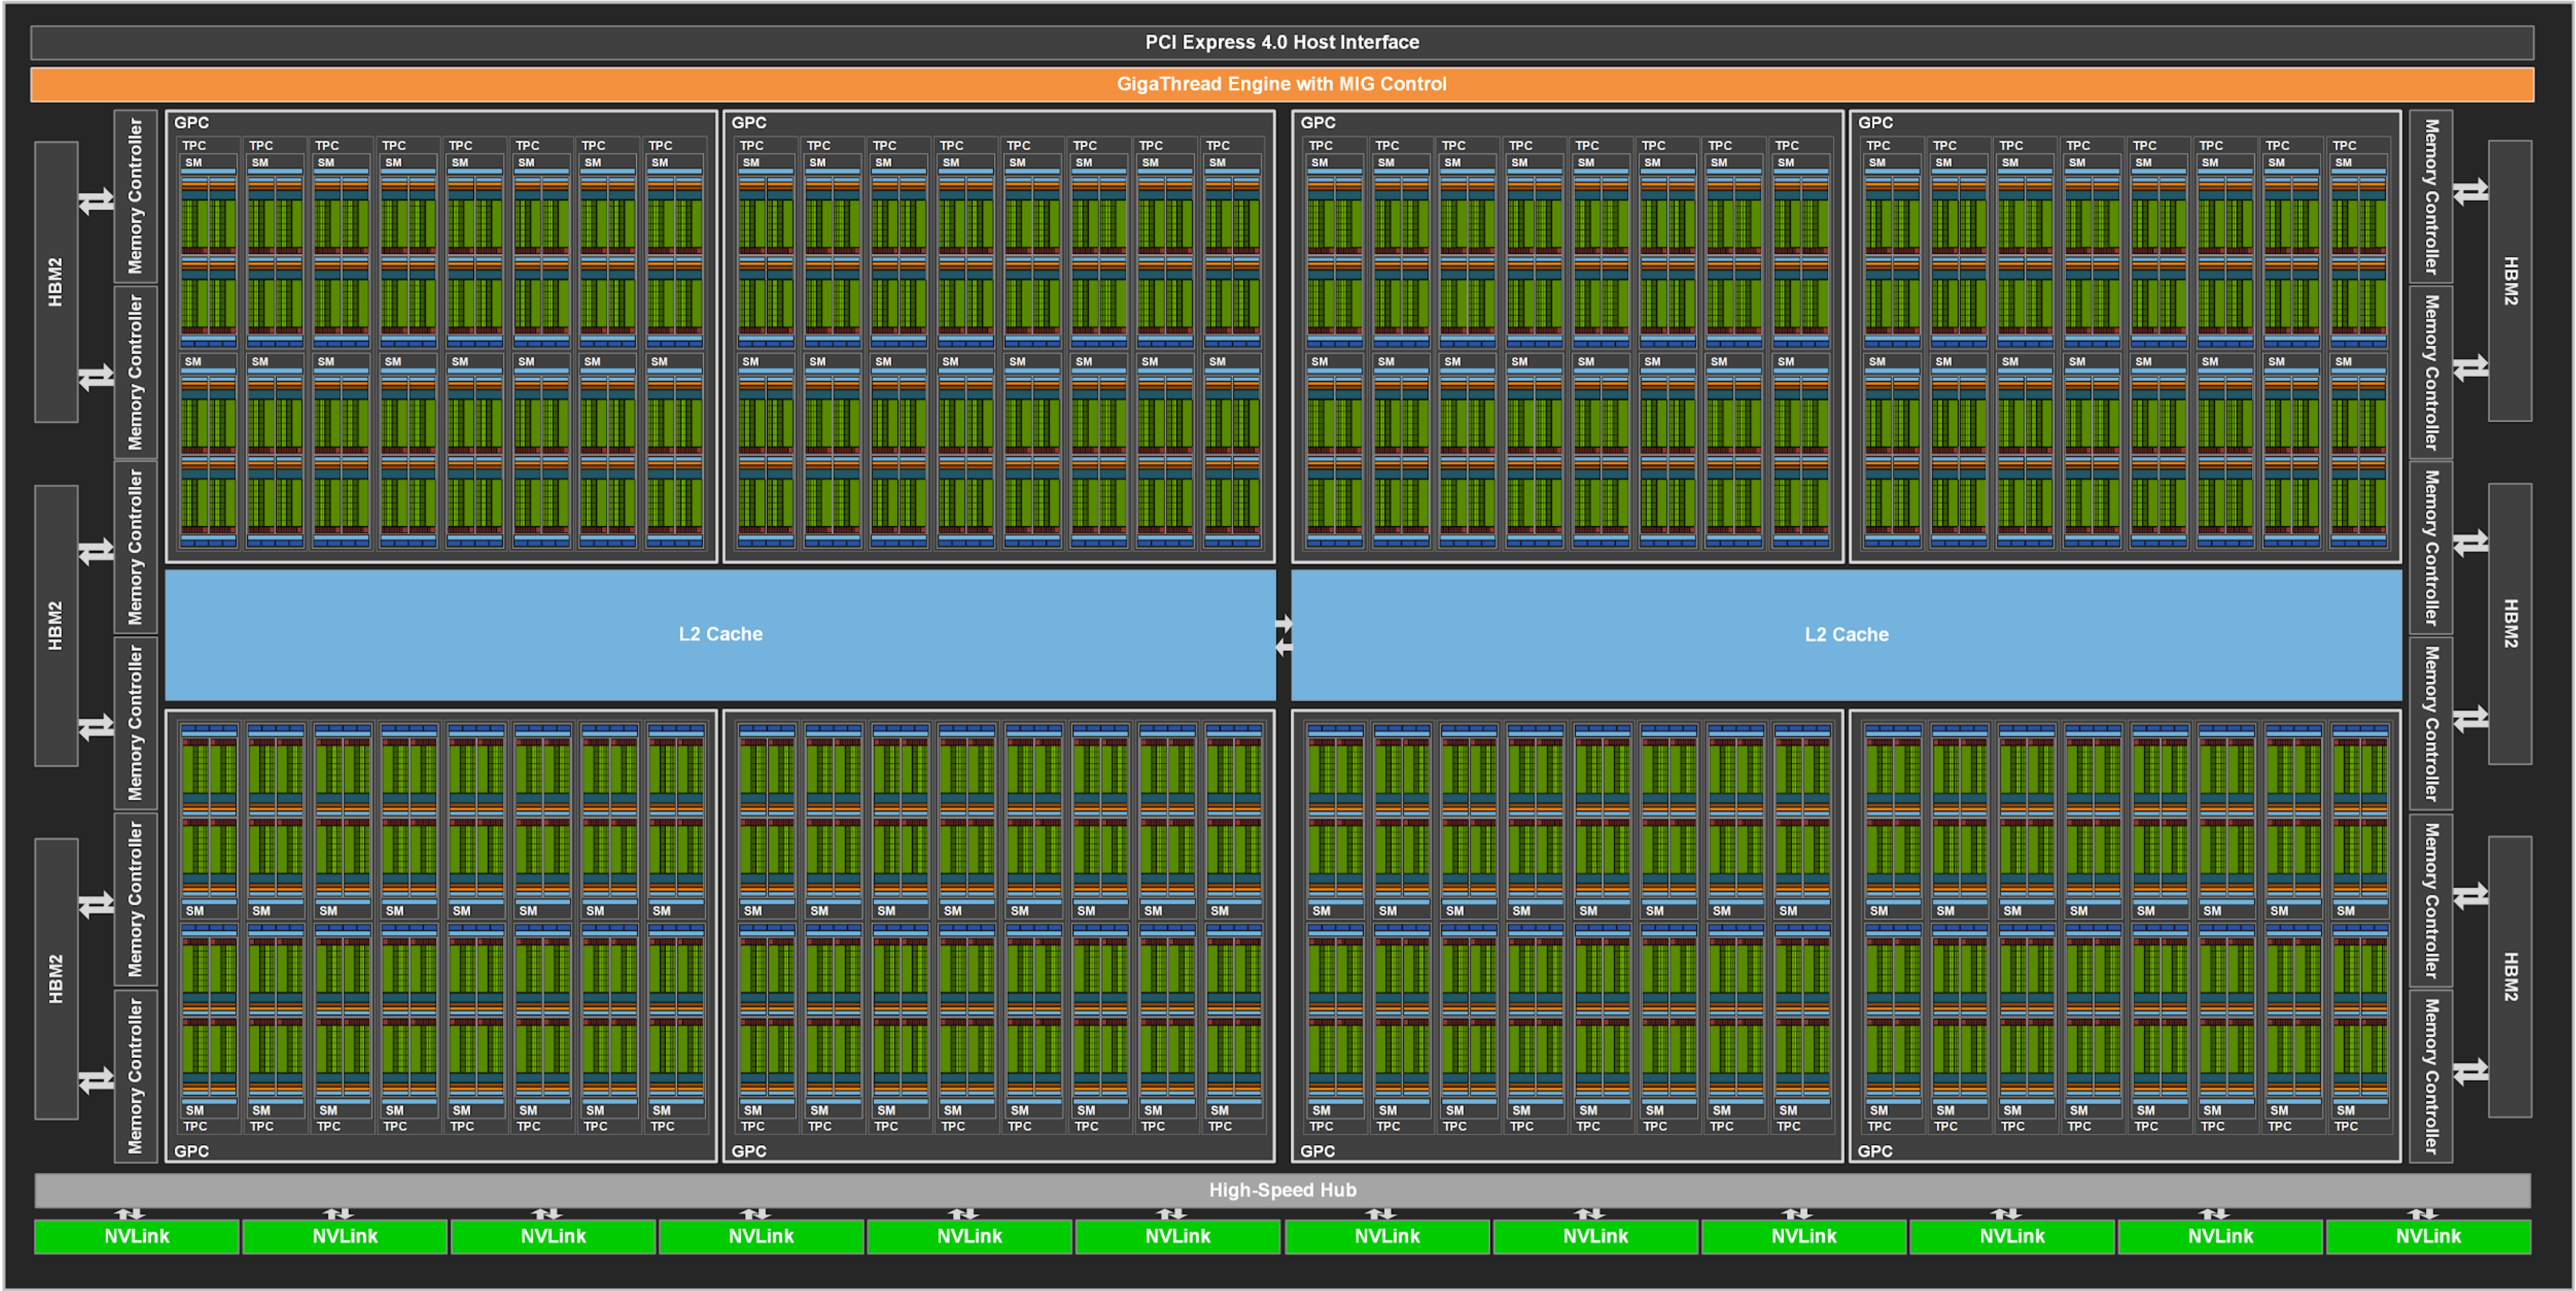
- **CUDA** (Compute Unified Device Architecture) has been introduced in 2006 with NVIDIA Tesla architecture. It offers a **C/C++-like language** which can be used to write programs leveraging GPUs as a processing resource to accelerate functions.
- CUDA's abstraction level closely **match** the **characteristics** and **organization** of NVIDIA GPUs

How Can We Program GPUs?

- **CUDA** (Compute Unified Device Architecture) has been introduced in 2006 with NVIDIA Tesla architecture. It offers a **C/C++-like language** which can be used to write programs leveraging GPUs as a processing resource to accelerate functions.
- CUDA's abstraction level closely **match** the **characteristics** and **organization** of NVIDIA GPUs

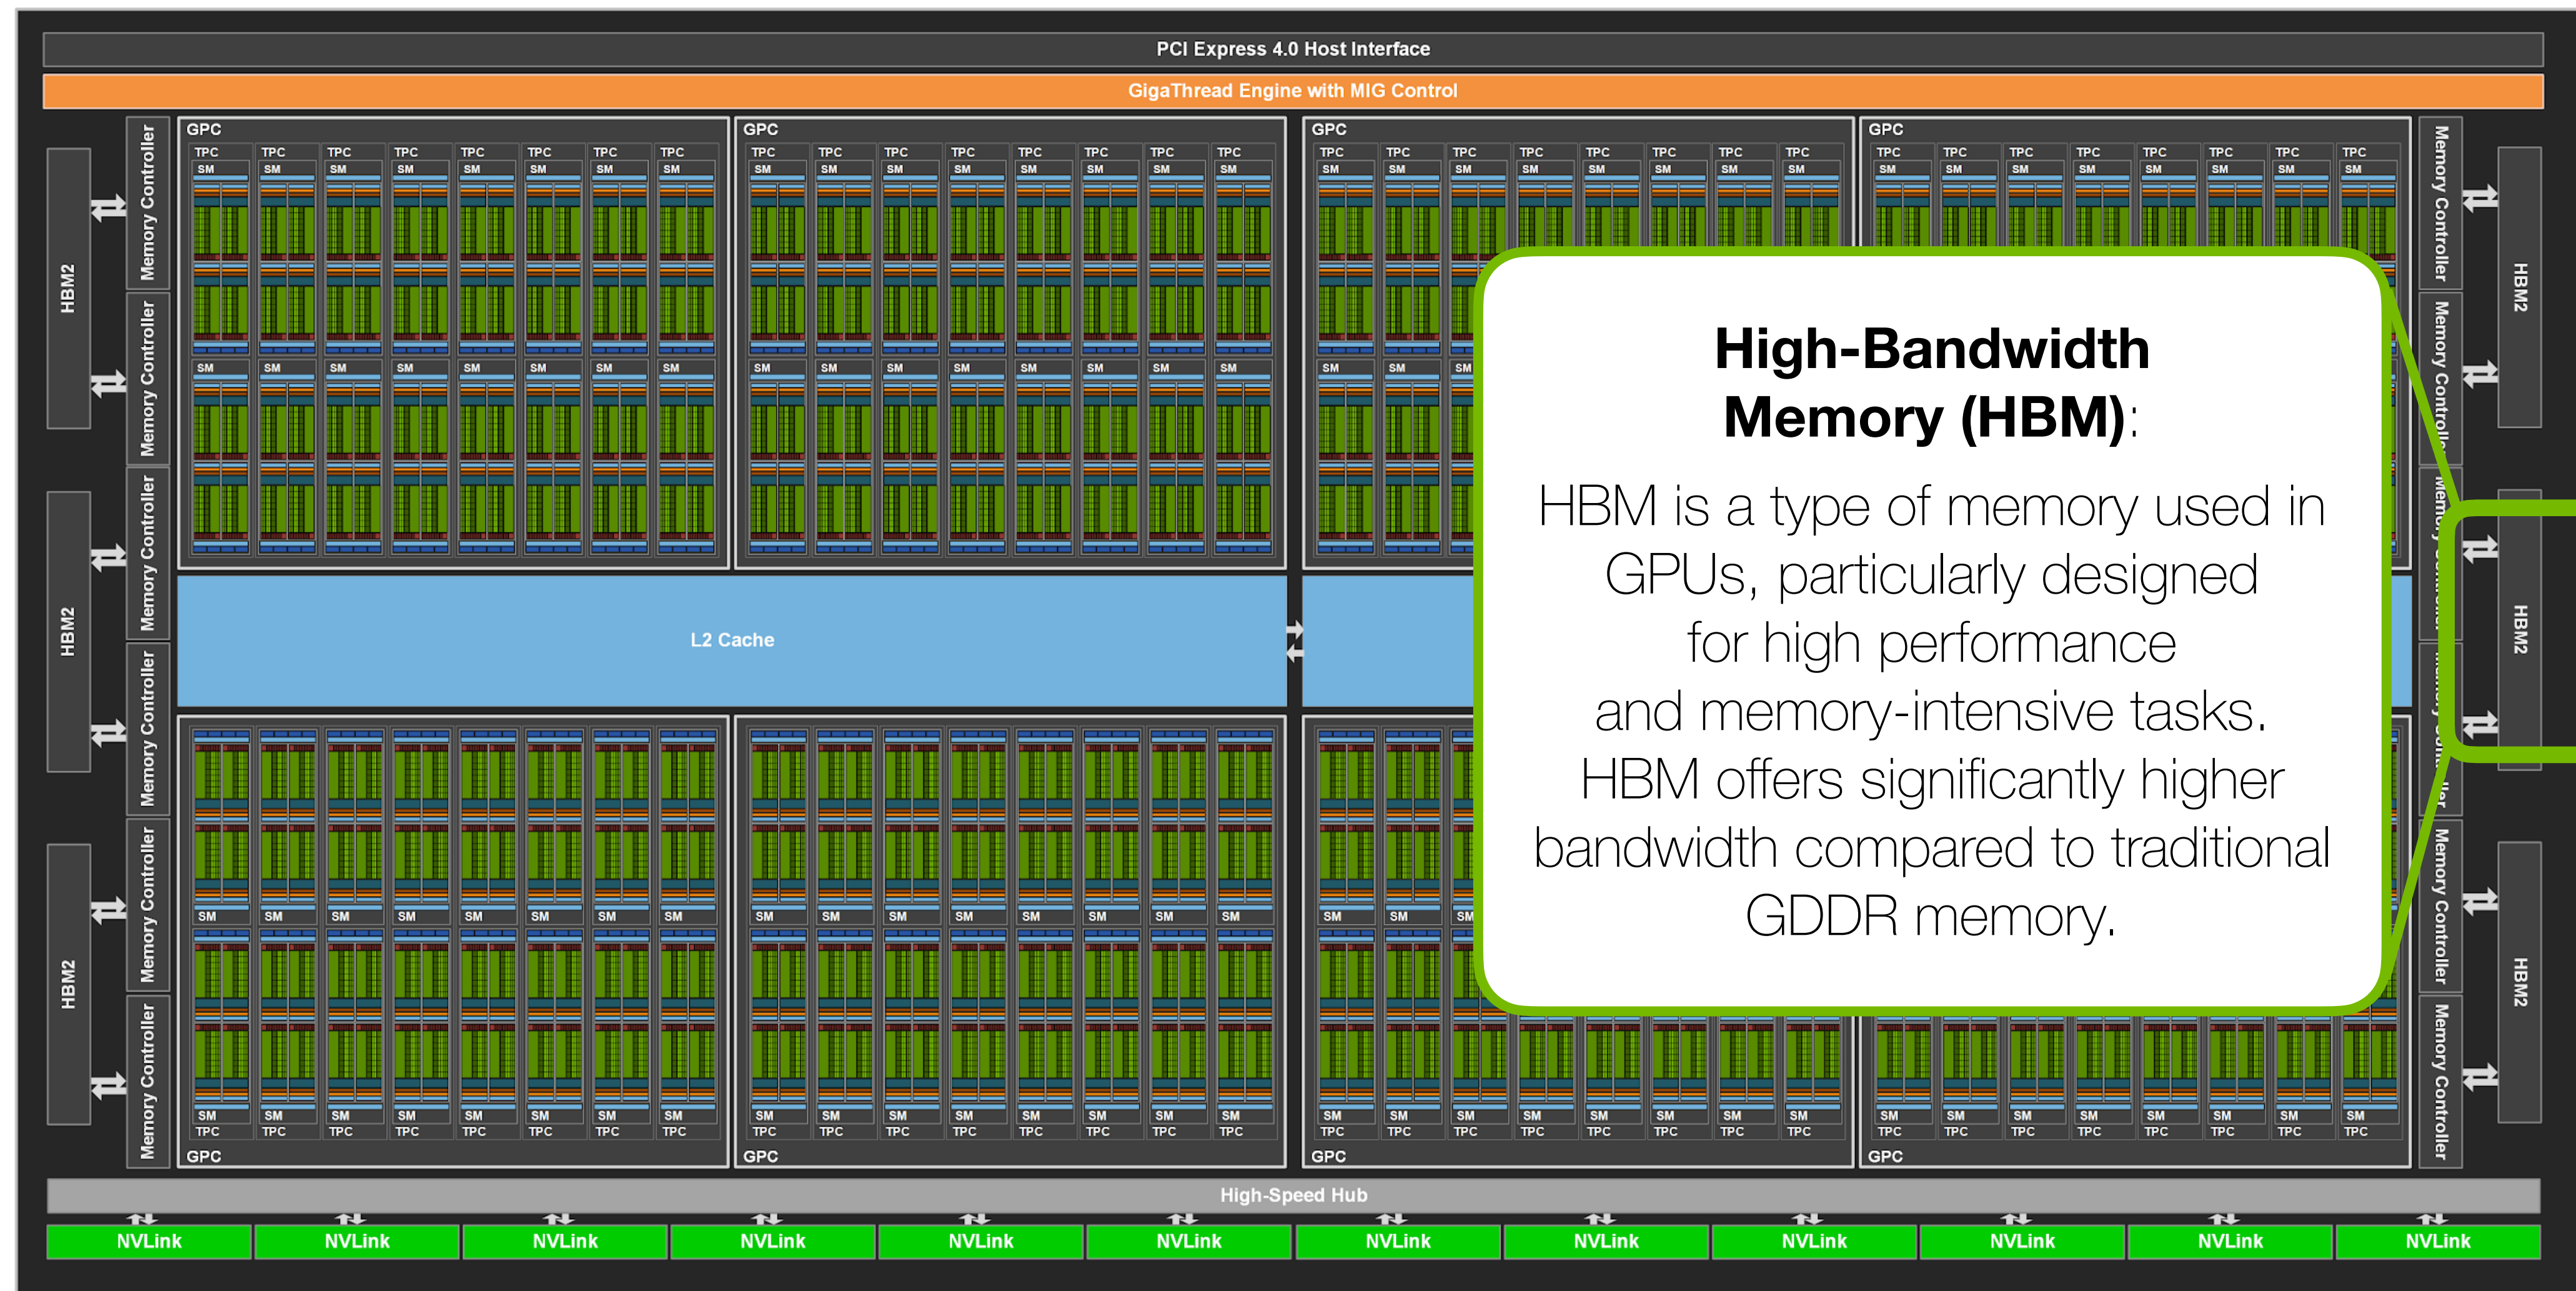
It is paramount to have a deep understanding of the device architecture to exploit GPUs efficiently

Recap: Device Structure



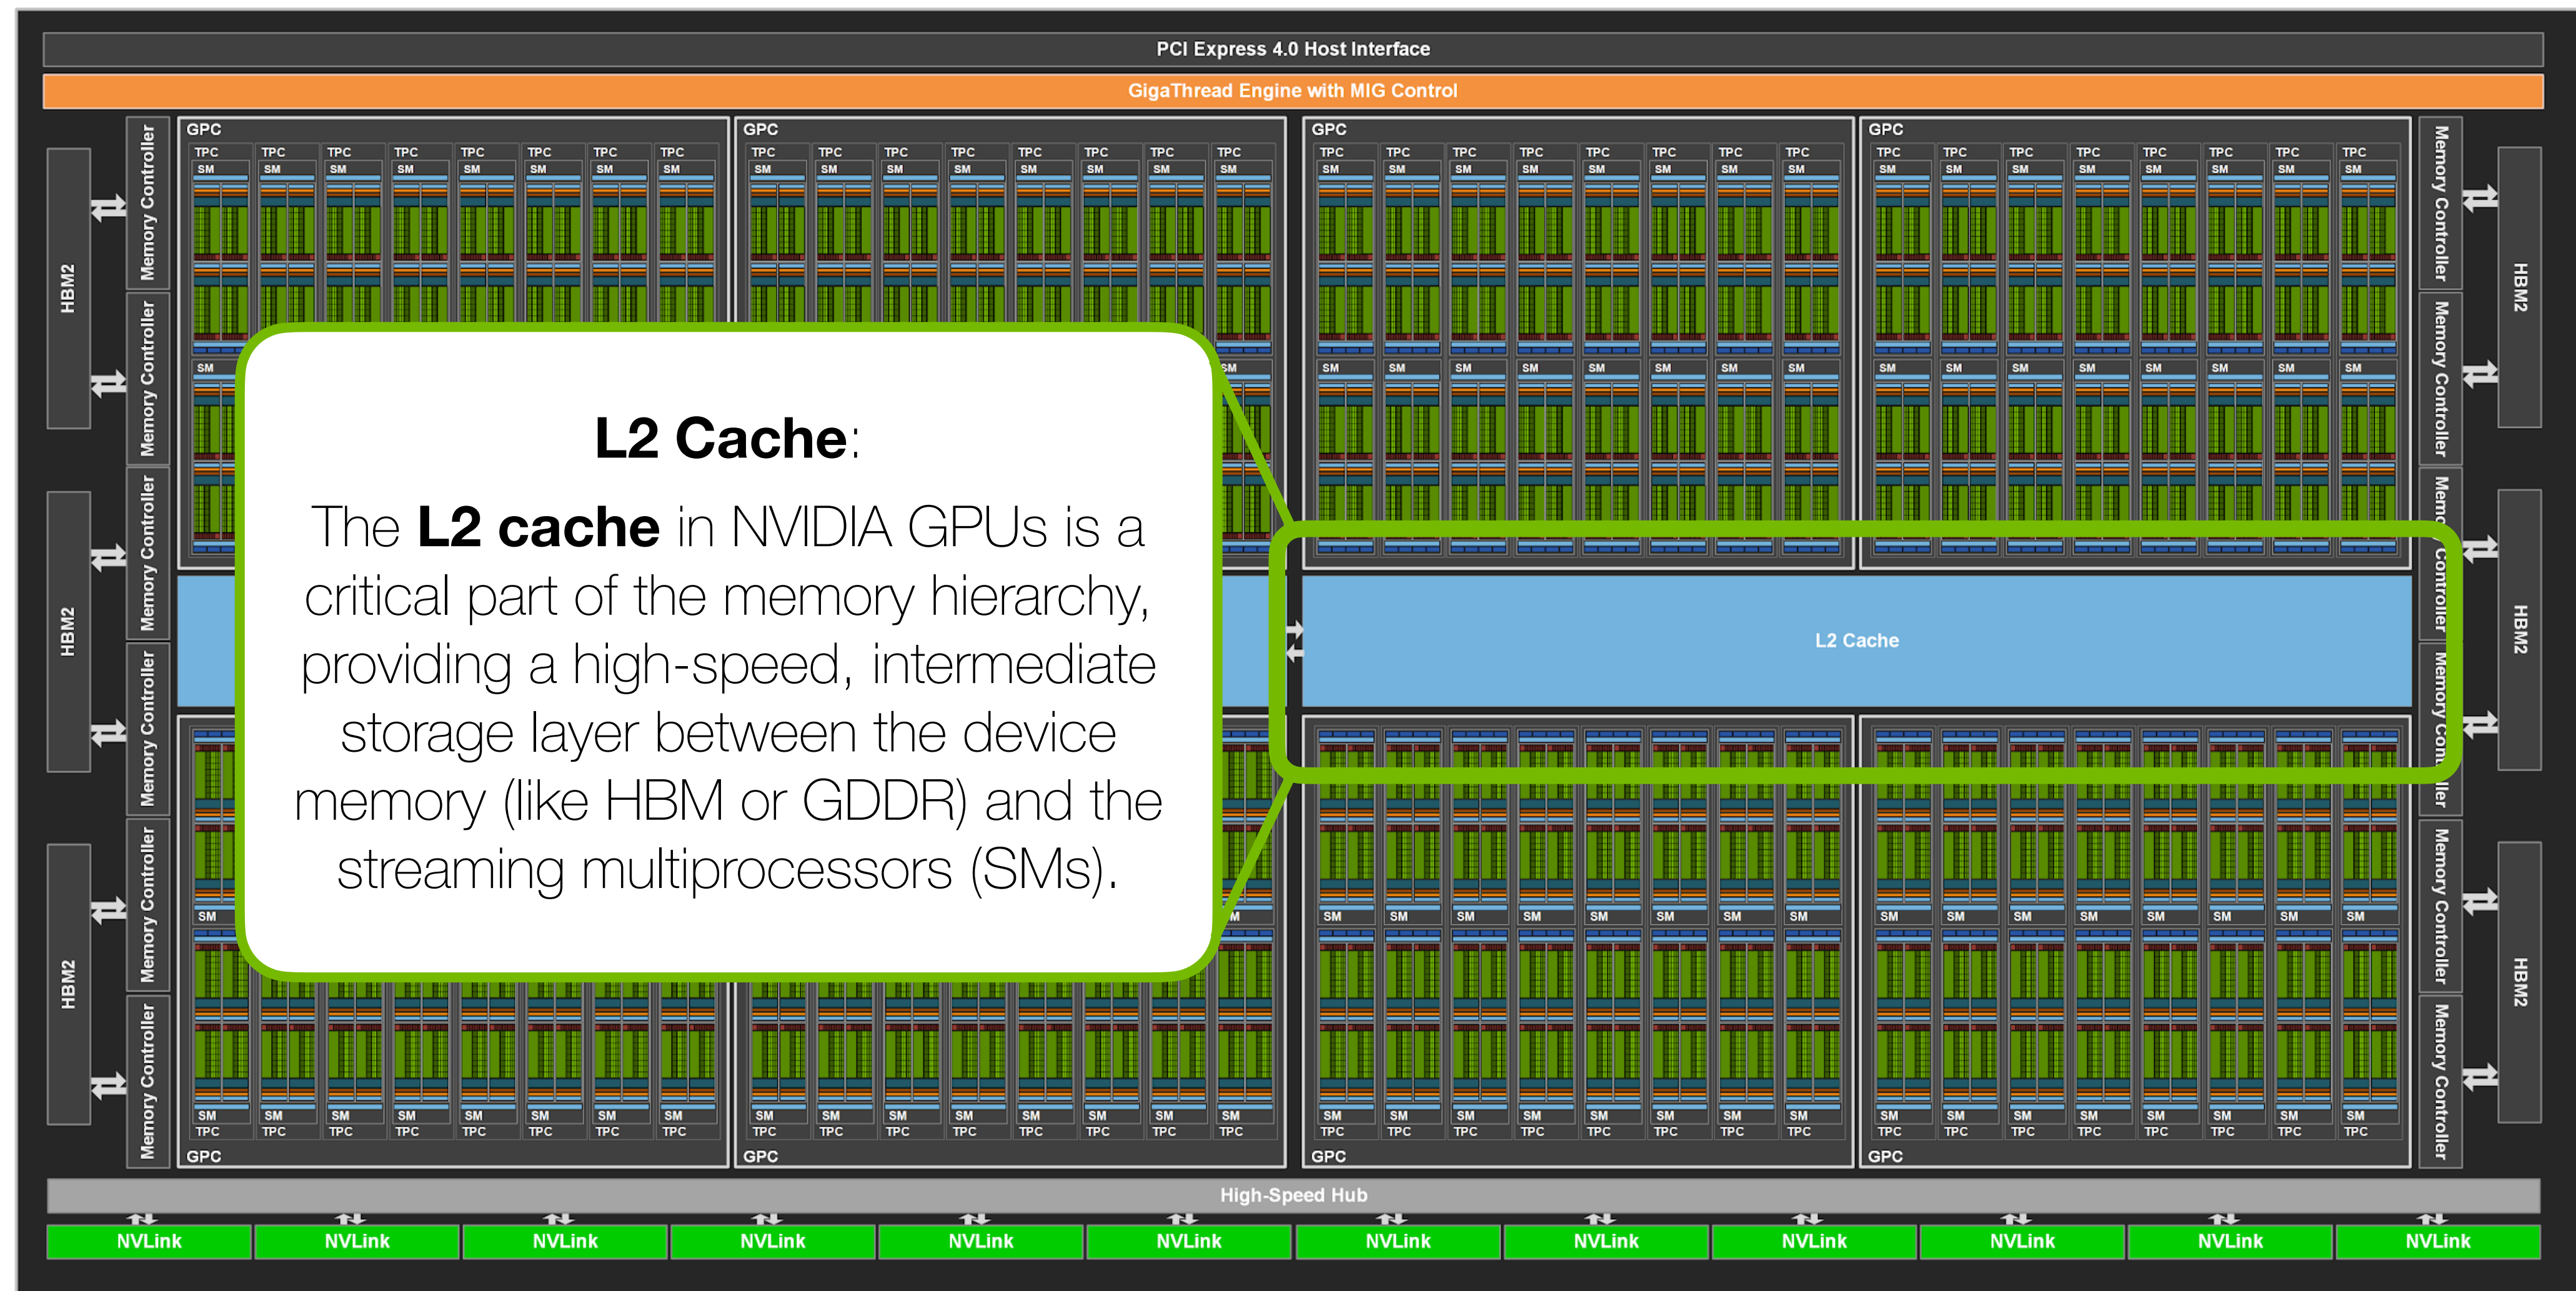
<https://developer.download.nvidia.com/devblogs/ga100-full-gpu-128-smx.png>

Recap: Device Structure

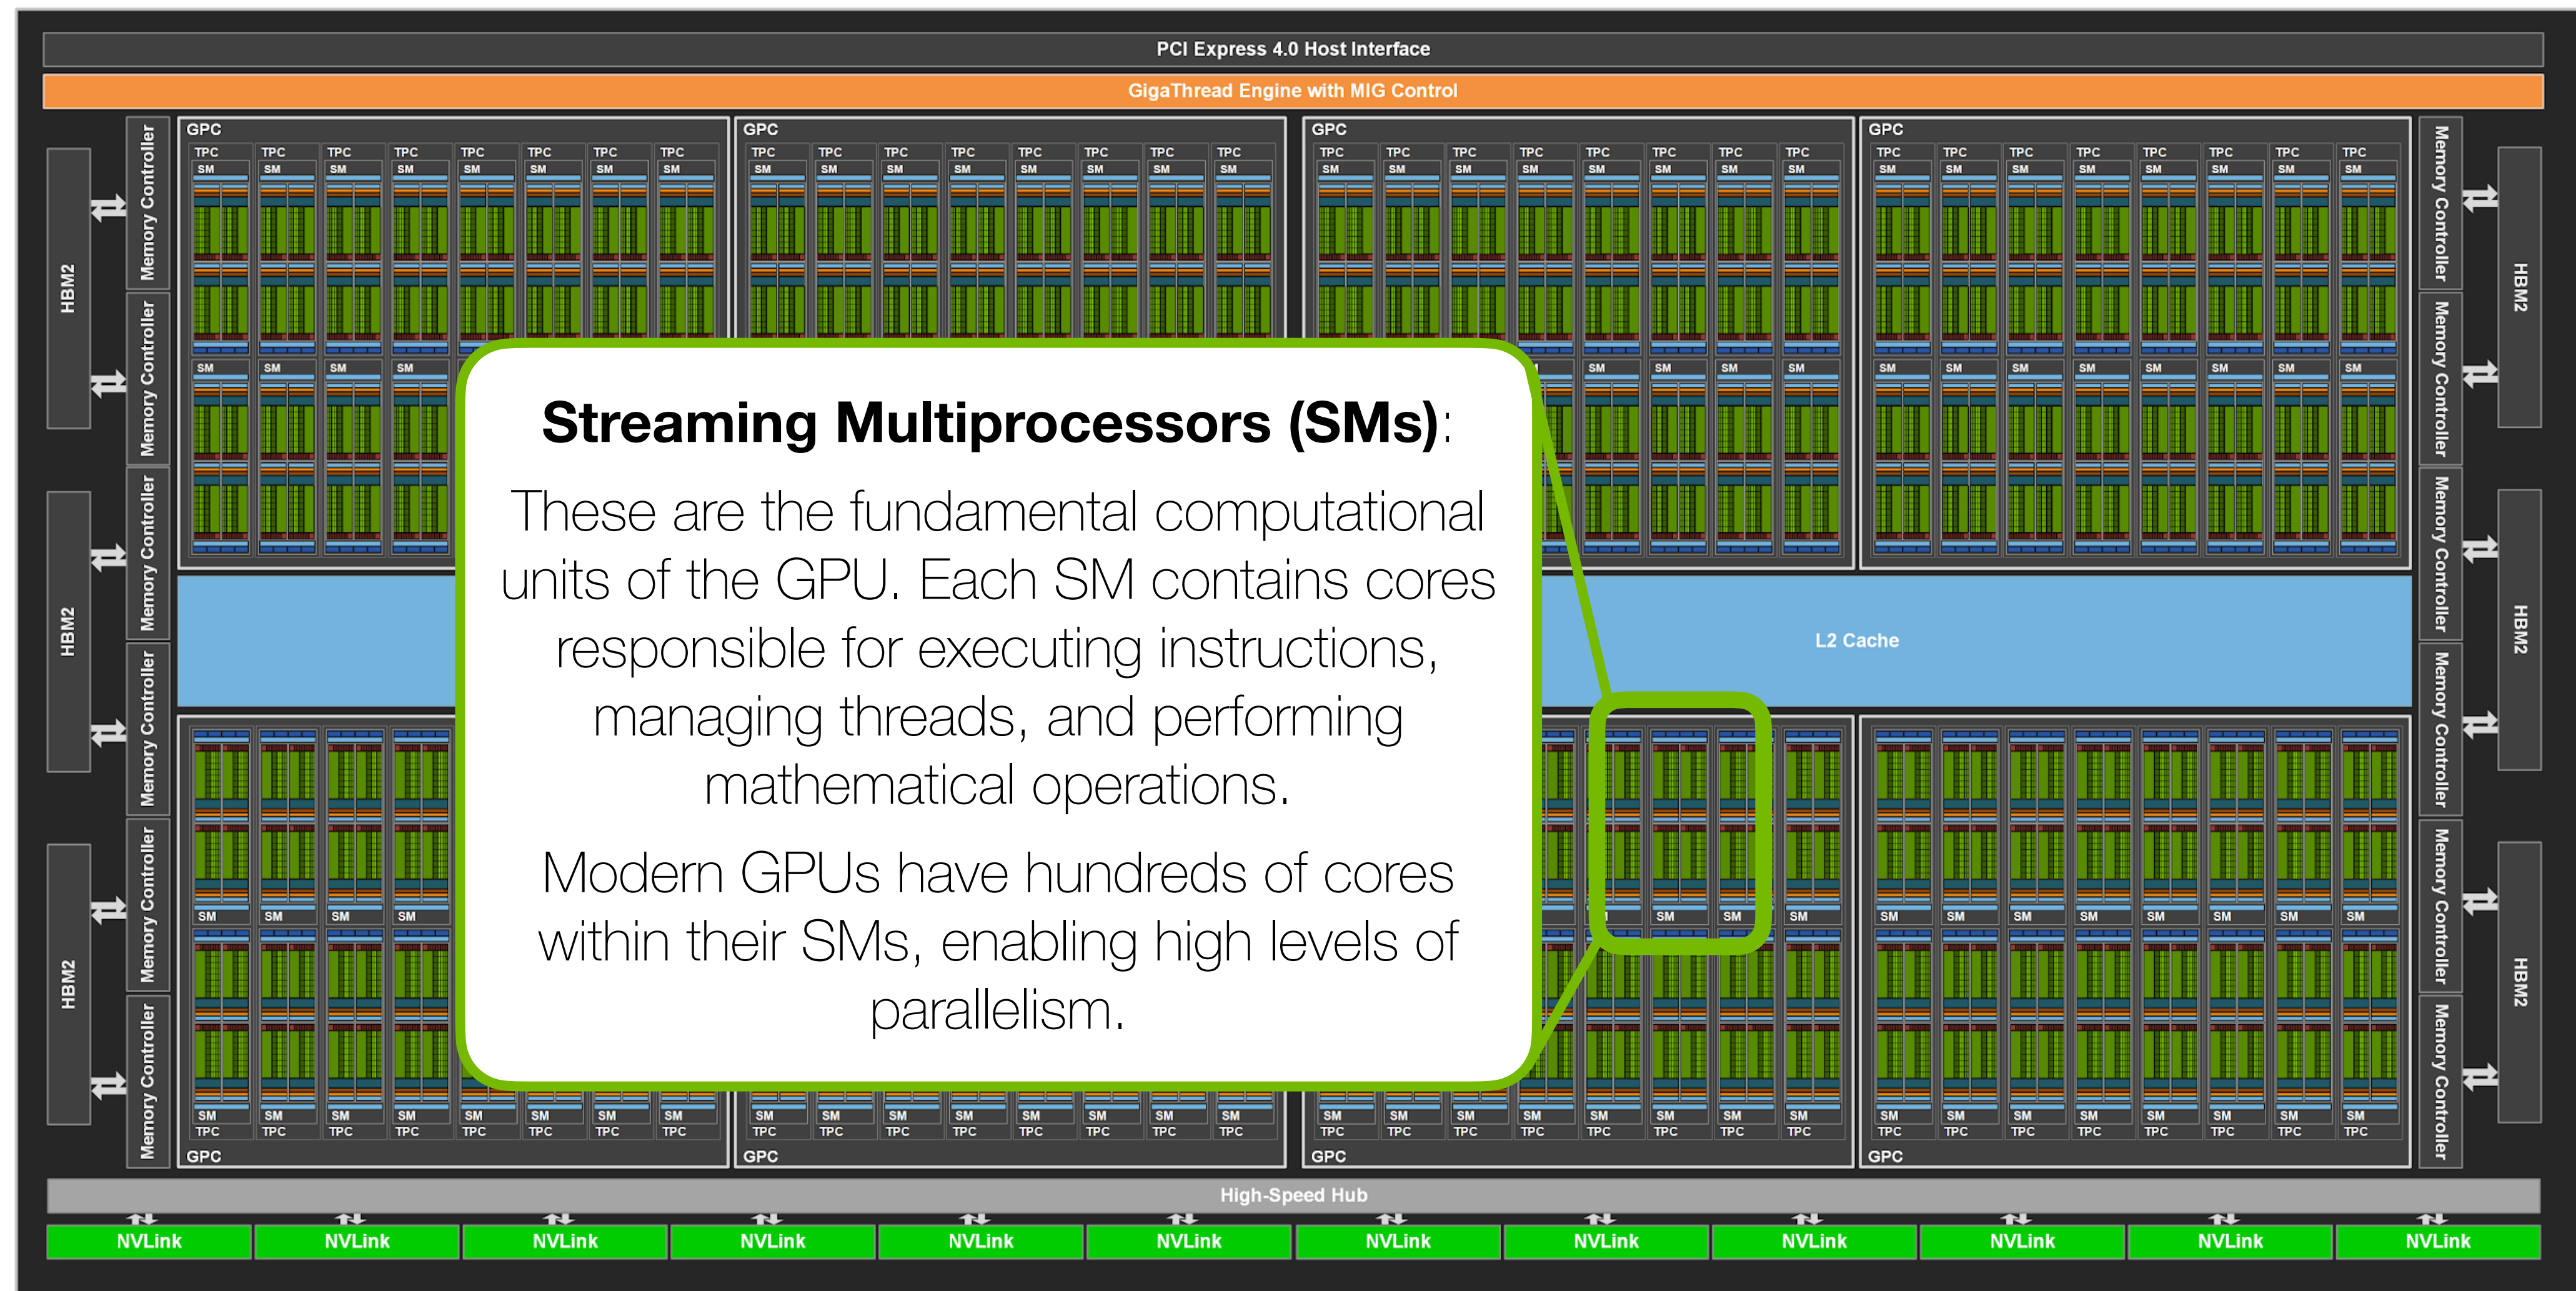


<https://developer.download.nvidia.com/devblogs/ga100-full-gpu-128-smx.png>

Recap: Device Structure



Recap: Device Structure



Recap: Device Structure



SMs feature:

- **CUDA Cores:** Responsible for general-purpose computations, including floating-point and integer arithmetic.
- **Specialized units:** These handle specific tasks, like:
 - **Tensor Cores:** Present in recent architectures like Volta, Turing, and Hopper, Tensor Cores accelerate deep learning operations by performing matrix multiplications at high speeds.

Recap: Device Structure



SMs feature:

- **L1 Cache**: a smaller, faster cache compared to the L2 cache, located within each SM. It primarily stores data that the SMs access frequently, reducing the need for repeated access to slower global memory.
- **Registers**: fastest form of memory available to each SM. They play a crucial role in storing data that threads frequently access, such as variables and intermediate results during computations.

CUDA

A **platform** designed jointly at **software** and **hardware levels** to make use of GPUs in **general-purpose computing**:

- **Software** –CUDA allows to program the GPU with minimal extensions to enable heterogeneous programming and attain an efficient and scalable execution
- **Firmware** – CUDA offers a driver oriented to GPGPU programming and APIs to manage devices, memory, etc.
- **Hardware** – CUDA exposes GPU parallelism for general-purpose computing via a number of multiprocessors endowed with cores and a memory hierarchy

CUDA

From an application perspective, CUDA defines:

- **Architecture model**
 - with many processing cores grouped in multiprocessors who share a SIMT control unit
- **Programming model**
 - based on **massive data parallelism** and fine-grained parallelism
 - **scalable**: the code is executed on a different number of cores without recompiling it
- **Memory model**
 - more **explicit** to the programmer, where **caches are not transparent** anymore

CUDA

From an application perspective, CUDA defines:

- **Architecture model**

- with many processing cores grouped in multiprocessors who share a SIMT control unit

- **Programming model**

- based on **massive data parallelism** and fine-grained parallelism
- **scalable**: the code is executed on a different number of cores without recompiling it

- **Memory model**

- more **explicit** to the programmer, where **caches are not transparent** anymore

CUDA Programming Model

It **abstracts** the **complexities of GPU hardware**, allowing developers to write code that leverages the massive parallelism provided by GPUs.

CUDA Programming Model

It **abstracts** the **complexities of GPU hardware**, allowing developers to write code that leverages the massive parallelism provided by GPUs.

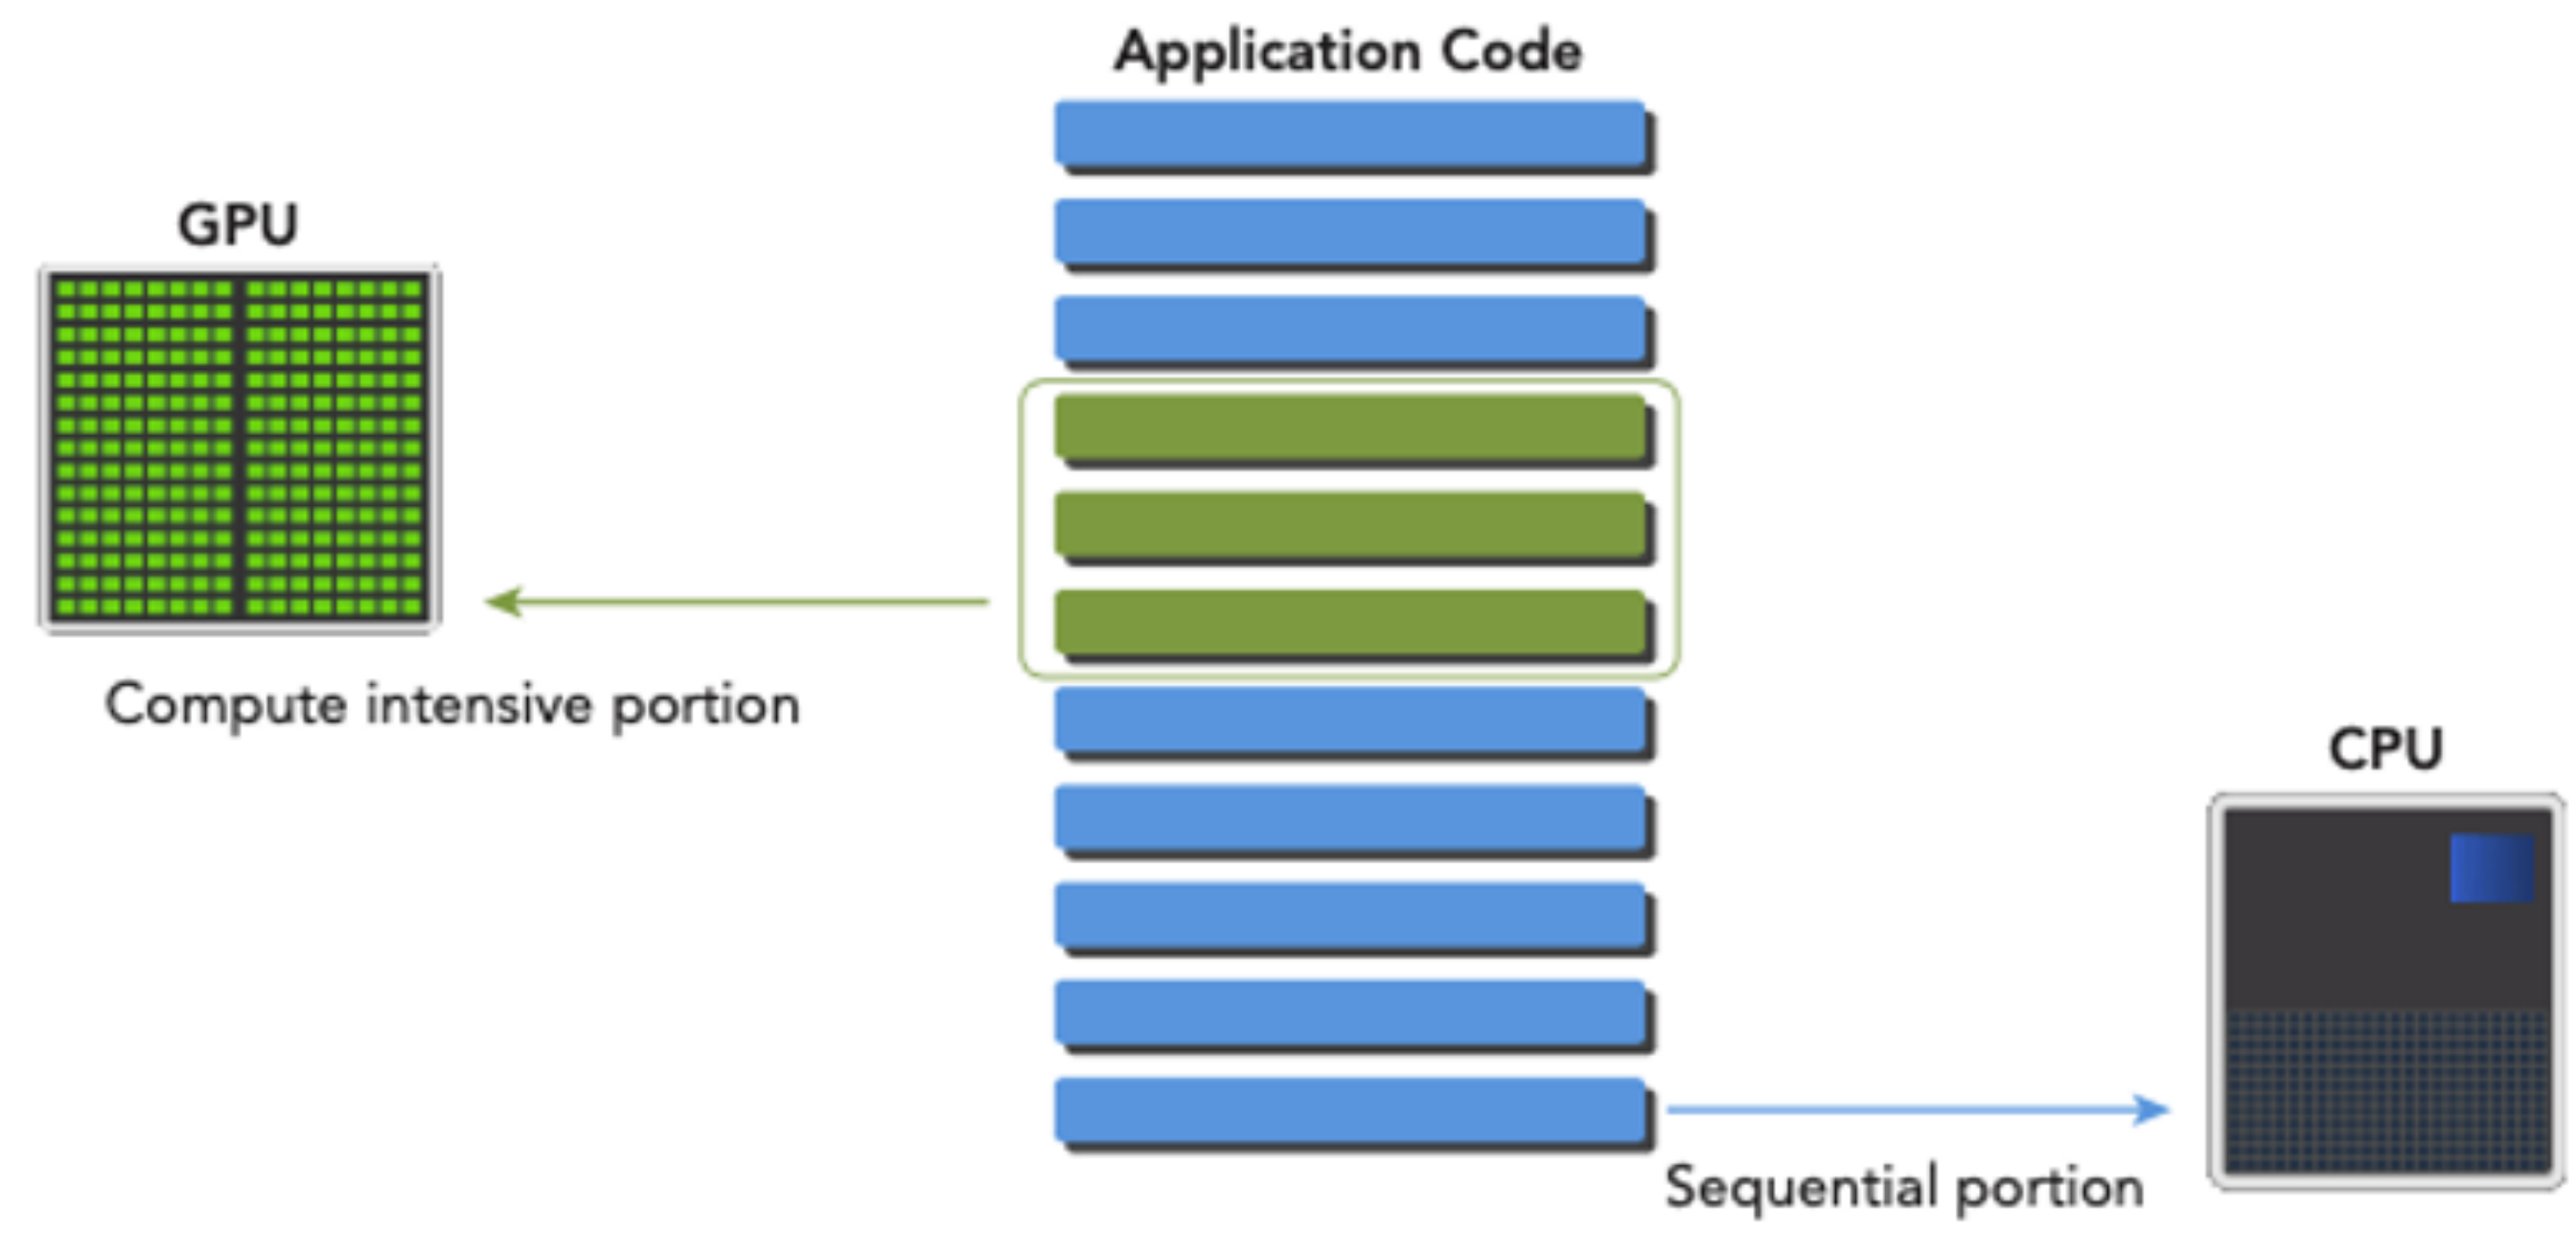
CUDA operates in a **heterogeneous computing environment**, where an application runs on two types of devices:

- **Host**: Typically the **CPU**, which runs the main application logic and controls the execution.
- **Device**: The **GPU**, which performs highly parallel tasks delegated by the host.

CUDA Program

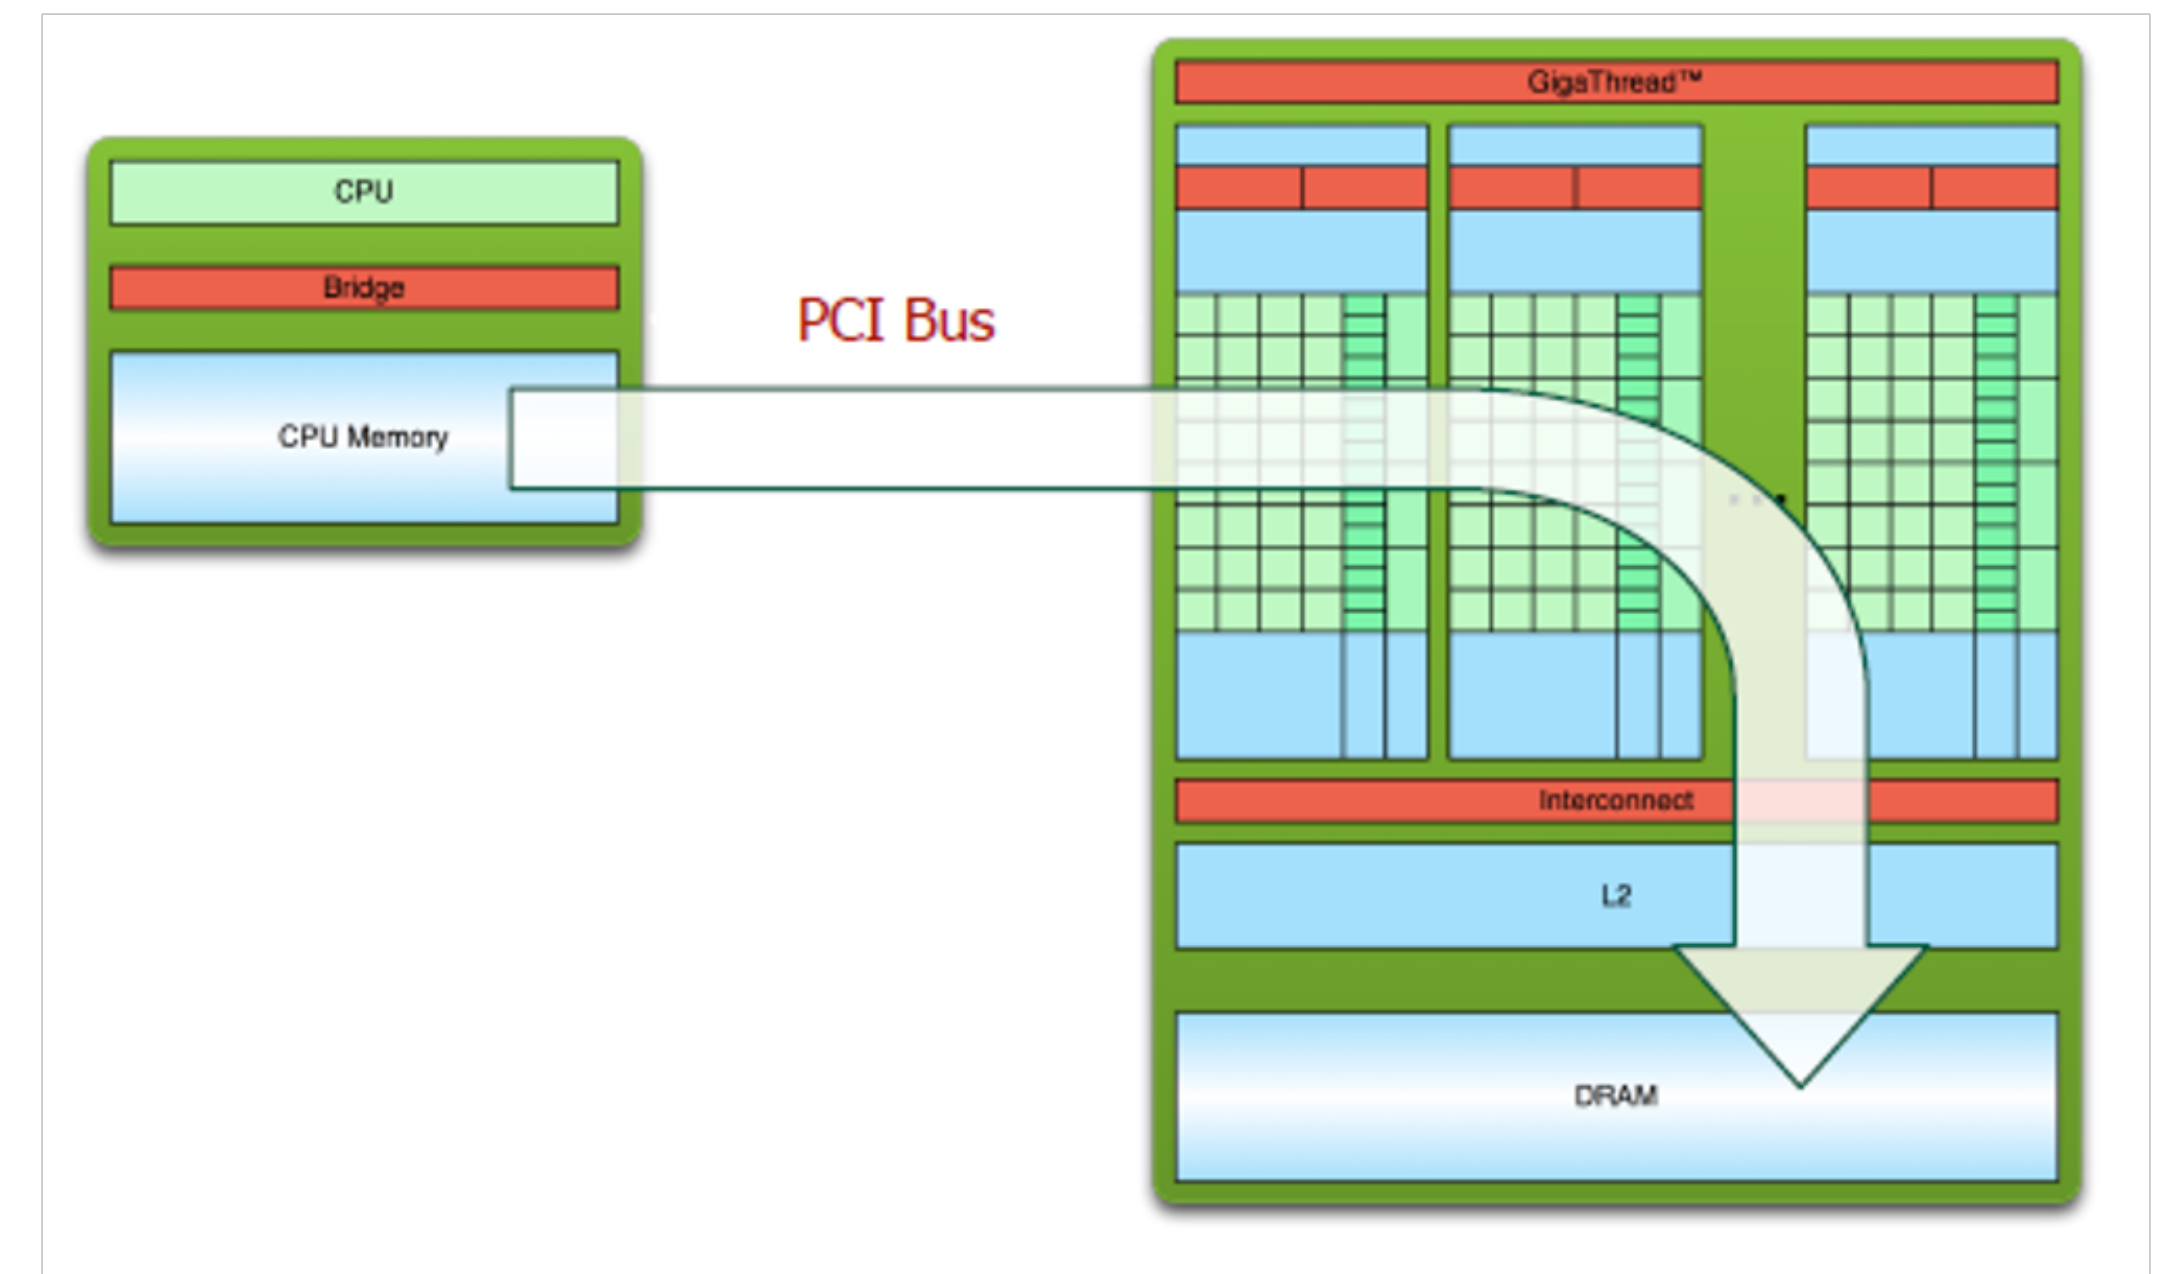
A CUDA program is organized in

- The **serial code** executed on the host (i.e., the CPU)
- One or many independent **parallelized** functions (called **kernels**) running on the **device** (i.e., the GPU)




Program Execution Flow

- Transfer data from the host to the device



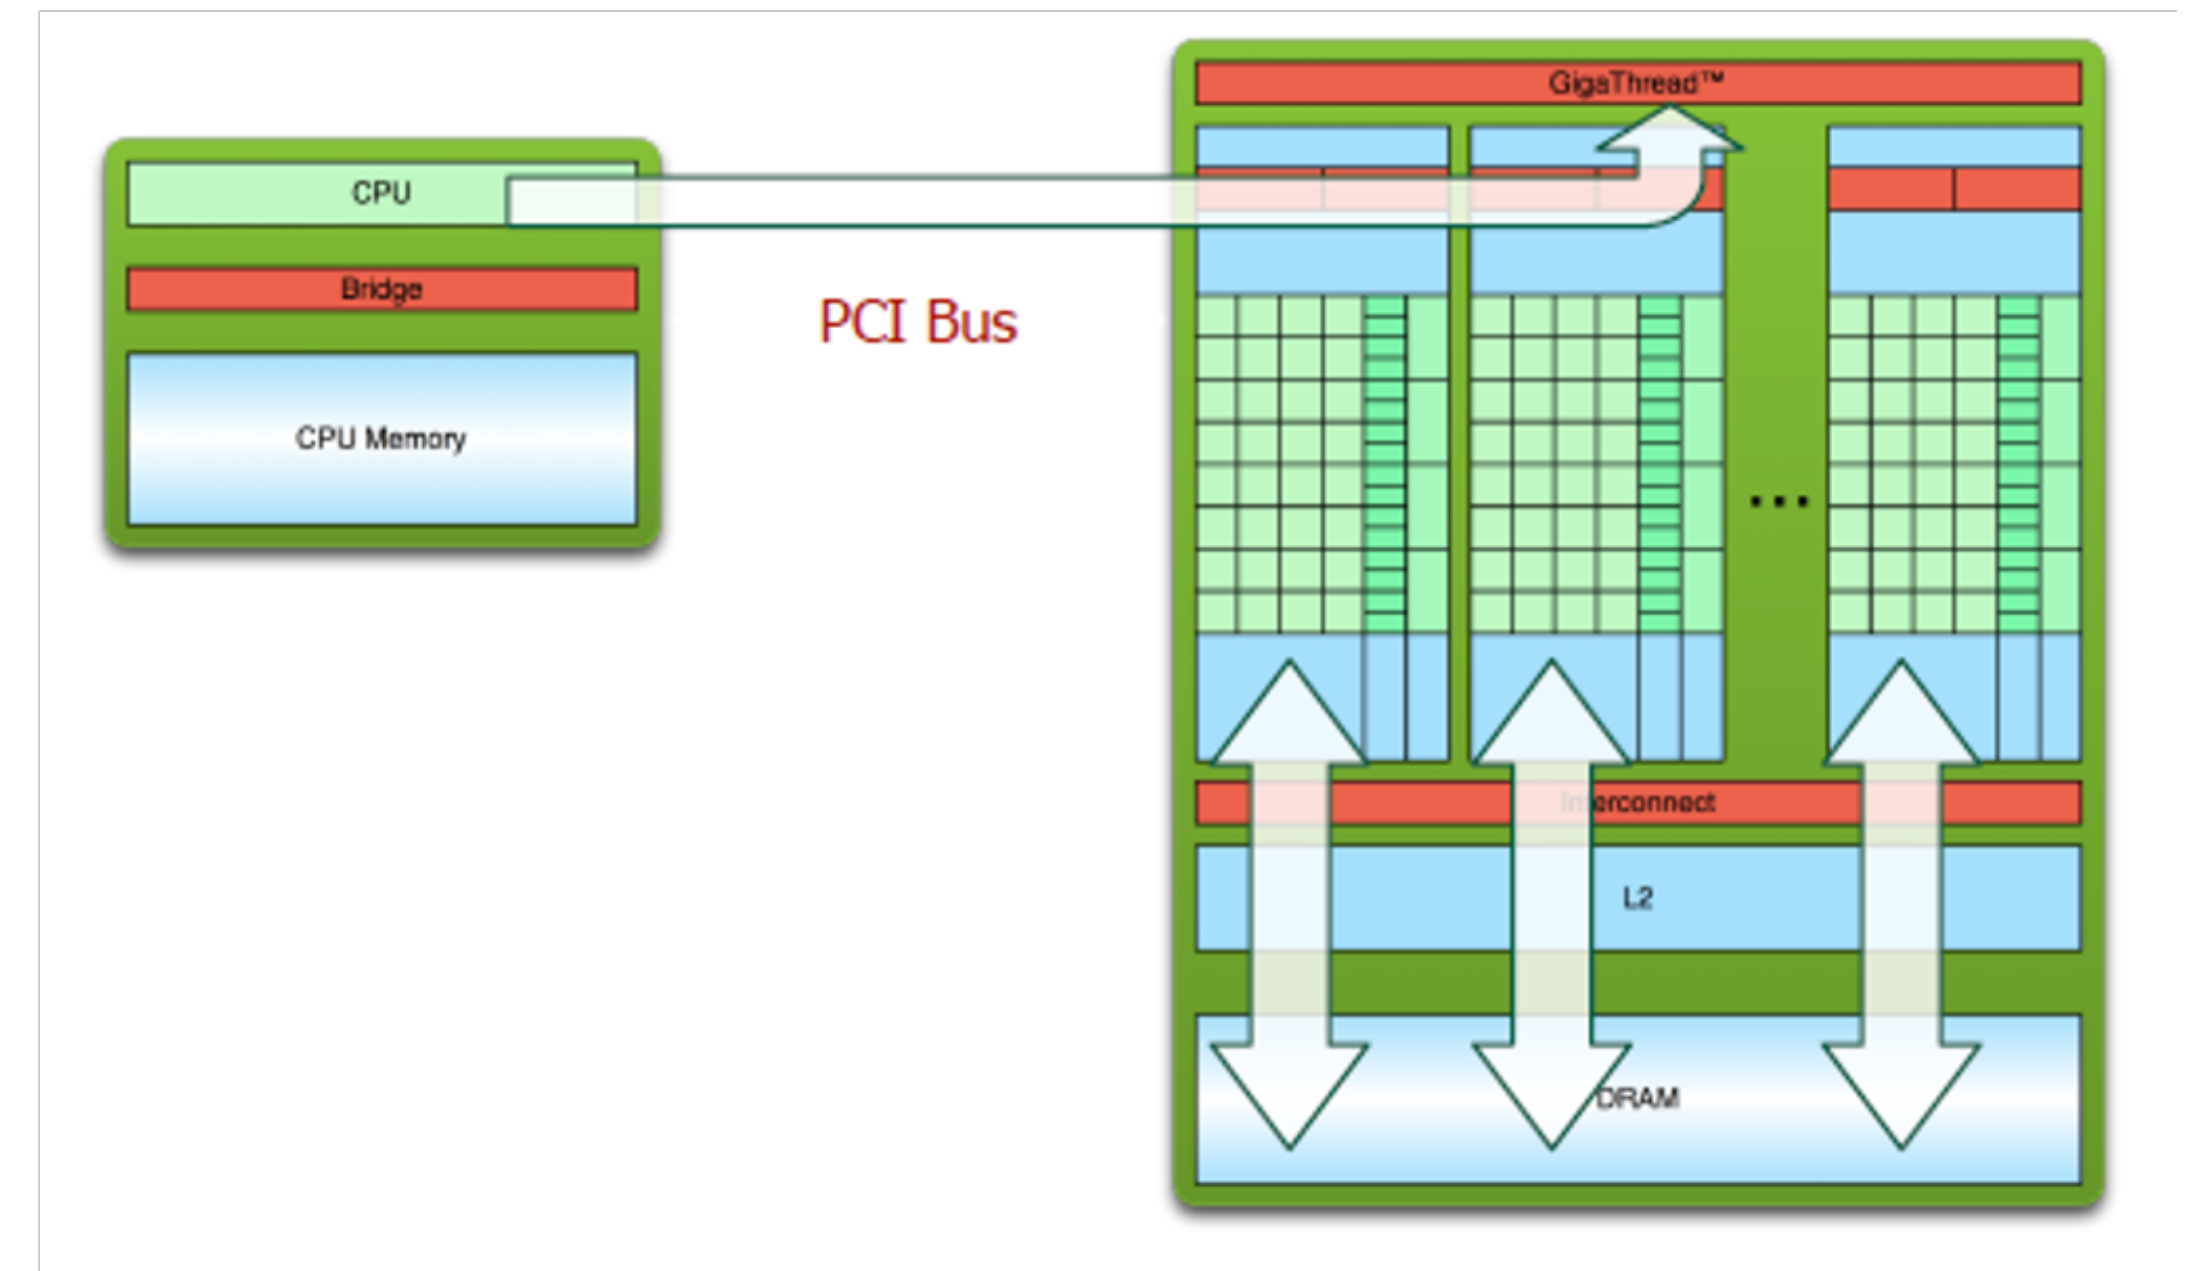
Program Execution Flow

- Transfer data from the host to the device
- Launch the kernel



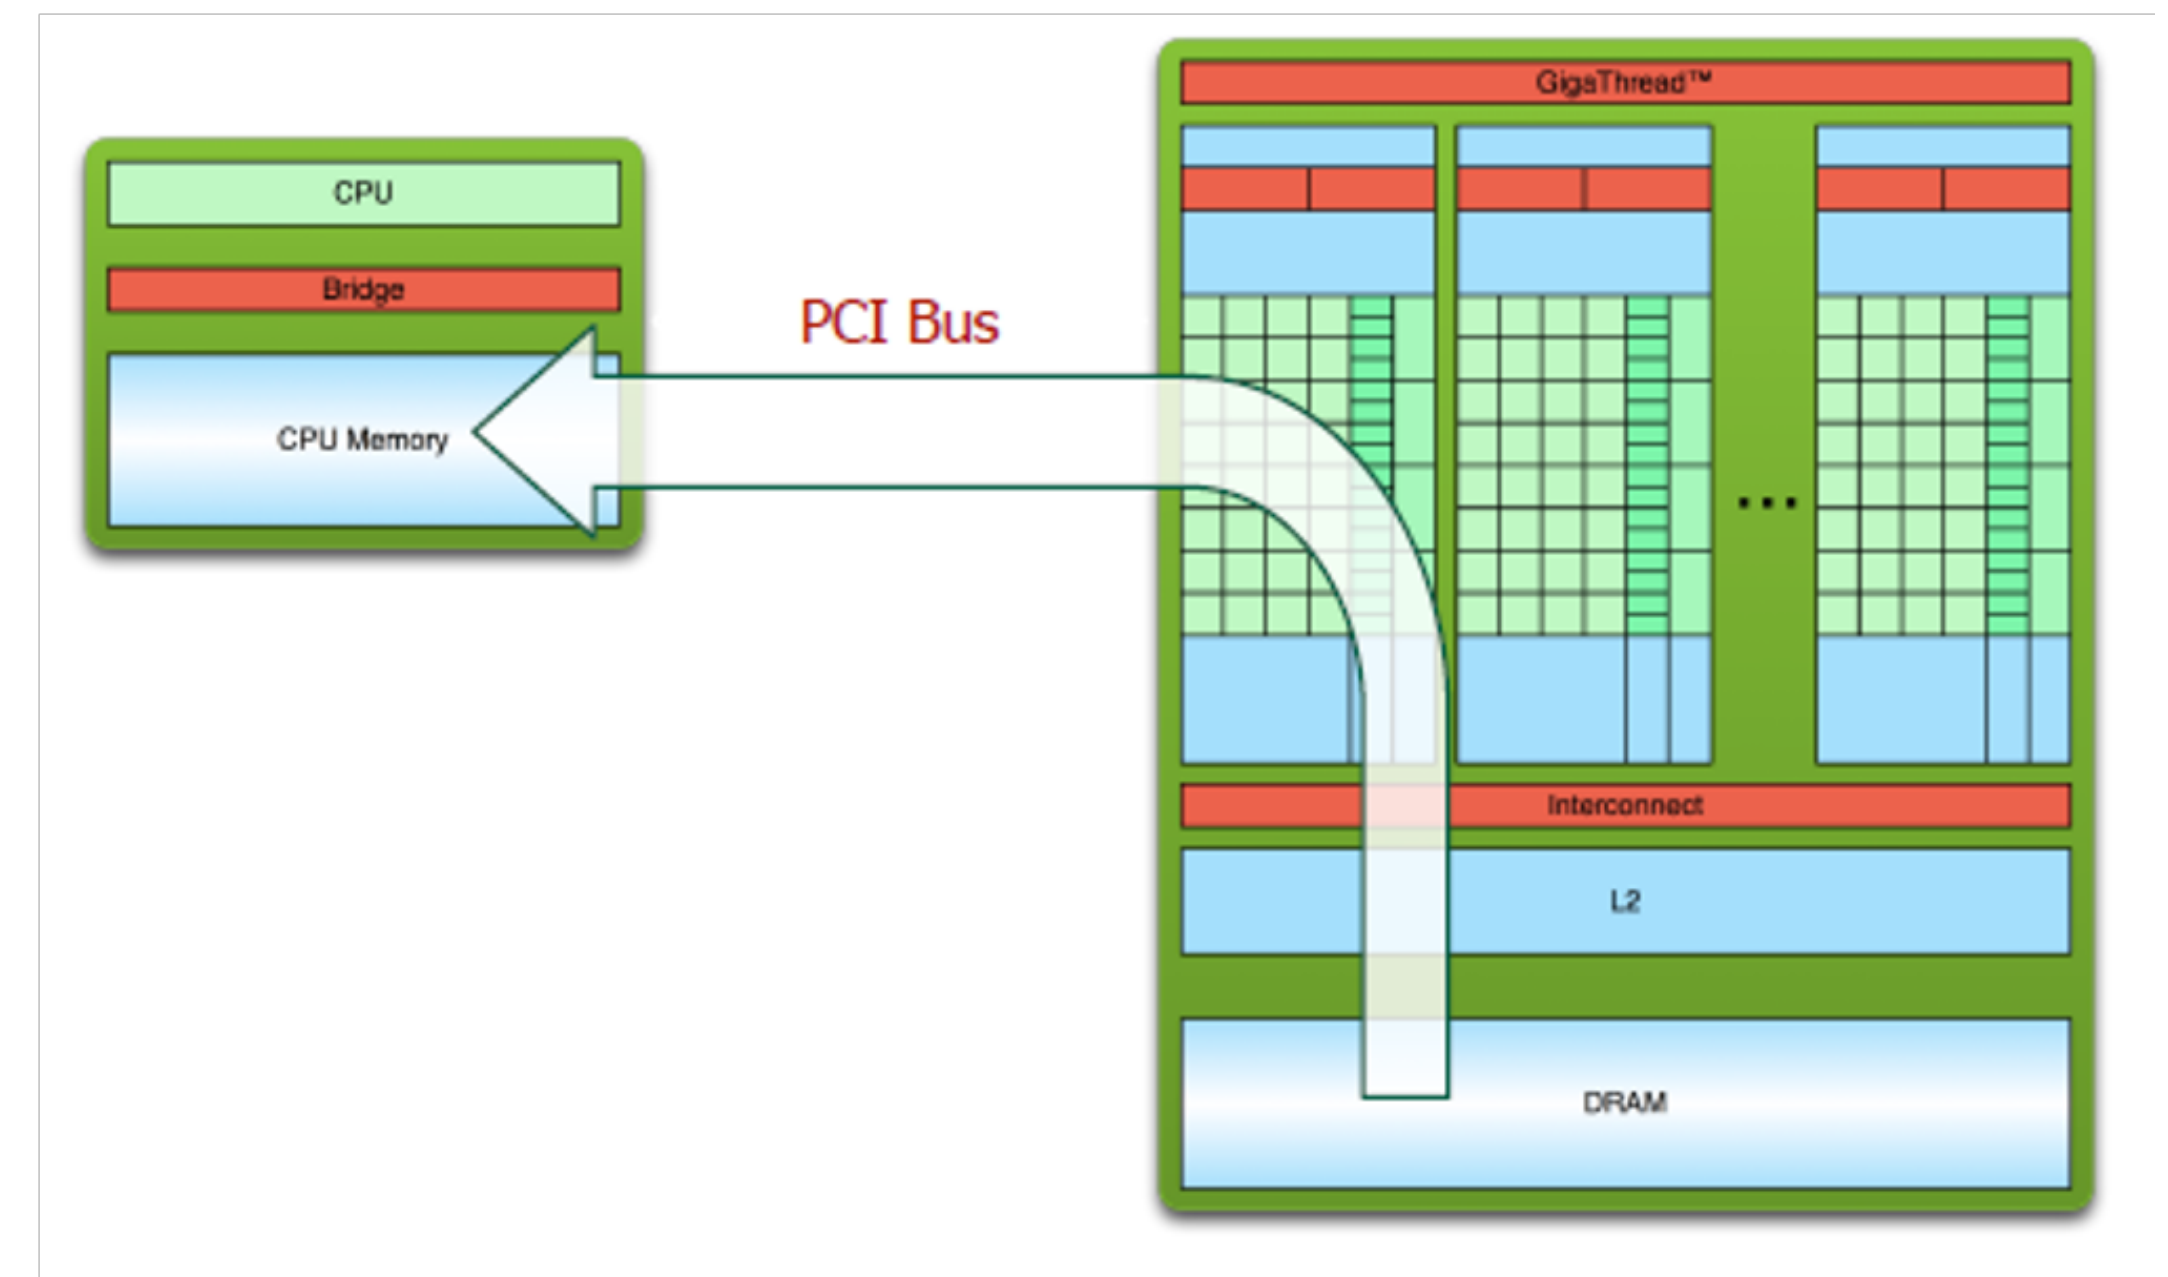
Program Execution Flow

- Transfer data from the host to the device
- Launch the kernel
- Synchronize to ensure all threads complete



Program Execution Flow

- Transfer data from the host to the device
- Launch the kernel
- Synchronize to ensure all threads complete
- Transfer results back to the host



CUDA Programming Model

It **abstracts** the **complexities of GPU hardware**, allowing developers to write code that leverages the massive parallelism provided by GPUs.

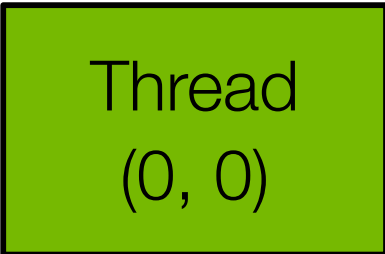
CUDA operates in a **heterogeneous computing environment**, where an application runs on two types of devices:

- **Host:** Typically the **CPU**, which runs the main application logic and controls the execution.
- **Device:** The **GPU**, which performs highly parallel tasks delegated by the host.

**In the CUDA programming model,
the function running on GPU is called kernel**

CUDA Programming Model

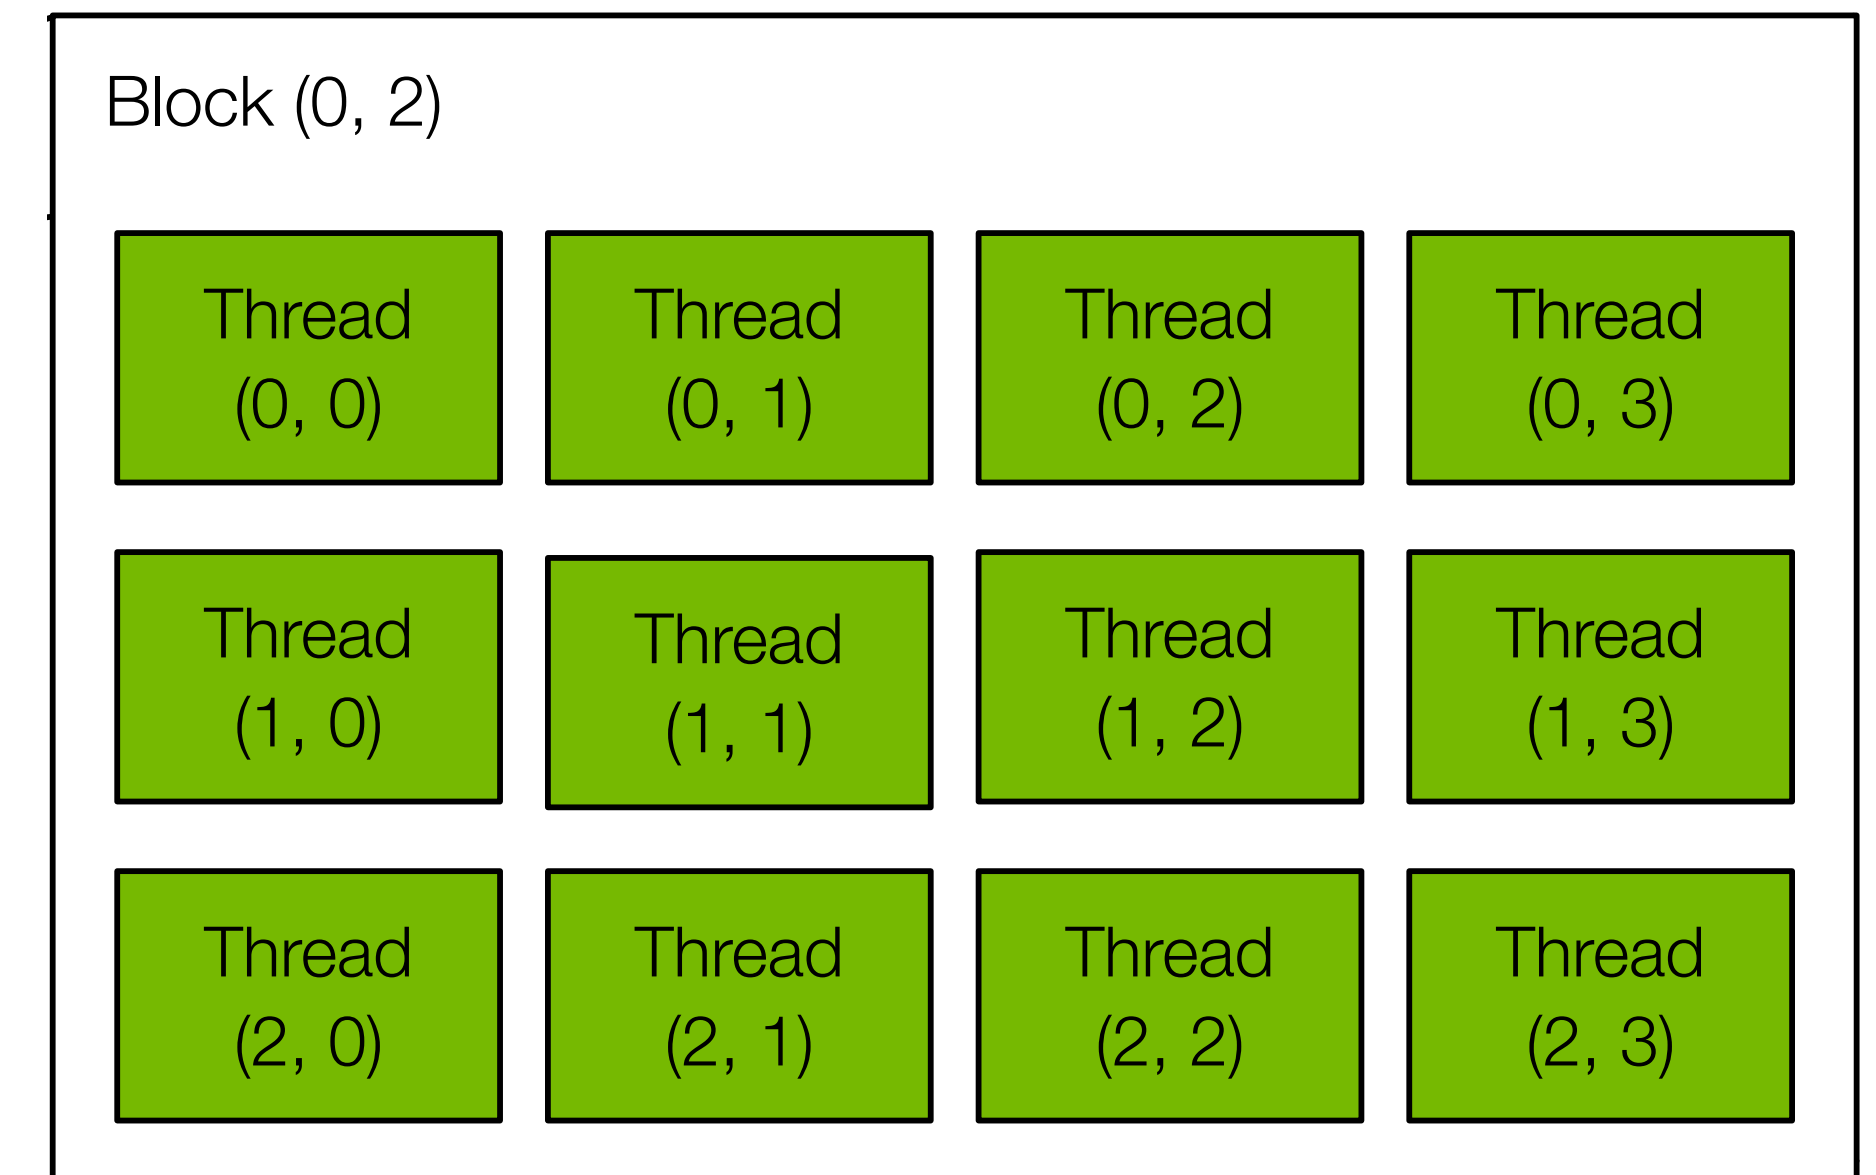
The **smallest unit of parallel execution** in CUDA. Each thread executes the kernel code independently and has its own set of registers and local memory.



Thread
(0, 0)

CUDA Programming Model


Threads are grouped into **thread blocks**. All threads in a block can communicate and synchronize on tasks. Threads in a block execute concurrently on a SM. Block dimensions can be 1D, 2D, or 3D depending on the problem at hand.



CUDA Programming Model

Threads are grouped into **thread blocks**. All threads in a block can communicate and synchronize on tasks. Threads in a block execute concurrently on a SM. Block dimensions can be 1D, 2D, or 3D depending on the problem at hand.

When a CUDA kernel is launched, threads within each thread block are organized into groups of 32 threads called **warps**. A warp **always** contains **32 threads**, regardless of the GPU architecture. All threads in a warp execute the **same instruction at the same time (SIMT paradigm)**, but each thread operates on its own data element.

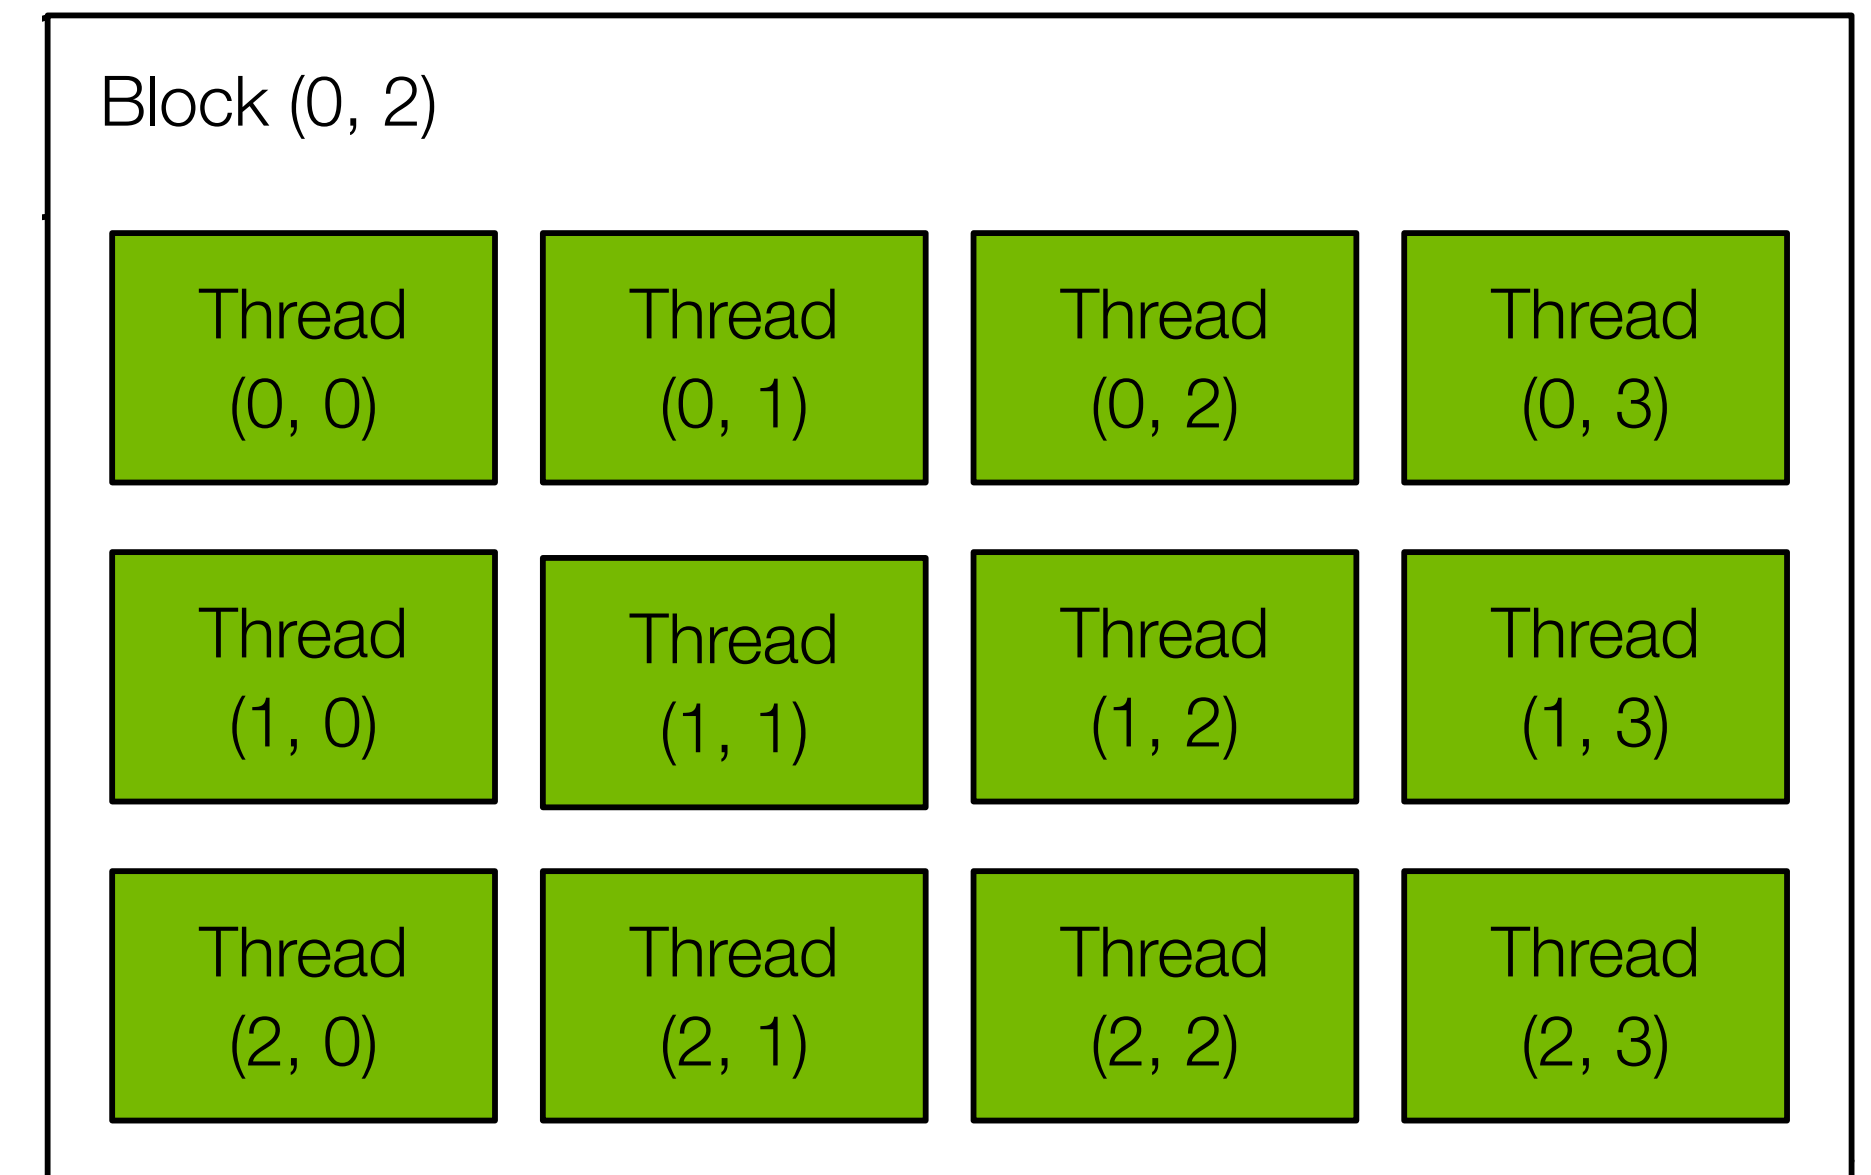


CUDA Programming Model

Threads are grouped into **thread blocks**. All threads in a block can communicate and synchronize on tasks. Threads in a block execute concurrently on a SM. Block dimensions can be 1D, 2D, or 3D depending on the problem at hand.

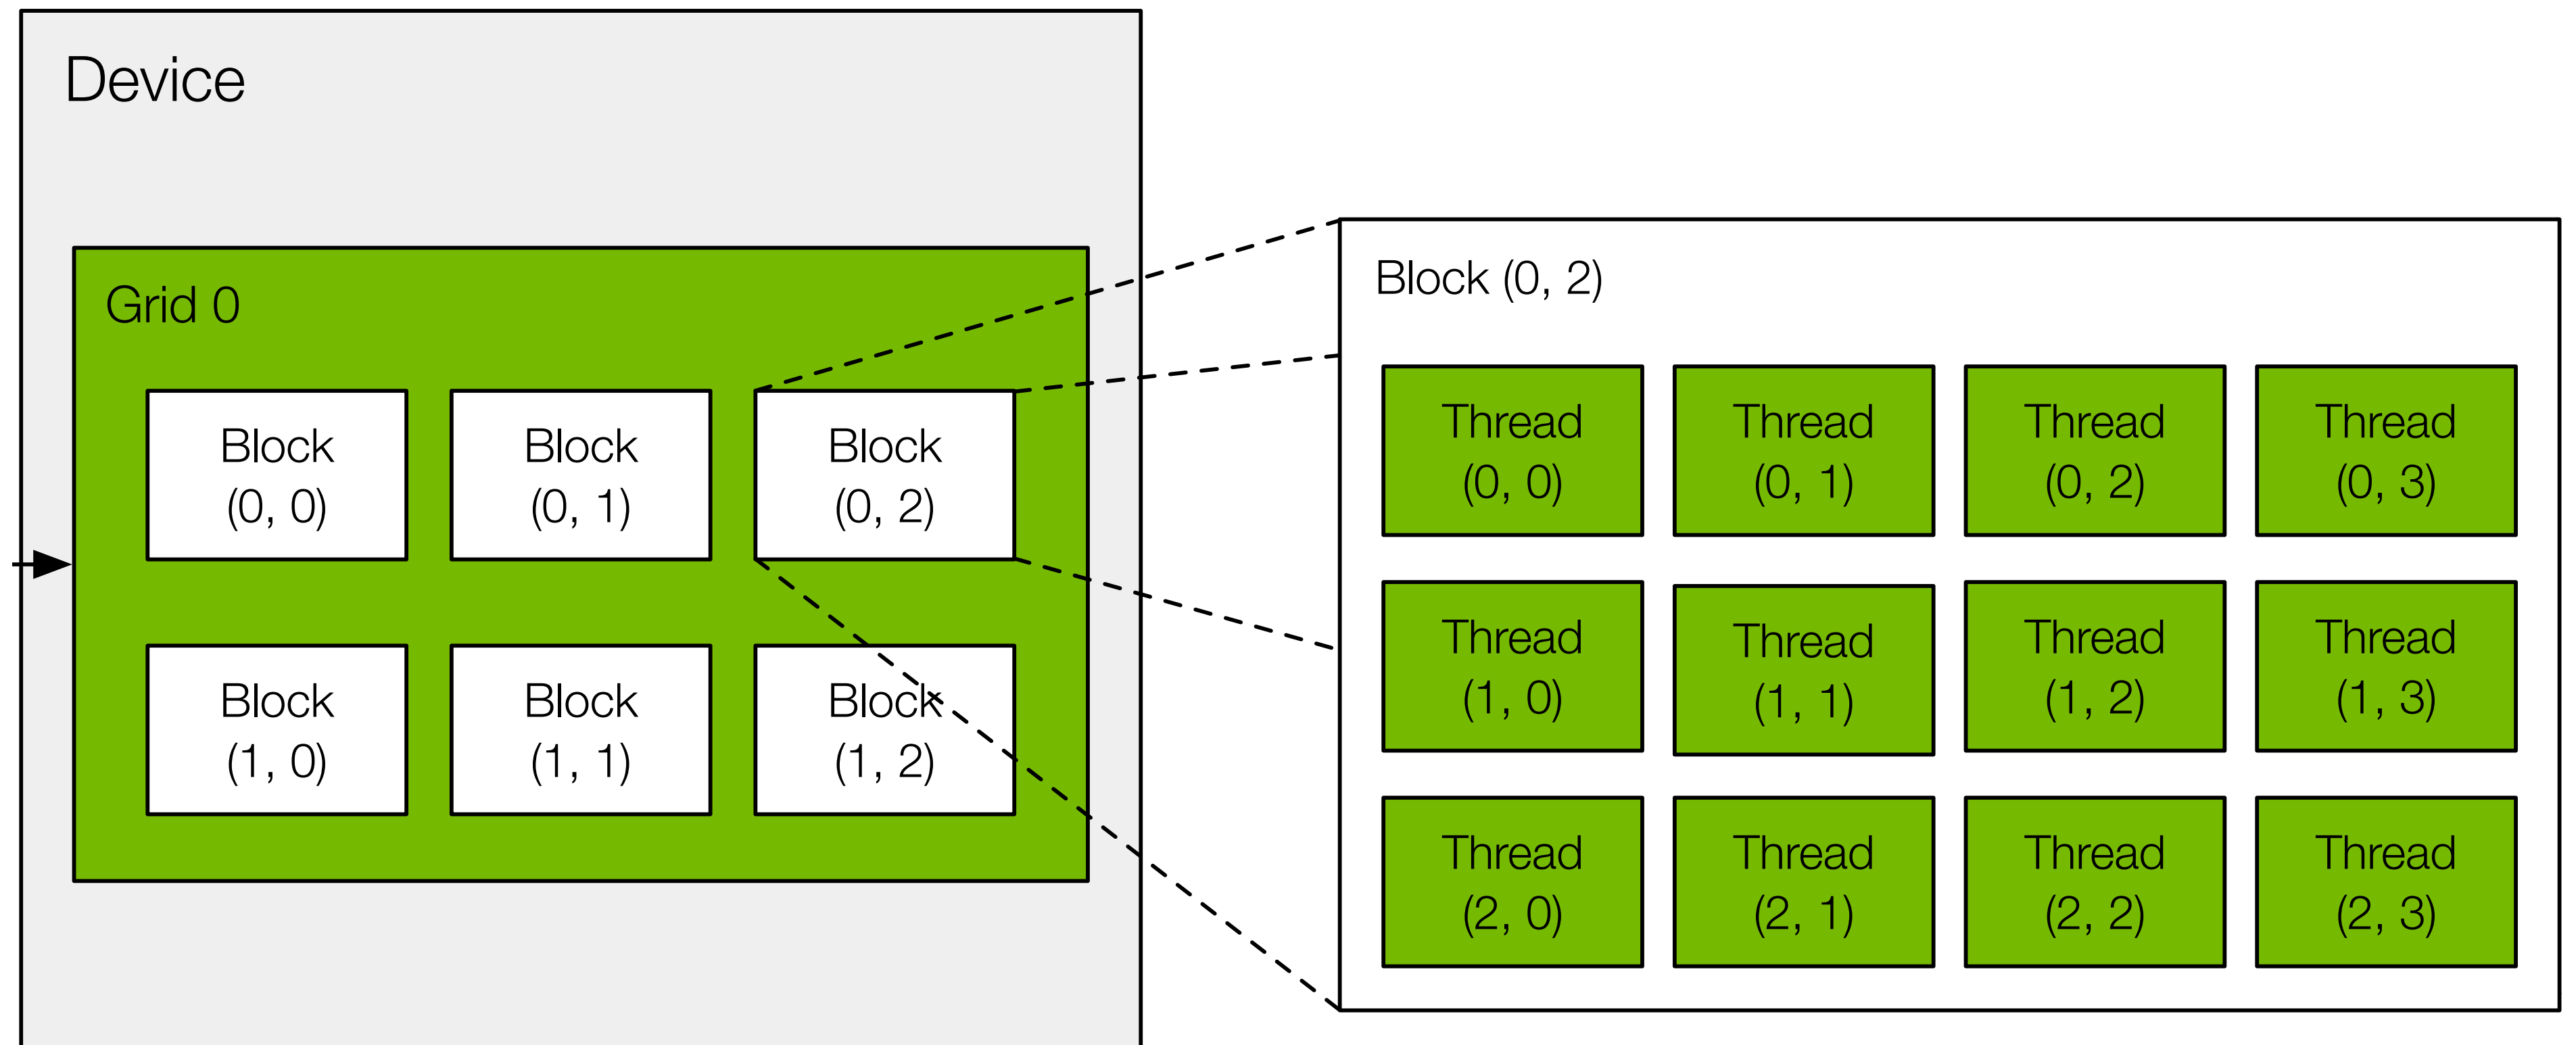
When a CUDA kernel is launched, threads within each thread block are organized into groups of 32 threads called **warps**. A warp **always** contains **32 threads**, regardless of the GPU architecture. All threads in a warp execute the **same instruction at the same time (SIMT paradigm)**, but each thread operates on its own data element.

 **Beware of divergence!**

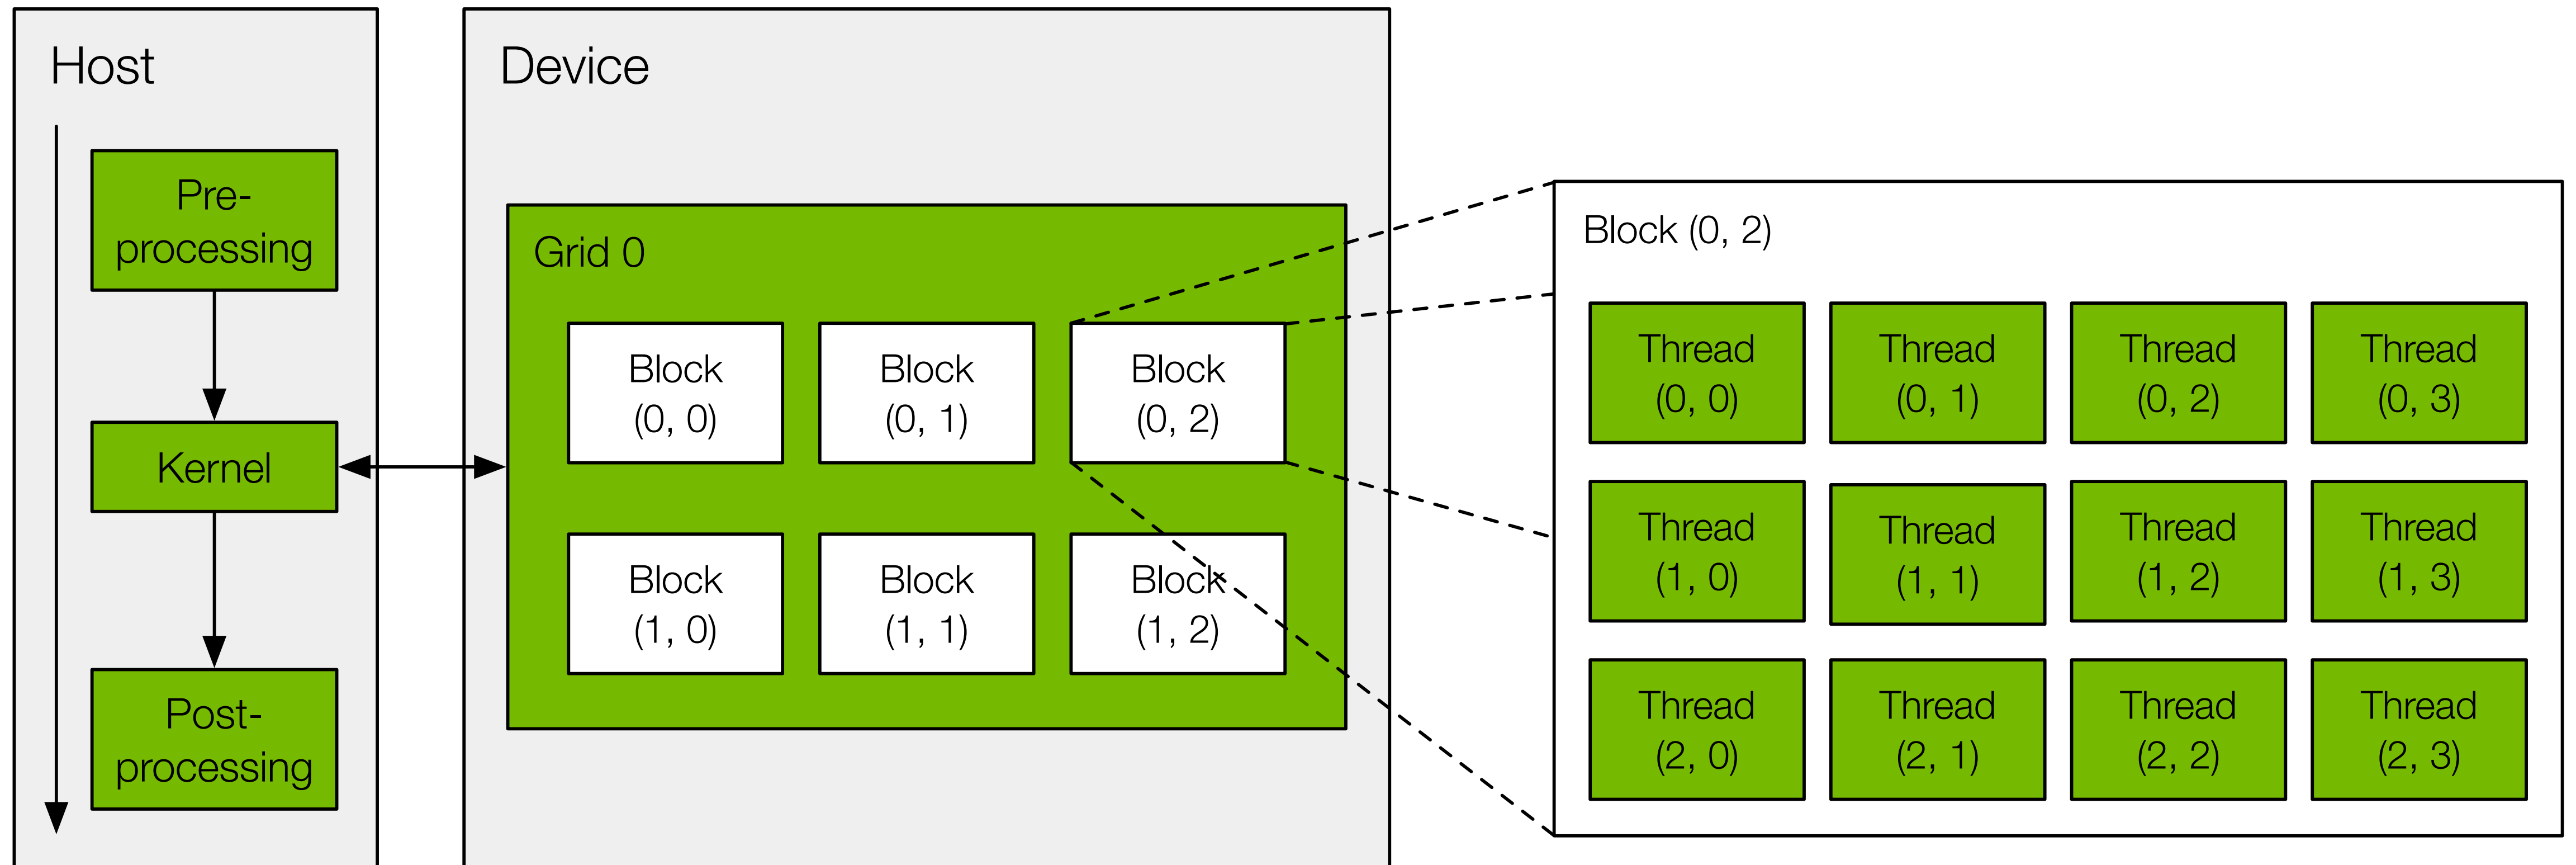


CUDA Programming Model

A collection of blocks forms a **grid**.
Blocks in the grid are independent of each other, meaning there's no direct synchronization or communication between them*.



CUDA Programming Model

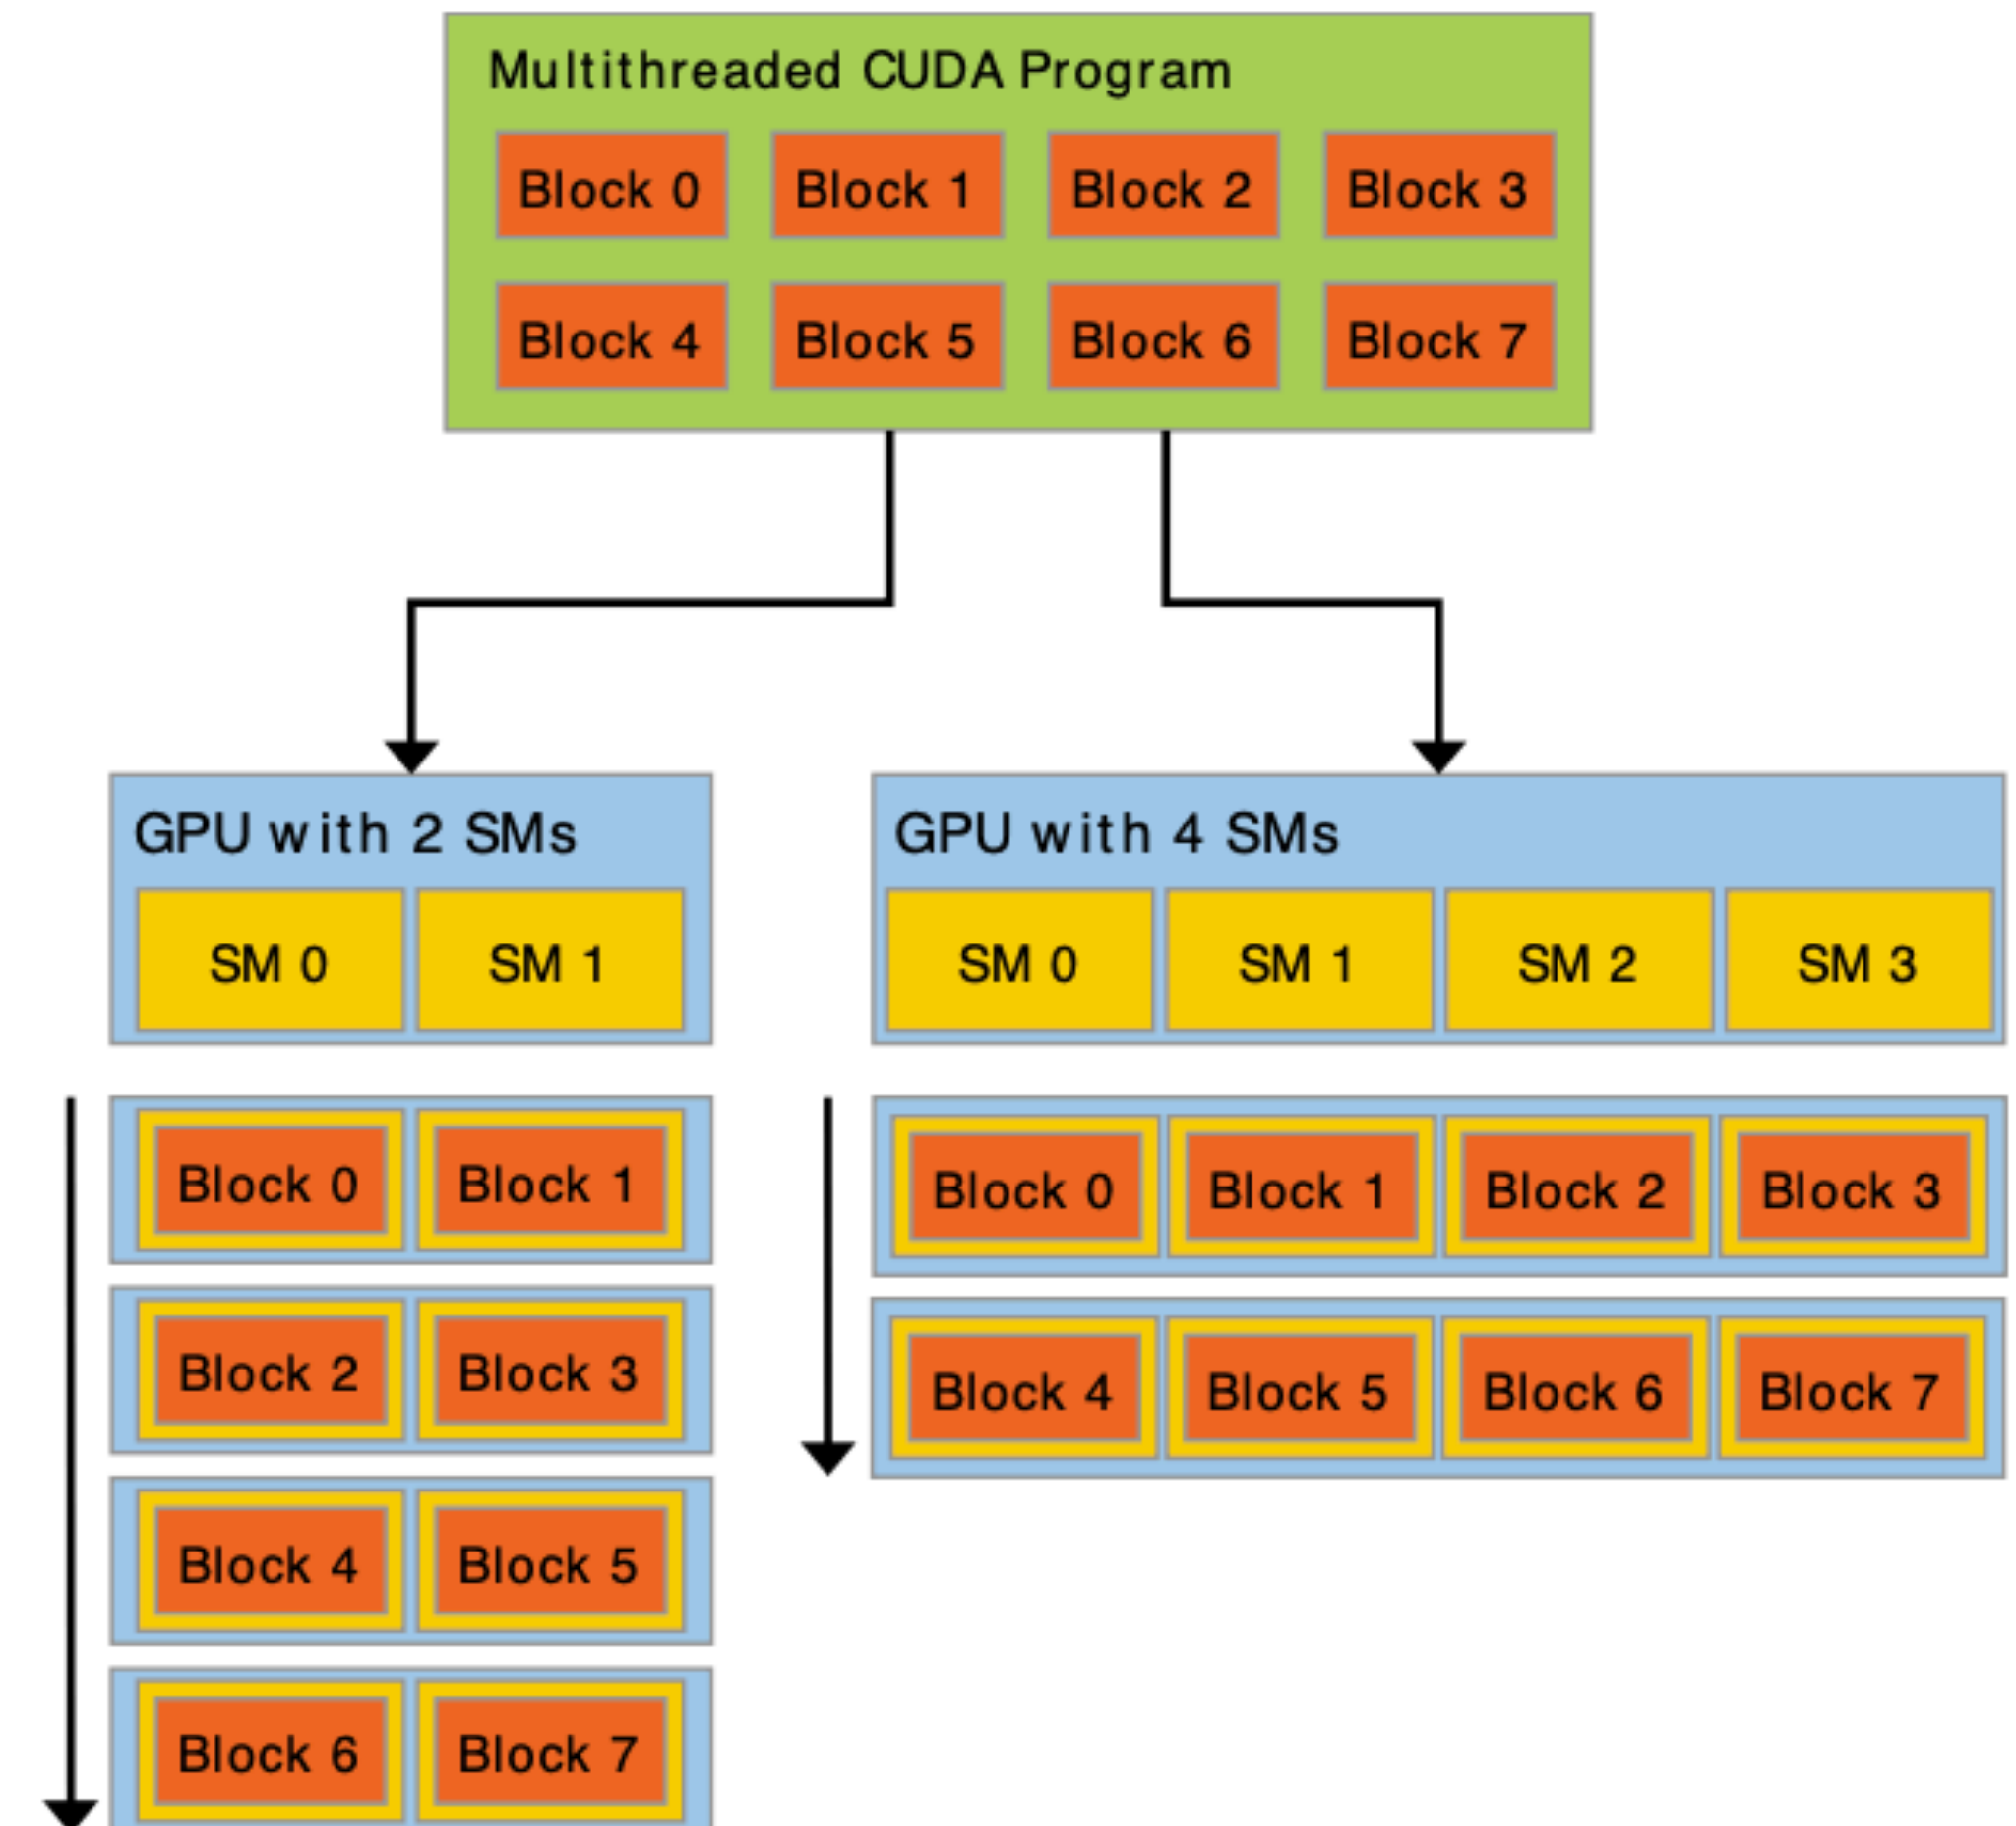


Block Dispatching

Each block is assigned to a single SM (no relocation!)

Multiple blocks may be assigned to a single SM.

Blocks are executed in **any order independently** from each other




Block Dispatching

A block is assigned to an SM only **if there are enough resources** and those resources are **reserved for the block for its entire execution**

Each SM has predefined resources:


- max number of threads
- max number of blocks
- number of registers
- shared memory size



Warp Scheduling

The SM divides the block in **warps** (32 threads each). A warp is executed with a **SIMT** approach: all threads in a warp execute the same instruction.

To improve throughput, the scheduler interleaves the execution of multiple warps.




If block size is not multiple of 32, **inactive threads are added to the last warp**

Warp Scheduling

Each active warp may be classified as:

- **Selected warp** – if currently executed
- **Stalled warp** – if not ready to be executed
- **Eligible warp** – if ready to be executed

Warp schedulers will select from **eligible** warps



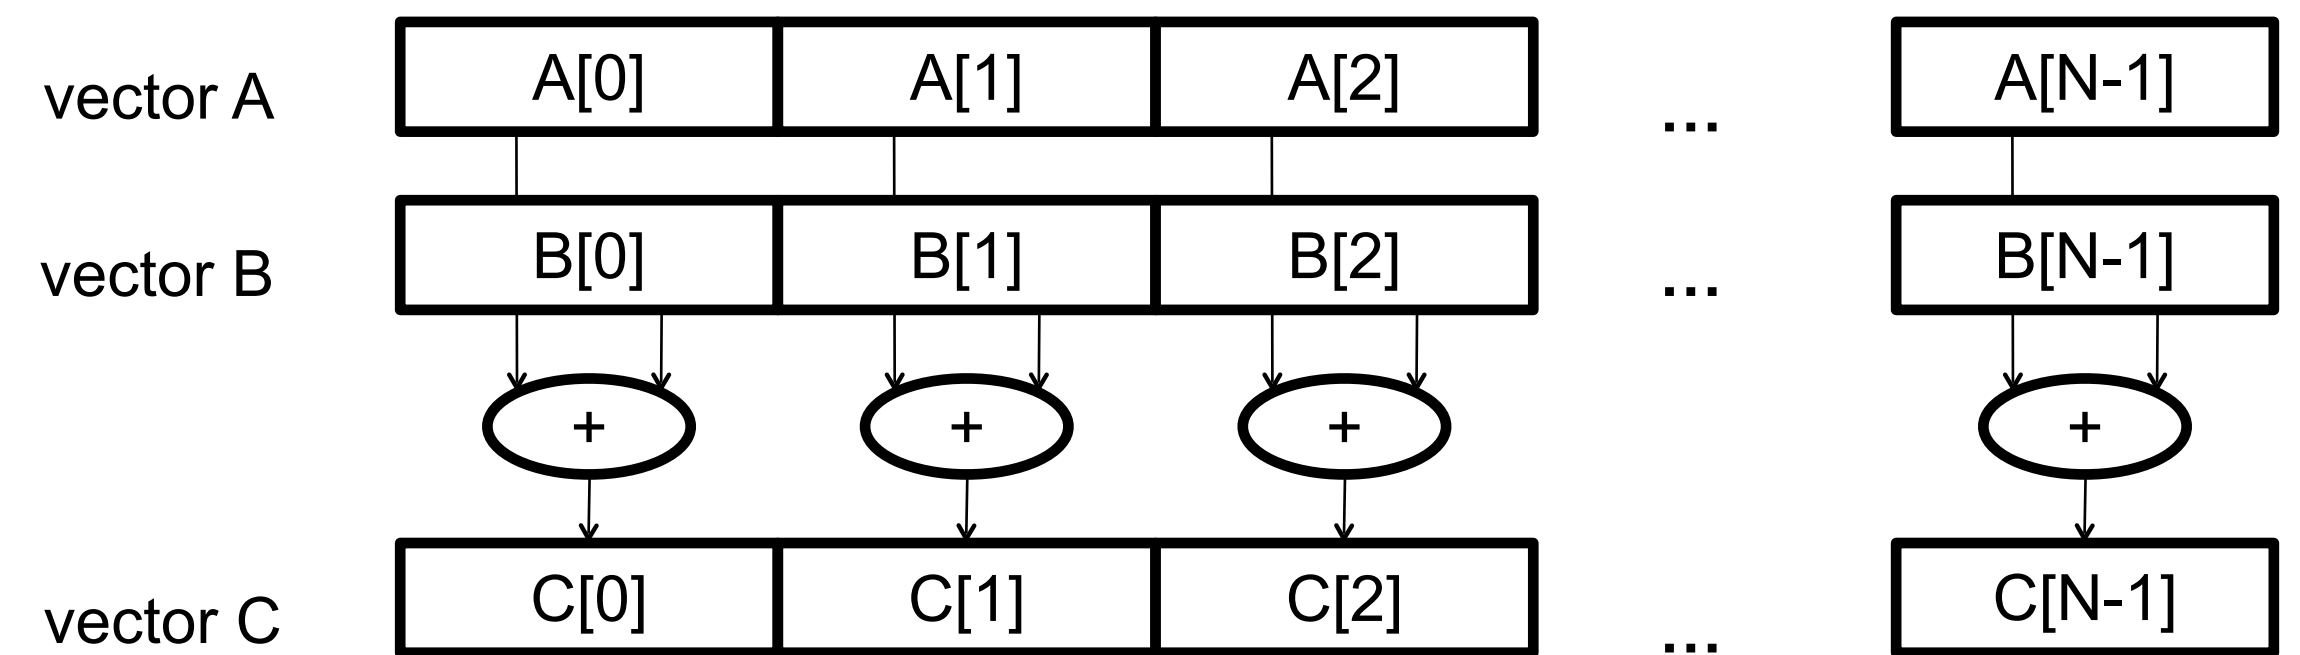
Vector Addition Example

C program computing the sum of two integer vectors:

```
#define N 1024

void vsum(int* a, int* b, int* c){
    int i;
    for (i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}

int main(){
    int va[N], vb[N], vc[N];
    /*... input data acquisition ...*/
    vsum(va, vb, vc);
    /*... output data visualization ...*/
}
```



Structure of a CUDA Program

```
#include <cuda_runtime.h>
#define N 1024

/* kernel function*/

int main(){
    int h_va[N], h_vb[N], h_vc[N];
    /*... input data acquisition ...*/
    /* device memory allocation */
    /* CPU->GPU data transmission */
    /* kernel launch */
    /* GPU->CPU data transmission */
    /* device memory freeing */
    /*... output data visualization ...*/
}
```

Structure of a CUDA Program

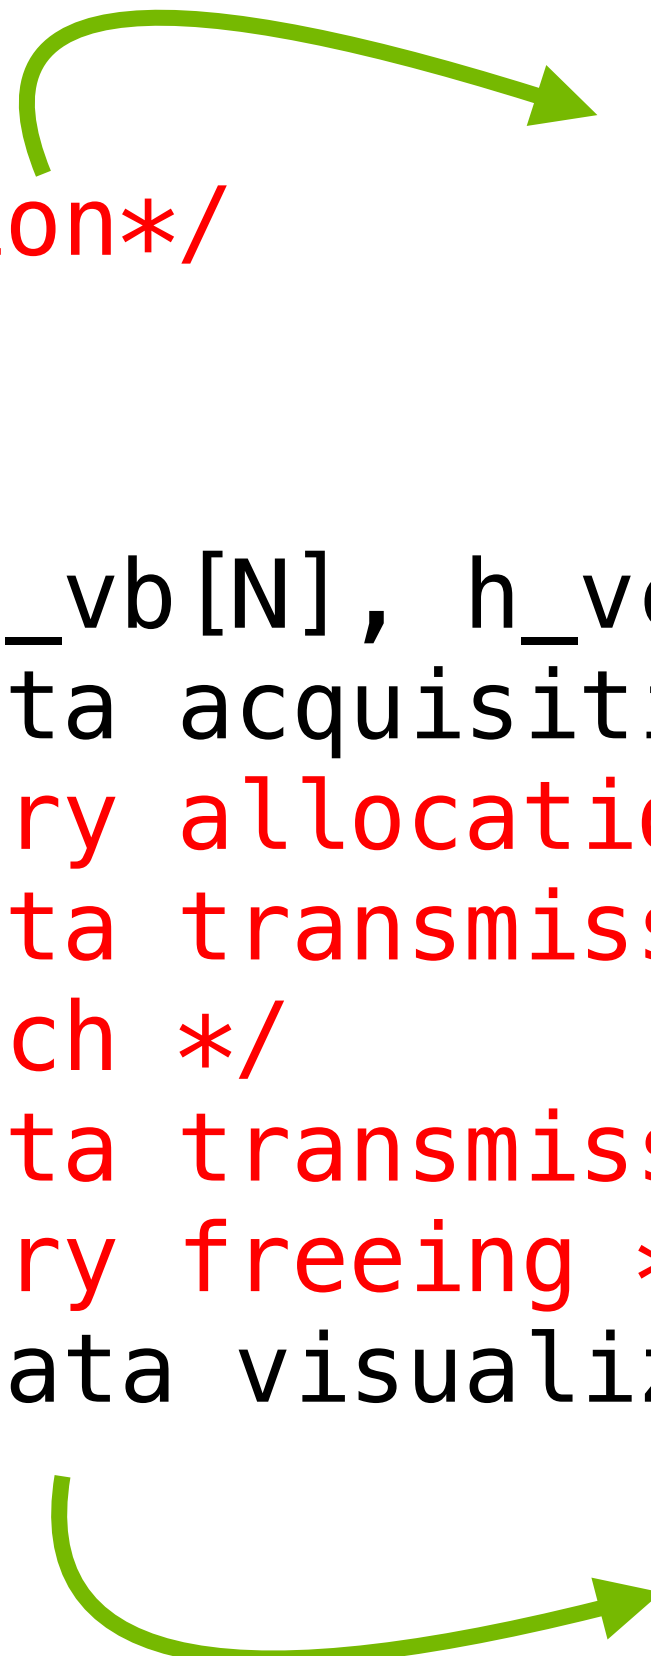
```
#include <cuda_runtime.h>
#define N 1024

/* kernel function*/

int main(){
    int h_va[N], h_vb[N], h_vc[N];
    /*... input data acquisition ...*/
    /* device memory allocation */
    /* CPU->GPU data transmission */
    /* kernel launch */
    /* GPU->CPU data transmission */
    /* device memory freeing */
    /*... output data visualization ...*/
}
```

Device code

Host code



- It is a C/C++ program with CUDA extensions
- The kernel function has some restrictions
 - Access to device memory only
 - **void** return type
 - Asynchronous behavior
 - No variable number of parameters
 - No static variables
 - No function pointers

Kernel Function

```
#include <cuda_runtime.h>
#define N 1024

__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
        + threadIdx.x;

    c[i] = a[i] + b[i];
}
```


Kernel Function

```
#include <cuda_runtime.h>
#define N 1024
```

```
__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
          + threadIdx.x;

    c[i] = a[i] + b[i];
}
```

Each GPU function has a **function qualifier**

	Executed on the:	Only callable from the:
__device__	device	device
__global__	device	host
__host__	host	host

Kernel Function

```
#include <cuda_runtime.h>
#define N 1024

__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
        + threadIdx.x;

    c[i] = a[i] + b[i];
}
```

Input to the `__global__` kernel are
pointers to the device memory

- Host and device codes have **separate scopes**
- The memory locations pointed by the parameters are accessed by **all the threads**

Kernel Function

```
#include <cuda_runtime.h>
#define N 1024

__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
        + threadIdx.x;

    c[i] = a[i] + b[i];
}
```

Local variables are
per-thread private

Kernel Function

```
#include <cuda_runtime.h>
#define N 1024
```

```
__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
          + threadIdx.x;

    c[i] = a[i] + b[i];
}
```

- **blockIdx**: id of the block in the grid
- **blockDim**: size of the block
- **threadIdx**: id of the thread in the block
- **gridDim**: size of the grid

CUDA built-in struct variables with three fields: x, y, z (3D grid!)

Kernel Function

```
#include <cuda_runtime.h>
#define N 1024
```

```
__global__ void vsumKernel(
    int* a,
    int* b,
    int *c)
{
    int i = blockIdx.x * blockDim.x
          + threadIdx.x;

    c[i] = a[i] + b[i];
}
```

- **Mapping of each thread to the single data element**
- **Elaboration on the data element**

Device Memory Allocation

```
#include <cuda_runtime.h>
#define N 1024

int main(){
    int h_va[N], h_vb[N], h_vc[N];
    int *d_va, *d_vb, *d_vc;
    /* device memory allocation */
    cudaMalloc(&d_va, N*sizeof(int));
    cudaMalloc(&d_vb, N*sizeof(int));
    cudaMalloc(&d_vc, N*sizeof(int));
    /* ... */
}
```

↓
Pointer storing the
address of the
allocated memory

↘
Size in bytes to
be allocated

Device memory has to be allocated
similarly to host memory (heap)

**Device memory is separate
from the host one!**

Pointers to device memory
CANNOT be used in the host
program

Data Transfer to the Device

```
int main(){
    /*...*/
    /* CPU->GPU data transmission */
    cudaMemcpy(d_va, h_va, N*sizeof(int),
               cudaMemcpyHostToDevice);
    cudaMemcpy(d_vb, h_vb, N*sizeof(int), cudaMemcpyHostToDevice);
    /* ... */
}
```

Destination
memory

Source
memory

Size in bytes
to be
transmitted

Transmission
direction

**Data must be transferred
explicitly to the device
memory**

- cudaMemcpyDeviceToHost
- cudaMemcpyHostToDevice
- cudaMemcpyDeviceToDevice
- cudaMemcpyHostToHost

Kernel Launch

```
/*...*/  
int main(){  
    /*...*/  
    /* kernel launch */  
    dim3 blocksPerGrid(N/256, 1, 1);  
    dim3 threadsPerBlock(256, 1, 1);  
    vsumKernel<<<blocksPerGrid, threadsPerBlock>>>(d_va, d_vb, d_vc);  
    /* ... */  
}
```

Execution configuration parameters:
block and grid dimension (dim3 is a
struct with three fields: x, y, z)

Kernel Launch

```
/*...*/  
int main(){  
    /*...*/  
    /* kernel launch */  
    dim3 blocksPerGrid(N/256, 1, 1);  
    dim3 threadsPerBlock(256, 1, 1);  
    vsumKernel<<<blocksPerGrid, threadsPerBlock>>>(d_va, d_vb, d_vc);  
    /* ... */  
}
```

The parameters are the
pointers to device memory

Data Transfer to the Host

```
/*...*/
```

**Data must be transferred
explicitly to the host memory**

```
int main(){  
    /*...*/  
    /* GPU->CPU data transmission */  
    cudaMemcpy(h_vc, d_vc, N*sizeof(int), cudaMemcpyDeviceToHost);  
    /* ... */  
}
```

Diagram illustrating the parameters of the `cudaMemcpy` function call:

- `h_vc`: Destination memory
- `d_vc`: Source memory
- `N*sizeof(int)`: Size in bytes to be transmitted
- `cudaMemcpyDeviceToHost`: Transmission direction

Device Memory Freeing

```
/*...*/
```

```
int main(){  
    /*...*/  
    /* device memory freeing */  
    cudaFree(d_va);  
    cudaFree(d_vb);  
    cudaFree(d_vc);  
    /* ... */  
}
```

Device memory has to be **explicitly released** by the application via `cudaFree()`

`cudaDeviceReset()` may be called to reset all resources of the device used by the process

Complete Code

```
#include <cuda_runtime.h>
#define N 1024

__global__ void vsumKernel(int* a, int* b, int *c){
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    c[i] = a[i] + b[i];
}

int main(){
    int h_va[N], h_vb[N], h_vc[N];
    int *d_va, *d_vb, *d_vc;

    /*... input data acquisition ...*/

    cudaMalloc(&d_va, N*sizeof(int));
    cudaMalloc(&d_vb, N*sizeof(int));
    cudaMalloc(&d_vc, N*sizeof(int));

    cudaMemcpy(d_va, h_va, N*sizeof(int),
               cudaMemcpyHostToDevice);
    cudaMemcpy(d_vb, h_vb, N*sizeof(int),
               cudaMemcpyHostToDevice);

    dim3 blocksPerGrid(N/256, 1, 1);
    dim3 threadsPerBlock(256, 1, 1);
    vsumKernel<<<blocksPerGrid, threadsPerBlock>>>
        (d_va, d_vb, d_vc);

    cudaMemcpy(h_vc, d_vc, N*sizeof(int),
               cudaMemcpyDeviceToHost);

    cudaFree(d_va);
    cudaFree(d_vb);
    cudaFree(d_vc);

    /*... output data visualization ...*/
    return 0;
}
```


CUDA

From an application perspective, CUDA defines:

- **Architecture model**

- with many processing cores grouped in multiprocessors who share a SIMT control unit

- **Programming model**

- based on **massive data parallelism** and fine-grained parallelism
- **scalable**: the code is executed on a different number of cores without recompiling it

- **Memory model**

- more **explicit** to the programmer, where **caches are not transparent** anymore

CUDA Memory Model

From an application perspective, CUDA defines:

- **Memory model**
 - more **explicit** to the programmer, where **caches are not transparent** anymore

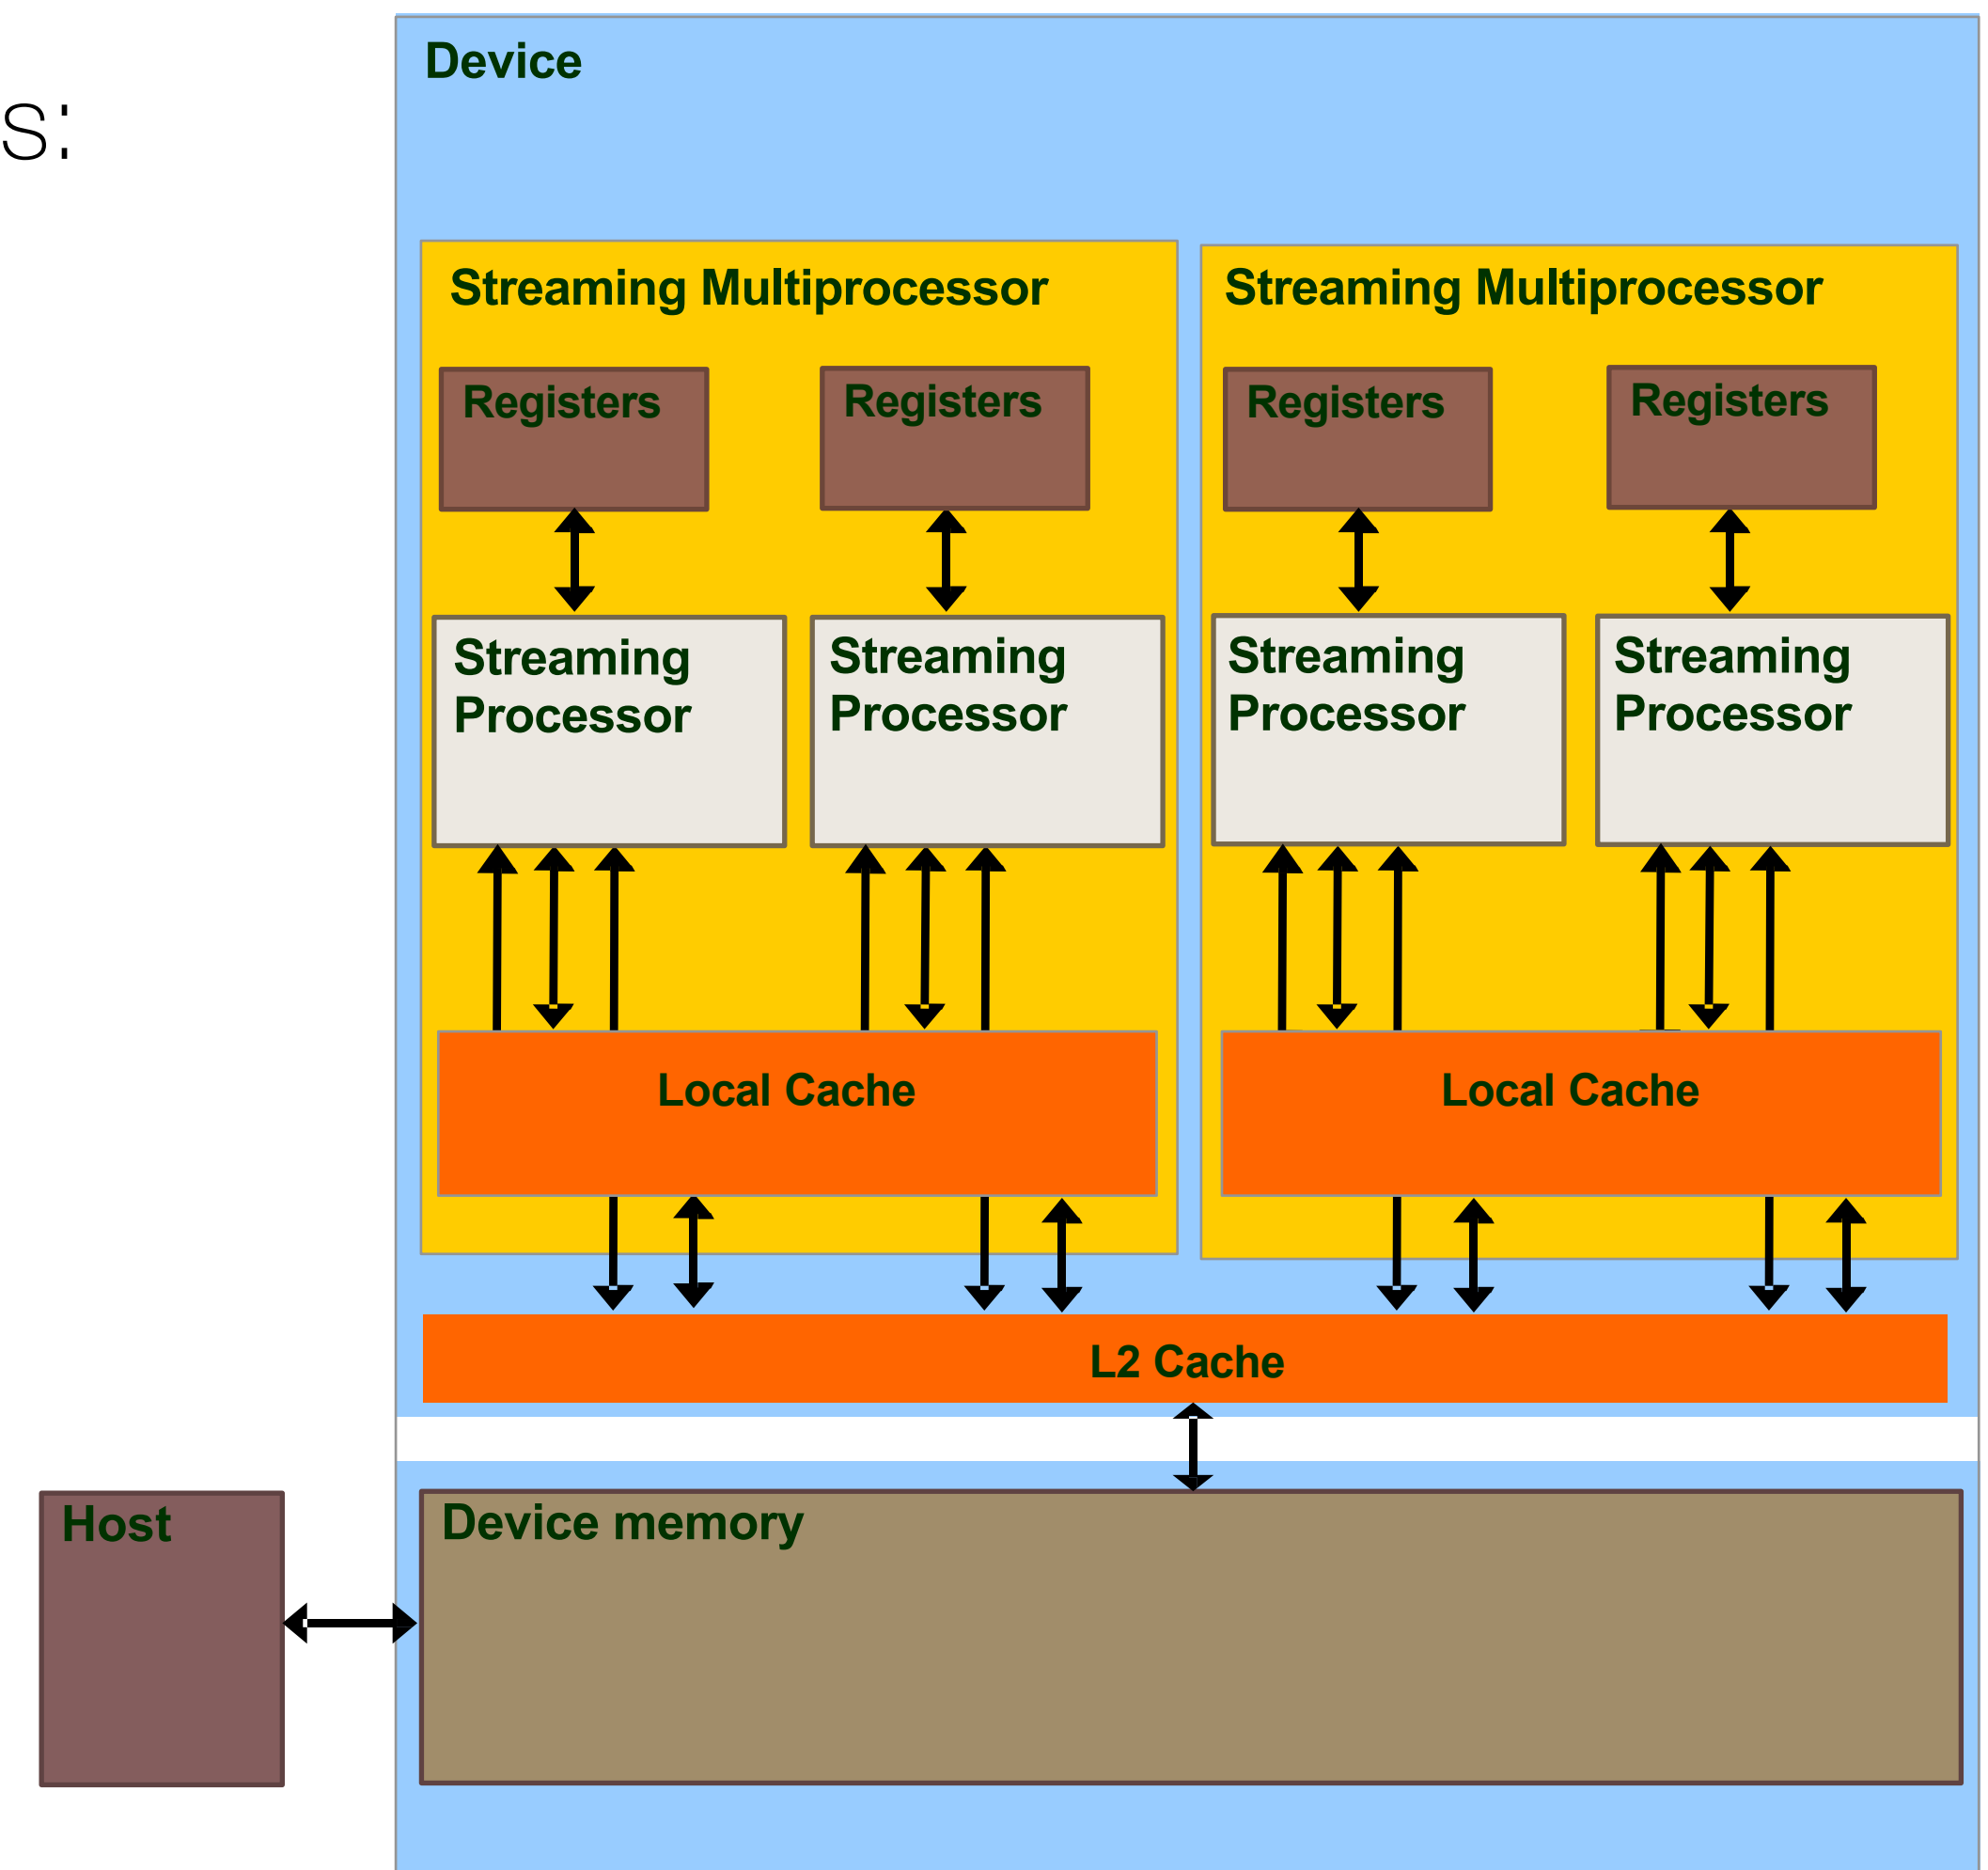
The **CUDA memory model** defines various types of **memory** available in a CUDA-enabled GPU, and it specifies how data is shared and accessed between threads, thread blocks, and the host (CPU).

This memory hierarchy is designed to **optimize** the **performance** and **manage** the **bandwidth** constraints of GPU-based parallel processing.

GPU Memory Hierarchy

GPU memory hierarchy is composed of three elements:

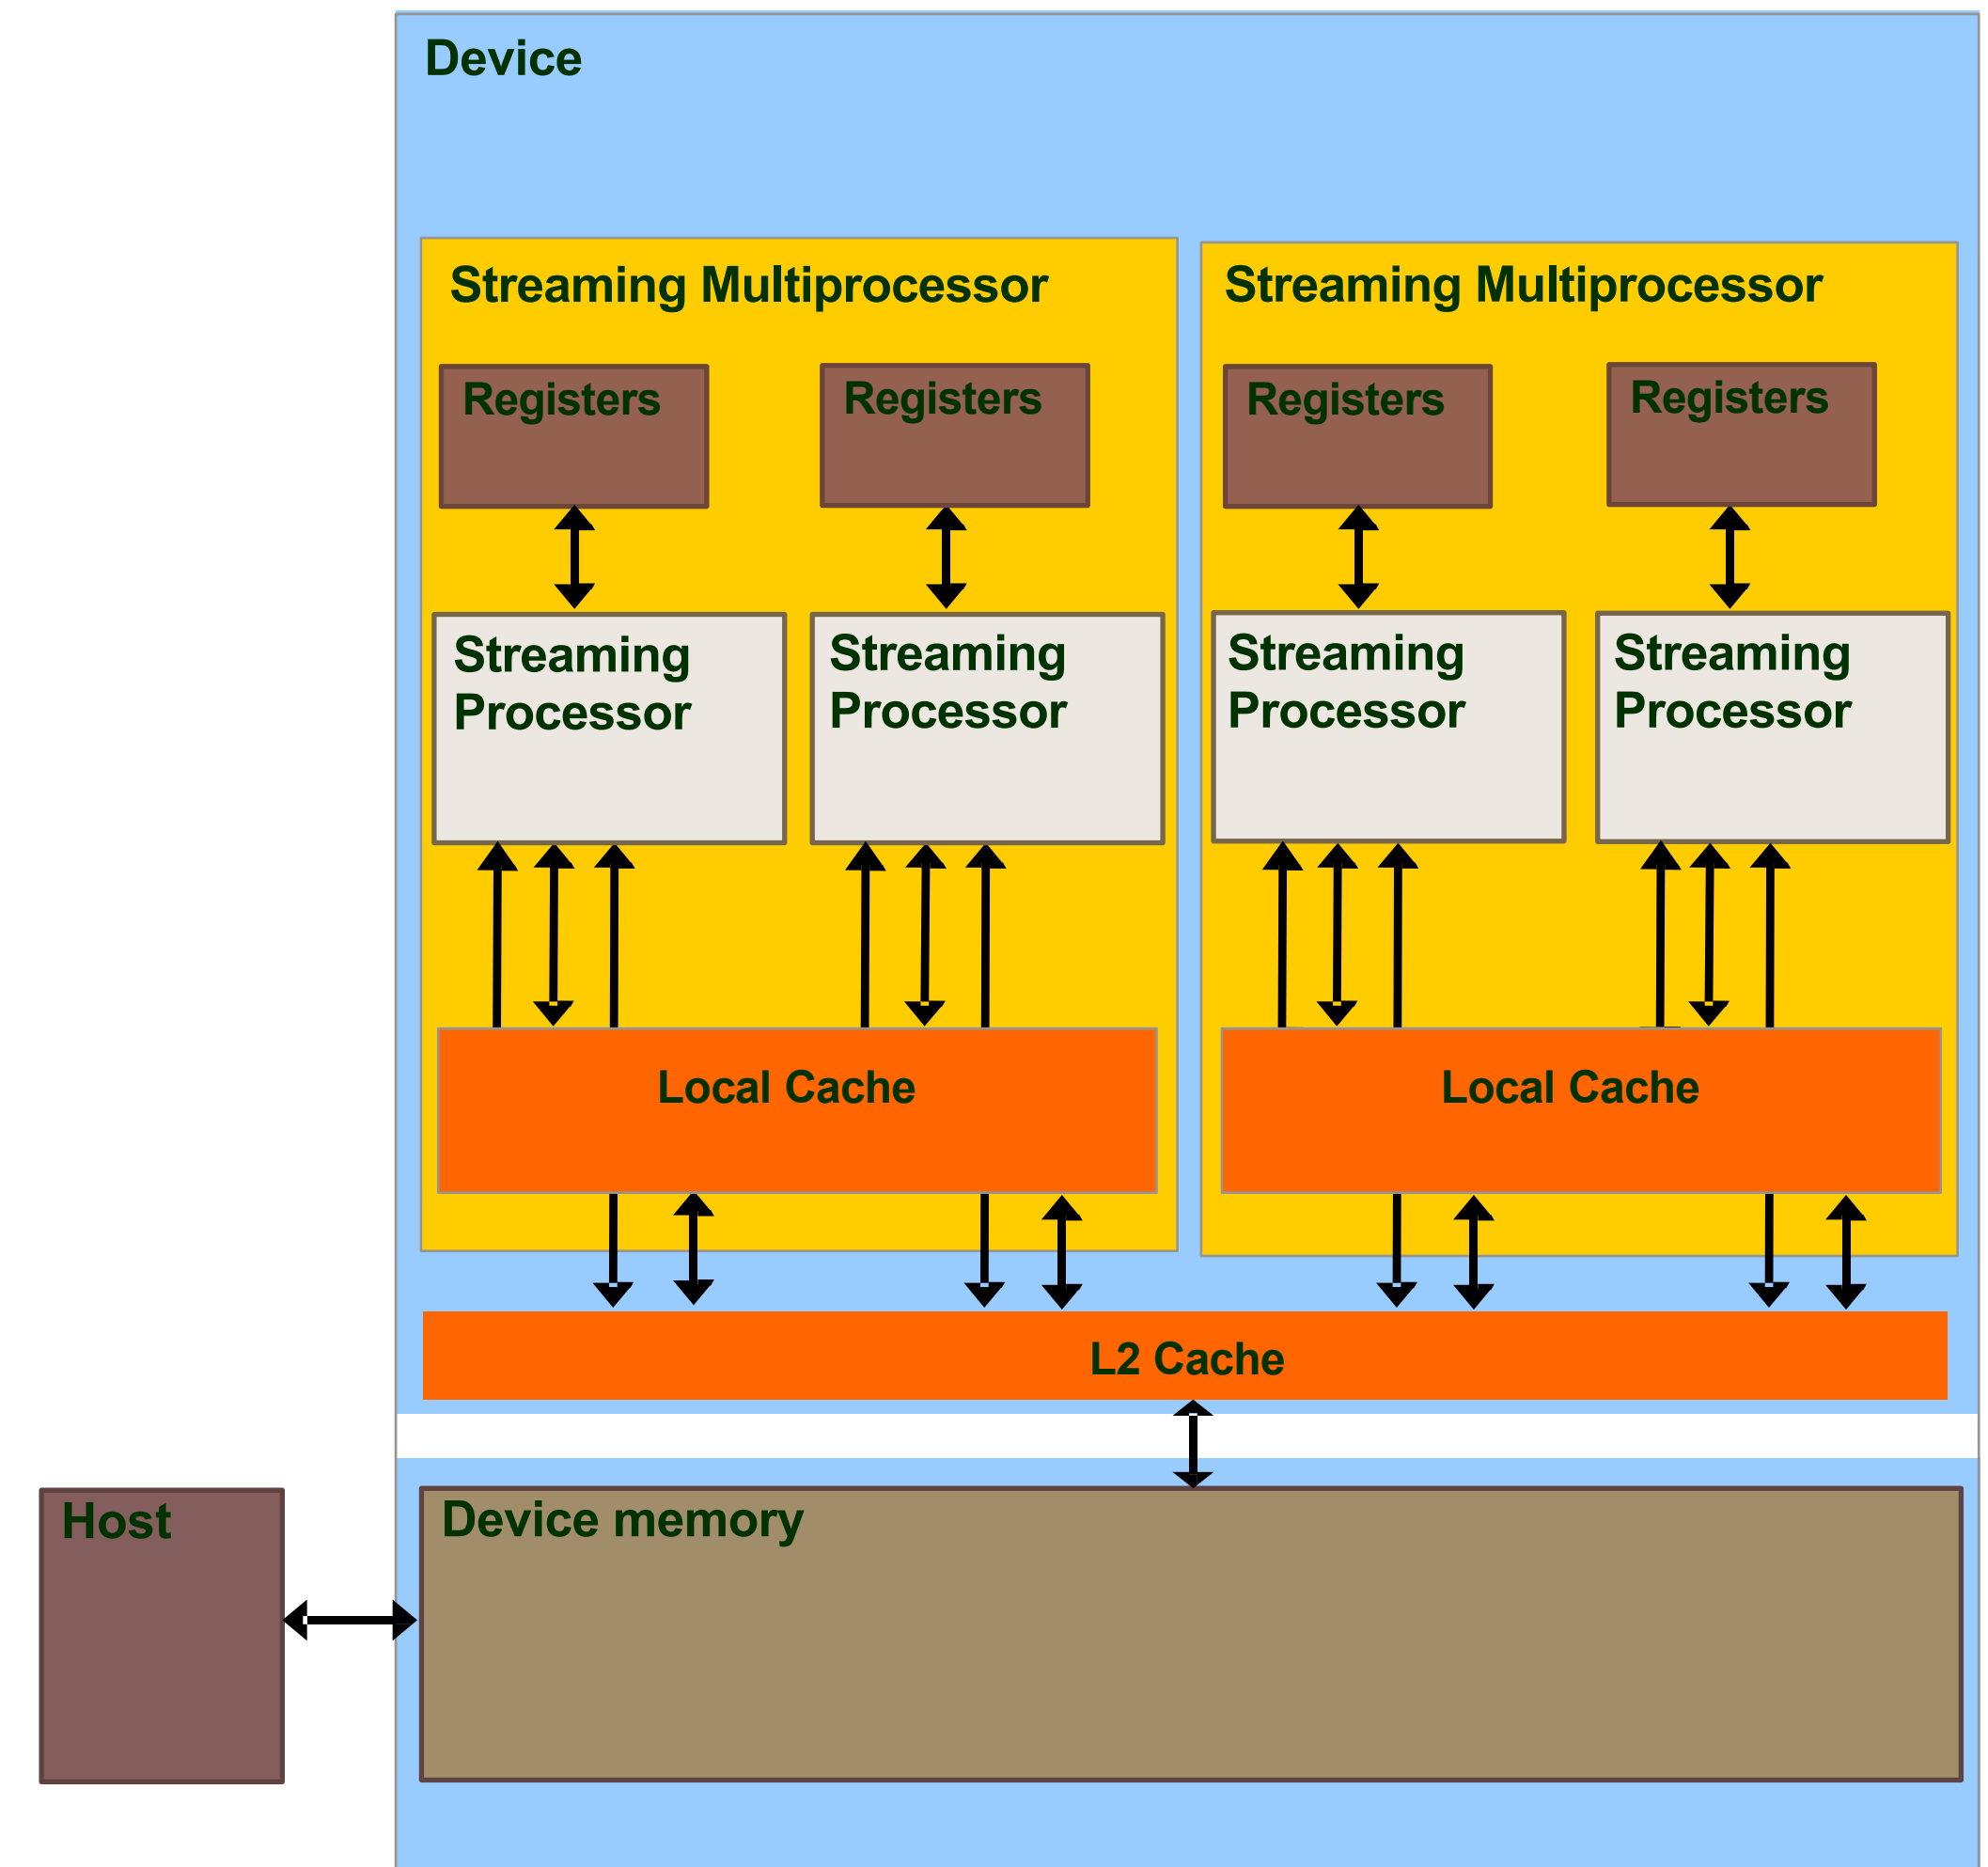
1. Registers
2. Caches
3. Device memory



GPU Memory Hierarchy

Registers:


- ▶ **Private** register file within each SM
- ▶ Partitioned among threads running in the SM
- ▶ Not enough to store all private data of each thread



GPU Memory Hierarchy

Caches:

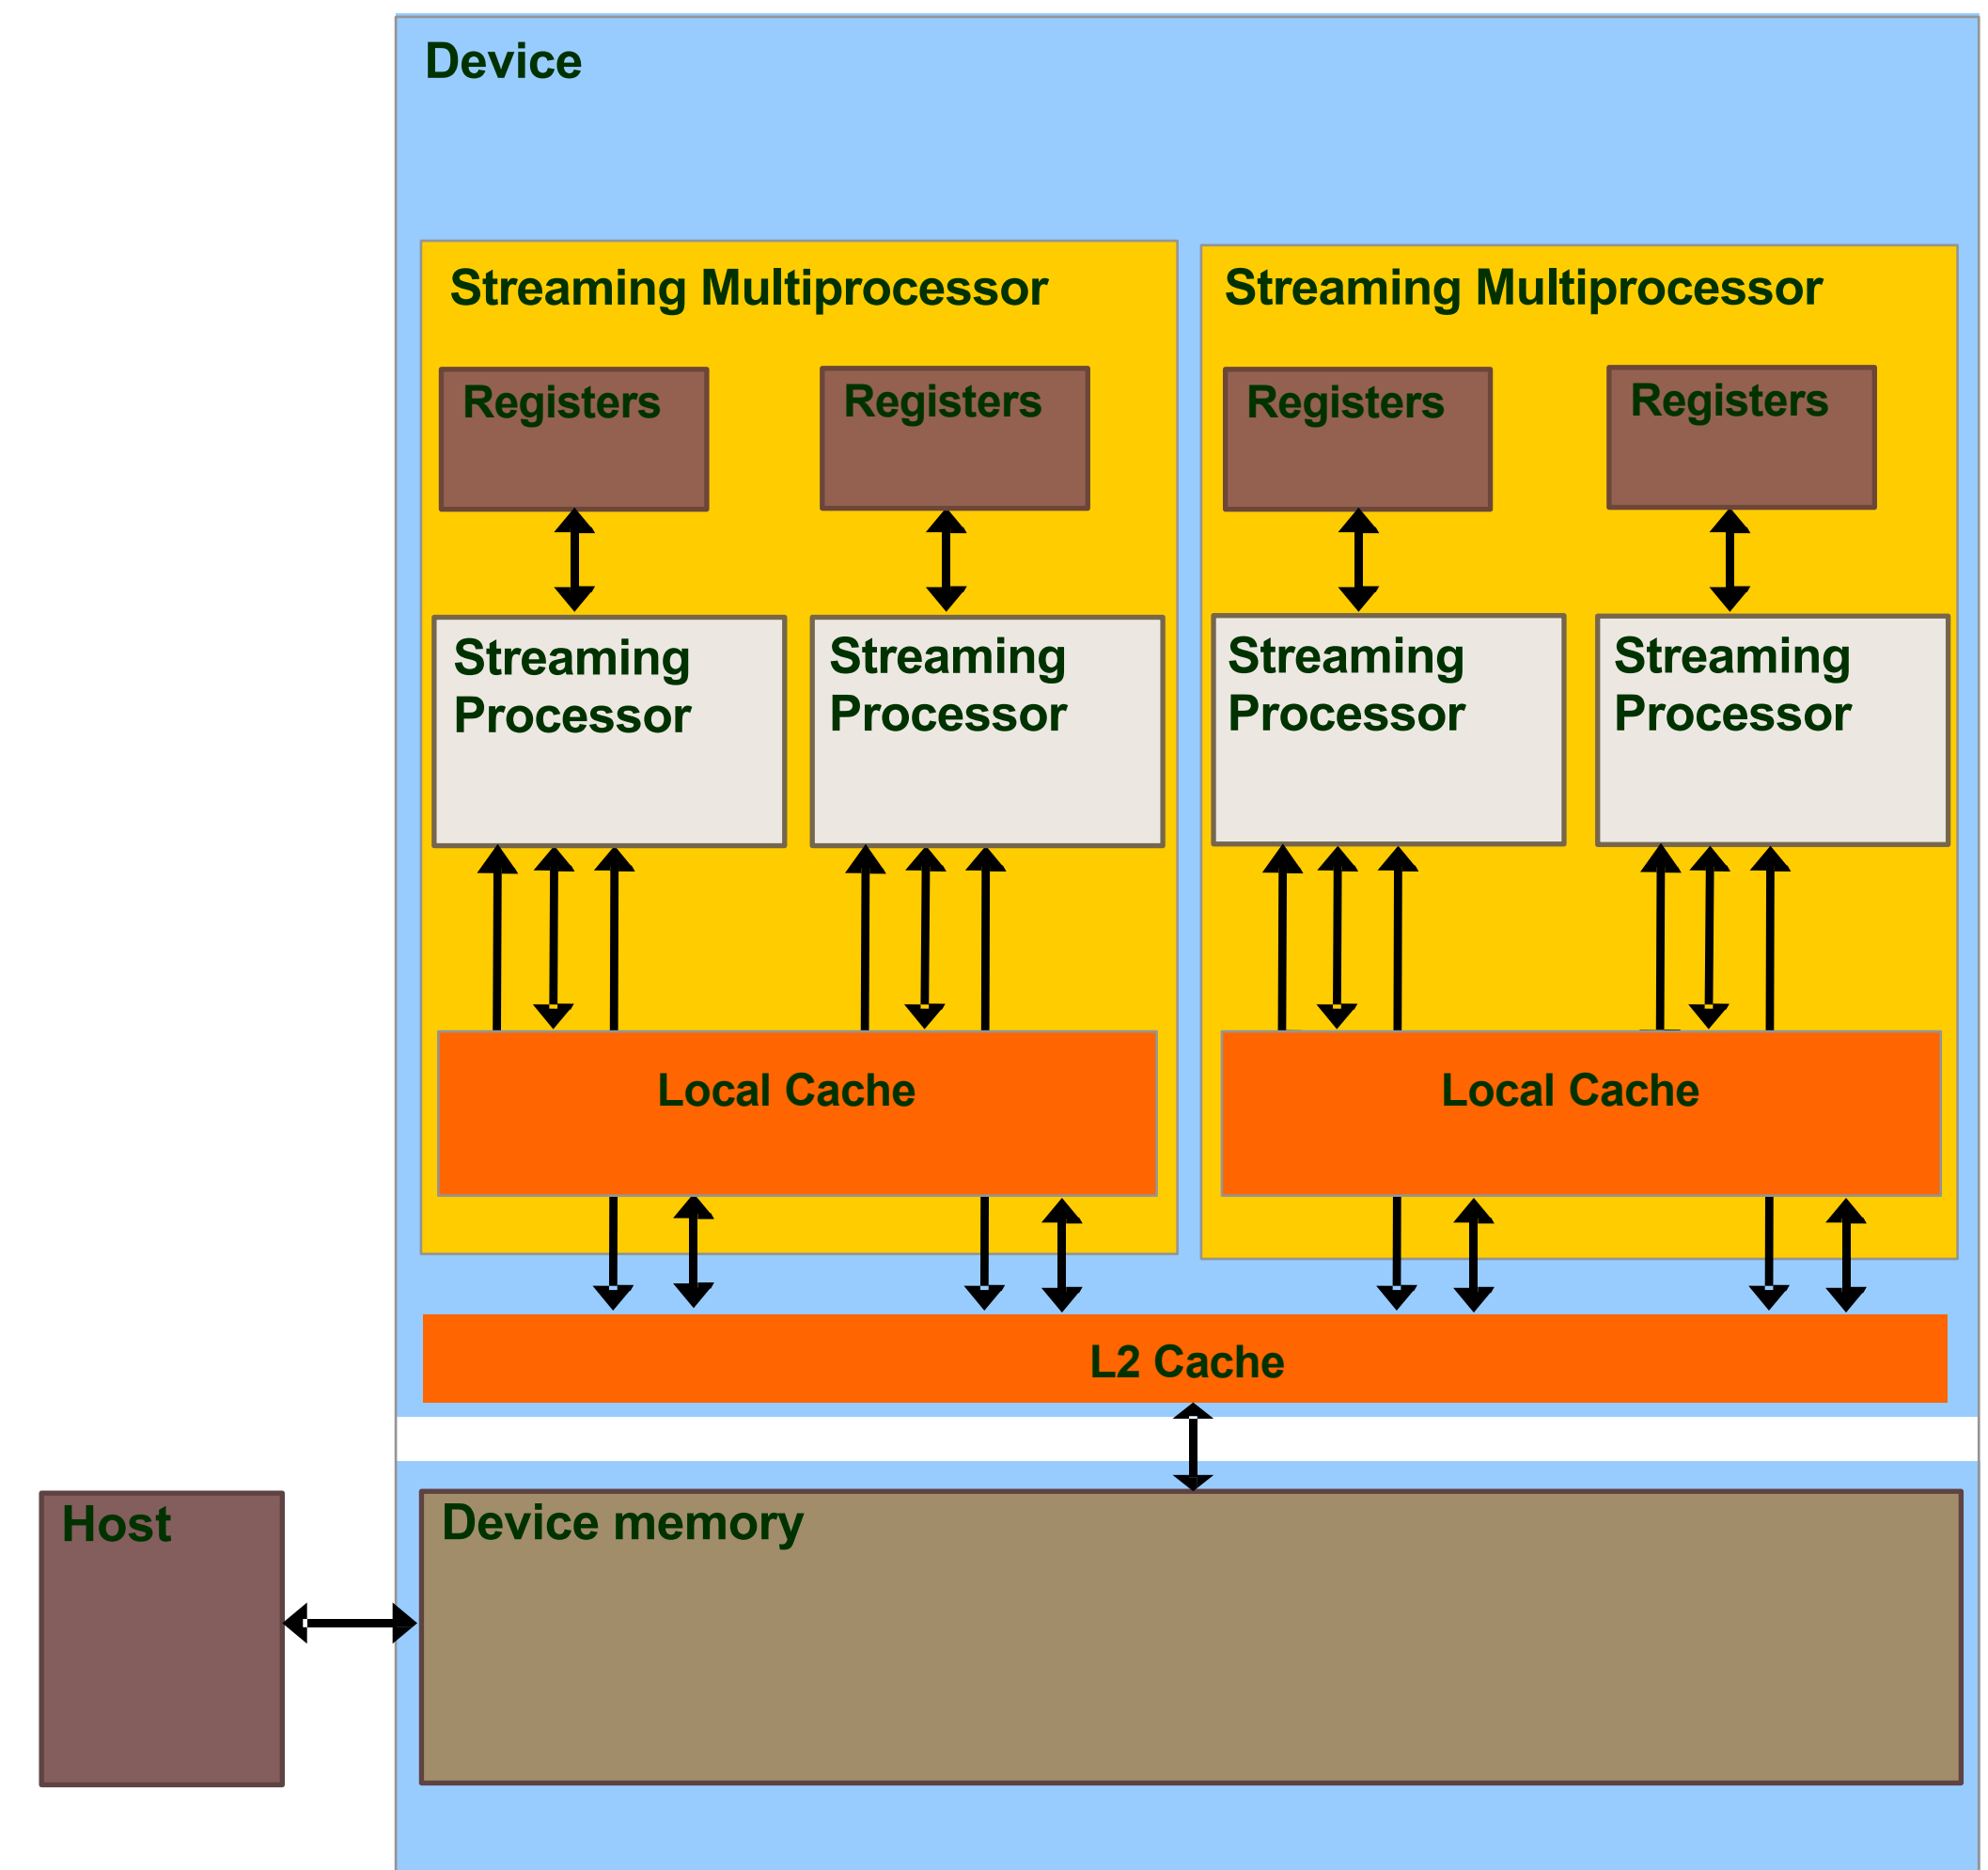
- ▶ Global **L2** cache
- ▶ Per-SM **L1** cache
- ▶ Per-SM **constant** cache
- ▶ Per-SM **read-only/texture** cache
- ▶ Per-SM **shared** memory
- ▶ Access depends on the type of cache



GPU Memory Hierarchy

Device memory:


- ▶ **Off-chip** memory
- ▶ Accessible by all threads on SMs
- ▶ Accessible by the host



GPU Memory Hierarchy

Memory latencies:

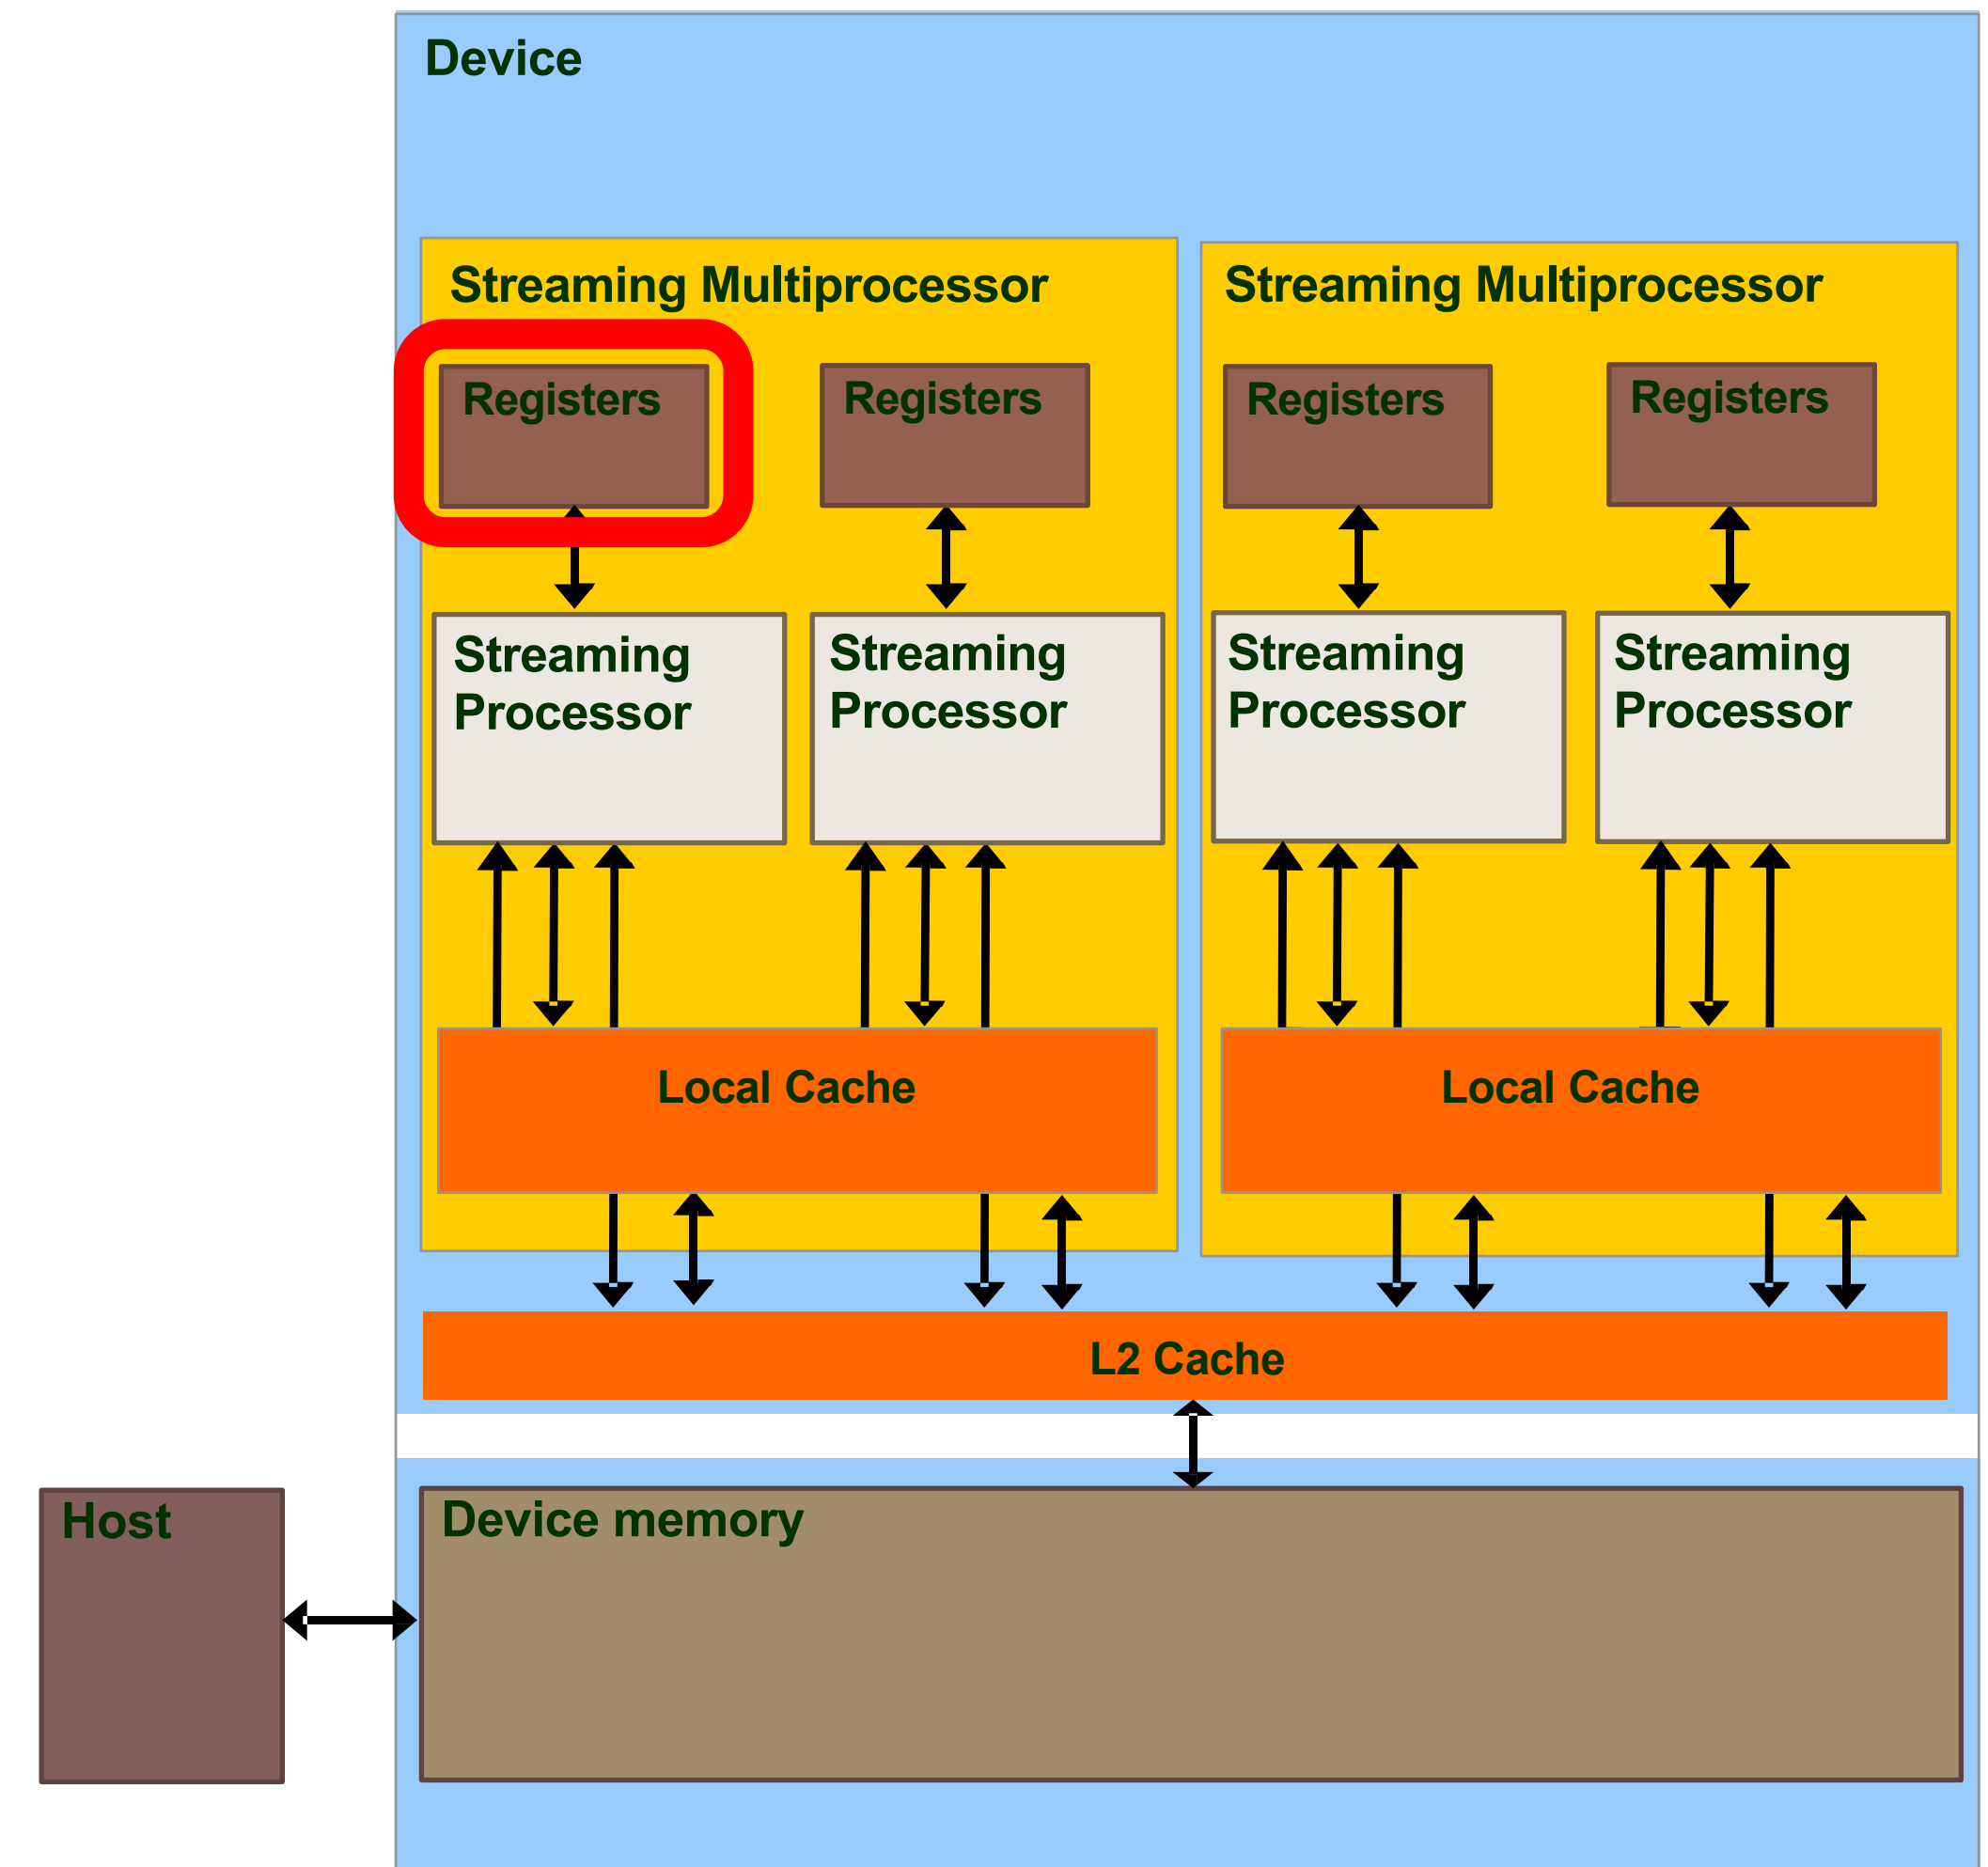
	Cost (cycles)
Registers	1
Device memory	200-800
Shared memory	~1
L1	1
Constant cache	~1
Read-only	1



CUDA Memory Model

Key components:

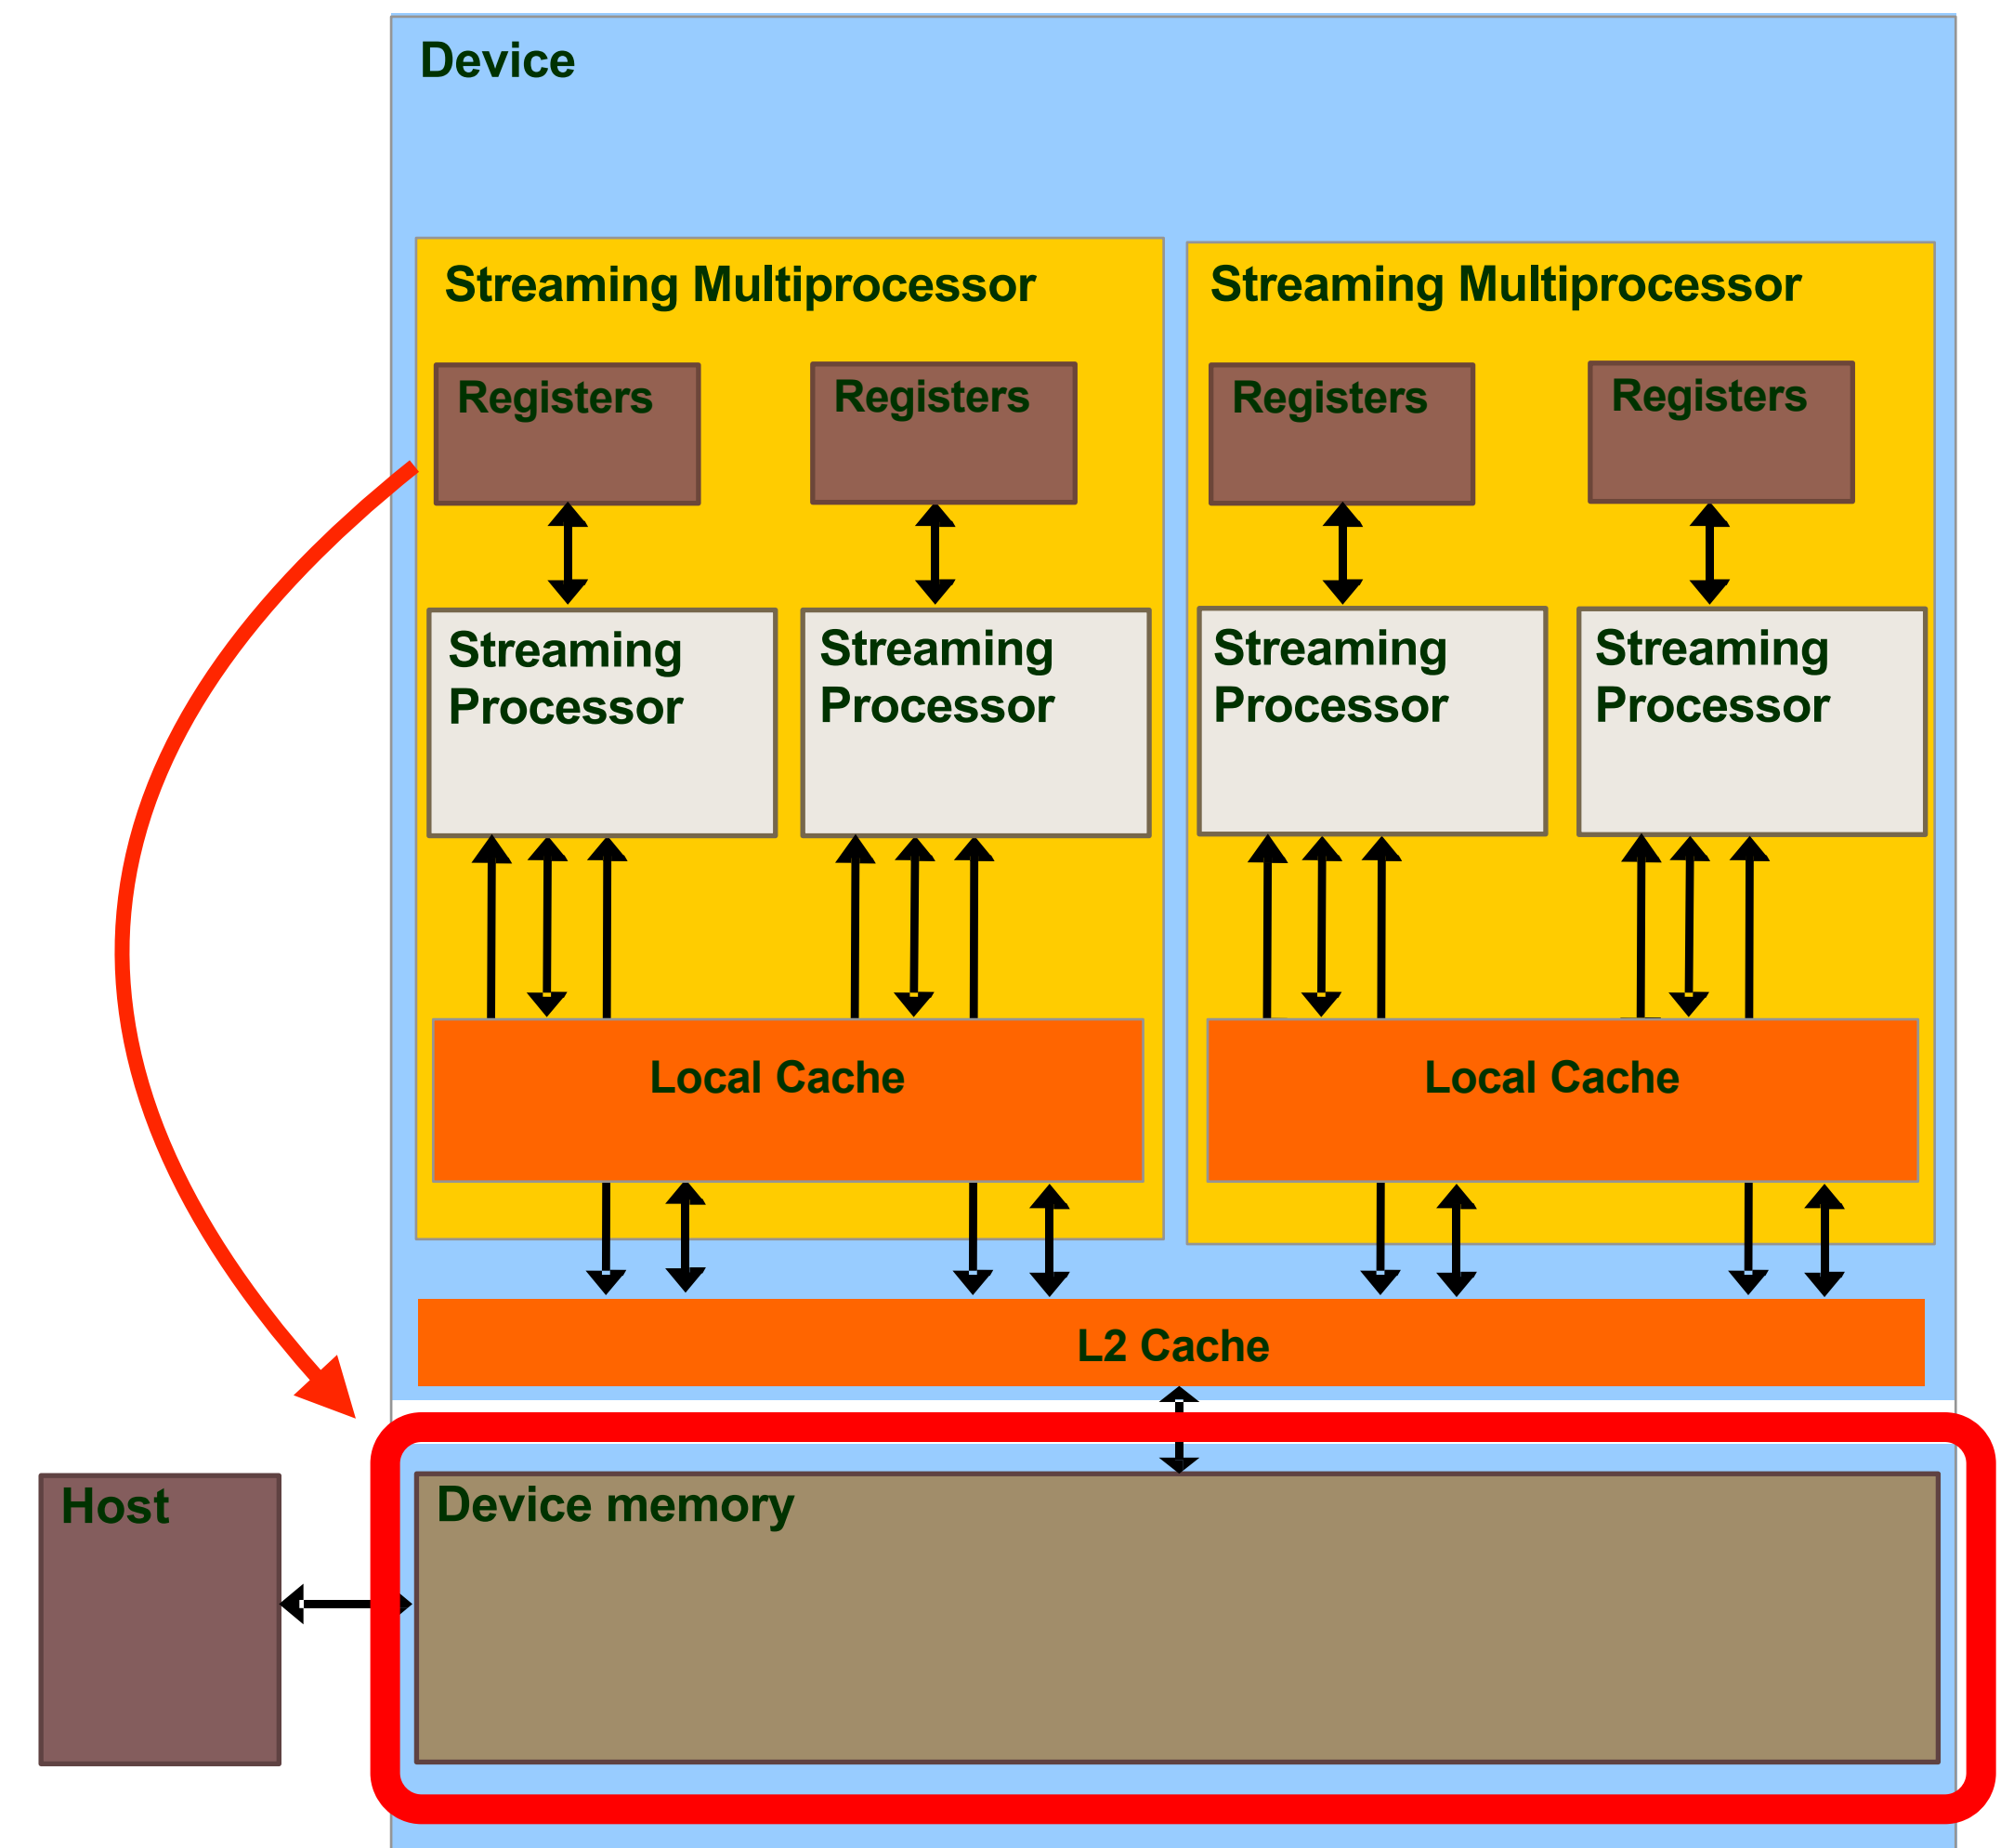
- **Registers (per-thread)**: fastest form of memory and are private to each thread, they store local variables used within a thread, making access almost instantaneous. The number of registers per thread is limited, and excessive register usage can lead to **register spilling** into slower local memory (see next slide).



CUDA Memory Model

Key components:

- **Registers (per-thread)**: fastest form of memory and are private to each thread, they store local variables used within a thread, making access almost instantaneous. The number of registers per thread is limited, and excessive register usage can lead to **register spilling** into slower local memory.
- **Local memory (per-thread)**: resides in device memory rather than being on-chip. It's used when there are not enough registers for all thread-local data. Local memory is primarily accessed when *register spilling* occurs, which should be minimized for performance reasons.



CUDA Memory Model

Key components:


- **Registers (per-thread)**: fastest form of memory and are private to each thread, they store local variables used within a thread, making access almost instantaneous. The number of registers per thread is limited, and excessive register usage can lead to **register spilling** into slower local memory.
- **Local memory (per-thread)**: resides in device memory rather than being on-chip. It's used when there are not enough registers for all thread-local data. Local memory is primarily accessed when register spilling occurs, which should be minimized for performance reasons.
- **Shared memory (per-block)**

Shared Memory

The shared memory is a very fast **scratchpad** memory integrated in the SM (**L1 cache**)

- ▶ Intra-block data sharing and thread cooperation
- ▶ Reduce accesses to the global memory

Together with registers, shared memory represents a constraint on the number of blocks assignable to a SM



Shared Memory

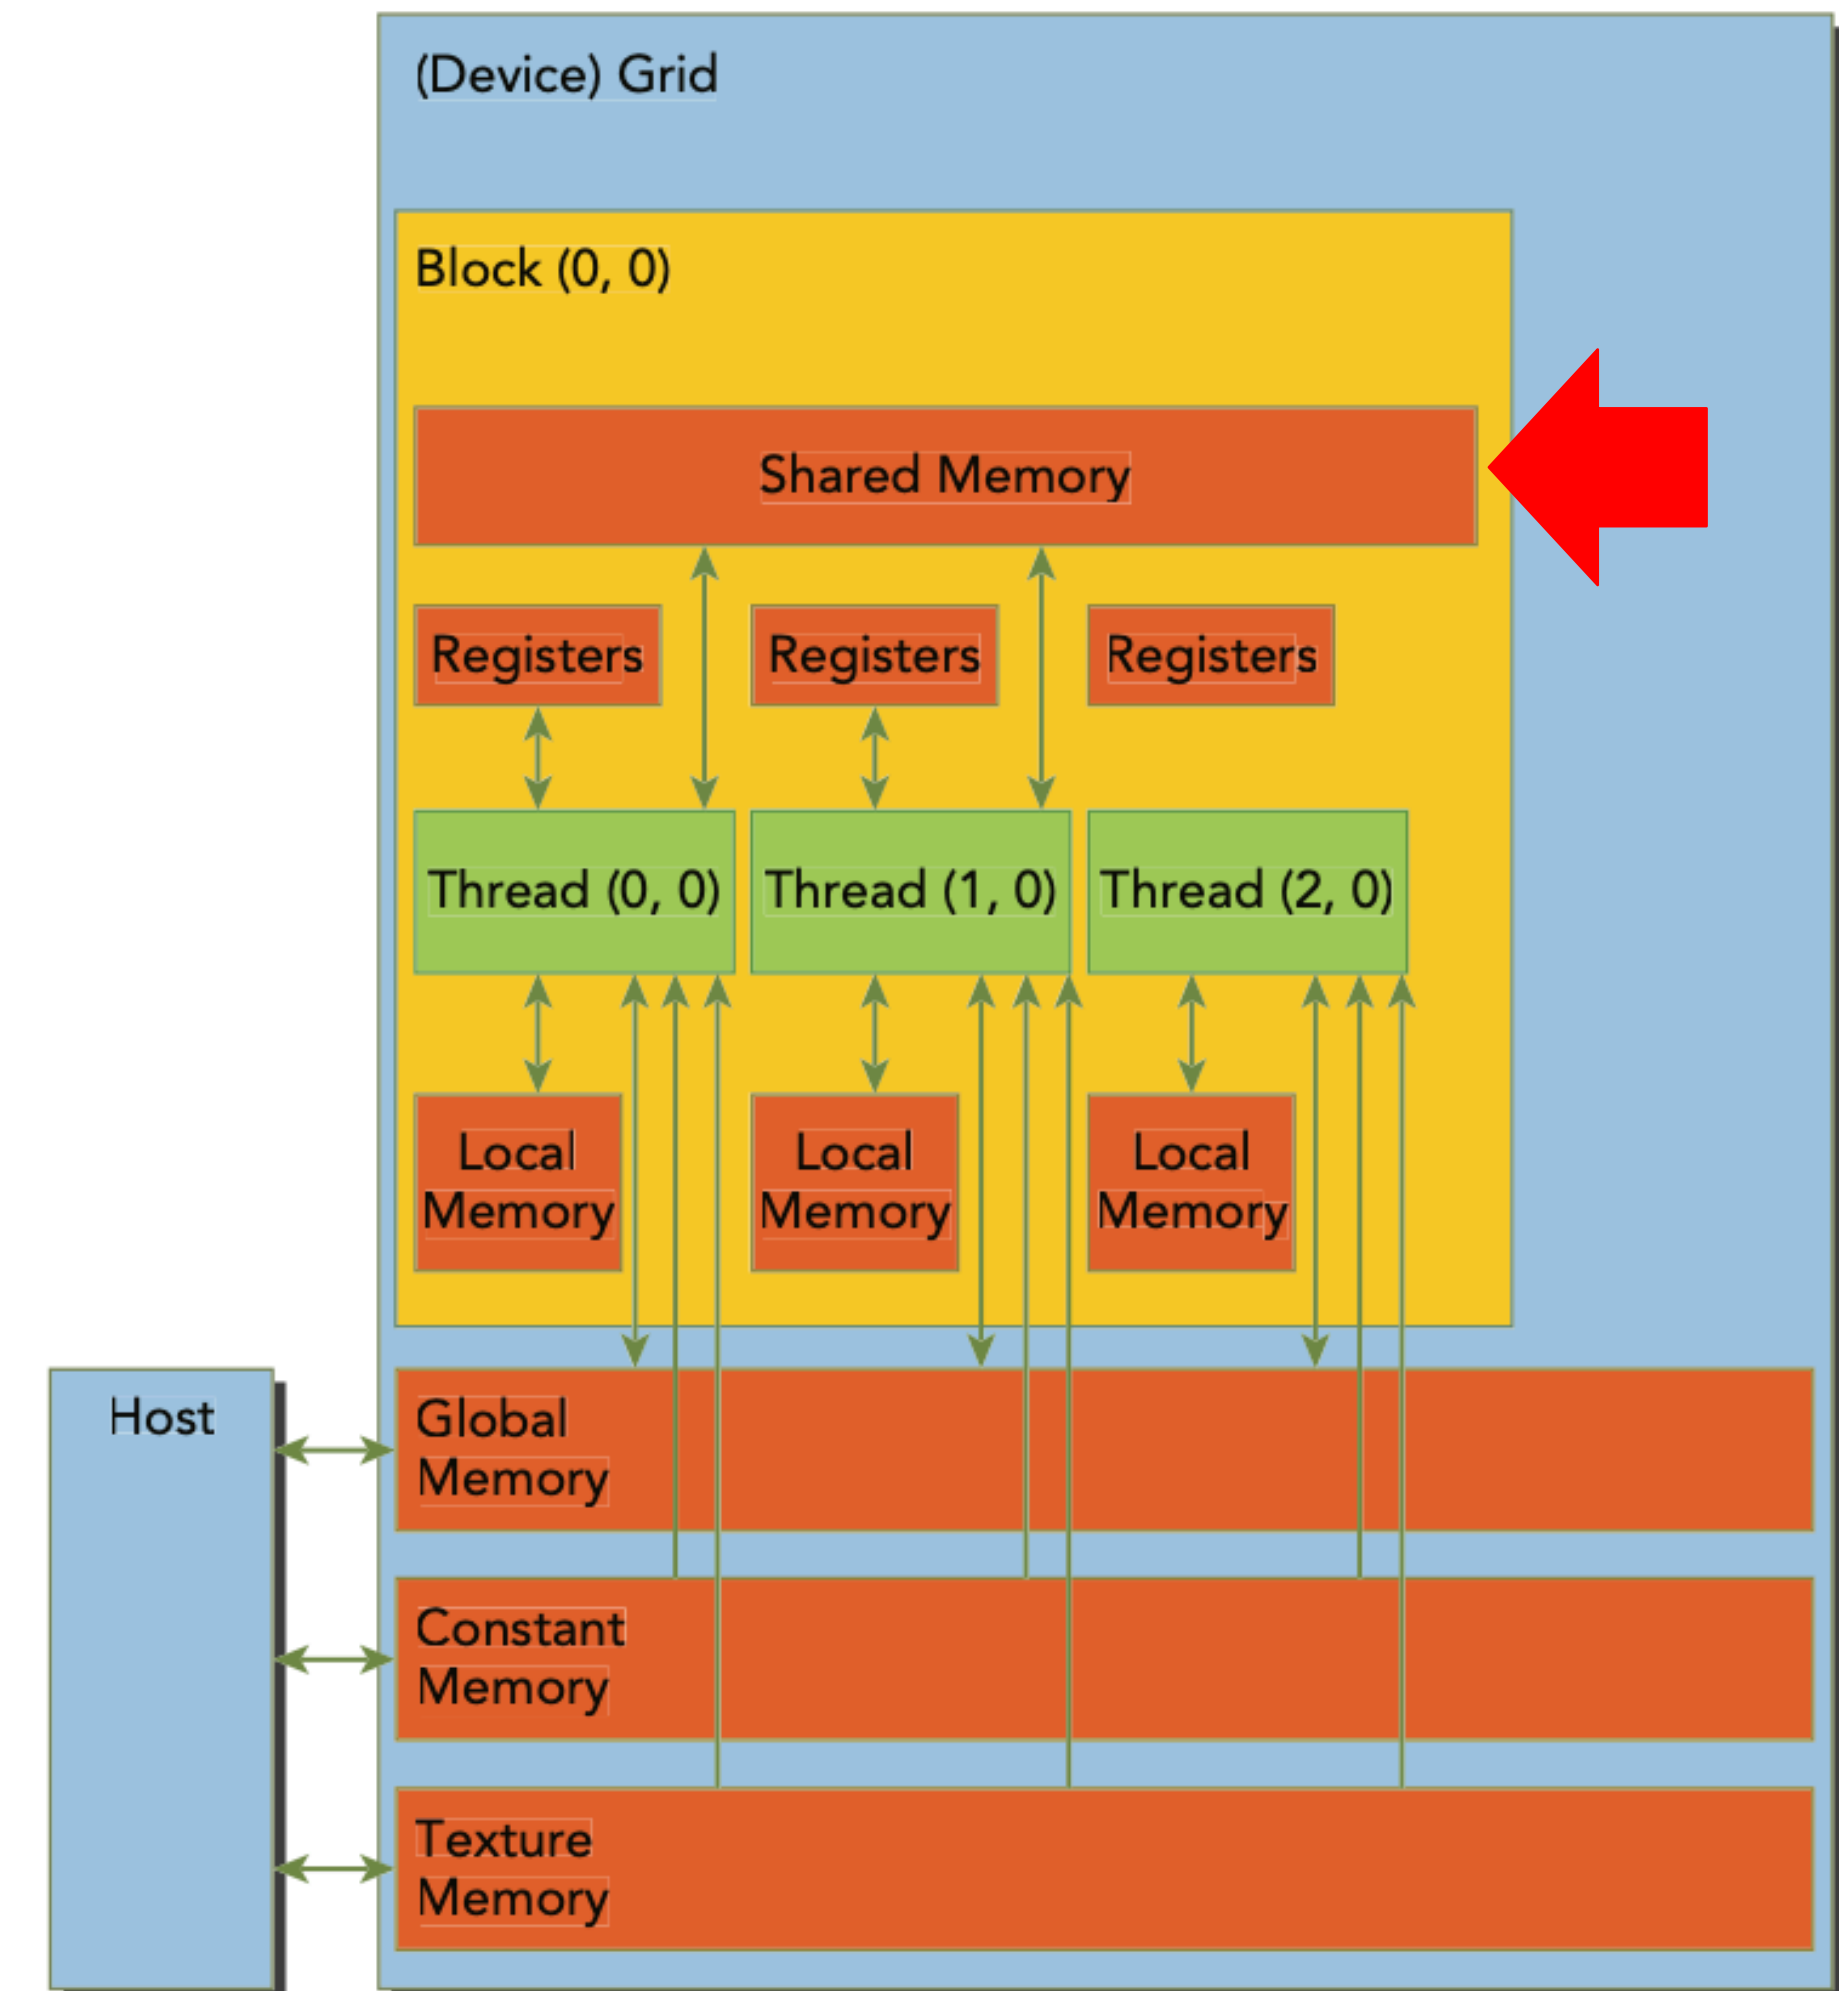
The shared variable is indicated by the **__shared__** qualifier

Two possible ways to allocate shared variables

- Static allocation
- Dynamic allocation

Shared variables are declared in the kernel code or globally (i.e., declared for all the kernels)

- Access: read/write for all block threads
- Scope: block
- Lifetime: block

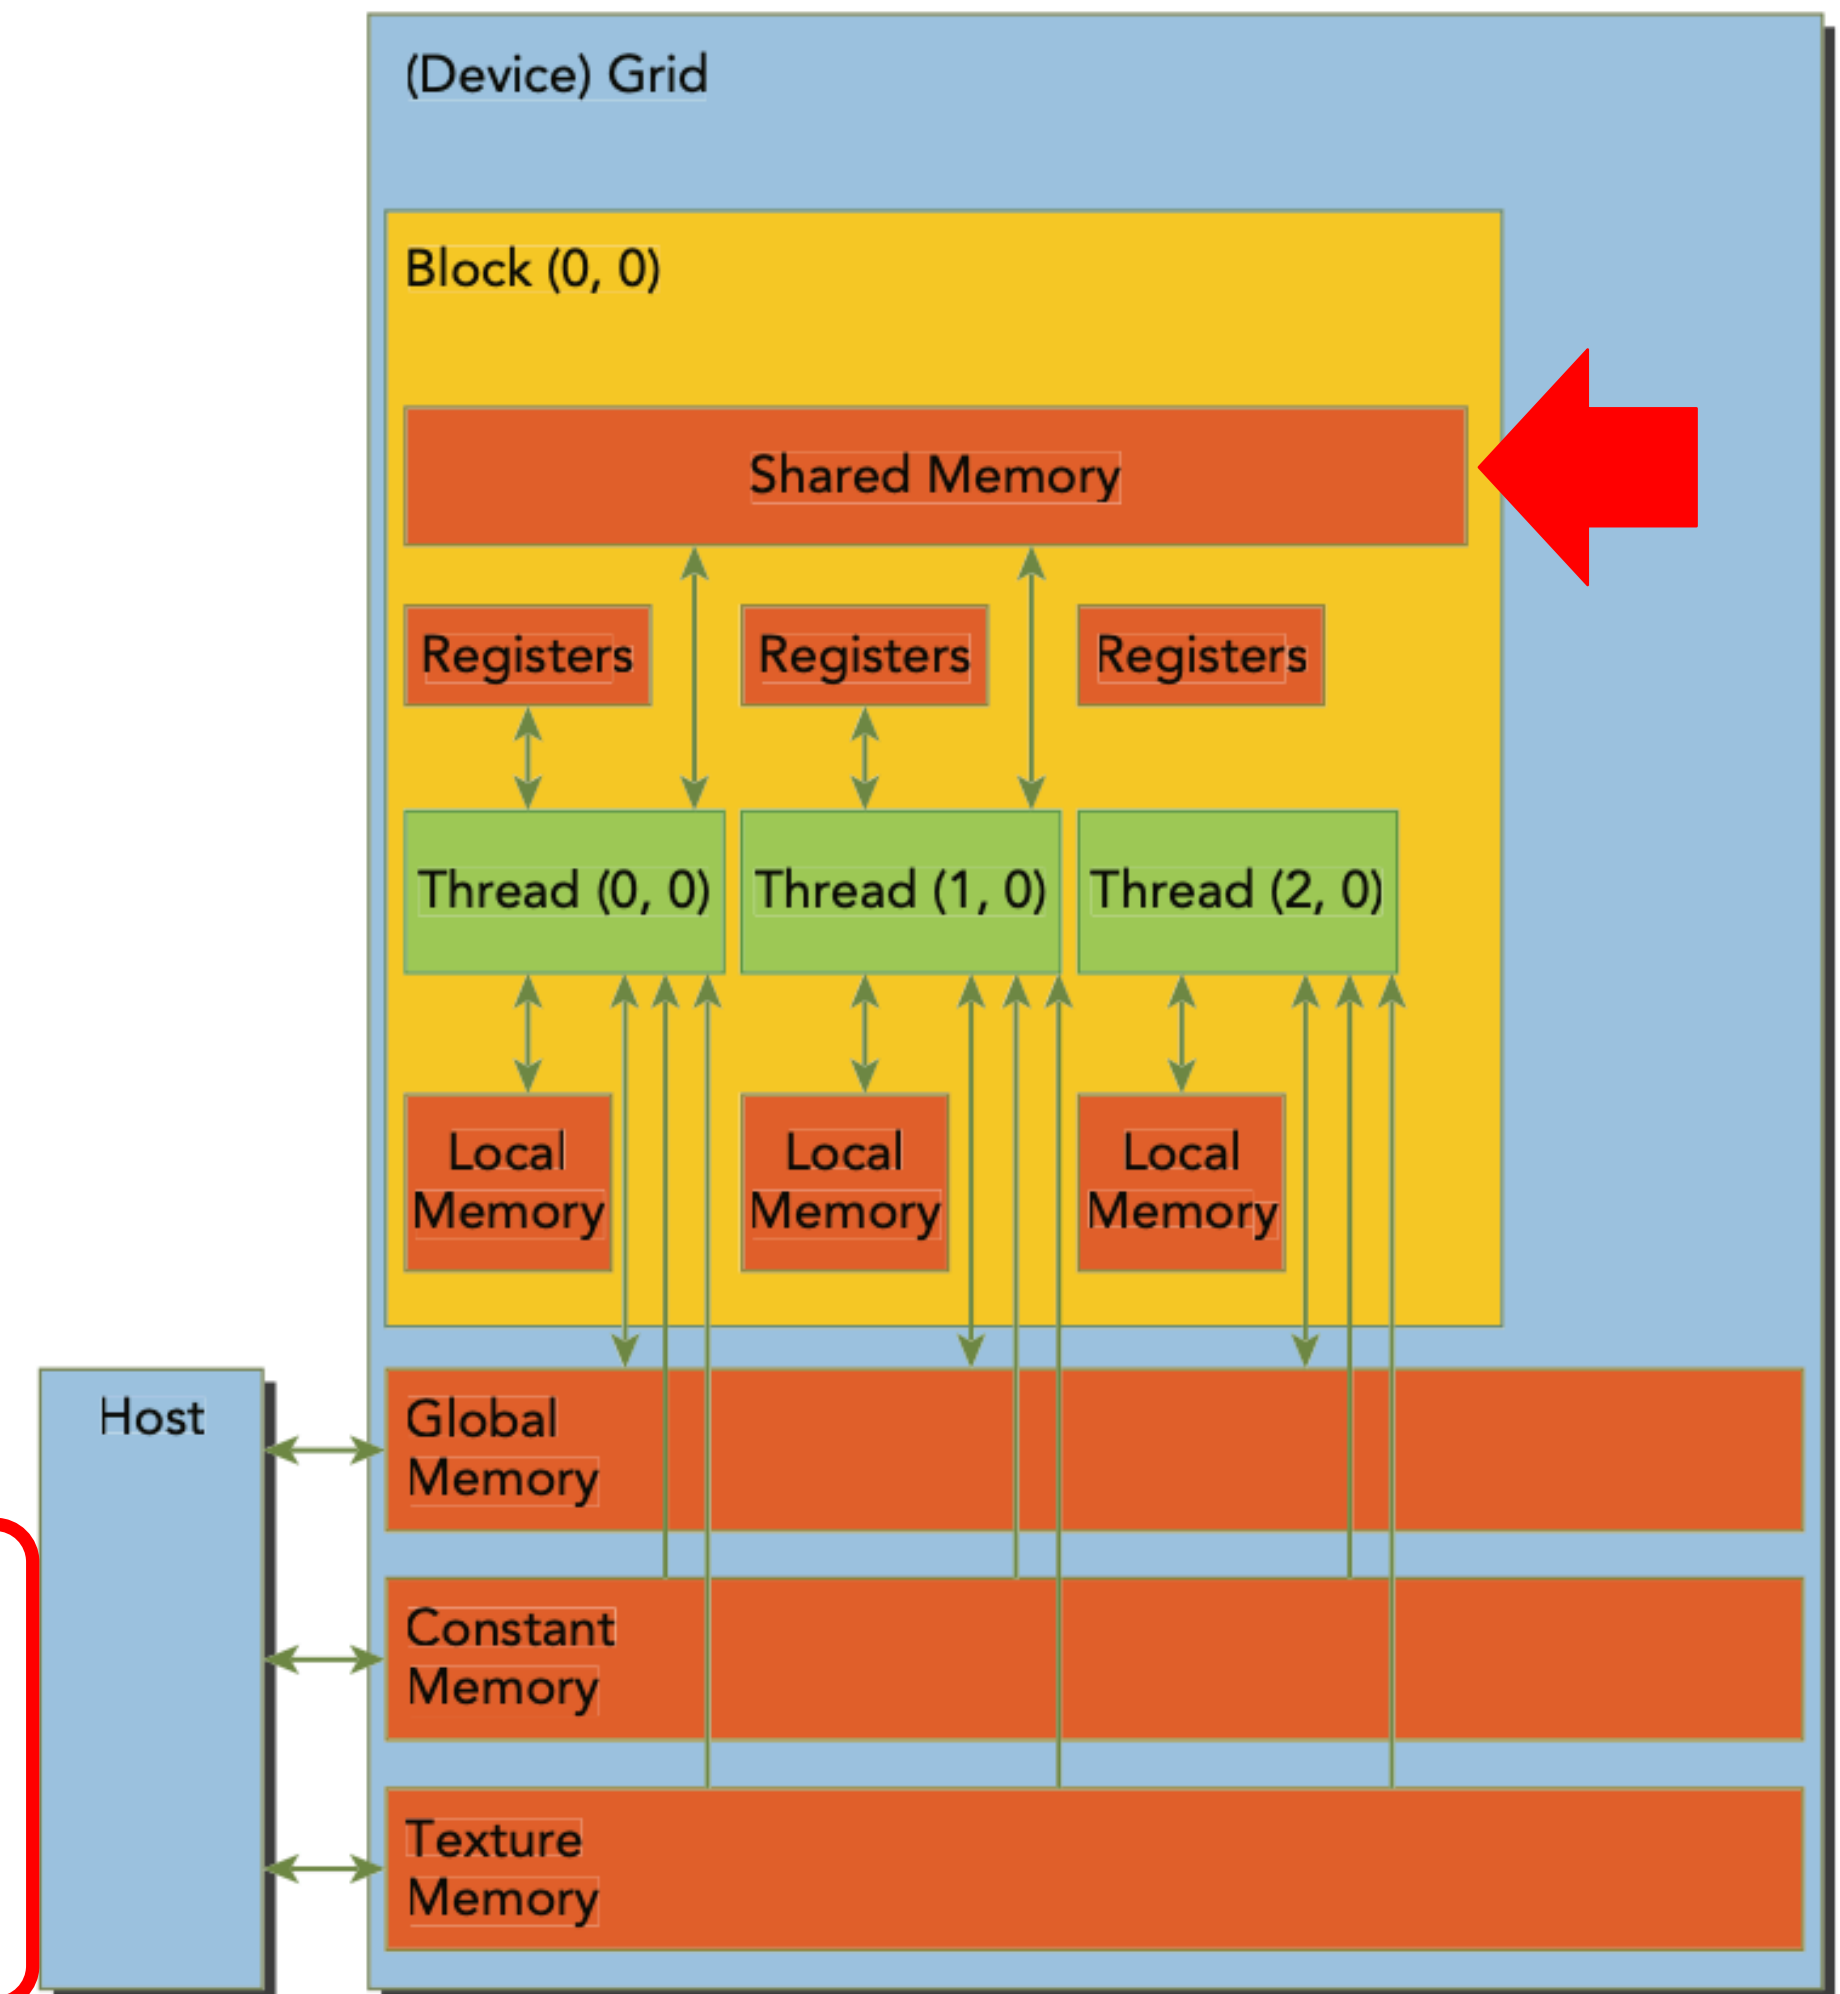


Shared Memory

Static allocation:

```
__global__ void foo(int* a, int* b) {  
    __shared__ int sv[N];  
    int i = blockIdx.x * blockDim.x +  
           threadIdx.x;  
    sv[i] = a[i];  
    __syncthreads();  
    /*...elaborate...*/  
    __syncthreads();  
    b[i] = sv[i];  
}
```

Multiple accesses of threads of the same block to shared memory must be **synchronized** through barriers (`__syncthreads()`)

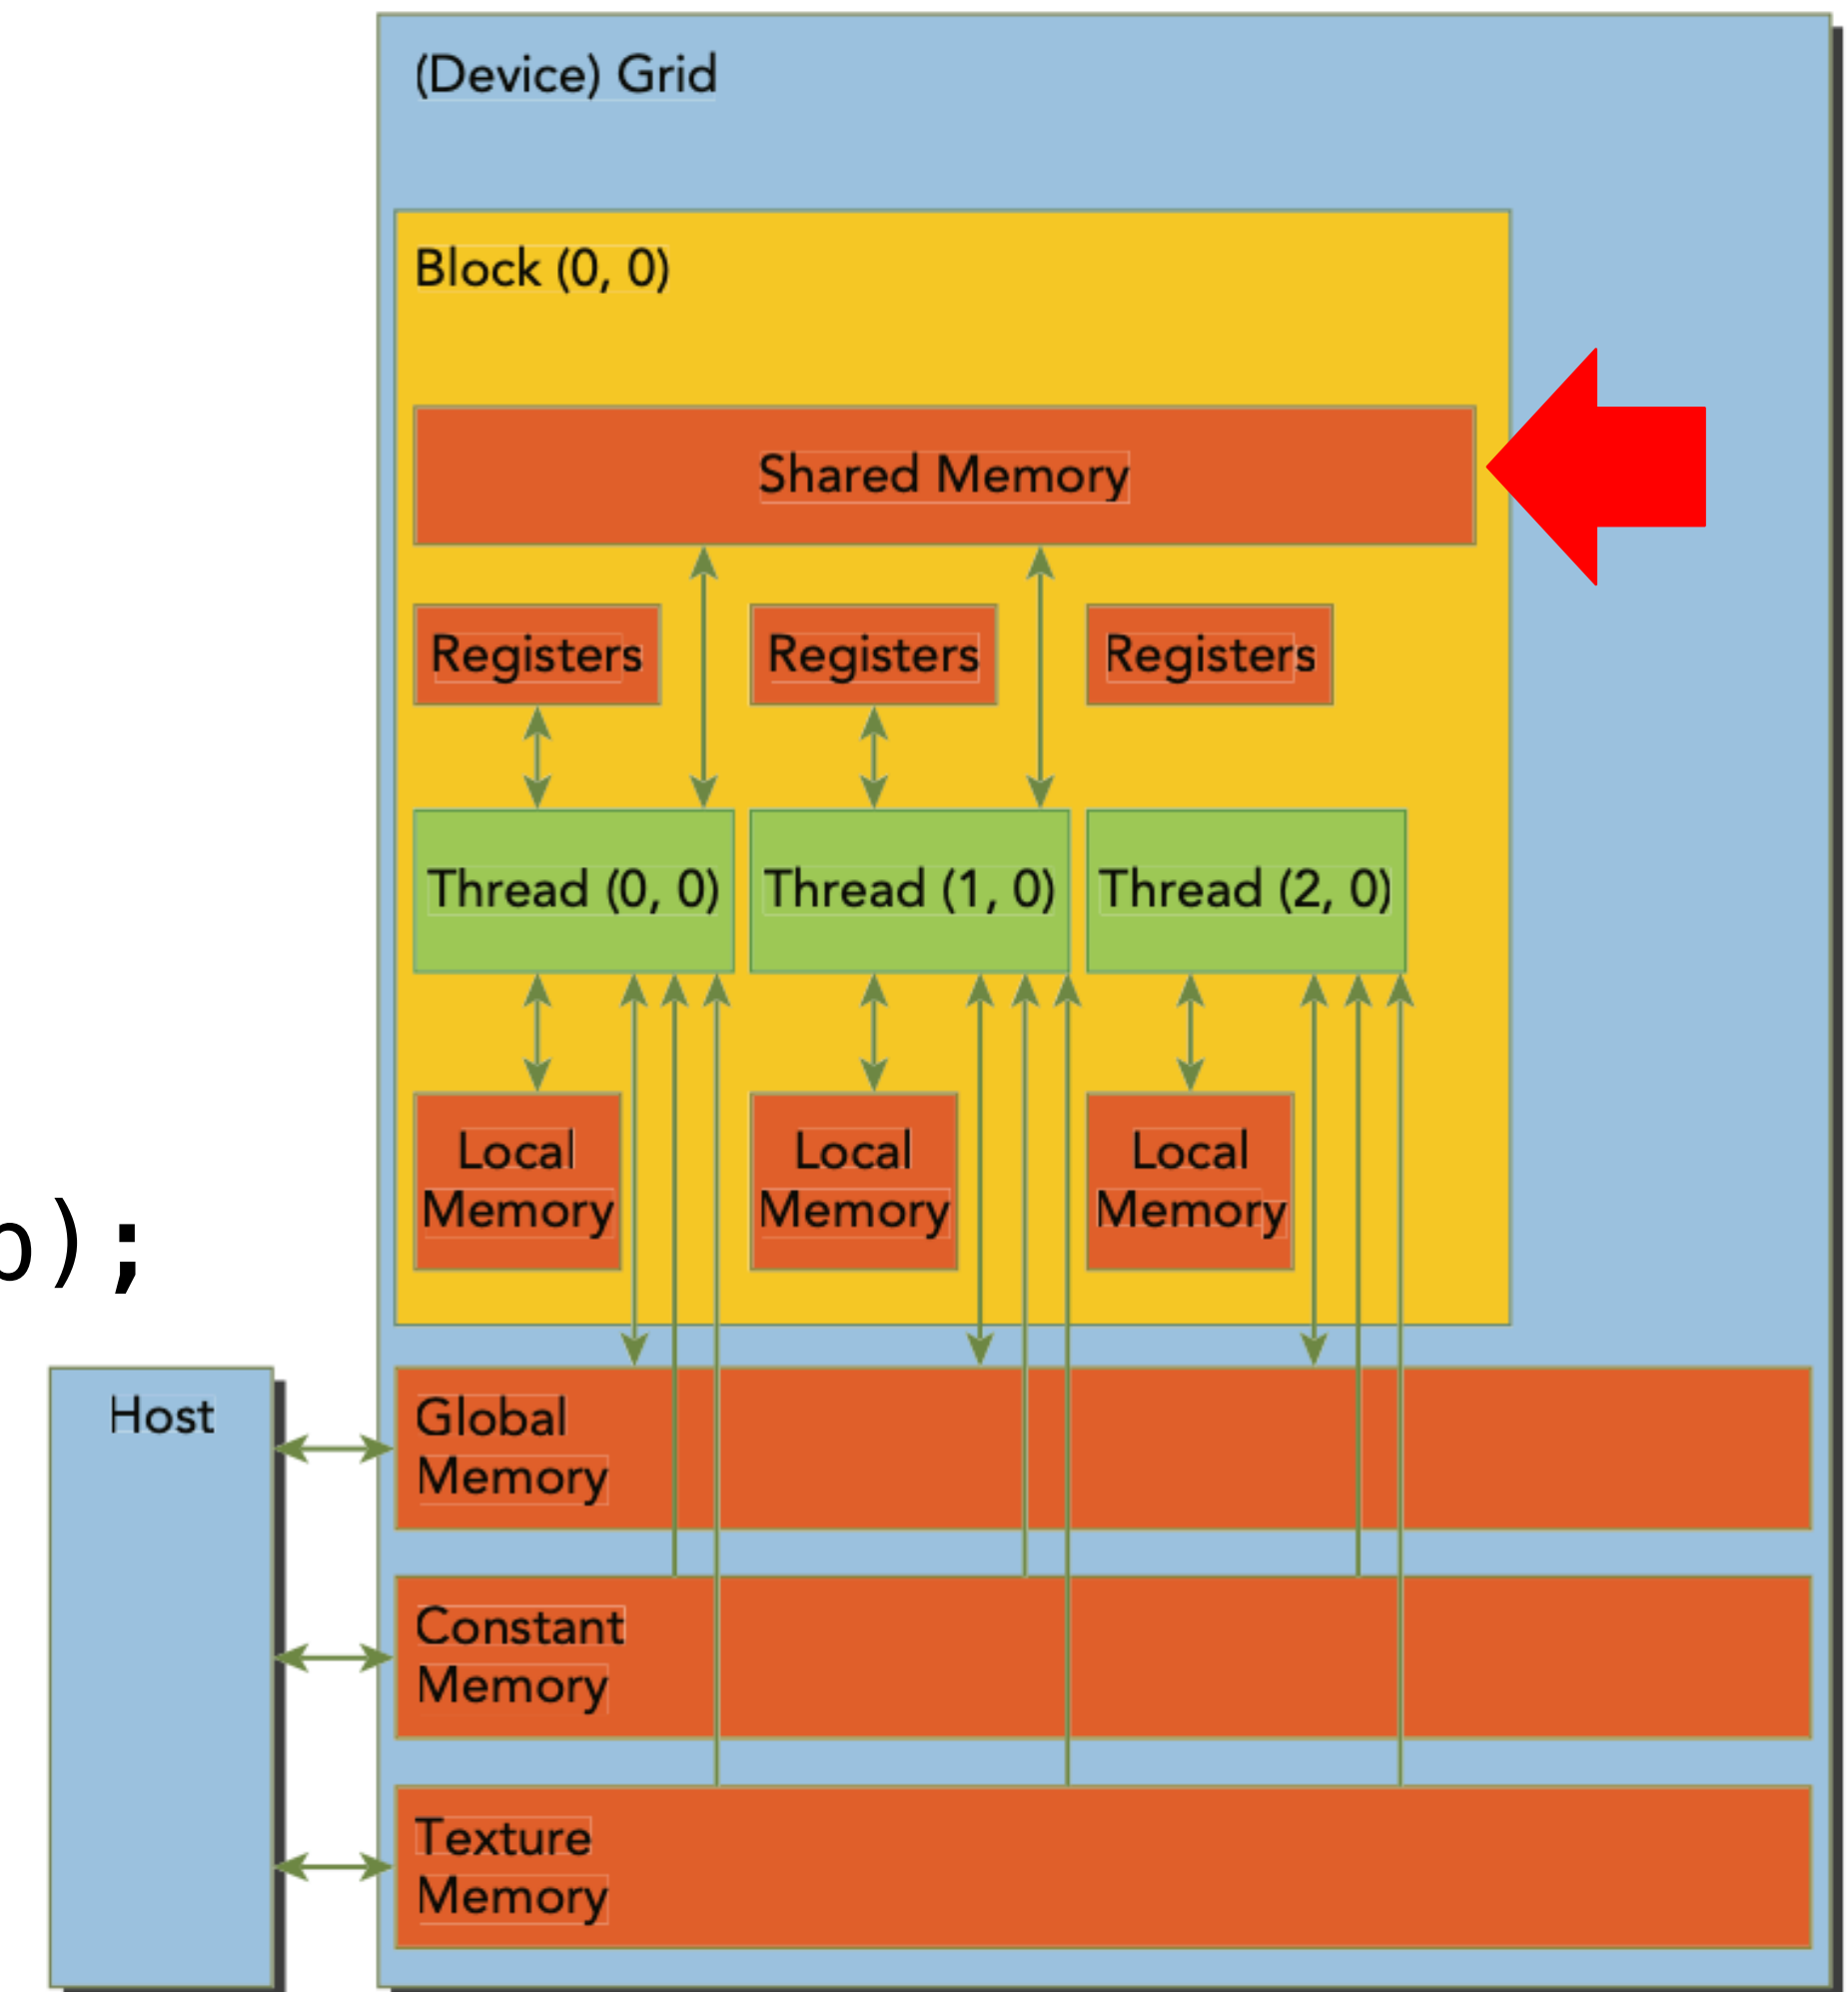


Shared Memory

Dynamic allocation:

```
__global__ void foo(int* a, int* b) {  
    extern __shared__ int sv[];  
    /*...*/  
}  
  
int main(){  
    /*...*/  
    foo<<<gdim, bdim, sizeof(int)*N>>>(va, vb);  
}
```

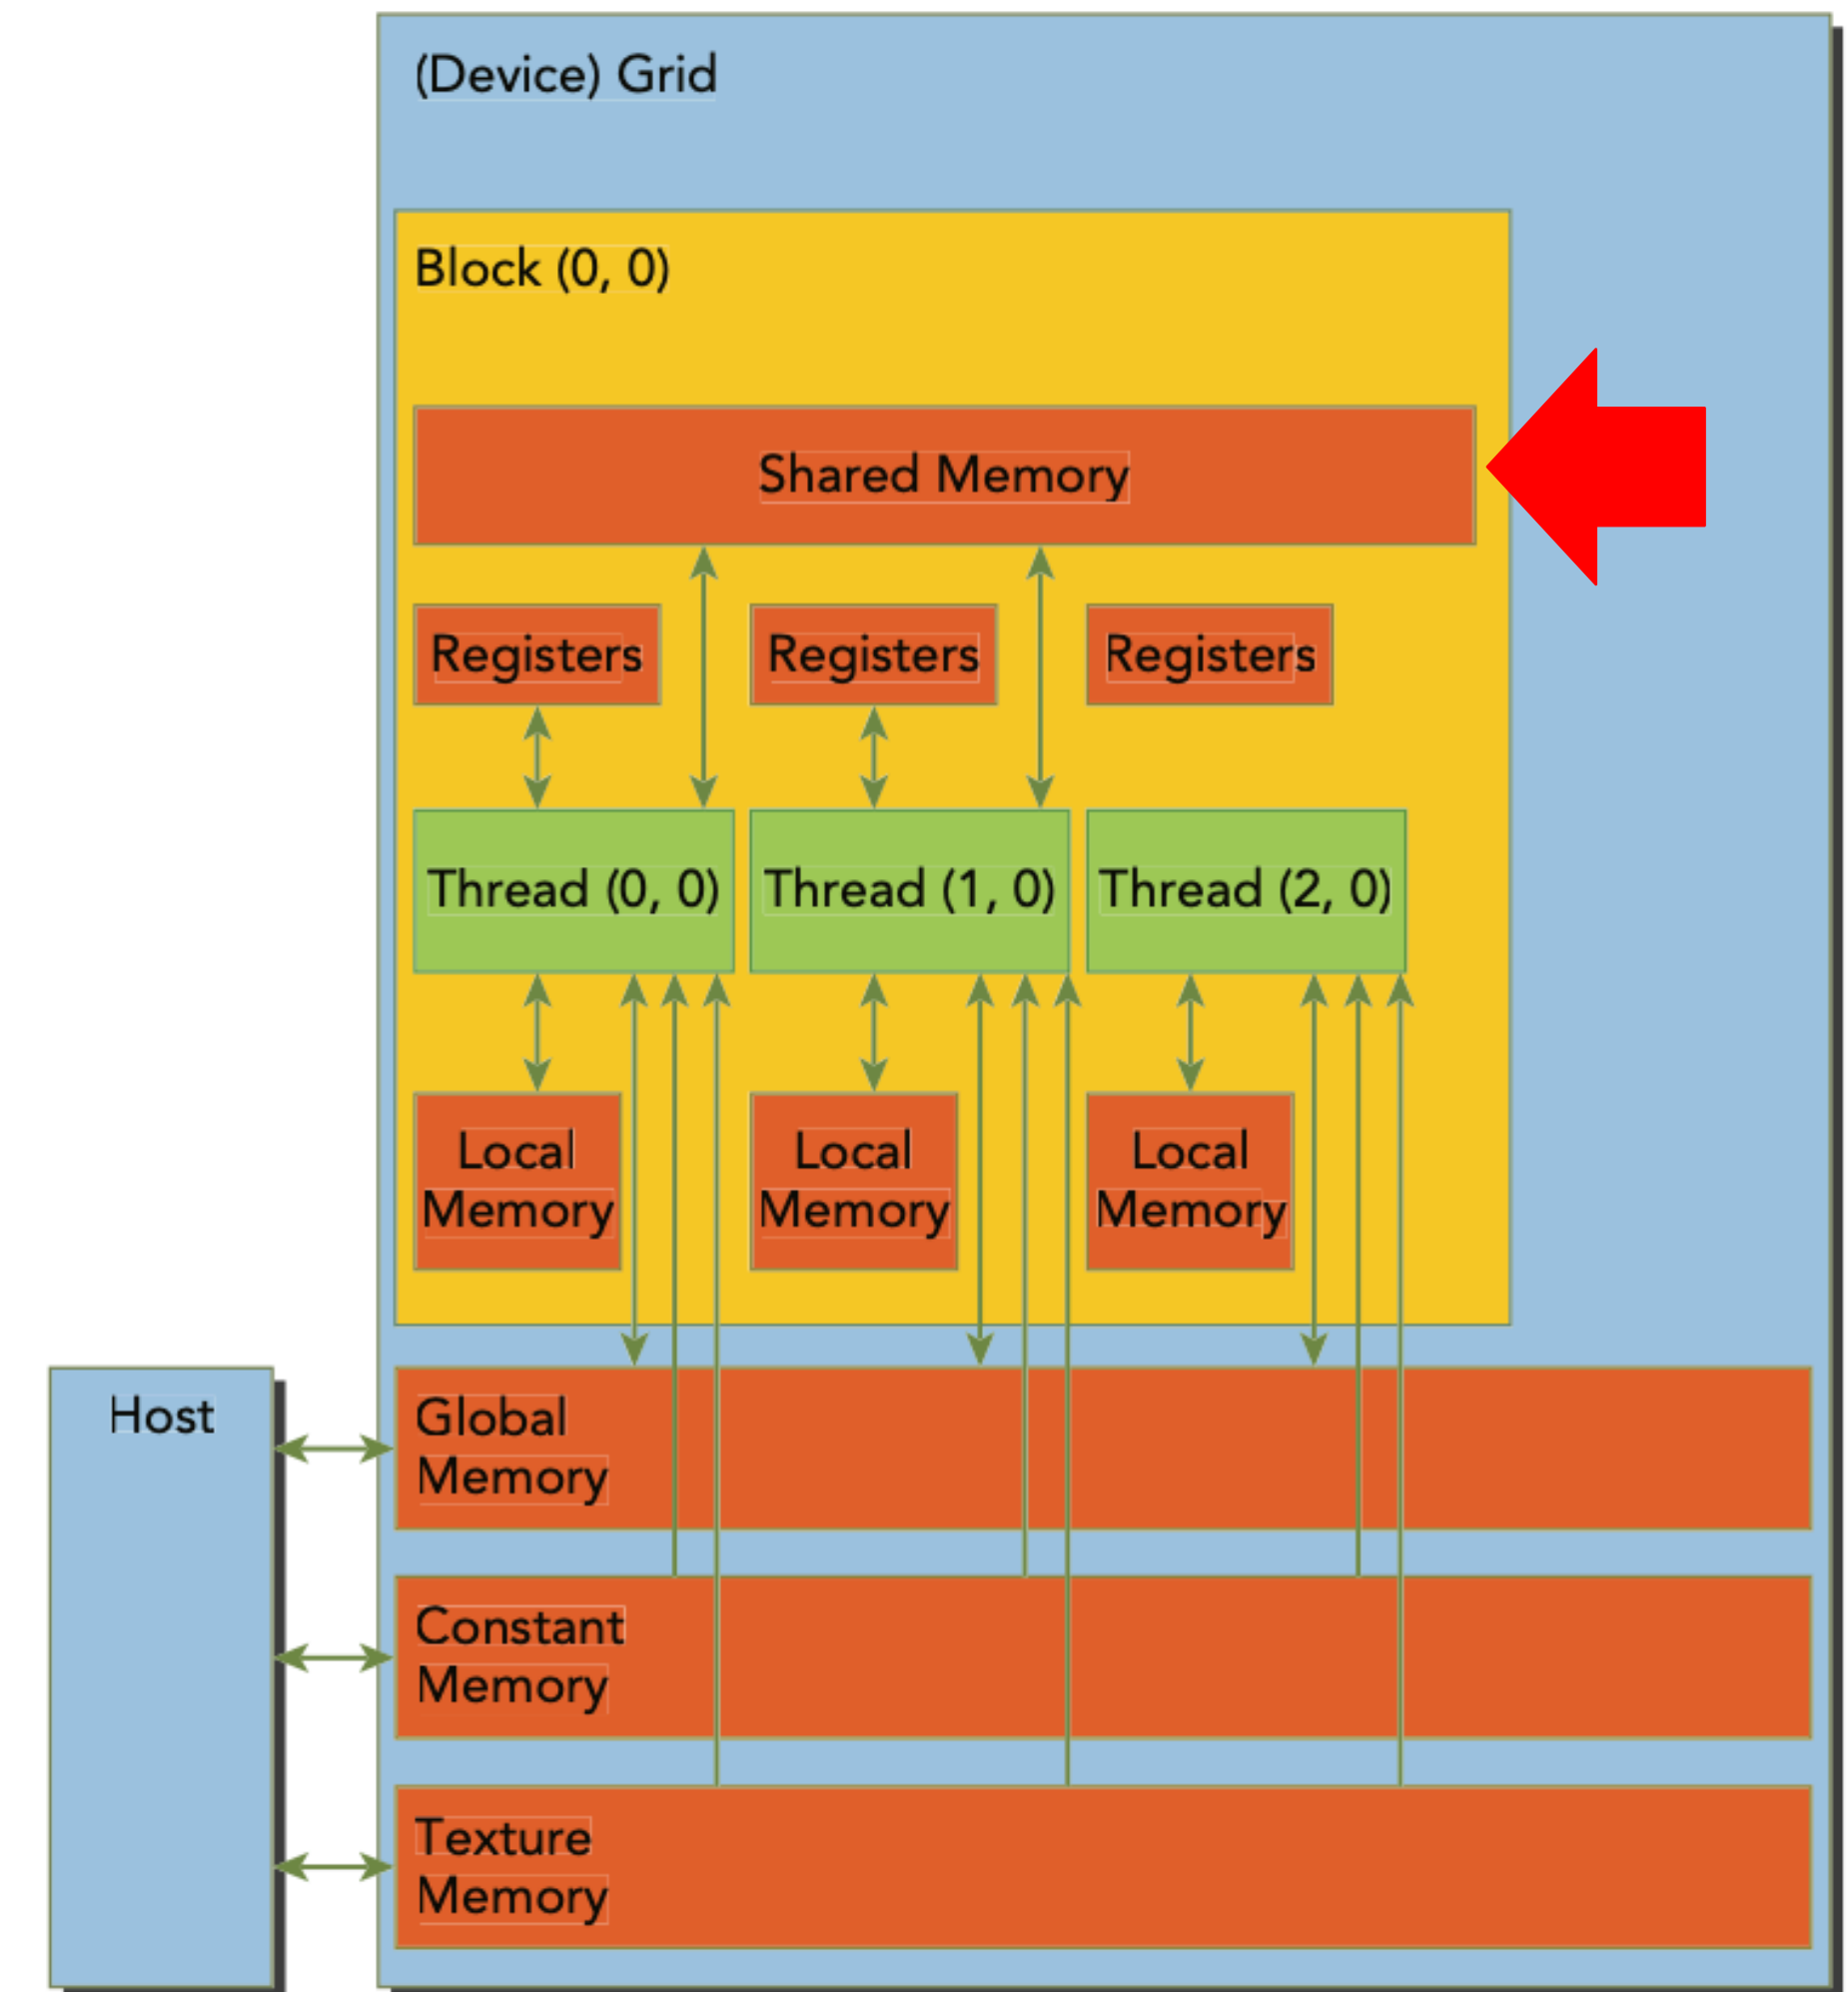
CUDA supports dynamic allocation of only 1D array
Allocation of multiple variables is tricky



Shared Memory

Typical shared memory usage:

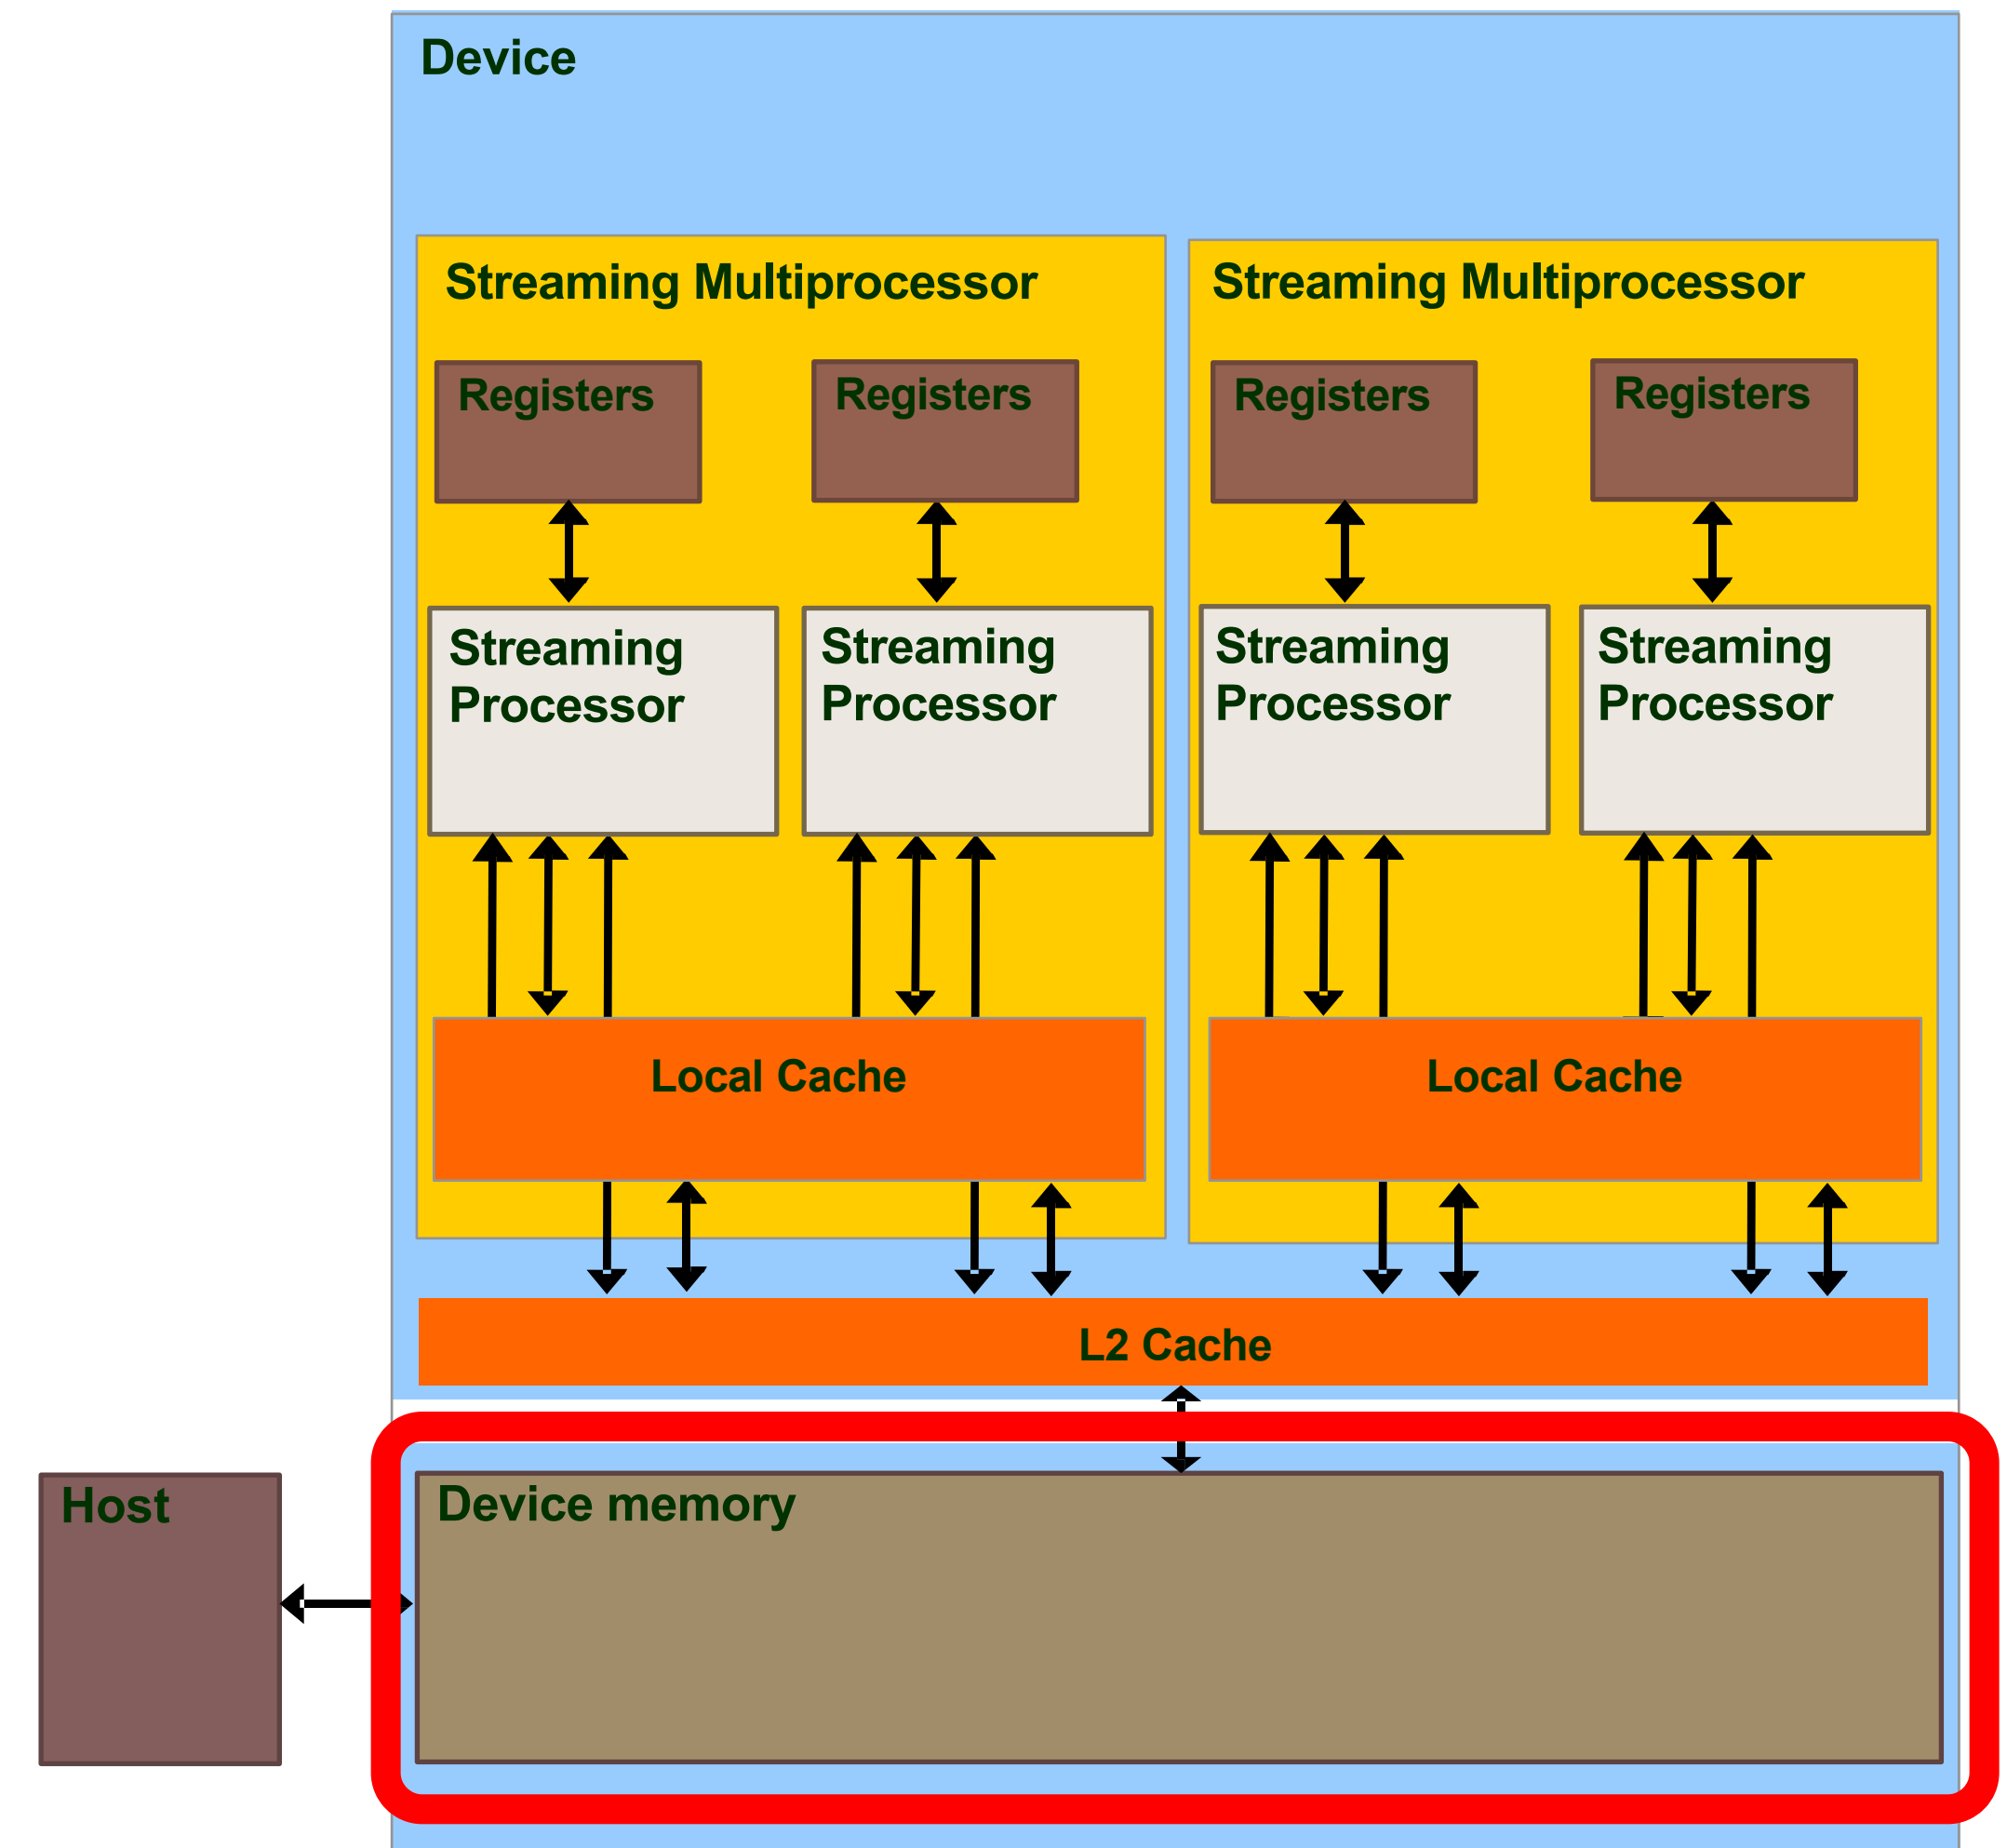
- Parallel load of data from global memory to shared memory
- Synchronization of block threads
- Elaboration of block threads on shared memory data
- Parallel store of data from shared memory to global memory



CUDA Memory Model

Key components:


- **Global Memory (device-wide): main memory** accessible by all threads across all blocks on the GPU, it has the largest capacity but is slower than registers and shared memory. It resides in device memory (HBM/DDR) and it is used for host/device communication. Coalesced Access: Access patterns in global memory can significantly affect performance. Coalesced memory accesses—where threads in a warp access consecutive memory addresses—optimize memory bandwidth.



CUDA Memory Model

Key components:

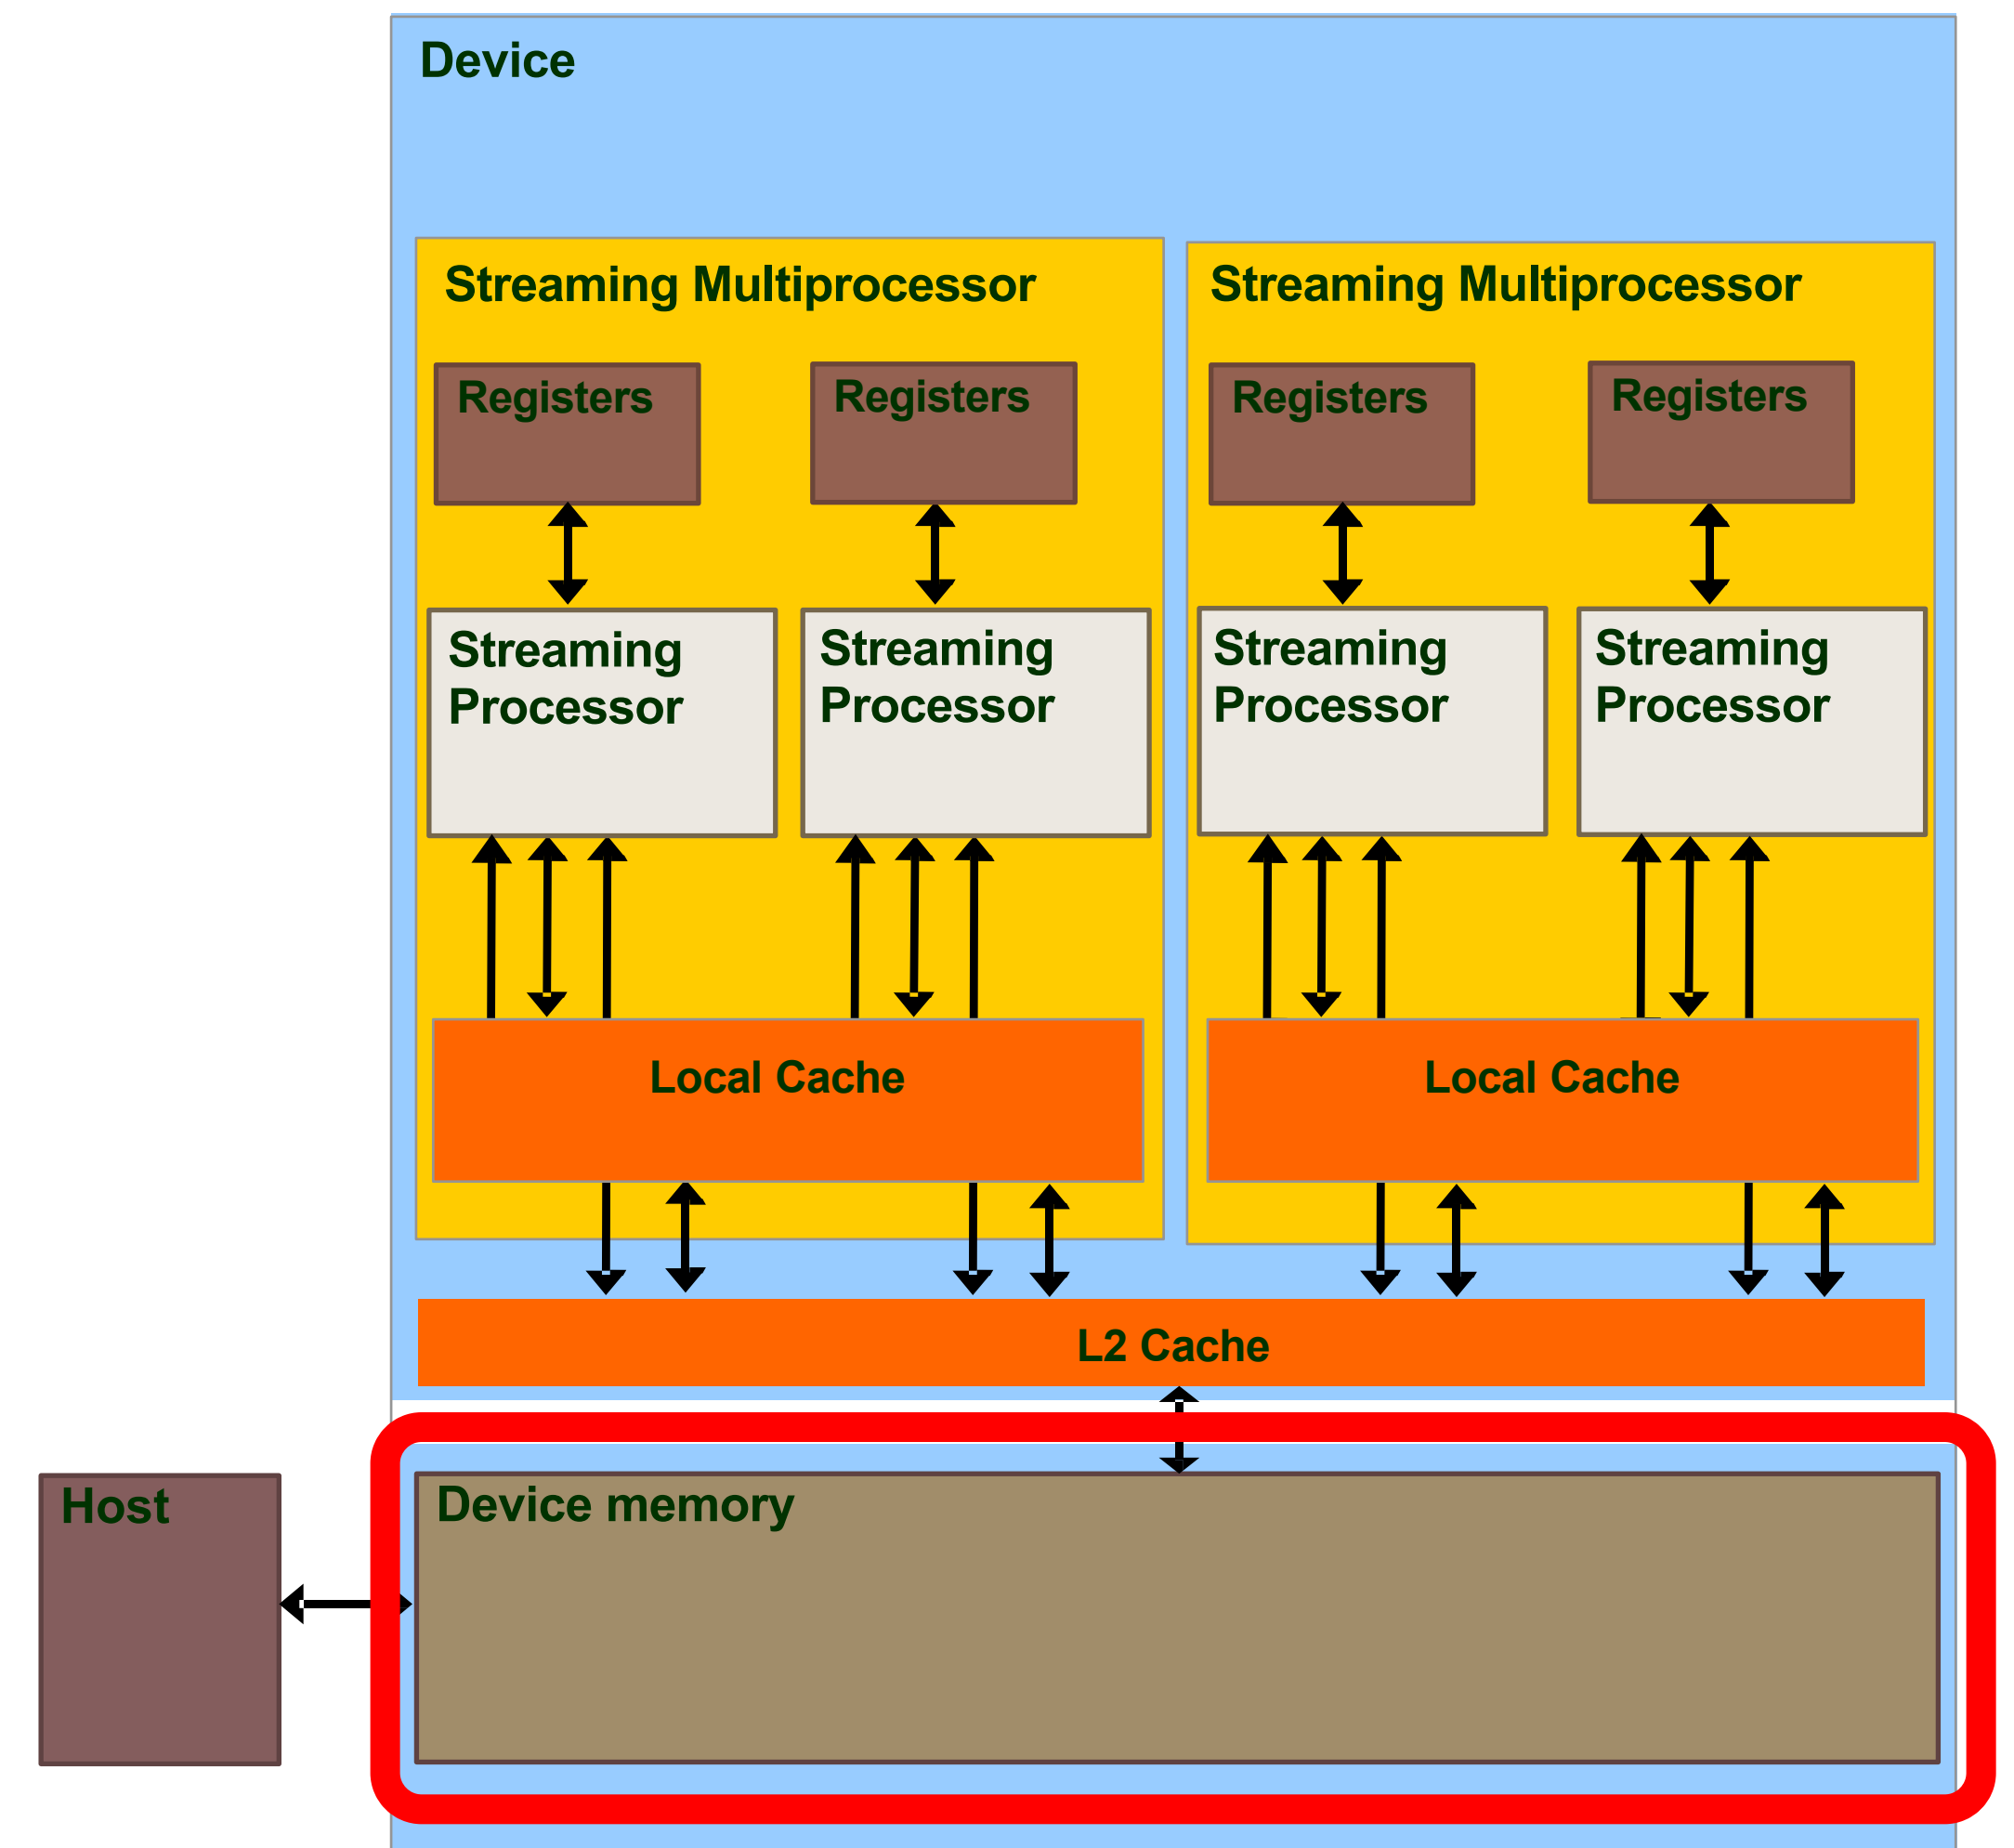
- **Constant Memory (device-wide, read-only)**: Constant memory is a **read-only** memory type that is cached, and it's accessible by all threads. It's mainly used to store data that remains constant throughout the execution of a kernel, such as lookup tables. Access to constant memory is optimized when all threads in a warp access the **same memory location**, as it avoids redundant fetches.



CUDA Memory Model

Key components:

- **Texture memory (device-wide, read-only):** read-only memory optimized for **spatial locality** and **2D locality**. It's cached, making it more efficient for applications requiring non-uniform access patterns, such as image processing. Texture memory is accessed through the **texture cache**, which is optimized for spatial locality and can handle non-coalesced access patterns better than global memory.



Credits

Slides mainly based on:

- D. B. Kirk, W.-m. W. Hwu, **Programming Massively Parallel Processors: A Hands-on Approach**
- J. Chen, M. Grossman, T. Mckercher, **Professional Cuda C Programming**
- A. Miele, G. Accordi, A. Zeni, **GPUs and Heterogeneous Systems (programming models and architectures)**
- **NVIDIA GPU Teaching Kit**

Thanks for your attention!

Beatrice Branchini
beatrice.branchini@polimi.it

Advanced Computer Architecture Course
T.1.1 - 5/5/2024