



# POLITECNICO DI MILANO

NECST  
laboratory

---

## Advanced Computer Architectures

# *Branch Hazards and Static Branch Prediction Techniques*

Politecnico di Milano

V0

Francesco Peverelli <francesco.peverelli@polimi.it>

Marco D. Santambrogio <marco.santambrogio@polimi.it>

Politecnico di Milano

---

# ACA

---


# Outline

---

- Dealing with Branches in the Processor Pipeline
- The Problem of Branch Hazards
- Branch Prediction Techniques
  - ▶ Static Branch Prediction

```
If (COND){  
    true statements;  
    [...]  
}else{  
    false statements;  
    [...]  
}
```

What it is all about



```
If (COND){  
    true statements;  
    [...]  
}else{  
    false statements;  
    [...]  
}
```

```
If (COND){  
PC → true statements;  
    [...]  
}else{  
    false statements;  
    [...]  
}
```


```
If (COND){  
  PC → true statements;  
    [...]  
  BTA → }else{  
    false statements;  
    [...]  
  }
```

```
If (COND){  
  PC → true statements;  
    [...]  
  BTA → }else{  
    false statements;  
    [...]  
  }
```



If **THAT'S THE TOUGH DECISION!!!**

```
PC → true  
    [...]  
    }  
BTA → else {  
        false  
        [...]   
    }
```



# Conditional Branch Instructions

---

- **Conditional Branch Instruction:** the branch is taken only if the **condition** is satisfied.  
The **branch target address (BTA)** is stored in the **Program Counter (PC)** instead of the address of the next instruction in the sequential instruction stream.

# Conditional Branch Instructions

- **Conditional Branch Instruction:** the branch is taken only if the **condition** is satisfied.  
The **branch target address (BTA)** is stored in the **Program Counter (PC)** instead of the address of the next instruction in the sequential instruction stream.
- Examples of branches for MIPS processor:  
**beq** (*branch on equal*) and **bne** (*branch on not equal*)  
`beq $s1, $s2, L1       # go to L1 if ($s1 == $s2)`  
`bne $s1, $s2, L1       # go to L1 if ($s1 != $s2)`

# Execution of conditional branches for 5-stage MIPS pipeline

**beq \$x,\$y,offset**

Instr. Fetch & PC Increm.	Register Read \$x e \$y	ALU Op. (\$x-\$y) & (PC+4+offset)	Write of PC
------------------------------	----------------------------	--------------------------------------	----------------

- Instruction fetch and PC increment
- Registers read (**\$x** and **\$y**) from Register File.
- ALU operation to compare registers (\$x and \$y) to derive **Branch Outcome** (branch taken or branch not taken).
  - Computation of **Branch Target Address** (PC+4+offset): the value (PC+4) is added to the least significant 16 bit of the instruction after sign extension
- The result of registers comparison from ALU is used to decide the value to be stored in the PC: (PC+4) or (PC+4+offset).

# Execution of conditional branches for 5-stage MIPS pipeline

IF	ID	EX	ME	WB
Instruction Fetch	Instruction Decode	Execution	Memory Access	Write Back

**beq \$x,\$y,offset**

Instr. Fetch & PC Increm.	Register Read \$x e \$y	ALU Op. (\$x-\$y) & (PC+4+offset)	Write of PC
------------------------------	----------------------------	--------------------------------------	----------------

- Branch Outcome and Branch Target Address are ready at the end of the EX stage (3<sup>th</sup> stage)
- Conditional branches are solved when PC is updated at the end of the ME stage (4<sup>th</sup> stage)

# The Problem of Control Hazards

---

- **Control hazards:** Attempt to make a decision on the next instruction to fetch before the branch condition is evaluated.
- Control hazards arise from the pipelining of conditional branches and other instructions changing the PC.
- Control hazards reduce the performance from the ideal speedup gained by the pipelining since they can make it necessary to stall the pipeline.

# Performance of Branch Schemes

---

- What is the performance impact of conditional branches?

$$\begin{aligned}\text{Pipeline Speedup} &= \frac{\text{Pipeline Depth}}{1 + \text{Pipe Stall Cycles per Instruction due to Branches}} \\ &= \frac{\text{Pipeline Depth}}{1 + \text{Branch Frequency} \times \text{Branch Penalty}}\end{aligned}$$


# Branch Hazards

---

- To feed the pipeline we need to fetch a new instruction at each clock cycle, but the branch decision (to change or not change the PC) is taken during the MEM stage.
- This delay to determine the correct instruction to fetch is called **Control Hazard or Conditional Branch Hazard**
- If a branch changes the PC to its target address, it is a **taken branch**
- If a branch falls through, it is **not taken or untaken**.



# Branch Hazards: Example



- The branch instruction may or may not change the PC in MEM stage, but the next 3 instructions are fetched and their execution is started.
- If the branch is **not taken**, the pipeline execution is OK
- If the branch is **taken**, it is necessary to *flush* the next 3 instructions in the pipeline and fetched the **lw** instruction at the branch target address (L1)

# Branch Stalls without Forwarding

---

```
beq $1, $3, L1  
and $12, $2, $5  
  or $13, $6, $2  
add $14, $2, $2
```

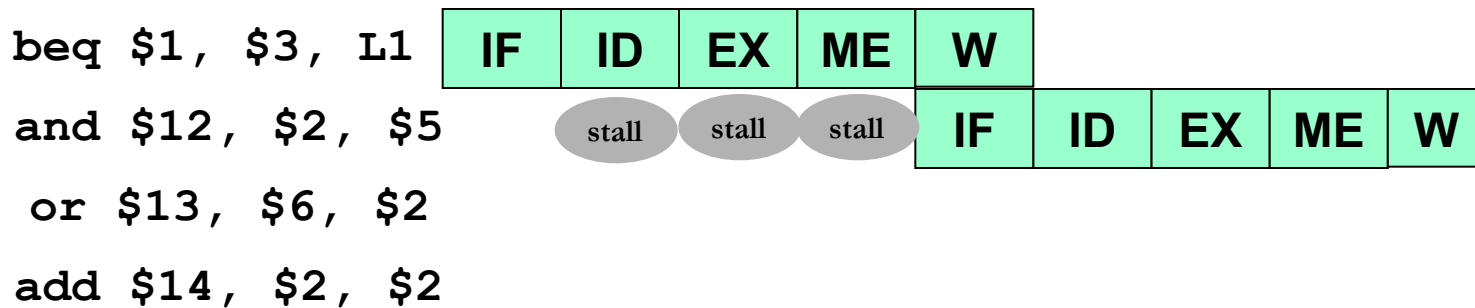
# Branch Stalls without Forwarding

---

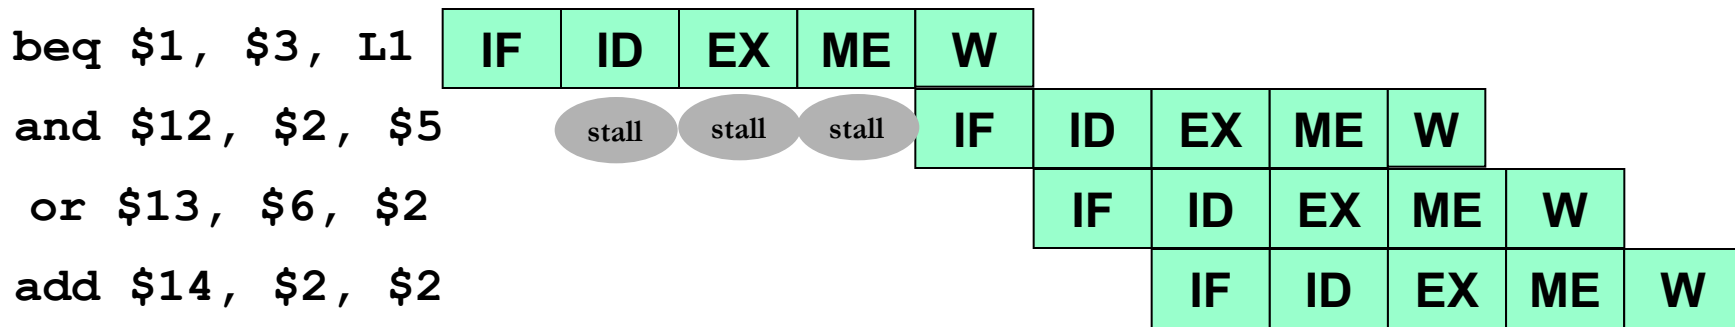
beq \$1, \$3, L1  
and \$12, \$2, \$5  
or \$13, \$6, \$2  
add \$14, \$2, \$2

IF	ID	EX	ME	W
----	----	----	----	---

# Branch Stalls without Forwarding



# Branch Stalls without Forwarding



# Branch Stalls with Forwarding

---

beq \$1, \$3, L1

IF

ID

EX

ME

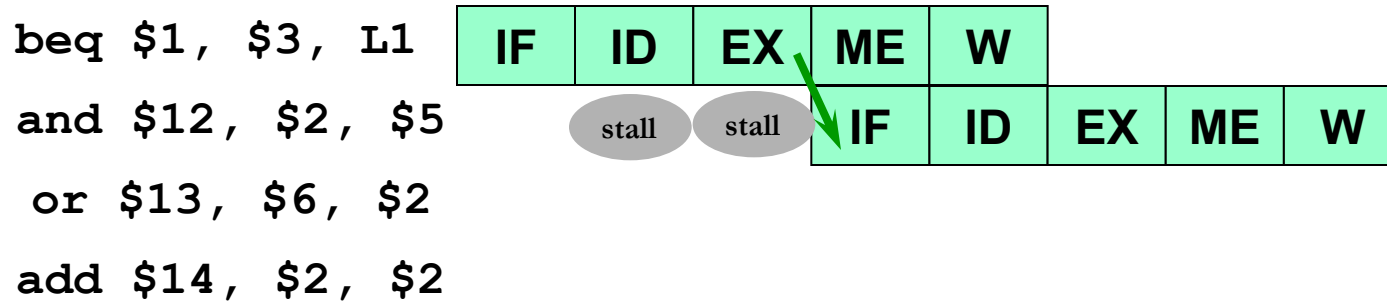
W

and \$12, \$2, \$5

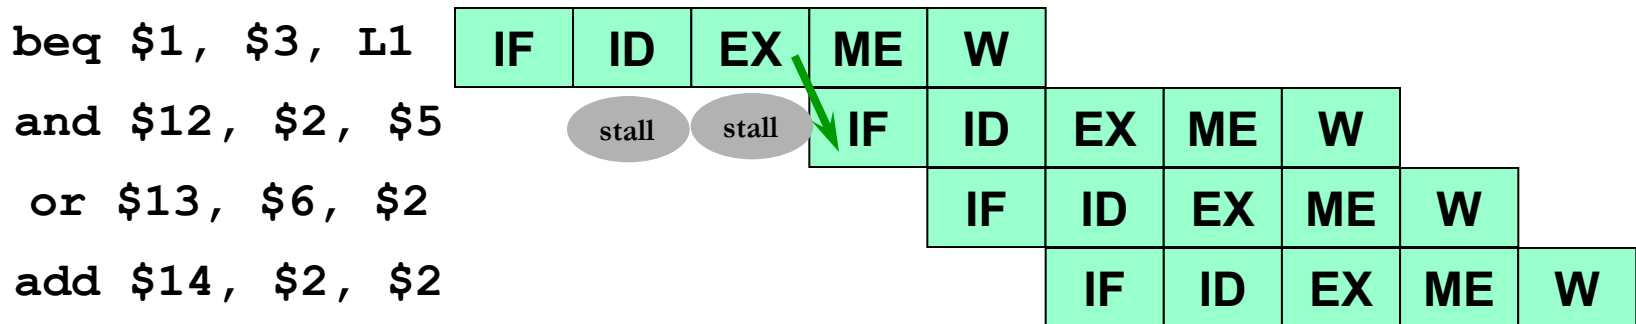
or \$13, \$6, \$2

add \$14, \$2, \$2

# Branch Stalls with Forwarding



# Branch Stalls with Forwarding





# Branch Hazards: Solutions

- To stall the pipeline until the branch decision is taken (**stalling until resolution**) and then fetch the correct instruction flow.
  - ▶ Without forwarding : for **three** clock cycles
  - ▶ With forwarding: for **two** clock cycles
- If the branch is *not taken*, the three cycles penalty is not justified  $\Rightarrow$  throughput reduction.
- We can assume the *branch not taken*, and **flush** the next 3 instructions in the pipeline only if the branch will be taken.

# Branch Hazards: Solutions

- To stall the pipeline until the branch decision is taken (**stalling until resolution**) and then fetch the correct instruction flow.
  - ▶ Without forwarding : for **three** clock cycles
  - ▶ With forwarding: for **two** clock cycles
- If the branch is *not taken*, the three cycles penalty is not justified  $\Rightarrow$  throughput reduction.
- We can assume the *branch not taken*, and **flush** the next 3 instructions in the pipeline only if the branch will be taken.

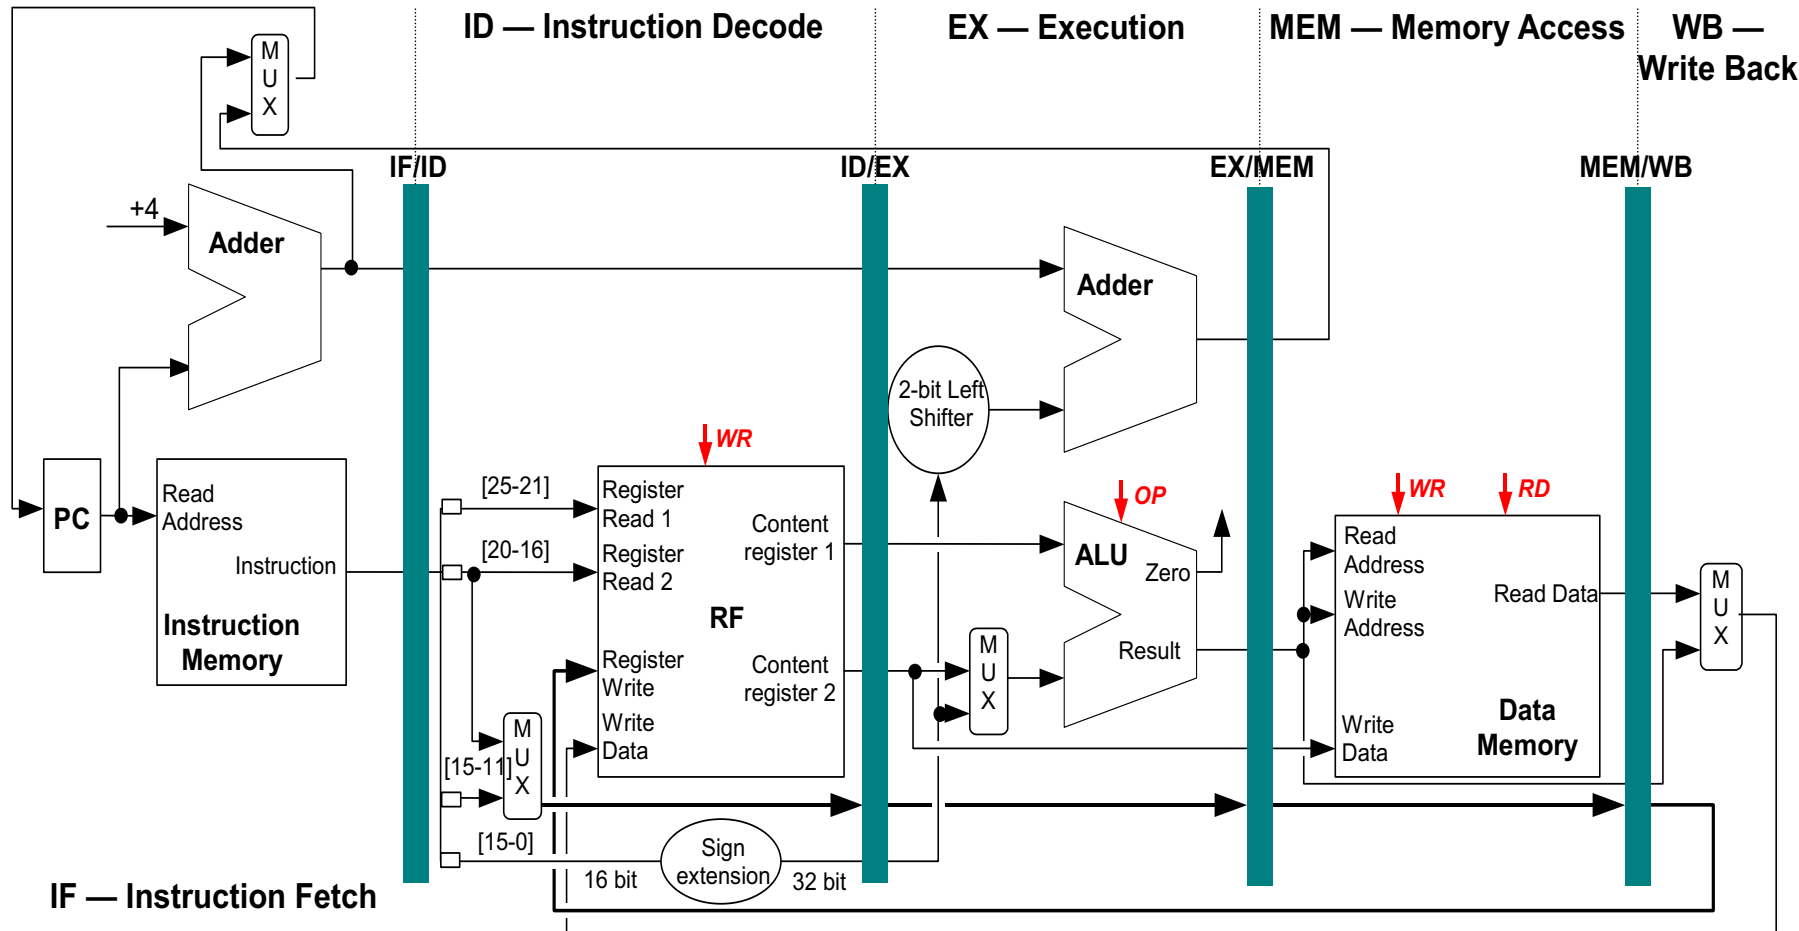
Can we do better?

# Early Evaluation of the PC

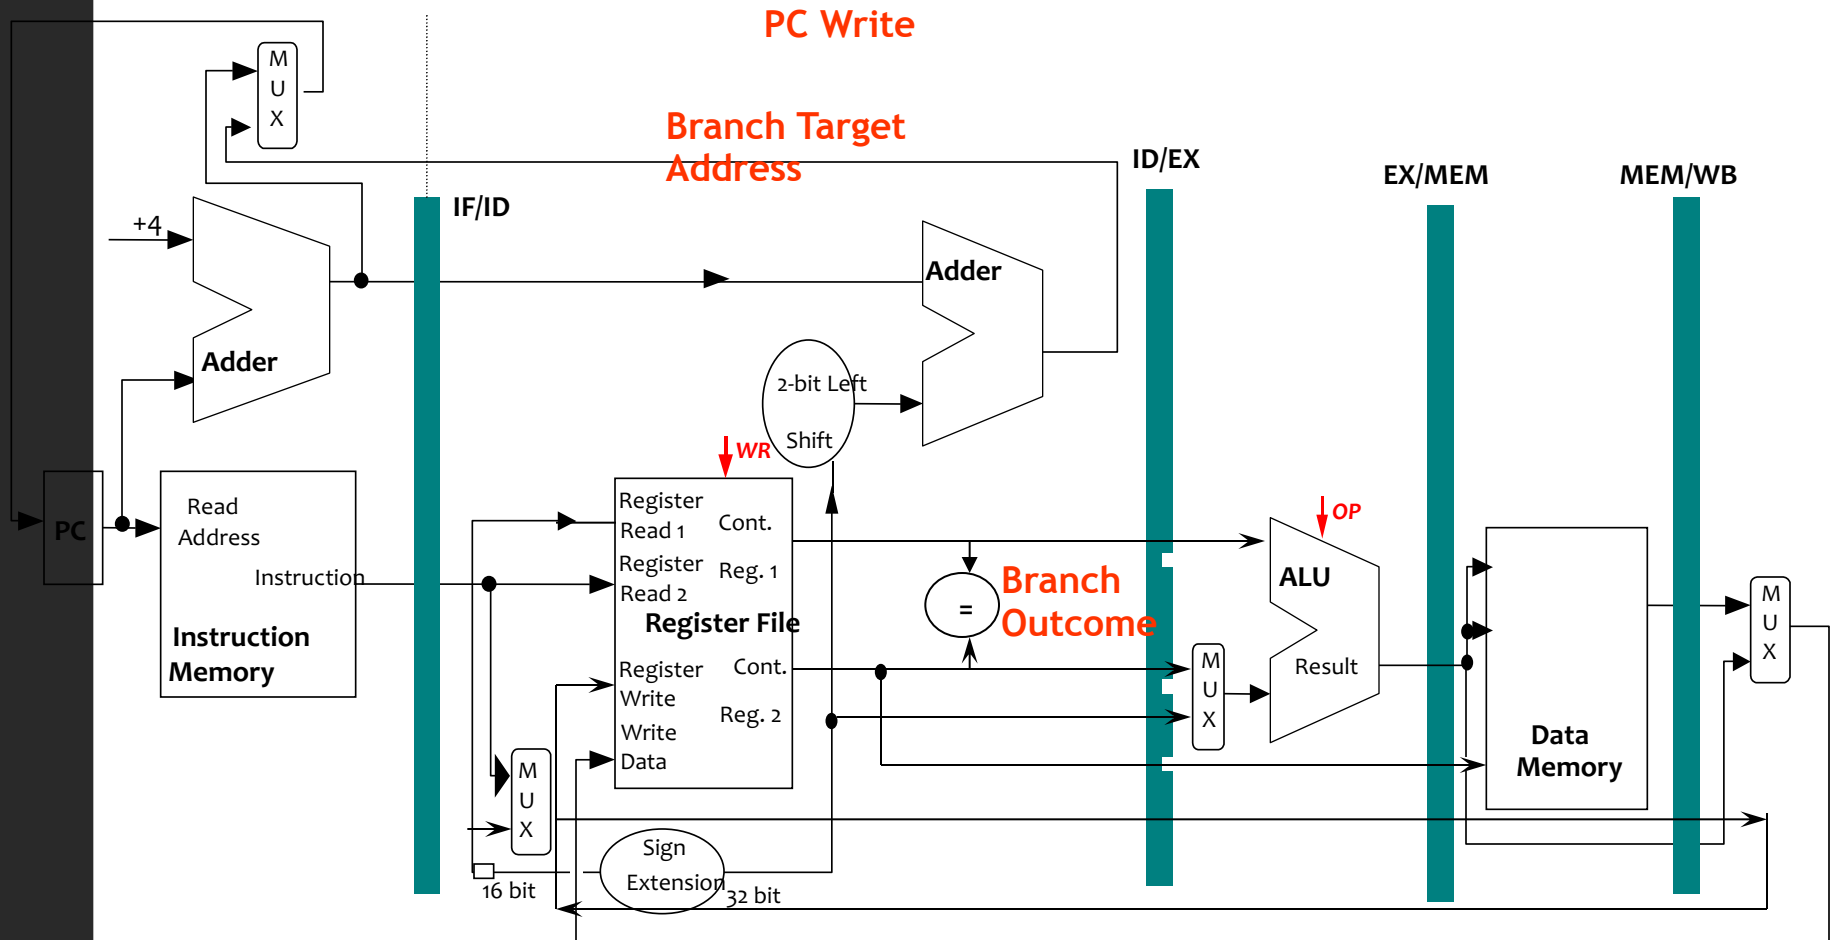
---

- To improve performance in case of branch hazards, we need to add hardware resources to:
  1. Compare registers to derive branch outcome
  2. Compute branch target address
  3. Update the PC register**as soon as possible in the pipeline.**
- MIPS processor compares registers, computes branch target address and updates PC **during ID stage.**

# Implementation of MIPS pipeline



# MIPS Processor: Early Evaluation of the PC



# MIPS Processor: Early Evaluation of the PC

---

- Each branch costs **one stall** to fetch the correct instruction flow: (PC+4) or branch target address

```
beq $1, $3, L1  
and $12, $2, $5  
or $13, $6, $2  
add $14, $2, $2
```

# MIPS Processor: Early Evaluation of the PC

- Each branch costs **one stall** to fetch the correct instruction flow: (PC+4) or branch target address

beq \$1, \$3, L1

IF	ID	EX	ME	W
----	----	----	----	---

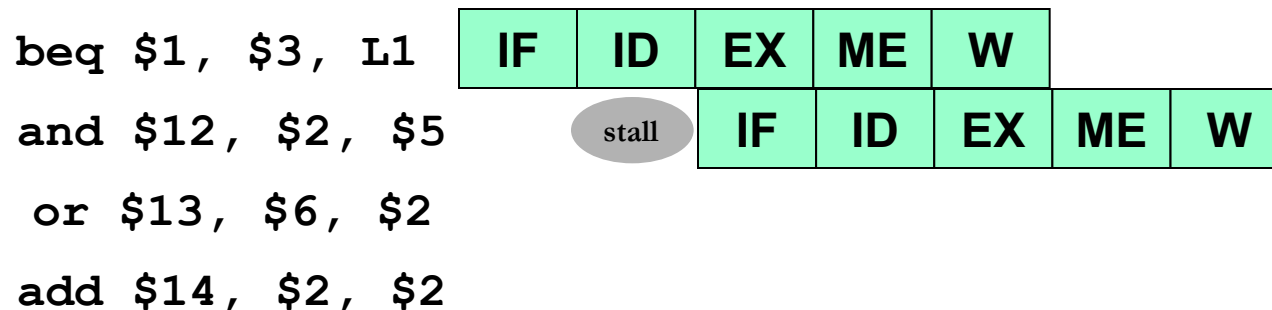
and \$12, \$2, \$5

or \$13, \$6, \$2

add \$14, \$2, \$2

# MIPS Processor: Early Evaluation of the PC

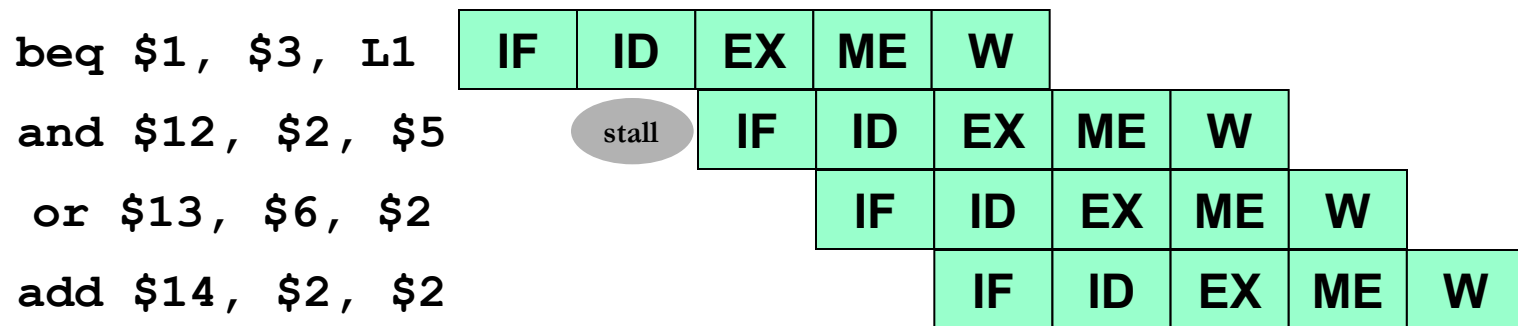
- Each branch costs **one stall** to fetch the correct instruction flow: (PC+4) or branch target address





# MIPS Processor: Early Evaluation of the PC

- Each branch costs **one stall** to fetch the correct instruction flow: (PC+4) or branch target address



## Early Evaluation of the PC: an “*issue*”

---

```
addi $1, $1, 4  
beq  $1, $6, L1
```

## Early Evaluation of the PC: an “issue”

---


addi \$1, \$1, 4  
beq \$1, \$6, L1

IF	ID	EX	ME	W
----	----	----	----	---

## Early Evaluation of the PC: an “issue”


---

addi \$1, \$1, 4  
beq \$1, \$6, L1




## Early Evaluation of the PC: an “issue”

addi \$1, \$1, 4  
beq \$1, \$6, L1




## Early Evaluation of the PC: an “issue”

addi \$1, \$1, 4  
beq \$1, \$6, L1




## Early Evaluation of the PC: an “issue”

addi \$1, \$1, 4  
beq \$1, \$6, L1



## Early Evaluation of the PC: an “issue”

addi \$1, \$1, 4  
beq \$1, \$6, L1





## MIPS Processor: Early Evaluation of the PC

---

- With the branch decision made during ID stage, there is a reduction of the cost associated with each branch (**branch penalty**):
  - ▶ We need only **one-clock-cycle stall** after each branch
  - ▶ Or a **flush** of only **one** instruction following the branch
- One-cycle-delay for every branch still yields a performance loss of 10% to 30% depending on the branch frequency:
- Pipeline Stall Cycles per Instruction due to Branches = Branch frequency x Branch Penalty
- We will examine some techniques to deal with this performance loss.


## MIPS Processor: Early Evaluation of the PC

---

- With the branch decision made during ID stage, there is a reduction of the cost associated with each branch (**branch penalty**):
  - ▶ We need only **one-clock-cycle stall** after each branch
  - ▶ Or a **flush** of only **one** instruction following the branch
- One-cycle-delay for every branch still yields a performance loss of 10% to 30% depending on the branch frequency:
- Pipeline Stall Cycles per Instruction due to Branches = Branch frequency x Branch Penalty
- We will examine some techniques to deal with this performance loss.

# Branch Prediction Techniques

---



# Branch Prediction Techniques

---

- There are many methods to deal with the performance loss due to branch hazards:
  - ▶ **Static Branch Prediction Techniques:** The actions for a branch are fixed for each branch during the entire execution. The actions are fixed at compile time.
  - ▶ **Dynamic Branch Prediction Techniques:** The decision causing the branch prediction can change during the program execution.
- In both cases, care must be taken not to change the processor state until the branch is definitely known.

# Static Branch Prediction Techniques

---

- Static Branch Prediction is used in processors where the expectation is that the branch behavior is highly predictable at compile time.
- Static Branch Prediction can also be used to assist dynamic predictors.

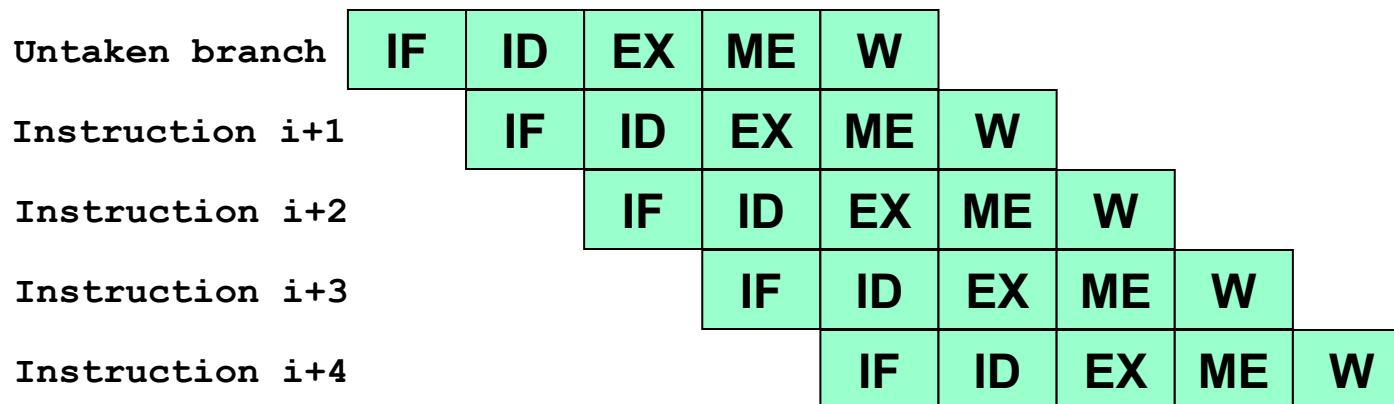
# Static Branch Prediction Techniques

---

- Branch Always Not Taken (Predicted-Not-Taken)
- Branch Always Taken (Predicted-Taken)
- Backward Taken Forward Not Taken (BTFNT)
- Profile-Driven Prediction
- Delayed Branch

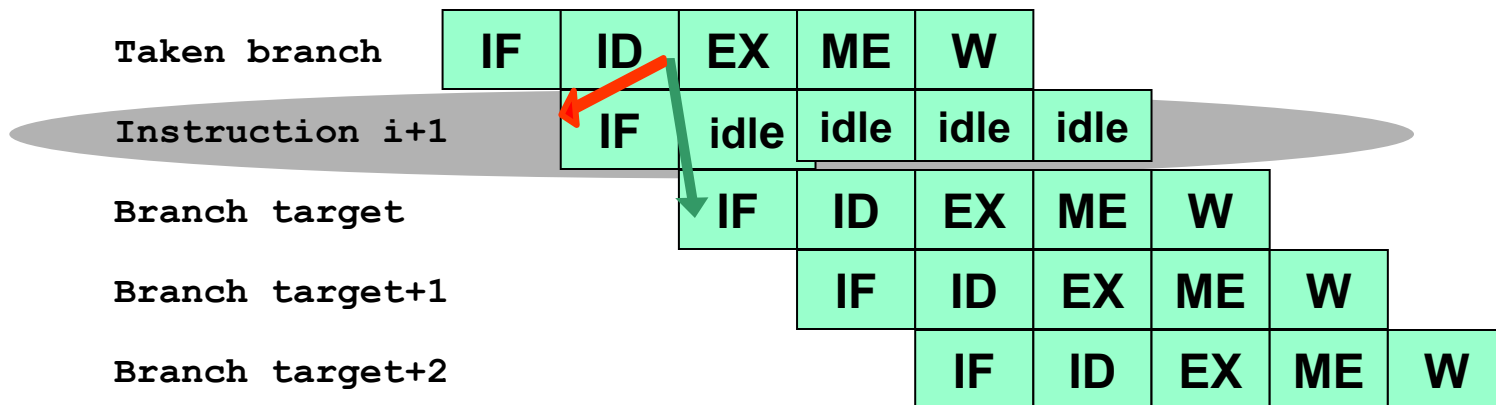
# Branch Always Not Taken

- We assume the **branch will not taken**, thus the sequential instruction flow we have fetched can continue as if the branch condition was not satisfied.
- If the condition in stage ID will result not satisfied (**the prediction is correct**), we can preserve performance.



# Branch Always Not Taken

- If the condition in stage ID will result satisfied (the prediction is incorrect), the branch is taken:
  - ▶ We need to flush the next instruction already fetched (the fetched instruction is turned into a **nop**) and we restart the execution by fetching the instruction at the branch target address  $\Rightarrow$  **One-cycle penalty**





# Branch Always Taken

---

- An alternative scheme is to consider every branch as taken: as soon as the branch is decoded and the branch target address is computed, we assume the branch to be taken and we begin fetching and executing at the target.

# Branch Always Taken

---

- An alternative scheme is to consider every branch as taken: as soon as the branch is decoded and the branch target address is computed, we assume the branch to be taken and we begin fetching and executing at the target.
- The predicted-taken scheme makes sense for pipelines where the branch target is known before the branch outcome.
- In MIPS pipeline, we don't know the branch target address earlier than the branch outcome, so there is no advantage in this approach for this pipeline.

## Backward Taken Forward Not Taken (BTFNT)

---

- The prediction is based on the branch direction:
  - ▶ Backward-going branches are predicted as taken
    - Example: the branches at the end of loops go back at the beginning of the next loop iteration  
⇒ we assume the backward-going branches are always taken.
  - ▶ Forward-going branches are predicted as not taken

# Profile-Driven Prediction

---

- The branch prediction is based on profiling information collected from earlier runs.
- The method can use compiler hints.


# Delayed Branch Technique

---

- The compiler statically schedules an independent instruction in the **branch delay slot**.
- The instruction in the branch delay slot is executed whether or not the branch is taken.
- If we assume a branch delay of one-cycle (as for MIPS)  
⇒ we have only **one-delay slot**
  - ▶ Although it is possible to have for some deeply pipeline processors a branch delay longer than one-cycle  
⇒ almost all processors with delayed branch have a single delay slot (since it is usually difficult for the compiler to fill in more than one delay slot).

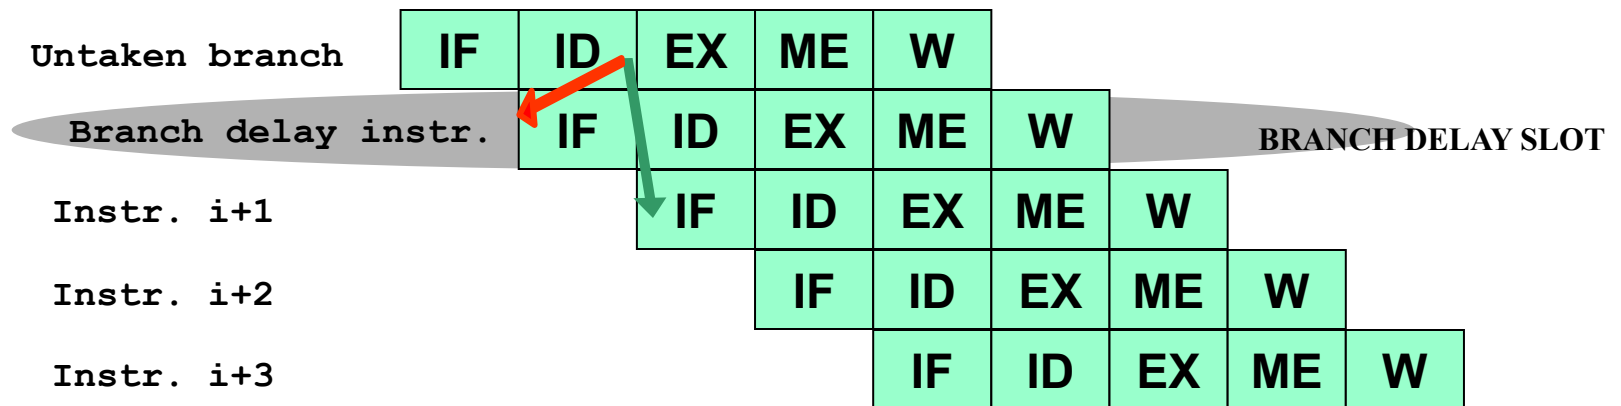
# Delayed Branch Technique

- The MIPS compiler always schedules a branch independent instruction after the branch.
- Example: A previous `add` instruction without any effects on the branch is scheduled in the **Branch Delay Slot**



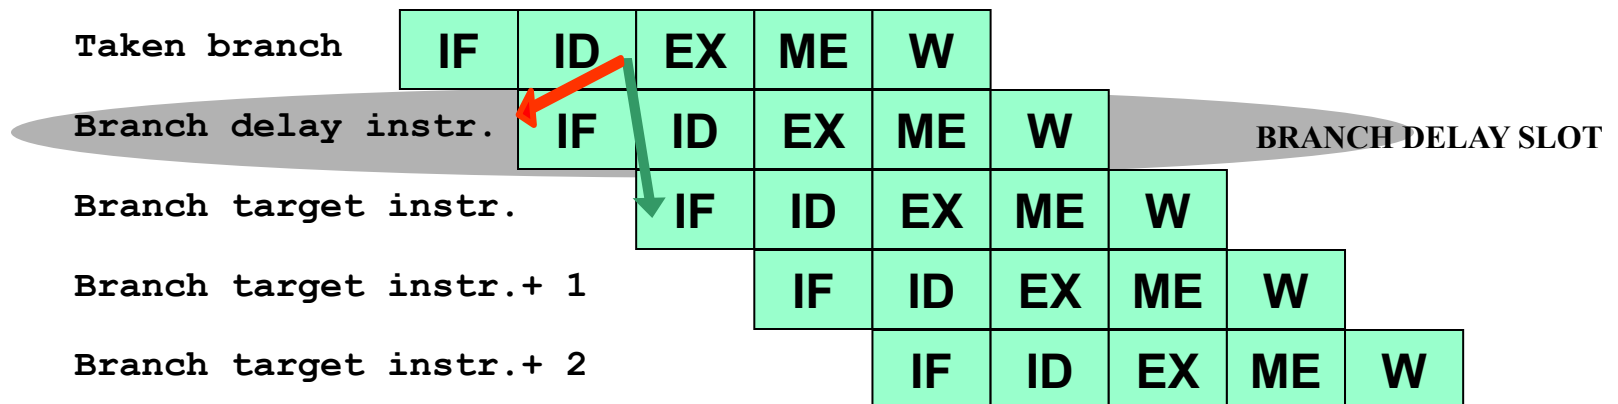
# Delayed Branch Technique

- The behavior of the delayed branch is the same whether or not the branch is taken.
  - ▶ If the branch is **untaken**  $\Rightarrow$  execution continues with the instruction after the branch



# Delayed Branch Technique

- If the branch is **taken**  $\Rightarrow$  execution continues at the branch target






# Delayed Branch Technique

---

- The job of the compiler is to make the instruction placed in the branch delay slot valid and useful.
- There are three ways in which the branch delay slot can be scheduled:
  1. **From before**
  2. **From target**
  3. **From fall-through**


# Delayed Branch Technique: From Before

- The branch delay slot is scheduled with an independent instruction from before the branch
- The instruction in the branch delay slot is **always executed** (whether the branch is taken or untaken).




# Delayed Branch Technique: From Target

- The use of \$1 in the branch condition prevents **add** instruction (whose destination is \$1) from being moved after the branch.
- The branch delay slot is scheduled from the target of the branch (usually the target instruction will need to be copied because it can be reached by another path).
- This strategy is preferred when the branch is taken with high probability, such as loop branches (**backward branches**).



# Delayed Branch Technique: From Fall-Through

- The use of \$1 in the branch condition prevents add instruction (whose destination is \$1) from being moved after the branch.
- The branch delay slot is scheduled from the not-taken fall-through path.
- This strategy is preferred when the branch is not taken with high probability, such as **forward branches**.



# Delayed Branch Technique

---

- To make the optimization legal for the target an fall-through cases, it must be OK to execute the moved instruction when the branch goes in the unexpected direction.
- By OK we mean that the instruction in the branch delay slot is executed but the work is wasted (the program will still execute correctly).
- For example, if the destination register is an unused temporary register when the branch goes in the unexpected direction.

# Delayed Branch Technique

---

- In general, the compilers are able to fill about 50% of delayed branch slots with valid and useful instructions, the remaining slots are filled with **nops**.
- In deeply pipeline, the delayed branch is longer than one cycle: many slots must be filled for every branch, thus it is more difficult to fill all the slots with useful instructions.

# Delayed Branch Technique

---

- The main limitations on delayed branch scheduling arise from:
  - ▶ The restrictions on the instructions that can be scheduled in the delay slot.
  - ▶ The ability of the compiler to statically predict the outcome of the branch.



# Questions

---

