

Advanced Computer Architectures

(High Performance Processors and Systems)

Hardware Based Speculation

Politecnico di Milano

v1

Alessandro Verosimile <alessandro.verosimile@polimi.it>

Marco D. Santambrogio <marco.santambrogio@polimi.it>

EXE ON APRIL 9
LIVE @25.S.2

Outline

- HW-based speculation main idea
- Reorder Buffer
- An example

HW-based Speculation

- HW-based Speculation combines 3 ideas:
 - **Dynamic Branch Prediction** to choose which instruction to execute
 - **Dynamic Scheduling** supporting out-of-order execution but in-order commit to prevent any irrevocable actions (such as register update or taking exception) until an instruction commits
 - **Speculation** to execute instructions before control dependences are resolved

HW-based Speculation

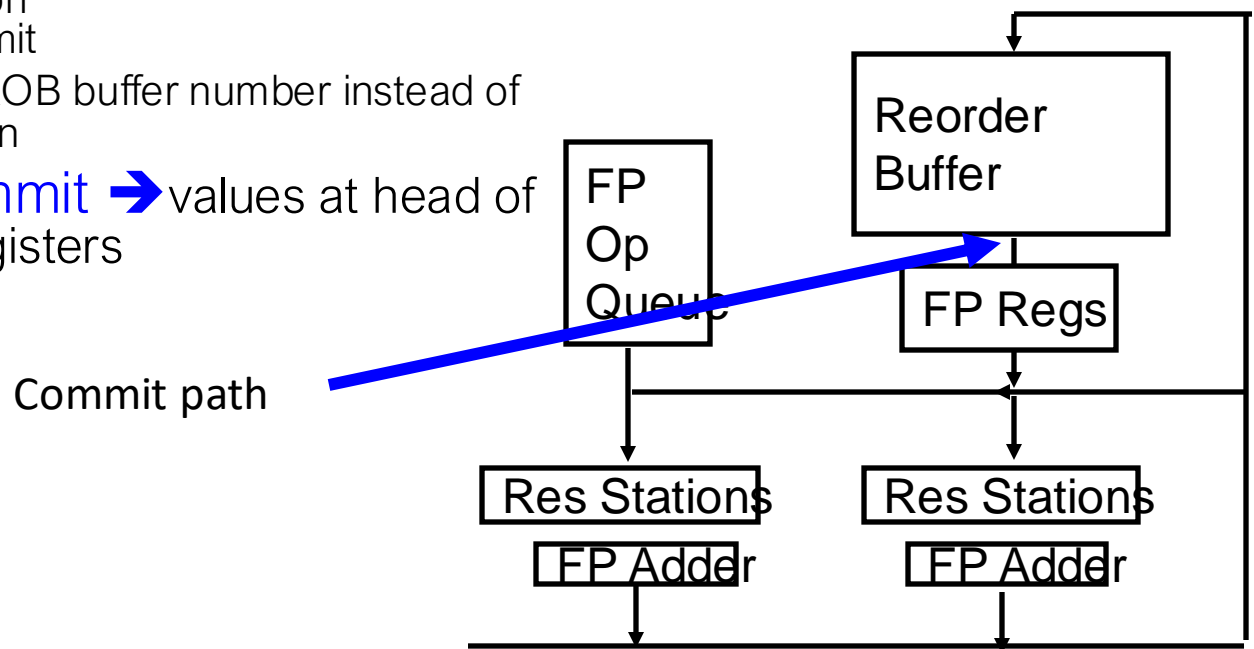
- The key idea behind speculation is:
 - to issue and execute instructions dependent on a branch before the branch outcome is known;
 - to allow instructions to execute out-of-order but to force them to commit in-order;
 - to prevent any irrevocable action (such as updating state or taking an exception) until an instruction commits;
- When an instruction is no longer speculative, we allow it to update the register file or memory (instruction commit).

HW-based Speculation

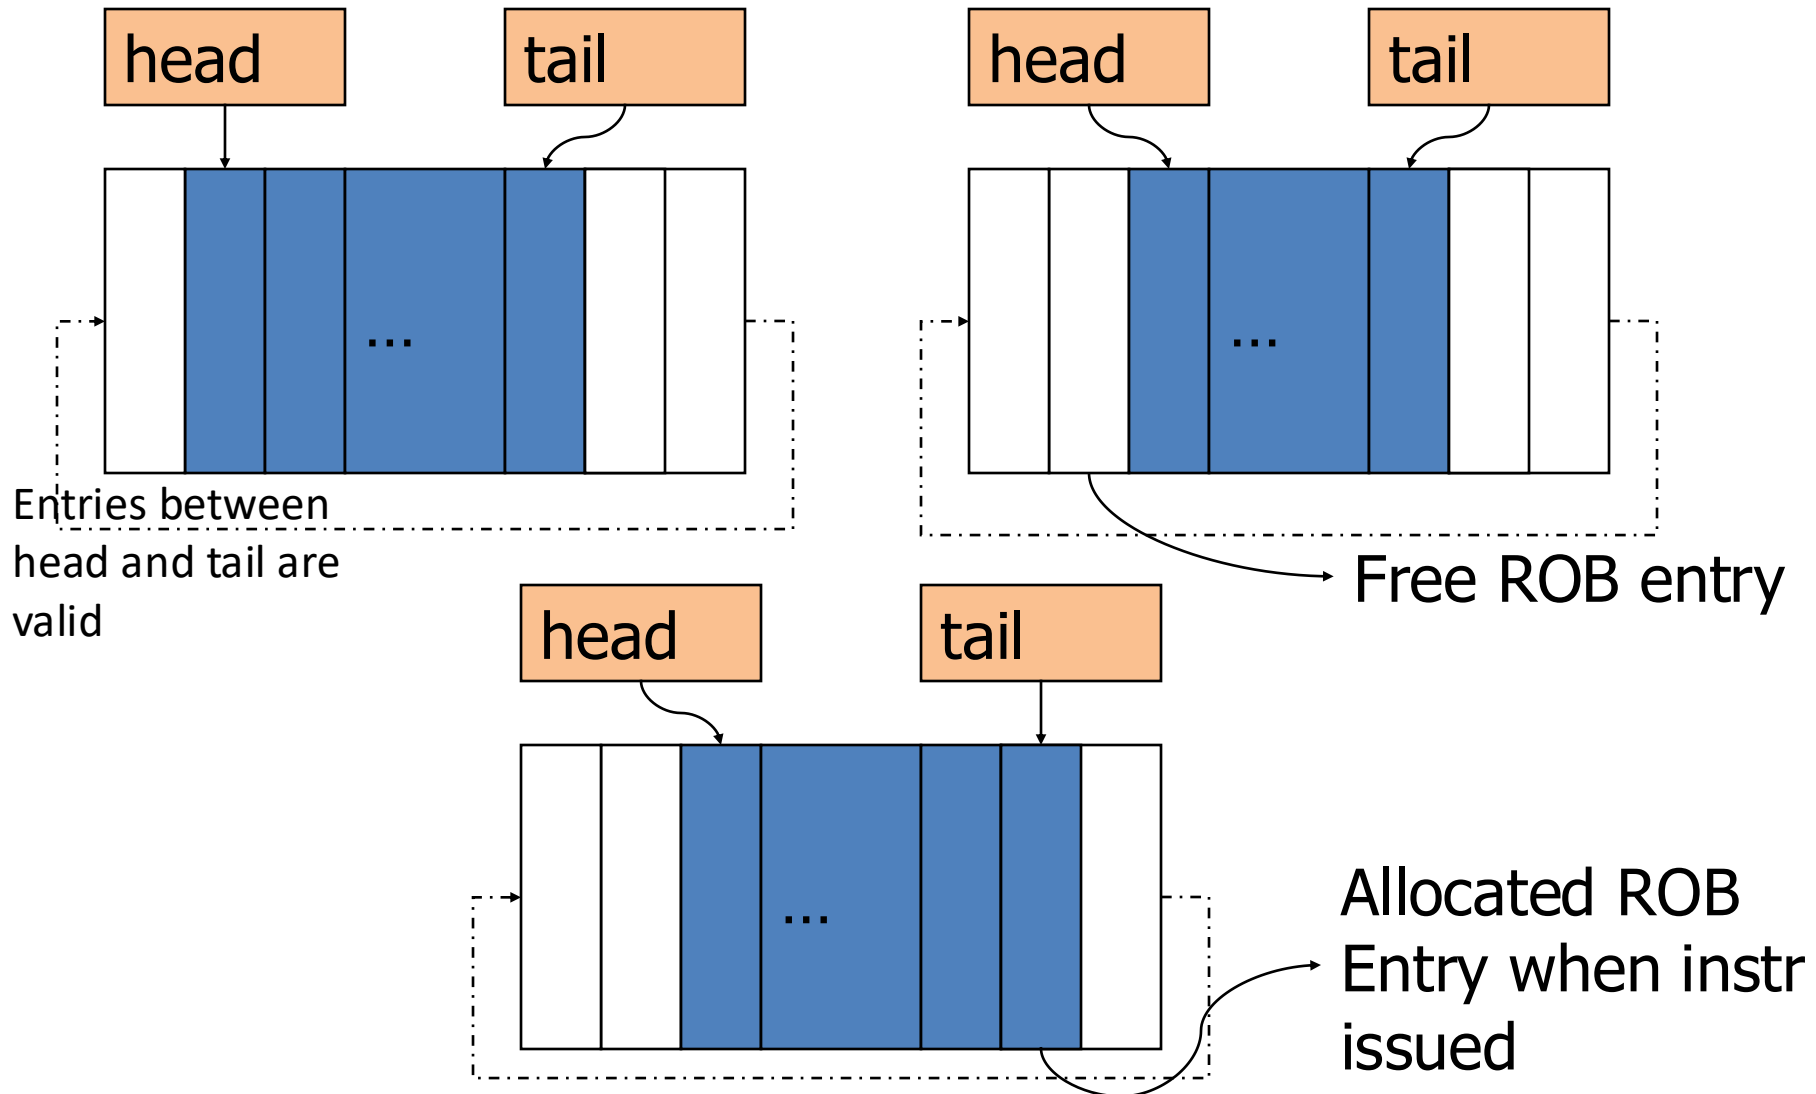
- Outcome of branches is speculated and program is executed as if speculation was correct (without speculation, the simple dynamic scheduling would only fetch and decode, not execute!)
- Mechanisms are necessary to handle incorrect speculation – hardware speculation extends dynamic scheduling beyond a branch (i.e. behind the basic block)

Reorder Buffer

- Holds instructions in FIFO order, exactly as issued
- When instructions complete, results placed into ROB
 - Supplies operands to other instruction between execution complete & commit
 - Tag results with ROB buffer number instead of reservation station
- Instructions **commit** → values at head of ROB placed in registers



ROB: Circular Buffer with Head/Tail Pointers



Separating Completion from Commit


- Re-order buffer holds register results from completion until commit
 - Entries allocated in program order during decode
 - Buffers completed values and exception state until in-order commit point
 - Completed values can be used by dependents before committed (bypassing)
 - Each entry holds program counter, instruction type, destination register specifier and value if any, and exception status (info often compressed to save hardware)
- Memory reordering needs special data structures
 - Speculative store address and data buffers
 - Speculative load address and data buffers

Relationship between precise interrupts and speculation

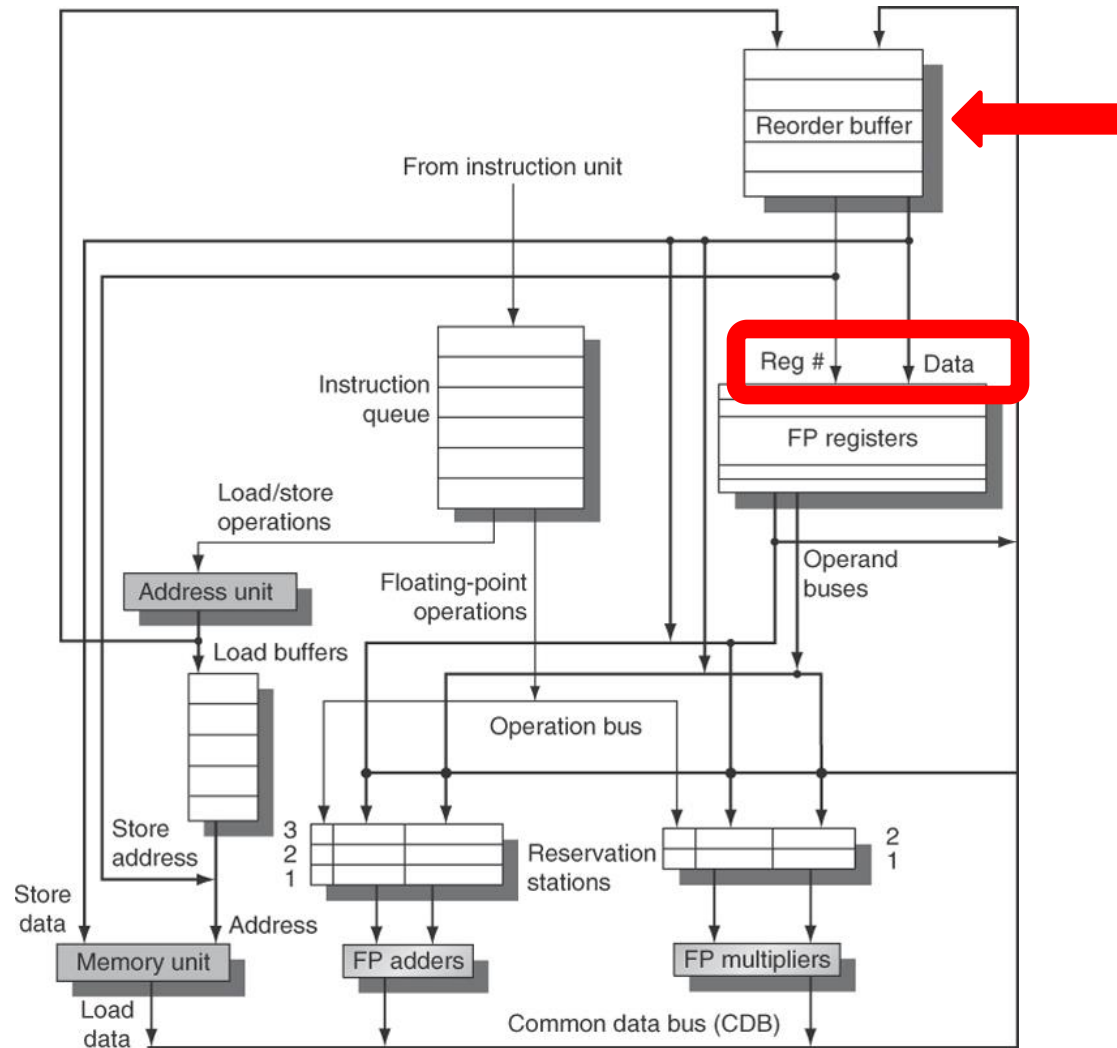
- Speculation: guess and check
- Important for branch prediction:
 - Need to “take our best shot” at predicting branch direction
- If we speculate and are wrong, need to back up and restart execution to point at which we predicted incorrectly:
 - This is exactly the same as precise exceptions!
- Technique for both precise interrupts/exceptions and speculation: *in-order completion or commit*

MORE ON **EXCEPTIONS/INTERUPTS** @ VIDEO **C9**

Tomasulo Algorithm for an FPU



Speculative Tomasulo Architecture for an FPU



Four Steps of Speculative Tomasulo Algorithm 1/2

1. Issue: get instruction from FP Op Queue

If reservation station **and reorder buffer slot** free, issue instr & send operands **& reorder buffer no. for destination** (this stage sometimes called “dispatch”)

2. Execution: operate on operands (EX)

When both operands ready then execute; if not ready, watch CDB for result; when both in reservation station, execute; checks RAW (sometimes called “issue”)

3. Write result: finish execution (WB)

Write on Common Data Bus to all awaiting FUs
& reorder buffer; mark reservation station available.

4. Commit: update register with reorder result

When instr. at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer. Mispredicted branch flushes reorder buffer (sometimes called “graduation”)

Steps of Speculative Tomasulo's Algorithm 2/2

4. **Commit**: 3 different possible sequences:
 1. **Normal commit**: instruction reaches the head of the ROB, result is present in the buffer. Result is stored in the register, instruction is removed from ROB;
 2. **Store commit**: as above, but memory rather than register is updated;
 3. **Instruction is a branch with incorrect prediction**: it indicates that speculation was wrong. ROB is flushed ("graduation"), execution restarts at correct successor of the branch.
If the branch was correctly predicted, branch is finished

Speculative Tomasulo's Algorithm

- Tomasulo's "Boosting" needs a buffer for uncommitted results (reorder buffer).

Each entry in ROB contains **four fields**:

- **Instruction type field** – indicates whether instruction is a branch (no destination result), a store (has memory address destination), or a load/ALU (register destination)
- **Destination field**: supplies register number (for loads and ALU instructions) or memory address (for stores) where results should be written;
- **Value field** (used to hold value of result until instruction commits)
- **Ready field**: indicates that instruction has completed execution, value is ready

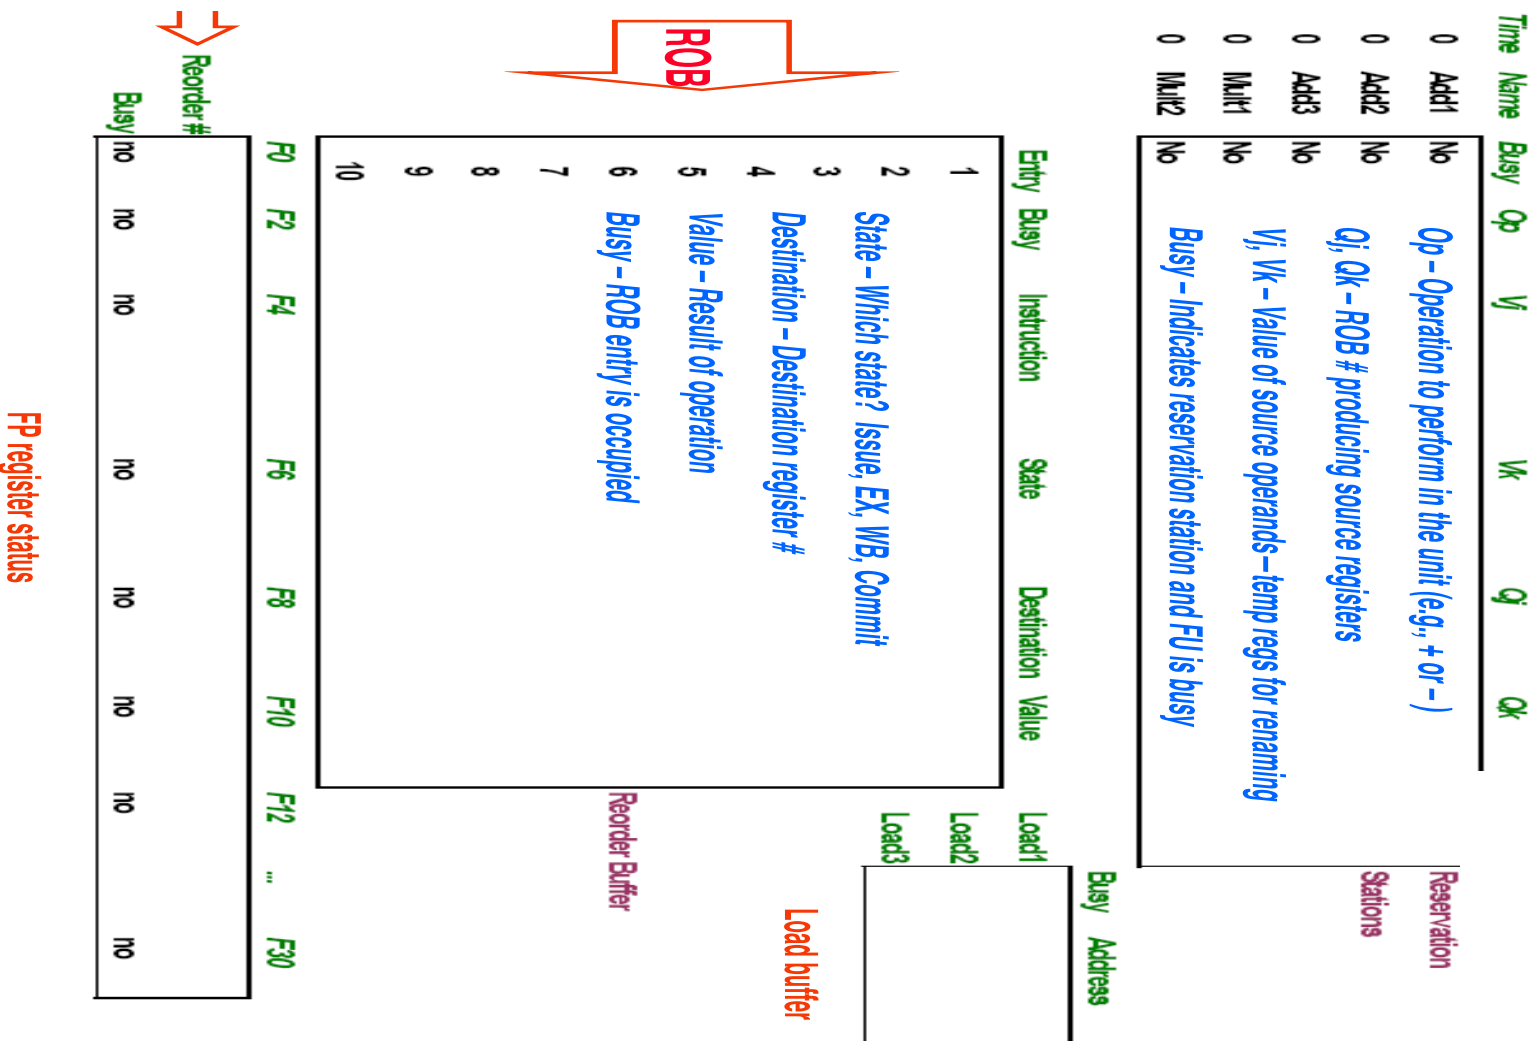
ReOrder Buffer Extension

- **ROB completely replaces store buffers:** stores execute in two steps, the second one when instruction commits;
- Renaming function of reservation stations completely replaced by ROB;
- Reservation stations now only queue operations (and operands) to FUs between the time they issue and the time they begin execution;
- Results are tagged with ROB entry number rather than with RS number \Rightarrow ROB entry assigned to instruction must be tracked in the reservation stations.

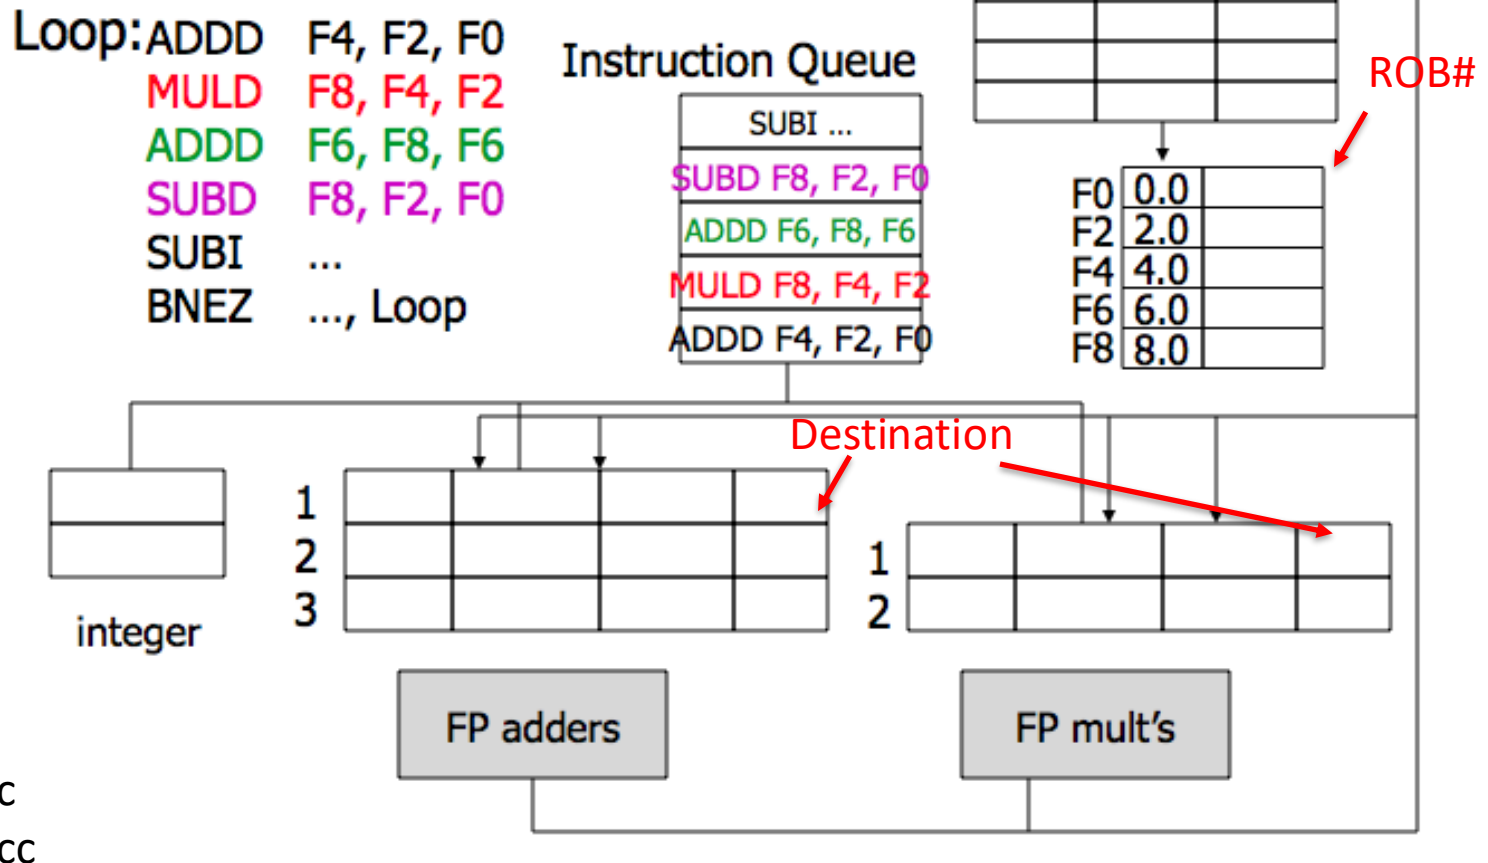
ReOrder Buffer Extension

- All instructions excluding incorrectly predicted branches (or incorrectly speculated loads) commit when reaching head of ROB;
- Incorrectly predicted branch reaches head of ROB
⇒ wrong speculation is indicated, ROB is flushed, execution restarts at correct successor of branch.
- Speculative actions are easily undone.
- Processors with ROB can dynamically execute while maintaining a precise interrupt model:
 - if instruction I_j causes interrupt, CPU waits until I_j reaches the head of the ROB and takes the interrupt, flushing all other pending instructions.

Tomasulo with ROB



ROB/Tomasulo – cycle 0



ADDD 3cc
 MULT 10cc

ROB/Tomasulo – cycle 1

ADDD 3cc
MULT 10cc

Loop:ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

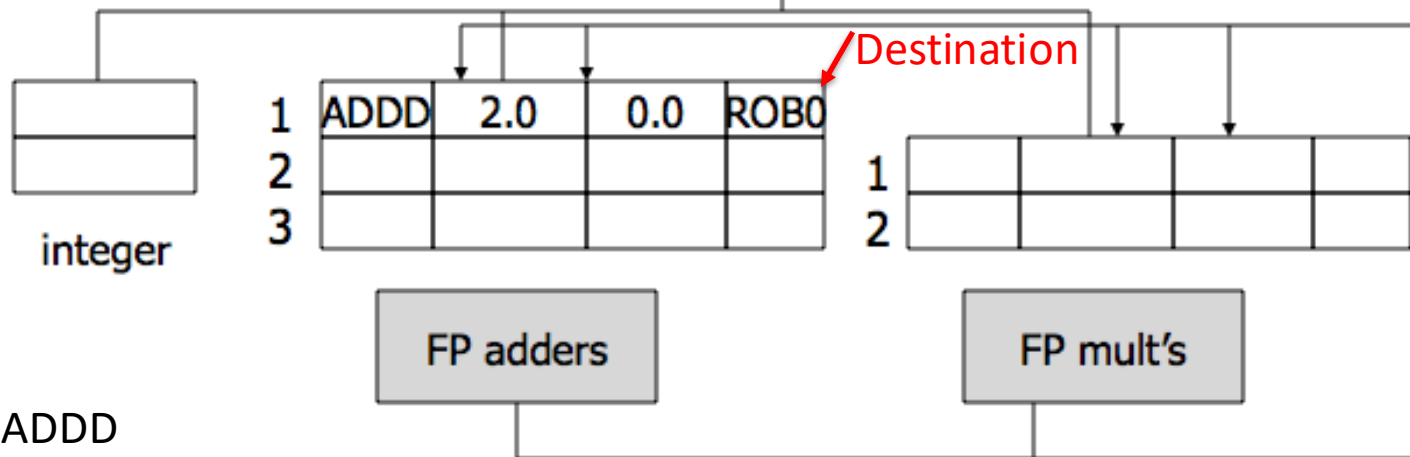
Instruction Queue

BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6
MULF F8, F4, F2

Destination Value

ROB			
0	ADDD	F4	-
1			
2			
3			
4			
5			
6			

F0	0.0	
F2	2.0	
F4	4.0	ROB0
F6	6.0	
F8	8.0	



ROB/Tomasulo – cycle 2

ADDD 3cc
MULT 10cc

Loop:ADDD F4, F2, F0
 MUL~~D~~ F8, F4, F2
 ADD~~D~~ F6, F8, F6
 SUB~~D~~ F8, F2, F0
 SUBI ...
 BNEZ ...

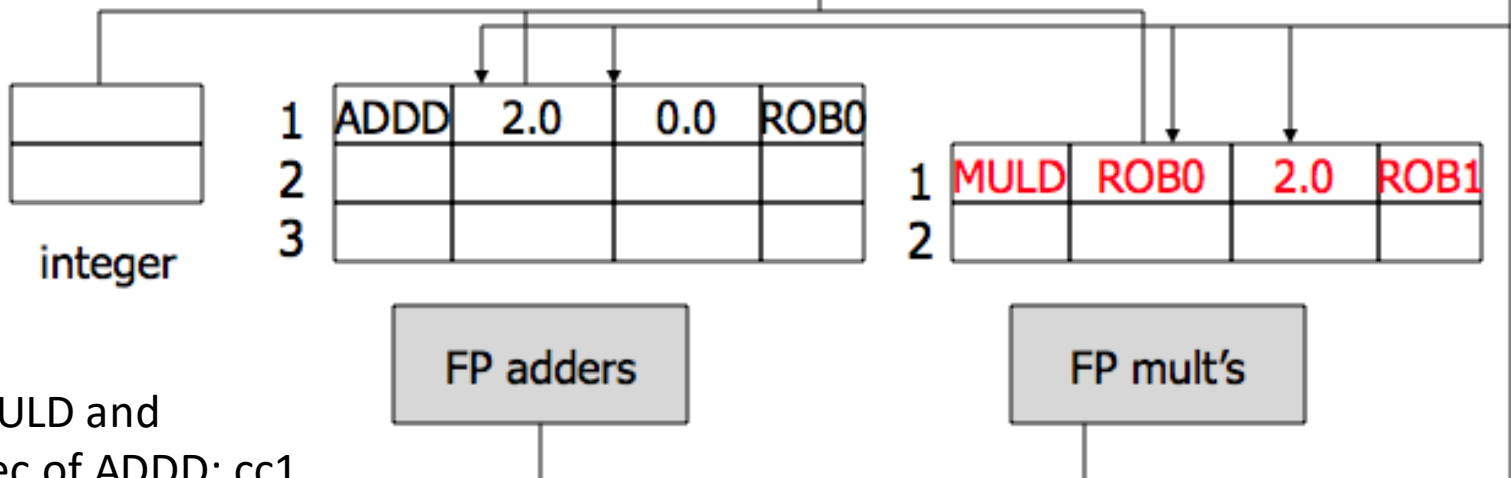
Instruction Queue

ADDD F4, F2, F0
BNEZ
SUBI
SUB D F8, F2, F0
ADD D F6, F8, F6

ROB

0	ADDD	F4	-	H
1	MUL D	F8	-	
2				
3				
4				
5				
6				

F0	0.0	
F2	2.0	
F4	4.0	ROB0
F6	6.0	
F8	8.0	ROB1



Issue MUL~~D~~ and
start exec of ADD~~D~~: cc1

ROB/Tomasulo – cycle 3

Loop: **ADDD** F4, F2, F0
MULD F8, F4, F2
ADDD F6, F8, F6
SUBD F8, F2, F0
SUBI ...
BNEZ ...

ADDD 3cc
 MULT 10cc

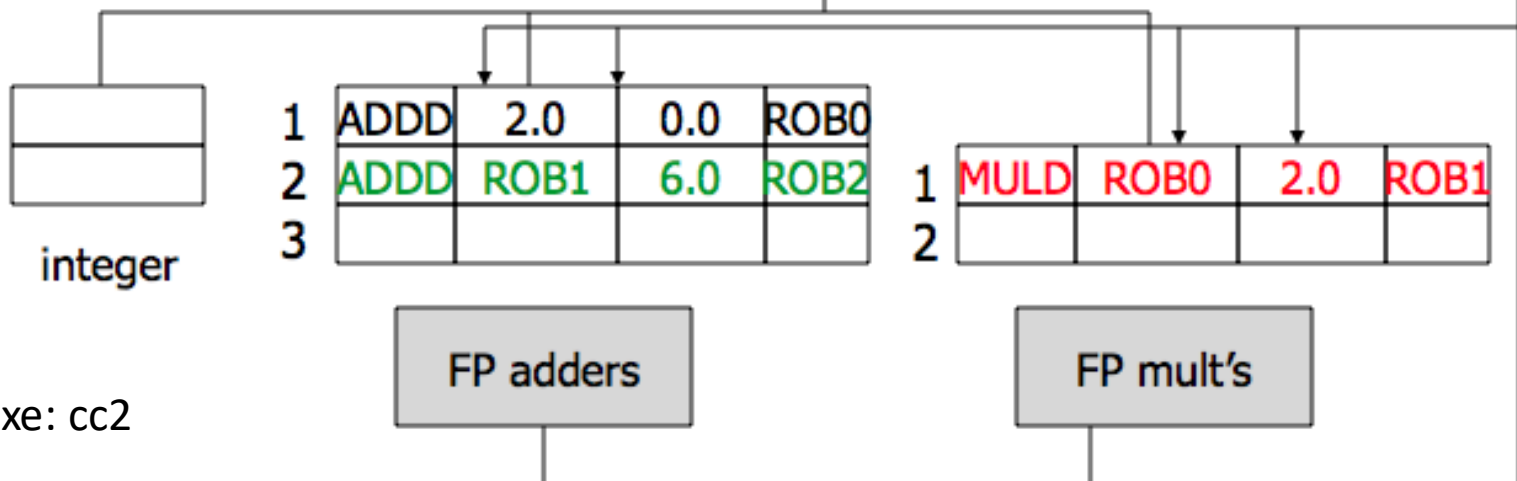
Instruction Queue

MULD F8, F4, F2
ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0

ROB

0	ADDD	F4	-	H
1	MULD	F8	-	
2	ADDD	F6	-	
3				
4				
5				
6				

F0	0.0	
F2	2.0	
F4	4.0	ROB0
F6	6.0	ROB2
F8	8.0	ROB1



ROB/Tomasulo – cycle 4

ADDD 3cc
MULT 10cc

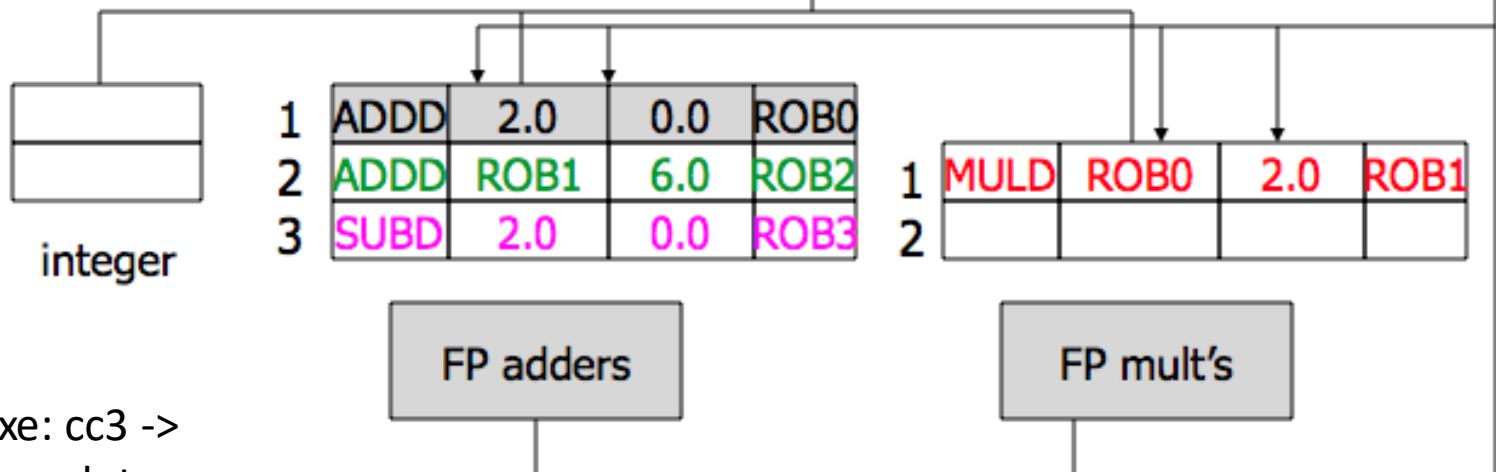
Loop:ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDF F6, F8, F6
 SUBF F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

ADDD F6, F8, F6
MULF F8, F4, F2
ADDD F4, F2, F0
BNEZ
SUBI

ROB			
0	ADDD	F4	-
1	MULF	F8	-
2	ADDD	F6	-
3	SUBF	F8	-
4			
5			
6			

F0	0.0	
F2	2.0	
F4	4.0	ROB0
F6	6.0	ROB2
F8	8.0	ROB3



ADDD exe: cc3 ->
ADDD completes

ROB/Tomasulo – cycle 5

Loop: ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

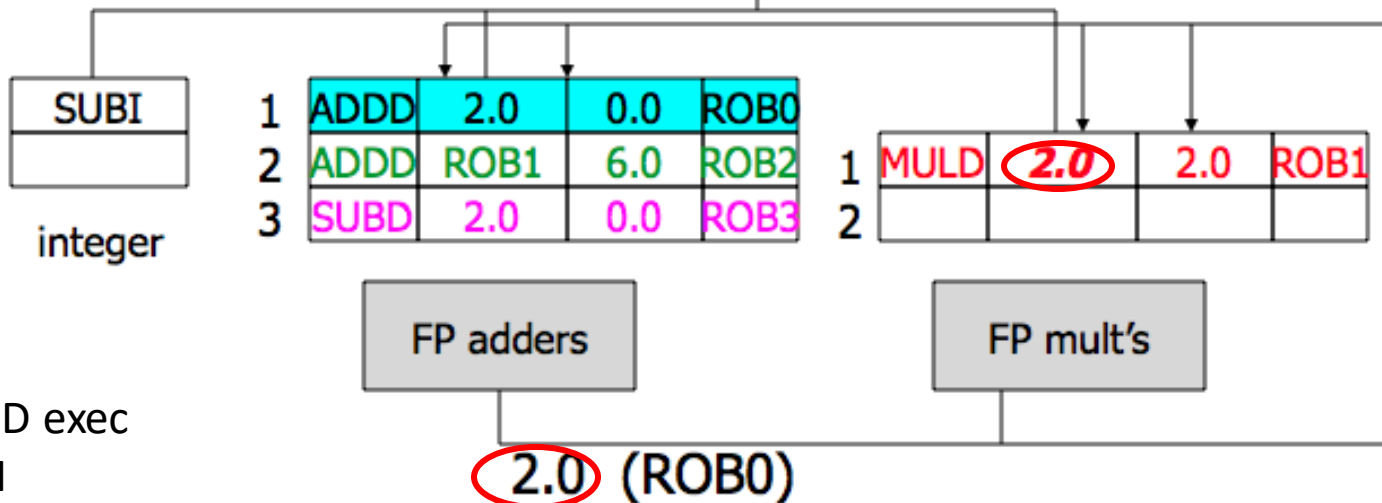
Instruction Queue

SUBD F8, F2, F0
ADDD F6, F8, F6
MULF F8, F4, F2
ADDD F4, F2, F0
BNEZ

ROB

0	ADDD	F4	2.0	H
1	MULF	F8	-	
2	ADDD	F6	-	
3	SUBD	F8	-	
4	SUBI			
5				
6				

F0	0.0	
F2	2.0	
F4	4.0	ROB0
F6	6.0	ROB2
F8	8.0	ROB3



Start SUBD exec

Exec SUBI

ROB/Tomasulo – cycle 6

Loop: ADDD F4, F2, F0
 MUL_D F8, F4, F2
 ADD_D F6, F8, F6
 SUB_D F8, F2, F0
 SUBI ...
 BNEZ ...

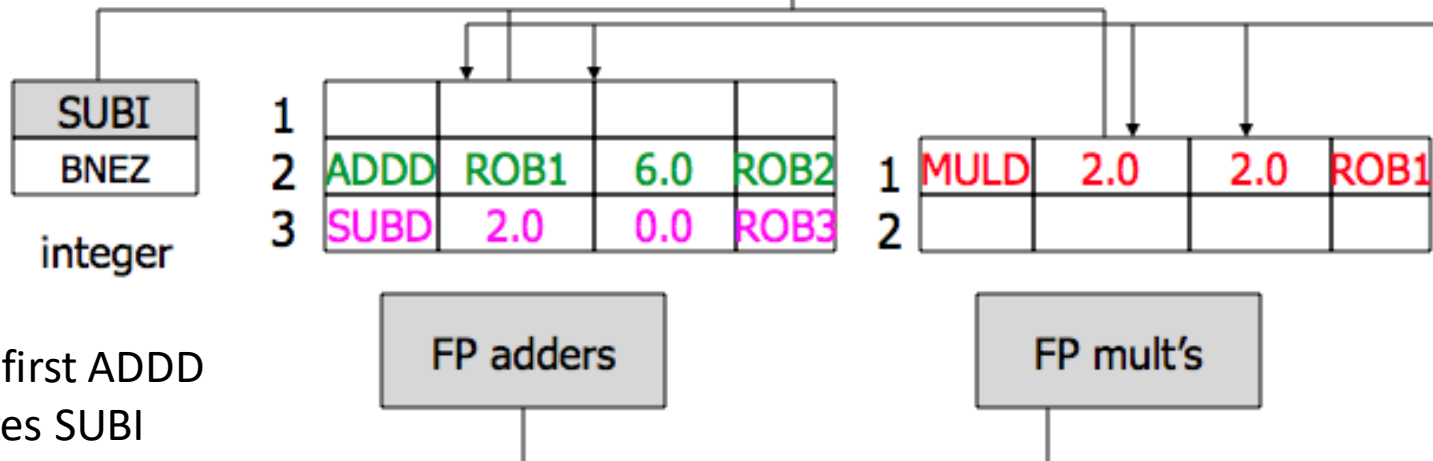
Instruction Queue

SUBI
SUB _D F8, F2, F0
ADD _D F6, F8, F6
MUL _D F8, F4, F2
ADD _D F4, F2, F0

ROB

0			
1	MUL _D	F8	-
2	ADD _D	F6	-
3	SUB _D	F8	-
4	SUBI		
5	BNEZ		
6			

F0	0.0	
F2	2.0	
F4	2.0	
F6	6.0	ROB2
F8	8.0	ROB3



Commit first ADD_D
 Completes SUBI
 Issue BNEZ

ROB/Tomasulo – cycle 7


Loop: ADDD F4, F2, F0
 MULD F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6
MULD F8, F4, F2

ROB		
0		
1	MULD	F8 -
2	ADDD	F6 -
3	SUBD	F8 -
4	SUBI	
5	BNEZ	
6	ADDD	F4

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB3



Issue second iteration ADDD
 Completes SUBD
 BNEZ not completed

ROB/Tomasulo – cycle 8

Loop: ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...


Instruction Queue

ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6

ROB

0	MULF	F8	-	H
1	MULF	F8	-	
2	ADDD	F6	-	
3	SUBD	F8	2.0	
4	SUBI			
5	BNEZ			
6	ADDD	F4		

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0



Issue second MULF
 Write SUBD
 BNEZ not completed

2.0 (ROB3)

ROB/Tomasulo – cycle 9

Loop: ADDD F4, F2, F0
 MULD F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6

ROB		
0	MULD	F8 -
1	MULD	F8 -
2	ADDD	F6 -
3	SUBD	F8 2.0
4	SUBI	
5	BNEZ	
6	ADDD	F4

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0

The 2nd **ADDD** iteration 2 cannot be issued!
 ROB is FULL!

ROB/Tomasulo – cycle 10

Loop: ADDD F4, F2, F0
 MULD F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

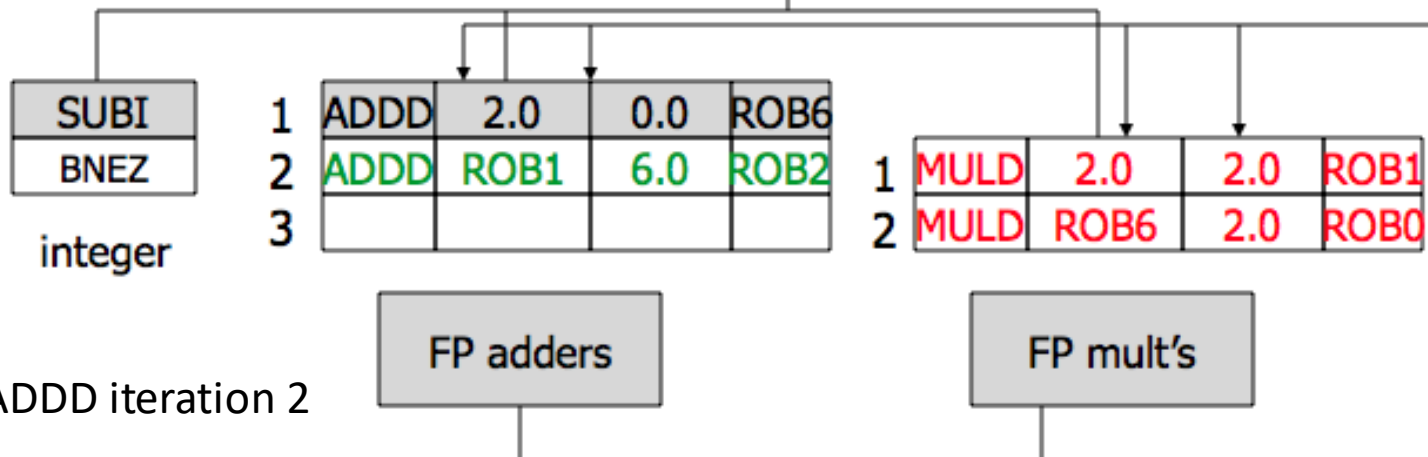
ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6

ROB

0	MULD	F8	-
1	MULD	F8	-
2	ADDD	F6	-
3	SUBD	F8	2.0
4	SUBI		
5	BNEZ		
6	ADDD	F4	

H

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0



Completes ADDD iteration 2

ROB/Tomasulo – cycle 11

Loop: ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

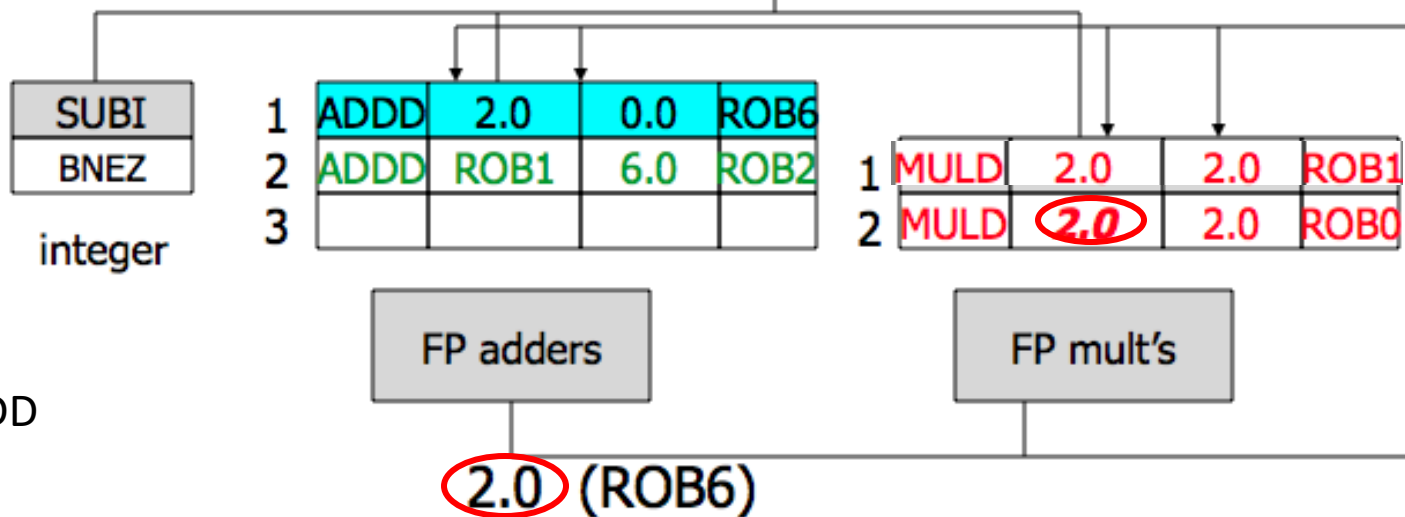
ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6

ROB

0	MULF	F8	-
1	MULF	F8	-
2	ADDD	F6	-
3	SUBD	F8	2.0
4	SUBI		
5	BNEZ		
6	ADDD	F4	2.0

H

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0



Write ADDD

ROB/Tomasulo – cycle 14

Loop: ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

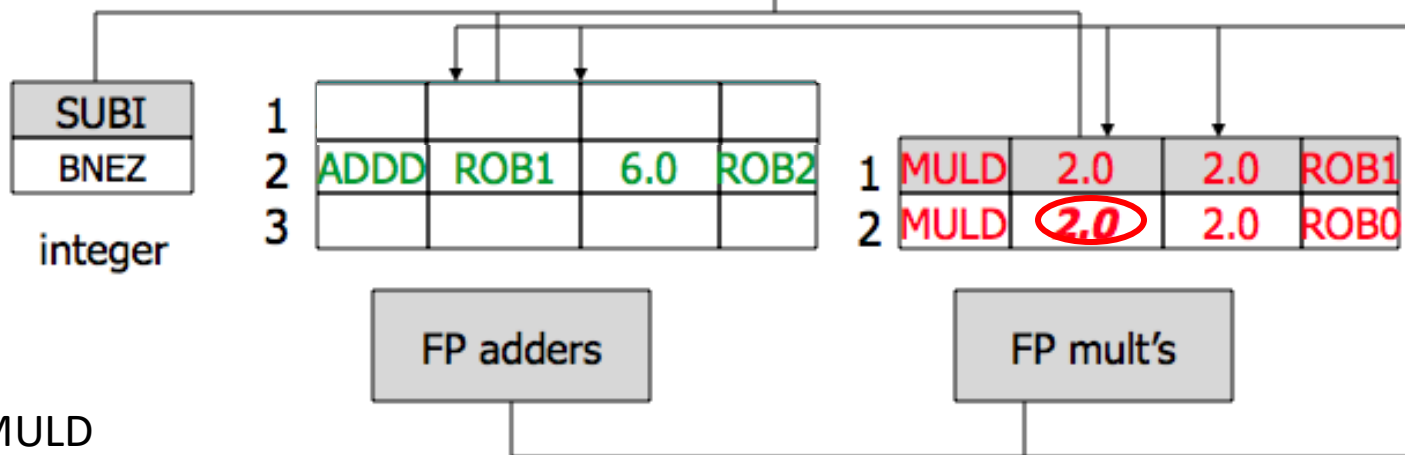
ADDD F4, F2, F0
BNEZ
SUBI
SUBD F8, F2, F0
ADDD F6, F8, F6

ROB

0	MULF	F8	-
1	MULF	F8	-
2	ADDD	F6	-
3	SUBD	F8	2.0
4	SUBI		
5	BNEZ		
6	ADDD	F4	2.0

H

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0



Completes MULF

ROB/Tomasulo – cycle 15

Loop: ADDD F4, F2, F0
 MUL_D F8, F4, F2
 ADD_D F6, F8, F6
 SUB_D F8, F2, F0
 SUBI ...
 BNEZ ...

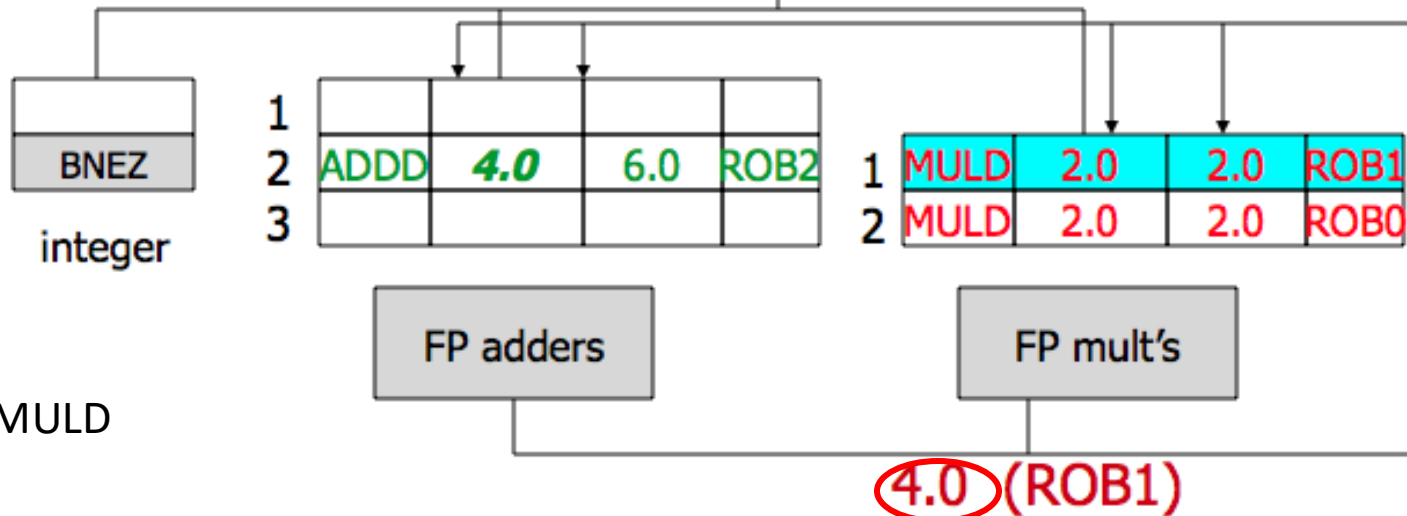
Instruction Queue

ADDD F4, F2, F0
BNEZ
SUBI
SUB _D F8, F2, F0
ADD _D F6, F8, F6

ROB

0	MUL _D	F8	-	
1	MUL _D	F8	4.0	H
2	ADD _D	F6	-	
3	SUB _D	F8	2.0	
4	SUBI		val	
5	BNEZ			
6	ADDD	F4	2.0	

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	8.0	ROB0



ROB/Tomasulo – cycle 16

Loop: ADDD F4, F2, F0
 MUL~~D~~ F8, F4, F2
 ADD~~D~~ F6, F8, F6
 SUB~~D~~ F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue


MUL D F8, F4, F2
ADDD F4, F2, F0
BNEZ
SUBI
SUB D F8, F2, F0

ROB

0	MUL D	F8	-
1	ADD D	F6	-
2	ADD D	F6	-
3	SUB D	F8	2.0
4	SUBI		val
5	BNEZ		
6	ADDD	F4	2.0

H

F0	0.0	
F2	2.0	
F4	2.0	ROB6
F6	6.0	ROB2
F8	4.0	ROB0



Need to undo 2nd iteration

ROB/Tomasulo – cycle 17

Loop: ADDD F4, F2, F0
 MULF F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

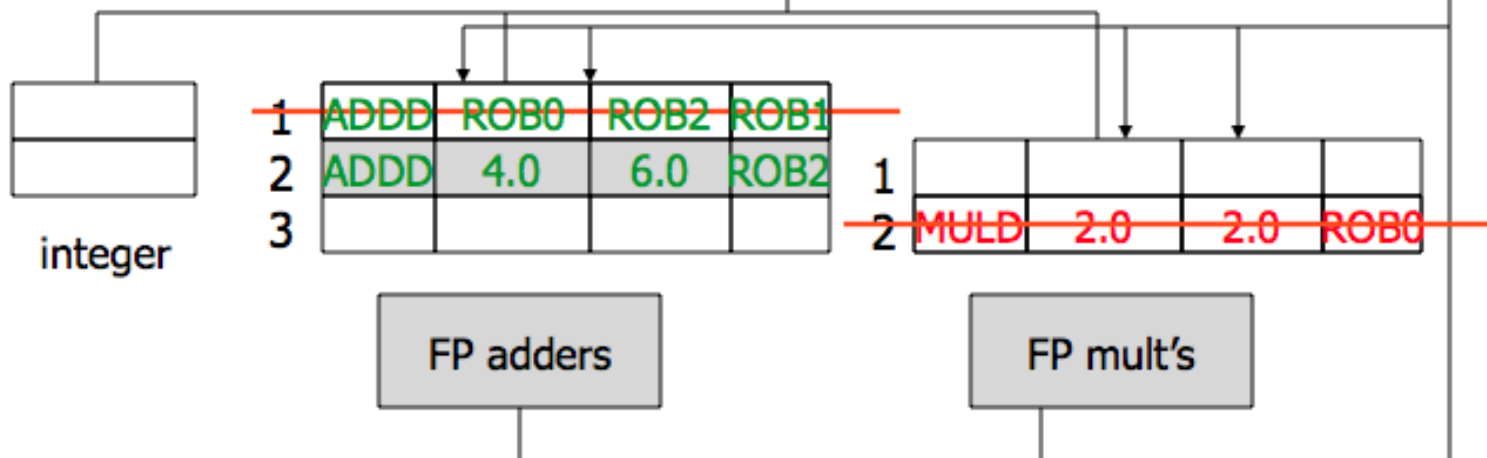
Instruction Queue

flushed
flushed
flushed
flushed
flushed

ROB		
0	flushed	
1	flushed	
2	ADDD	F6 -
3	SUBD	F8 2.0
4	SUBI	val
5	BNEZ	nt
6	flushed	

F0	0.0	
F2	2.0	
F4	2.0	
F6	6.0	ROB2
F8	4.0	ROB3


RST
entries
from
branch



ROB/Tomasulo – cycle 18

Loop: ADDD F4, F2, F0
 MULDD F8, F4, F2
 ADDD F6, F8, F6
 SUBD F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

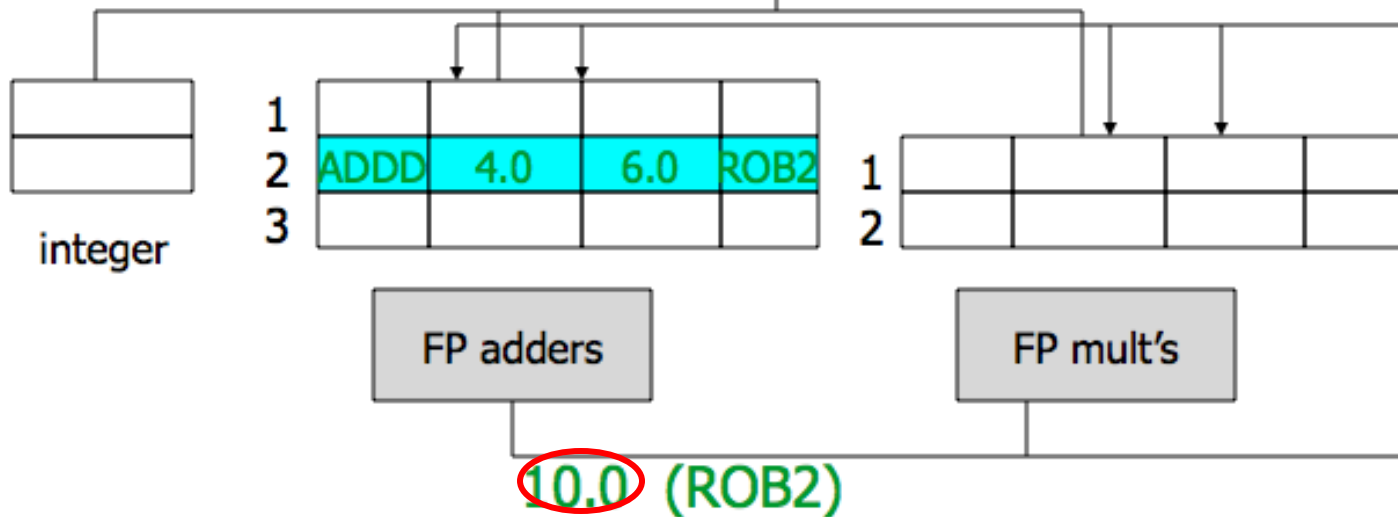


ROB

0			
1			
2	ADDD	F6	10.0
3	SUBD	F8	2.0
4	SUBI		val
5	BNEZ		nt
6			

H


F0	0.0	
F2	2.0	
F4	2.0	
F6	6.0	ROB2
F8	4.0	ROB3



ROB/Tomasulo – cycle 19

Loop:ADDD F4, F2, F0
 MUL~~D~~ F8, F4, F2
 ADD~~D~~ F6, F8, F6
 SUB~~D~~ F8, F2, F0
 SUBI ...
 BNEZ ...

Instruction Queue

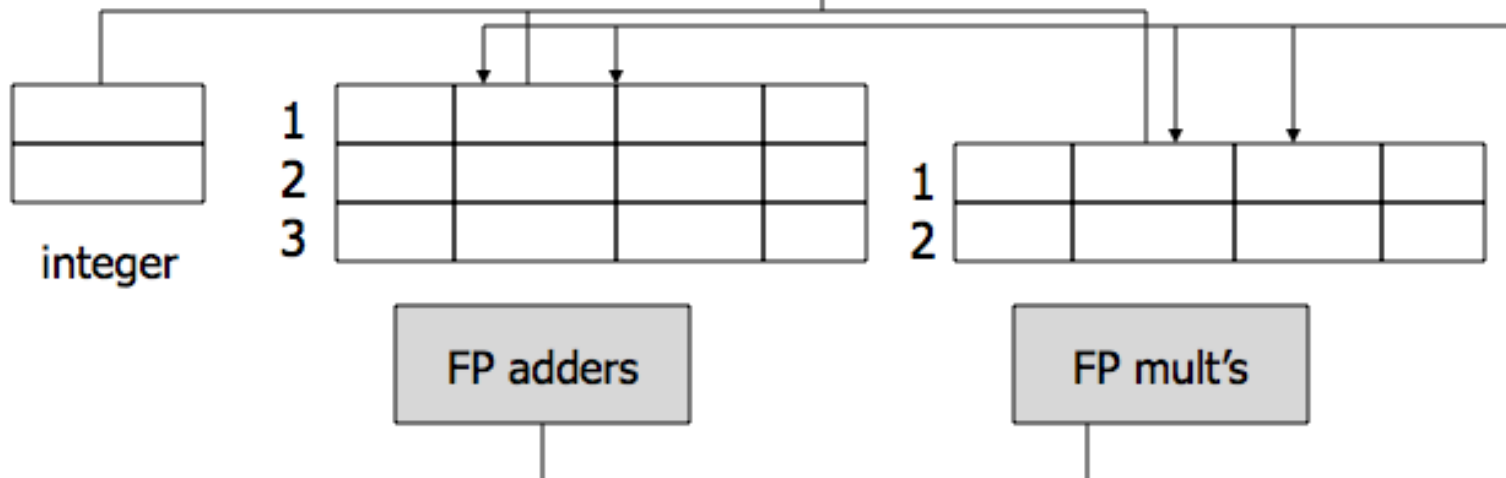


ROB

0			
1			
2			
3	SUBD	F8	2.0
4	SUBI		val
5	BNEZ		nt
6			

H


F0	0.0	
F2	2.0	
F4	2.0	
F6	10.0	
F8	4.0	ROB3



ROB/Tomasulo – cycle 20

Loop:ADDD F4, F2, F0
 MUL~~D~~ F8, F4, F2
 ADD~~D~~ F6, F8, F6
 SUB~~D~~ F8, F2, F0
 SUBI ...
 BNEZ ...

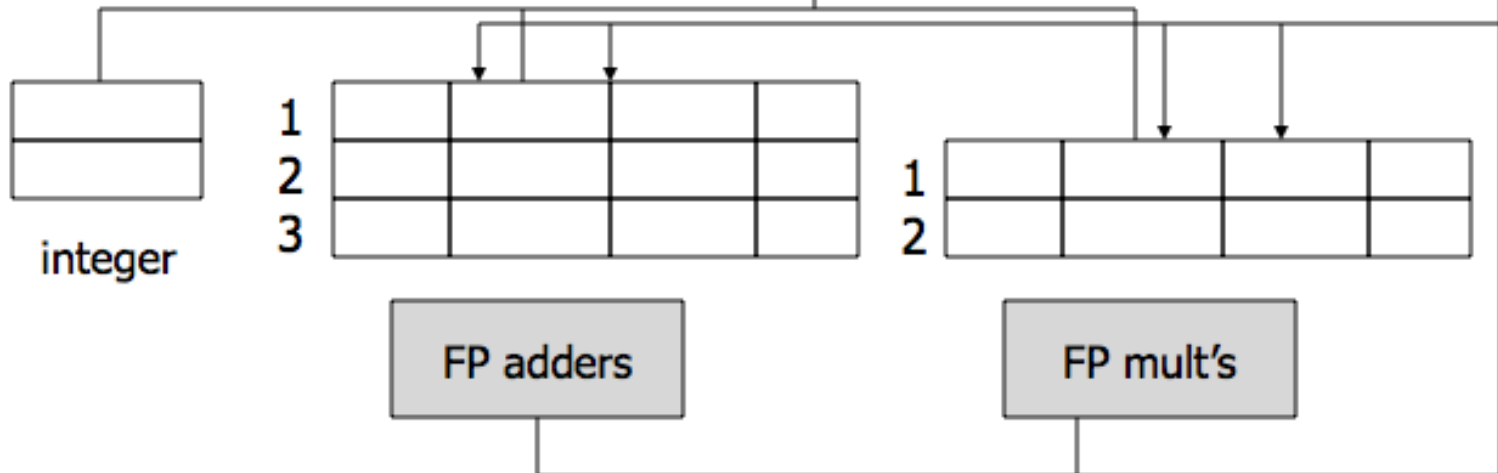
Instruction Queue



ROB		
0		
1		
2		
3		
4	SUBI	val
5	BNEZ	nt
6		

H

F0	0.0
F2	2.0
F4	2.0
F6	10.0
F8	2.0



Hardware Based Speculation

Politecnico di Milano

v1